

HS Data Science 26

Session 04: Scientific Working III — Sciometrics, Research Culture & Organisation

Prof. Dr. Michael Granitzer

19 May 2026

Session Overview

Sciometrics & Research Quality

- measuring science: h-index, impact factor
- ranking venues: CORE, Scimago, DBLP
- predatory publishers and how to spot them
- open access and reproducibility

Research Culture

- the replication crisis and what it means for ML
- open science principles
- research ethics and AI-specific concerns

Organising Yourself

- Getting Things Done (GTD) for researchers
- weekly review as a habit
- tools: Zotero, Obsidian, Git research diary
- time-blocking and deep work

Good science requires two things in equal measure: rigorous methods **and** sustainable work habits. Neither alone is enough.

Speaker notes

This session extends the research loop framework from S02/S03 with the external context: how is science measured and rewarded, and how do you organise yourself sustainably? Both topics are essential for PhD students and seminar participants doing serious empirical work.



Measuring Science — Sciometrics

Guiding Question: How is scientific impact measured — and why does it matter which venue you publish in?

The h-index

Source: Hirsch 2005; Garfield 1955

Definition

A researcher has h-index h if exactly h of their papers have been cited at least h times (Hirsch 2005).

Example

Paper	Citations
P1	42
P2	17
P3	9

Strengths

- combines productivity and impact in one number
- robust to one-off highly-cited outliers
- widely used by hiring committees and funding agencies

Limitations

- favours quantity over quality of individual papers
- heavily field-dependent (ML h-indices are high; humanities are low)



Paper	Citations
P4	6
P5	2

h-index =

4

(P1–P4 each have ≥ 4 cites)

Where to check

Google Scholar profile · Semantic Scholar ·
Publish or Perish (free)

- accumulates over time — disadvantages early-career researchers
- does not account for author order or self-citations

Never compare h-indices across disciplines. An h-index of 15 means very different things in physics vs. computer science vs. literary studies.

Speaker notes

Hirsch (2005) introduced the h-index. Today it is almost universally present in researcher profiles. Emphasise its field-dependence and the fact that it rewards sustained output over decades — so students should not feel behind.



Impact Factor & Journal Metrics

Impact Factor (IF)

Average citations per paper published in a journal over the last 2 years.

$$IF = \frac{\text{citations to papers from years } t-1, t-2}{\text{papers published in years } t-1, t-2}$$

Published annually by Clarivate (Journal Citation Reports).

CiteScore (Scopus)

Similar idea, but uses a 4-year window — somewhat more stable.

SJR (Scimago)

What it tells you

- a rough signal of journal prestige within a field
- useful when selecting where to submit

What it does NOT tell you

- quality of any individual paper
- relevance to your specific research question
- reproducibility or rigour of published work

For Computer Science

Conference proceedings often matter more than journals. Many top venues (NeurIPS, ICML, ACL, WWW) have no impact factor — rank instead by CORE ranking or community acceptance.

Impact Factor is a journal-level metric. It says nothing about whether a specific paper is worth

Weighted by prestige of the citing journal —
analogous to PageRank for journals.

reading.

Many students arrive believing IF is the ultimate quality signal. Correct this early. In CS and AI, a paper at NeurIPS or ICML (CORE A*) beats a mediocre journal article by any IF measure.

Ranking Research Venues — CORE

Source: Research & Australasia 2023; Beall 2012

CORE Conference Ranking

The Computing Research and Education Association (CORE) ranks CS conferences:

Rank	Meaning
A*	flagship, top tier worldwide
A	excellent, highly selective
B	good, respected
C	peer-reviewed, limited impact
Unranked	no quality signal

How to use it

When reading a paper: check venue rank to calibrate how much trust to extend before diving in.

When submitting: target A*/A venues even if rejection is likely — reviewers' feedback is educational and forces quality.

Warning: Predatory Venues

Signs of a predatory conference or journal:

- unsolicited invitation email promising fast review
- vague scope ("multidisciplinary", "all sciences")
- no indexing in DBLP, ACM DL, IEEE Xplore, Scopus
- excessive article processing charges with no clear affiliation
- check at: **beallslist.net**



Check at:

portal.core.edu.au/conf-ranks

For journals:

Scimago (SJR) Q1–Q4 quartiles per discipline

Publishing in a predatory venue harms your reputation and wastes time. Always verify before submitting.

Show students the CORE portal live if possible. Demonstrate checking a known good conference (NeurIPS: A*) and a suspicious-sounding one they might encounter. The predatory conference problem is severe in some sub-fields.

Open Science & Reproducibility

The Replication Crisis

Many published results — particularly in psychology and medicine — fail to replicate when experiments are repeated. ML is not immune:

- code not released → results unverifiable
- hyperparameter tuning hidden → baselines cherry-picked
- dataset leakage → inflated metrics
- random seeds not reported → variance hidden

Open Science Principles (UNESCO 2021)

1. Open data — share datasets where ethically possible



Practical Steps for AI/ML Papers

- release code on GitHub with a clear README
- report all hyperparameters and hardware specs
- use fixed random seeds, report variance across multiple runs
- publish to arXiv pre-print before or alongside journal submission
- use **Papers With Code** to link results to implementations
- use model cards and datasheets for datasets (Mitchell et al., 2019)

Reproducibility Checklists

NeurIPS and ICML now require reproducibility statements. Read them before writing — they are a free checklist of what reviewers will check.

2. Open code — publish reproducible implementations
3. Open methods — describe setup completely
4. Open peer review — transparency in the review process

Open, reproducible science is increasingly a requirement, not a bonus. Build the habit early.

Speaker notes

The Joelle Pineau NeurIPS reproducibility checklist (2018 onward) is worth examining. Discuss the "stochastic parrots" and LLM reproducibility challenges — especially relevant for students working on generative AI.



Research Culture & Ethics

Research Ethics in AI

Fundamental Principles

- **Honesty:** report results as they are, including negative results
- **Transparency:** describe methods fully enough for others to replicate
- **Accountability:** acknowledge limitations and failure modes
- **Non-maleficence:** consider harms of the research and its applications

Specific AI Concerns

Academic Integrity

Plagiarism:

using others' text or ideas without attribution — zero tolerance, automated detection (iThenticate, Turnitin).

Fabrication / falsification:

inventing or manipulating data — career-ending, potentially criminal.

Authorship:

every named author must have made a substantial intellectual contribution. Gift authorship and ghost authorship are both violations.

AI-assisted writing:

LLMs may be used as a writing aid, but you are responsible for every claim. Disclose use per



venue policy.

- bias in training data propagates to model outputs
- benchmark saturation creates false progress signals
- dual-use: the same model can help and harm
- environmental cost: large model training emits significant CO₂

When in doubt: cite, disclose, and describe. The cost of over-attribution is zero; the cost of under-attribution can end a career.

Discuss the 2023 ACL and NeurIPS policies on LLM use in paper writing as a concrete, recent example. Ask students: "If an LLM writes a paragraph that turns out to be factually wrong, who is responsible?" Answer: the authors.

Organising Yourself — GTD for Researchers

Guiding Question: How do you manage a research project that has no clear "done" and spans months?

Getting Things Done (GTD) — Core Workflow

Source: Allen 2001

The Five Steps

1. **Capture** — collect every commitment, idea, task, and worry into a trusted inbox (notebook, phone, email folder)
2. **Clarify** — process each item: is it actionable? what is the next physical action?
3. **Organise** — put it in the right bucket: next-actions
≡ list, project list, waiting-for, someday/maybe,

The Research-Specific Buckets

Bucket	Examples
Next Actions	"read Smith 2024 abstract", "run ablation on seed 42"
Projects	"seminar paper on RAG", "reproduce BEIR baseline"
Waiting For	"supervisor feedback on draft §3", "GPU cluster queue"
Someday/Maybe	"explore diffusion models for tabular data"
Reference	Zotero library, Obsidian notes, Git repo

calendar

4. **Reflect** — weekly review: scan all lists, empty inbox, check project progress
5. **Engage** — choose what to work on based on context, energy, and priority

Your brain is for thinking, not for storing. Capture everything externally and trust your system.

GTD (Allen, 2001) is especially powerful in research because the work is open-ended — there is no clear "done". The weekly review is the key habit: without it the system collapses within a week. Recommend 30-45 minutes every Friday afternoon.

The Weekly Review

Why it matters: Research has no natural stopping point. Without a regular review, tasks accumulate invisibly and projects stall.

Weekly Review Checklist

- ☐ Process and empty inbox (email, notes, paper pile)
- ☐ Review last week's calendar — anything that needs follow-up?
- ☐ Check next-actions lists — mark done, identify stuck items
- ☐ Review all active projects — does each have a  clear next action?

Time and place

- 30–45 minutes, same time every week
- Friday afternoon is ideal (fresh week starts clean on Monday)
- physical separation helps: close all tabs first

Tools that support this

Purpose	Tool
Task management	Obsidian Tasks, Notion, plain text
Literature	Zotero (with Obsidian plugin)
Code & experiments	Git (commit = natural checkpoint)
Notes	Obsidian, Logseq, Roam

- ☐ Review waiting-for list — send follow-up nudges if overdue
- ☐ Review someday/maybe — promote anything now relevant
- ☐ Check upcoming calendar — prepare for commitments
- ☐ Identify the one most important thing for next week

Purpose	Tool
Calendar	Google Calendar (block deep work)

The weekly review is a habit, not a method. It takes four weeks to feel natural. Do not skip it.

Speaker notes

Students often resist the weekly review because it feels like overhead. Reframe it: the review is when you make all difficult priority decisions in one go, so that during the week you can focus entirely on doing without deciding.



Deep Work and Research Time

Source: Newport 2016

Cal Newport's Deep Work Framework

- **Deep work:** cognitively demanding tasks that push the limits of ability — writing, proving, designing experiments
- **Shallow work:** logistical tasks that can be done while distracted — email, administrative forms, formatting references

Research value comes almost exclusively from deep work.

Key principles

Research Rhythms

Block	Example
Morning deep work	writing, analysis, proofs
Mid-day shallow	email, Slack, administrative
Afternoon deep work	coding, experiments
Evening	reading papers (lighter cognitive load)

Protecting Deep Work

- batch email to two fixed times per day
- say no to low-value meetings

1. Schedule deep work blocks (2–4 hours minimum) in the calendar
2. Eliminate distractions during deep work: phone off, notifications off
3. Measure deep work hours per week as a leading indicator
4. Rest deeply between blocks — shallow work in between is fine

- use "office hours" concept for collaboration requests
- log deep work hours (even a tally mark) — what gets measured improves

Most researchers get 1–2 hours of genuine deep work per day. Aim for 4–6. That gap is where careers diverge.

Speaker notes

Newport's "Deep Work" (2016) is worth recommending as supplementary reading. The key practical message: calendar blocking is more reliable than willpower. If deep work is not in the calendar, it will not happen.



Research Tools Stack

Literature Management — Zotero

- free, open-source reference manager
- browser plugin captures papers in one click
- syncs PDF highlights and annotations
- exports BibTeX directly for LaTeX
- Zotero Connector + Better BibTeX = seamless workflow

Note-Taking — Obsidian

- plain Markdown files, no lock-in
- bidirectional links: paper notes link to concept notes link to project notes
- Zotero plugin: pulls metadata and annotation highlights automatically
- graph view reveals connection patterns in your reading

The literature note template



Code & Experiments — Git as Research Diary

Every commit message is a lab notebook entry:

```
feat: add BM25 baseline on MSMARCO dev split
```

```
Scores: MRR@10 = 0.187, Recall@100 = 0.857  
Same as reported in Thakur et al. 2021 Table 3.  
Confirms our reimplementation is correct.  
Next: add dense retrieval (DPR).
```

ML Project Structure

```
project/  
├── data/           # raw (gitignored), processed  
├── notebooks/      # exploration only, not production  
├── src/            # importable modules  
├── experiments/    # one folder per experiment run  
│   └── 2026-05-12-bm25-baseline/  
│       ├── config.yaml  
│       ├── results.json  
│       └── README.md
```

Citation
Author (year). Title. Venue.

Core Claim
One sentence.

Evidence
How did they test it?

Limitations & Gap
What does this leave open?

Connection to My Work

tests/
README.md

Your future self is a stranger. Write notes and
commit messages for them.

Speaker notes

Walk through a live Zotero import if possible. The key insight is that the note-taking system is only valuable if it is used consistently — a perfect system used inconsistently is worse than a simple system used every day.



Summary

Sciometrics

- h-index = personal impact indicator (field-dependent!)
- Impact Factor = journal-level metric (not paper-level)
- CORE ranking = venue quality for CS conferences
- Predatory venues: always verify before submitting
- Open science: share code, data, methods

Research Ethics

- honesty, transparency, accountability
- reproducibility is a professional obligation
- AI-specific: bias, dual-use, environmental cost
- plagiarism, fabrication, authorship norms



Organisation

- GTD: Capture → Clarify → Organise → Reflect → Engage
- Weekly review: 30 min every Friday, non-negotiable
- Deep work: 4–6 hours/day is the goal
- Tools: Zotero + Obsidian + Git = complete research stack

The researcher's compact: Measure your impact without being ruled by it. Organise your work so that deep thinking gets protected time. Open your results so others can build on them.

Close by asking students: what is the one organisational change they will make this week? Capture commitments. Follow up next session.

Image Credits & References

Allen, David (2001). *Getting Things Done: The Art of Stress-Free Productivity*. Viking.

Beall, Jeffrey (2012). Predatory Publishers Are Corrupting Open Access. *Nature*.

Computing Research and Education Association of Australasia (2023). CORE Conference Ranking Portal. <https://portal.core.edu.au/conf-ranks/>

Garfield, Eugene (1955). Citation Indexes for Science: A New Dimension in Documentation Through Association of Ideas. *Science*.

Hirsch, Jorge E. (2005). An Index to Quantify an Individual's Scientific Research Output. *Proceedings of the National Academy of Sciences*.

Newport, Cal (2016). *Deep Work: Rules for Focused Success in a Distracted World*. Grand Central Publishing.