

Data Science

Lecture 01: Introduction

Prof. Dr. Michael Granitzer

Converted from LNKB-tex part-en-ds-chp1-introduction.tex. Examples and framing updated from the 2022 "big data" era to the current one; the argumentative spine (data does not speak for itself) is unchanged and is what the rest of the course rests on. No prerequisites.



Module 1

The Data Deluge

Guiding Question: *We generate more data than we can store. Does having more of it make us know more?*

How Much Data Exists?

Year	Global datasphere	Per person on earth
2002	5 exabytes of new stored information (Lyman & Varian 2003)	≈ 0.8 GB
2010	≈ 1.2 zettabytes (Gantz & Reinsel 2010)	≈ 170 GB
2020	≈ 64 ZB; IDC's 2012 <i>forecast</i> had said 40 ZB (Gantz & Reinsel 2012)	≈ 8,000 GB
2025	≈ 175 ZB — a 2018 <i>forecast</i> (Reinsel et al. 2018)	≈ 21,000 GB

In 2007 the amount of information produced first **surpassed available storage capacity**. Since then, most data that is generated is never stored.

The exact figures matter less than the two structural facts: growth is exponential, and generation outruns storage — so *selection* is unavoidable and someone is choosing. Ask who decides what gets kept; that is a data-science question before it is an engineering one.

The rows are not one measure: 2002 counts newly *stored* information (Lyman & Varian, Berkeley), the later rows IDC's "digital universe" of data created or replicated. 2020's 64 ZB is IDC's later estimate; its 2012 forecast had said 40 ZB, and 175 ZB for 2025 is still a forecast. Per person: 1.2 ZB / 6.9 bn $\approx$ 170 GB; 64 ZB / 7.8 bn $\approx$ 8,200 GB; 175 ZB / 8.1 bn $\approx$ 21,600 GB. The 2007 crossover is IDC's claim (Gantz et al. 2008).

Growth Is Not the Interesting Part

Volume

More bytes than anyone can read, index or curate.

Heterogeneity

Text, sensors, images, logs, transactions — rarely in a common schema.

Velocity

Streams that must be acted on before they can be stored.

The scarce resource is not data. It is **the ability to turn data into justified conclusions** — which is what this course is about.

Deliberately deflate the "big data" framing (P2): volume is a logistics problem. The course's subject is the inferential problem, which does not get easier with more rows — and sometimes gets harder, as the next unit shows.

Science at Scale

Large Hadron Collider

1 MB per event, 600 million events per second — roughly **18 zettabytes per year** if everything were kept.

A single tier-1 centre stores ~2.2 petabytes per year.

The filter is the experiment

Almost all events are **discarded in hardware**, in real time, by trigger rules written in advance.

What is kept is decided by a theory of what matters.

The LHC is the sharpest possible illustration that data is never raw: the trigger encodes physics theory, so the stored dataset is already theory-laden. This directly sets up the "end of theory" argument next.

Data at Web Scale

Search

Web-scale indexing processes petabytes per day and answers in milliseconds.

Platforms

Interaction logs — clicks, dwell time, scrolls — dwarf the content itself.

Models

Training corpora for language models are measured in trillions of tokens.

In every case the value comes from **aggregation across many weak signals**, not from any single record.

Data is never raw: at the LHC a physics theory decides what is recorded, and on the web value comes from aggregating signals that are individually worthless. **Someone always chose what the dataset contains.**

Speaker notes

Ask what a single click is worth versus a billion clicks — the answer motivates statistics as the discipline of extracting signal from individually uninformative observations. The takeaway is the bridge to Module 2.



Module 2

Does Data Speak for Itself?

Guiding Question: *Twenty-nine teams analysed the same dataset and the same question. Why did they not get the same answer?*

"The End of Theory"

With enough data, the numbers speak for themselves. [...] Correlation supersedes causation, and science can advance even without coherent models, unified theories, or really any mechanistic explanation at all.

Chris Anderson, *Wired*, 2008 (Anderson 2008). The argument has returned, almost verbatim, in claims that sufficiently large models make domain theory unnecessary.

Is it right?

Guiding question of the module (P1). Do not answer yet. Ask for a show of hands and hold both camps — the next four slides supply the evidence, and students should feel the argument's pull before it is dismantled.



One Dataset, Many Answers

Twenty-nine research teams received **the same dataset** and the same question:

Are football referees more likely to give red cards to dark-skinned players than to light-skinned players?

The result

Odds ratios ranged from **0.89 to 2.93**.
Twenty teams found a significant effect;
nine did not.

The reason

Not mistakes: the spread persisted among
analyses peer-rated as sound. Teams
made **different defensible choices** —
covariates, exclusions, model.

The data did not determine the answer. The **analysis** did. (Silberzahn et al. 2018)

Speaker notes

The most important slide of the lecture. Emphasise that this is not a story about incompetence or fraud — it is what "researcher degrees of freedom" means. Ask what the honest reporting practice would be (pre-registration; publishing the specification curve), and note that this is exactly what makes P3 (epistemic modelling) a professional obligation rather than a philosophical nicety.



Correlation Is Not Causation

Spurious correlations are easy to find and get **easier as datasets grow wider**: with enough variables, some pair will track another closely by chance alone.

More columns: chance

Testing many pairs makes an extreme correlation likely *somewhere*, even when every pair is independent.

More rows: no cure for confounding

A confounded correlation does not go away with more rows. The interval shrinks around a real but **non-causal** number: more precise, not more valid.

Speaker notes

The two boxes are different failures. Chance correlations grow with the number of *variables* and fade with more rows; a confounded correlation (ice cream and drownings, both driven by summer) survives any number of rows. The second box is the deep point and the one students miss: precision and validity are different axes. This returns in hypothesis testing, where a large n makes trivial effects "significant".



Summary Statistics Hide the Data

Datasets whose **summary numbers agree** to two decimal places — same averages, same spread, same strength of association — can look completely different when plotted: a dinosaur, a star, a set of lines (Matejka & Fitzmaurice 2017).

We will make "average", "spread" and "association" precise in Lecture 02, and meet this phenomenon again there as Anscombe's quartet.

Plot the data. Then plot it again, differently.

Forward *pointer*, not a forward dependency (P7): the phenomenon is shown visually and named in plain words; the statistics that make it precise are Lecture 02's job, where Anscombe's quartet returns as consolidation. Do not use the words variance or correlation as if defined here.

Where This Leaves Us

Data does not speak for itself

Every dataset was selected, filtered and encoded by someone.

Data cannot be read intuitively

Human intuition fails at scale, in high dimensions, and with base rates.

So theory is needed *more*

Models make the assumptions explicit enough to be criticised.

Big data does not end theory. It **raises the cost of implicit theory**, because assumptions that stay unstated now get automated and scaled.

Return to the guiding question and answer it explicitly. The last sentence is the one to leave on the board: the danger is not wrong theory, it is invisible theory. This is the thread the whole course pulls on.



Module 3

What Is Data Science?

Guiding Question: *If data science is not statistics, not machine learning and not engineering, what is left that is its own?*

A Definition

Data science is the study of the generalizable extraction of knowledge from data. — Vasant Dhar (Dhar 2013)

Course definition, after Dhar. Data science researches and applies technologies, systems and algorithms to gain knowledge from potentially large, heterogeneous and dynamically changing data sets — for a particular domain.

The phrase doing the work is "**for a particular domain**": without it, there is no criterion for which patterns count as knowledge.

Speaker notes

Math-first equivalent for a definitional slide (P5): show it, then interrogate it. Dhar's one line is quoted verbatim; the box is this course's working definition, which adds the domain. "Generalizable" in Dhar is the same demand as "valid" in the KDD criteria later in this lecture. Ask which word could be deleted without loss (probably "large") and which could not ("domain", "knowledge").



Neighbouring Fields

Statistics

Inference from samples: what may be concluded, and with what uncertainty.

Machine Learning

Learning from past experience to improve at a task in the future.

Artificial Intelligence

Building systems that behave intelligently.

Data Engineering

The twin field: managing, organising and providing access to data.

Data science is not a subset of any of them. It is the discipline of getting from a **domain question** to a **defensible answer**, using whichever of these tools fits.

Speaker notes

Students arrive with these words used interchangeably. The distinction that sticks: statistics and ML are methods, data engineering is an infrastructure, data science is an inquiry.



Three Kinds of Task

Explorative

Find patterns.

Which customers behave similarly?

Descriptive

Summarise what is there.

What does an average customer spend?

Predictive

Anticipate future outcomes.

Is this transaction fraudulent?

Naming the task first constrains everything else — the method, the evaluation measure, and what counts as success.

Data science is defined by its **question and its domain**, not by its tools. Explorative, descriptive and predictive tasks demand different methods and different evidence — so the first decision in any project is which one you are doing.

This triple is the organising map of the whole course: clustering and projections are explorative, statistics is descriptive, classification and regression are predictive. Point forward to where each returns.

Module 4

How Knowledge Is Acquired

Guiding Question: A pattern in data is not yet knowledge. What has to happen in between?

Three Modes of Reasoning

	From	To	What it guarantees
Deduction	premises	what must follow	certain — but adds no new content
Induction	observed patterns	a regularity , a hypothesis	nothing — content-adding, not guaranteed
Abduction	an observation	the best available explanation	nothing — only comparable against rivals

Machine learning lives in the bottom two rows.

Three modes, one criterion to compare them: what does each guarantee, and what does it buy? Deduction guarantees everything and discovers nothing; induction and abduction discover and guarantee nothing. Data science lives in the latter two, which is why validation is a permanent obligation. The next slides show the three modes taking turns on a real case.

The Cycle, on a Real Case: Uranus

DEDUCTION

Newton's laws plus the known planets give the orbit Uranus **must** follow.

INDUCTION

Six decades of observed positions show a **systematic** deviation — larger than any measurement error.

ABDUCTION

Best explanation: an **unseen planet** further out, pulling on Uranus.

Speaker notes

Uranus was discovered in 1781; by the 1840s its tables no longer matched the sky. The point of the slide is the last step: an explanation earns its keep by predicting something checkable — here, a position in the sky. A data-science project runs the same loop: choosing a model family is deduction from assumptions, fitting is induction, diagnosing why it fails on a subgroup is abduction. The success is famous. The next slide asks why it was not obvious.



The Cycle, on a Real Case: Uranus

DEDUCTION

Newton's laws plus the known planets give the orbit Uranus **must** follow.

INDUCTION

Six decades of observed positions show a **systematic** deviation — larger than any measurement error.

ABDUCTION

Best explanation: an **unseen planet** further out, pulling on Uranus.

The cycle turns: the abducted planet becomes a new premise. Le Verrier deduces where it must be; in 1846 Galle points a telescope there and finds **Neptune**, about a degree from the prediction.

An Anomaly Leaves Two Moves

SAVE THE THEORY

Add a hidden cause

The bet: the theory is right; our picture of the world is incomplete.

If wrong: we chase something that is not there.

CHANGE THE THEORY

Replace the premises

The bet: the theory fails here, however well it worked elsewhere.

If wrong: we throw away two centuries of correct predictions.

1859: Mercury's orbit turns about **43" per century** more than Newton allows. Which move would you make?

Collect votes before the next slide. After Neptune, almost everyone chooses "save the theory" — and that is the right answer *for 1859*: the same scientist had just been spectacularly vindicated by exactly that move. Make students commit; the next slide pays off only if they have.

Same Move, Opposite Outcomes

anomaly	the move	what happened
Uranus , 1840s	a hidden planet	Neptune found, 1846 ✓
Mercury , 1859	a hidden planet, "Vulcan" (Le Verrier again)	?

Speaker notes

Same scientist, same method, same kind of anomaly, opposite outcomes. Vulcan was searched for at eclipses for decades and a few sightings were claimed; none survived. Einstein's theory produced the 43" per century as a consequence, with nothing adjusted to fit it — which is why it replaced Newton rather than merely patching him.



Same Move, Opposite Outcomes

anomaly	the move	what happened
Uranus , 1840s	a hidden planet	Neptune found, 1846 ✓
Mercury , 1859	a hidden planet, "Vulcan" (Le Verrier again)	?

Vulcan is **never found** ✗ In 1915 **general relativity** explains the 43" with no new planet — here, changing the theory was right.

Same Move, Opposite Outcomes

anomaly	the move	what happened
Uranus , 1840s	a hidden planet	Neptune found, 1846 ✓
Mercury , 1859	a hidden planet, "Vulcan" (Le Verrier again)	?

Vulcan is **never found** ✗ In 1915 **general relativity** explains the 43" with no new planet — here, changing the theory was right.

Both decisions were **reasonable when they were made**. What separated them was evidence that did not exist yet.

Deciding Before the Evidence Is In

- **Trust earned.** After two centuries of success, absorbing one anomaly *into* the theory was the rational first move.
- **Evidence accumulates.** Every failed search made "a hidden planet" less plausible — confidence moved step by step, not in one jump.
- **Replace, don't just refute.** Newton fell when a rival explained *more*, not at the first failure.
- **Costs decide.** Chasing Vulcan forever and abandoning a good theory too early are both expensive; which move to make depends on what being wrong costs.

A hidden planet is an **assumption added to save a pattern**. It is worth exactly as much as the check it survives.

The data-science translation: a model that fails on new data can be saved by an extra assumption ("that month was unusual", "the new customers are different") or replaced. Both are legitimate first moves; neither is knowledge until something checkable has been predicted and observed. This is the course's thread in its oldest form — name the assumption, then check it.

Mini-Exercise

A dashboard shows: customers who use the mobile app **churn less**. Besides "the app keeps customers", give **one other explanation** — and one experiment that tells the two apart.

Speaker notes

Solution on the same slide, click-gated. The exercise runs the whole cycle in a business setting: induction reads the regularity off the dashboard, abduction supplies rival explanations, deduction turns one of them into a checkable prediction. The key move is the randomisation: a hypothesis earns its keep by predicting something its rival does not. This subscription business returns throughout the course.



Mini-Exercise

A dashboard shows: customers who use the mobile app **churn less**. Besides "the app keeps customers", give **one other explanation** — and one experiment that tells the two apart.

Loyal customers are the ones who **install** the app: same correlation, opposite direction.

Experiment: prompt a **random** subgroup to install it. If the app causes loyalty, their churn drops; if loyalty causes the app, it does not.

Takeaway

A pattern is not a proof. Induction and abduction **propose**; deduction turns a proposal into a **checkable prediction**; only the check makes it knowledge. When a check fails, saving the theory with a hidden cause or changing it is a **decision made before the evidence is in**.

What Counts as Knowledge?

Knowledge discovery in databases (KDD) is the non-trivial process of identifying patterns in data that are (Fayyad et al. 1996)

- **valid** — they hold on new data, beyond the sample,
- **novel** — they were not already known,
- **potentially useful** — they can support a decision or an action,
- **ultimately understandable** — a person can make sense of them.

A pattern that fails **any** of the four is not knowledge. Most patterns a large dataset offers fail at least one.

Speaker notes

Use the four criteria as a checklist against a real finding — say, "customers who buy nappies buy beer". Valid? Maybe. Novel? Perhaps. Potentially useful? Only if it changes a decision. Ultimately understandable? Only with a mechanism. This is the antidote to pattern-hunting. The wording is Fayyad et al.'s; "ultimately" matters — the pattern need not be obvious at first sight, but someone must be able to understand it in the end.



Five Steps from Data to Knowledge



Fayyad et al. draw the five steps in a line, with arrows back from every step: evaluation routinely sends you back to an earlier one.

Knowledge is a pattern that is **valid, novel, potentially useful and ultimately understandable**. Data mining is one step of five; it produces patterns, and only the four criteria turn one into knowledge.

Speaker notes

The figure students will meet in every textbook, drawn as a flow so the next unit can walk through it step by step. Point out that "data mining" — the part most people mean by the field — is only the fourth box.



Module 5

The Process and the Craft

Guiding Question: *Most of a data-science project is spent before any model is fitted. On what, and why does it matter?*

Step 1 — Selection of Data

Types of source

- **Structured:** relational databases, CSV tables
- **Semi-structured:** XML, JSON, web APIs
- **Unstructured:** documents, text, multimedia

Types of data

media · object-relational · spatial and spatio-temporal · graphs and networks · time series and sequences · data streams

Selection is where the **first and largest** bias enters. What was not selected cannot be discovered later.

Tie back to the LHC trigger: selection encodes a theory of relevance. Ask what is systematically missing from a dataset of app users (everyone who never installed it) — survivorship bias, before it has a name.



Steps 2–3 — Preprocessing and Transformation

Cleaned

errors and noise removed —
garbage in, garbage out

Integrated

combined with other sources,
resolving conflicts

Transformed

features constructed or
reduced into a representation
for analysis

These steps routinely consume most of a project's effort — and every choice in them is a modelling decision, not a clerical one.

Speaker notes

The "80% of the work" folklore is true but usually told as a complaint. Reframe it: preprocessing and transformation build the representation — which Lecture 04 will call α — and the representation determines what can be learned at all. Fayyad et al. count transformation as its own step (step 3); cleaning and integration are step 2.



Steps 4–5 — Mining and Interpretation, then Deployment

Data mining

Apply the algorithm — clustering, classification, rule mining — that fits the task.

Interpretation / evaluation

Test the patterns against the four criteria: valid, novel, potentially useful, ultimately understandable.

Then: deployment

Not one of Fayyad's five steps (CRISP-DM adds it): put the result where a decision is made, and monitor it there.

The process is a **loop**, not a pipeline. Evaluation routinely sends you back to selection.

Speaker notes

The linear diagram is the standard depiction and the standard lie. Ask where the arrows back go — and note that in production, deployment feeds new data into selection, which is how feedback loops (and their pathologies) arise.



What the Work Actually Looks Like

Ask

Turn a domain question into something a dataset could answer — and check that it is the question worth answering.

Build

Acquire, clean, model, evaluate. Iterate, mostly backwards.

Communicate

Deliver a result someone can act on, with its uncertainty attached.

Speaker notes

The first and last boxes are the ones curricula under-serve and employers complain about. Say plainly that the middle box is what this course teaches, and the other two are where the value is realised or lost.



Three Skill Areas

Domain knowledge

What the data means, which values are impossible, which differences matter.

Statistics and modelling

What may be inferred, with what uncertainty, under which assumptions.

Engineering

Getting the data, at scale, reproducibly, into production.

The interesting failures happen **between** the three: a technically flawless model of the wrong quantity, or a correct analysis nobody can deploy.

The work runs from a **question** to a result someone can **act on**, and it spans three skill areas that nobody holds equally. The characteristic failures happen at the seams — a flawless model of the wrong quantity, or a correct analysis nobody can deploy.

Resist the "unicorn" framing — nobody holds all three deeply. The realistic skill is knowing enough of the neighbouring areas to ask the right question of a colleague, which is also why the course insists on stating assumptions aloud.

Module 6

Applications, and Their Failures

Guiding Question: *What happens to a model's validity when the world it measures changes — sometimes because of models?*

Web and Social Media

Google Flu Trends

Search queries were used to nowcast influenza faster than public health reporting.

And then it failed

It overestimated flu levels — in 2012/13 by more than double. Causes named: overfitting to seasonal search terms, **changes to Google's own search** (e.g. suggested queries), and searches driven by media coverage of the flu season (Lazer et al. 2014).

The world a model measures **keeps changing** — sometimes because of other models. Validity is not a property a model keeps once earned.

The failure is more instructive than the success and belongs on the slide. Lazer et al. ("The Parable of Google Flu", Science 2014) call the two main causes *big data hubris* — fitting 50 million search terms to about a thousand data points picks up winter terms that merely co-vary with flu — and *algorithm dynamics*: Google kept changing its search engine, whose suggestions changed what people typed, i.e. one model altered the data another model relied on. It previews model monitoring, distribution shift, and the KDD loop from the previous unit.

Text and Language

Classic tasks

Summarisation, entity extraction, sentiment analysis, question answering — each once a research programme of its own.

Today

A single pretrained model handles all of them, with the task specified in words rather than in code.

What did **not** change: the need to evaluate on data the model has not seen, and to state what the output may be used for.

Speaker notes

The students' baseline experience is post-LLM, so the historical framing lands better as "these were separate fields" than as a list of tasks. The note is the continuity claim the course rests on.



Business and the Public Sector

Reported potentials from the era of the first big-data wave: large retail revenue gains, hundreds of billions in healthcare savings, comparable figures for e-government.

These are **projections from consultancy reports**, not measured outcomes. Treat them as claims to be interrogated — whose baseline, whose costs, and what happened afterwards?

Applications show what is possible; their failures show what must be checked. Ask of every deployed model: what changed in the world because it exists?

Speaker notes


Explicit critique practice (P2 and the AI-context section): give students a confident number and ask what evidence would be needed to believe it. Ending the lecture on scepticism sets the tone for the rest of the course.



The Road Ahead

Guiding Question: *Every lecture from here on brings a method. What should you ask of each one before you believe its patterns?*

Argument Map



Speaker notes

A recap of the lecture, shown once the room has argued it. Read it top-left to bottom. There is more data, but not more insight. Twenty-nine teams, one dataset and different answers: the analysis decided, not the data. So theory is needed more, not less, because an unstated assumption gets automated at scale. Data science is defined by its question, not by its tools. A pattern found by induction is a guess until it is checked, and it is knowledge only if it passes the four criteria. The loop, not the pipeline, is the honest picture.

Point at the strip: every later lecture introduces a method, and its map names the assumption that speaks for the data in it. The check follows in the same lecture. Every later lecture opens with its own map and the same strip.



The Road Ahead: Name the Assumption, Then Check It

L	The assumption that speaks for the data	The check
02 Chance has rules	a model of chance: outcomes, probabilities	would chance alone produce this?
03 Draw the dependences	a graph: a missing edge claims independence	does the data fit the factorisation?
04 Assume a form	a line, log-odds, independent attributes	error on unseen data, against a baseline
05 Learn the features	stacked units learn form and features	training vs validation error
06 Ask short questions	short trees first; a majority vote	pruning on held-out data
07 Look before you fit	a scale per column; why values are missing	plots, quality checks
08 Define similarity	a distance, a linkage, k	cluster validation
09 Beat the base rate	co-occurrence counts, thresholds	lift against independence
10 Choose what to keep	variance, distances or co-occurrence	explained variance, stress, biplot

Knowing a method's assumption tells you where its patterns stop being knowledge.


Speaker notes

Only now the detailed structure: every row is a method family the students will meet, read as the assumption it adds and the check that tells a pattern from knowledge. This is the course arc, "name the assumption, then check it", and the strip at the bottom of every argument map. Do not explain the rows; the words (independence, loss, lift, variance) are meant to be recognised when they return, not understood today.

Return to this table at the start of every later lecture: each lecture's argument map has the same strip, with its own chip highlighted.



The Lectures on the KDD Loop



The course does not walk the loop in order. Lecture 02 builds the language every check is written in, and Lectures 03–06 build models first, because judging data needs both. Lecture 07 then returns to the first steps, before the methods without labels (08–10).

Speaker notes

Say the order problem out loud rather than hiding it. Interpretation is where each lecture's check lives: the test in Lecture 02, measuring performance in Lecture 04, cluster validation in Lecture 08, lift in Lecture 09. Lecture 07 is the moment the course turns back to the table itself; from Lecture 08 on there is no label, and looking at the data becomes the main check left.



Image Credits & References

Dhar, Vasant (2013). Data science and prediction. *Communications of the ACM*.

Anderson, Chris (2008). The End of Theory: The Data Deluge Makes the Scientific Method Obsolete. *Wired*.

Fayyad, Usama and Piatetsky-Shapiro, Gregory and Smyth, Padhraic (1996). From Data Mining to Knowledge Discovery in Databases. *AI Magazine*.

Gantz, John and Reinsel, David (2010). The Digital Universe Decade -- Are You Ready?.

Gantz, John and Reinsel, David (2012). The Digital Universe in 2020: Big Data, Bigger Digital Shadows, and Biggest Growth in the Far East.

Lazer, David and Kennedy, Ryan and King, Gary and Vespignani, Alessandro (2014). The Parable of Google Flu: Traps in Big Data Analysis. *Science*.

Lyman, Peter and Varian, Hal R. (2003). How Much Information? 2003.

Matejka, Justin and Fitzmaurice, George (2017). Same stats, different graphs: generating datasets with varied appearance and identical statistics through simulated annealing. *Proceedings of the 2017 CHI Conference on Human Factors in Computing Systems*.

Reinsel, David and Gantz, John and Rydning, John (2018). The Digitization of the World: From Edge to Core.

Silberzahn, Raphael and Uhlmann, Eric L and Martin, Daniel P and others (2018). Many analysts, one data set: making transparent how variations in analytic choices affect results. *Advances in Methods and Practices in Psychological Science*.