

# Data Science

---

## Lecture 06: Decision Trees and Ensembles

Prof. Dr. Michael Granitzer

Converted from LNKB-tex part-en-ds-chp5-decision-tree-learning.tex. Builds on Lecture 04 (hypothesis space, inductive bias, bias–variance, cross-validation) and Lecture 02 (probabilities of class labels; entropy itself is introduced here, in Module 2). The arc runs from the most interpretable model in the course to a committee that gives that interpretability up.

# Where We Left Off

Lectures 04–05 gave us models that predict and cannot explain themselves — a few hundred weights, or millions. Today: a model you can read aloud, and what it costs to keep it readable.

## Lecture 05

accurate, opaque



## Modules 1–4

one tree: readable, unstable



## Module 5

a forest: stable, opaque again


"Ultimately understandable" was one of Lecture 01's four criteria for knowledge. This lecture is the first to trade it away on purpose.

Speaker notes

Name the trade at the start so that the students watch for it: every module until the forest buys readability; Module 5 sells it for stability and accuracy. The churn tree at the end of the lecture makes the trade concrete.



# Argument Map



A tree is a sequence of questions; the only real decision is the next split, which information gain scores. ID3's greed carries a bias: short trees, and attributes with many values win. Growing to purity memorises, so prune, and judge the pruning on unseen data. The instability of a single tree becomes the raw material of bagging and random forests. The bottom row states the trade the students must be able to name.

# Module 1

---

## Trees as Hypotheses

**Guiding Question:** *A linear model hides its reasoning in 200 weights. What would a model look like that can be read aloud?*

# A Tree Is a Sequence of Questions

Each inner node tests one attribute, each branch is an answer, each leaf is a class.

| Example | Sky   | Temperature | Humidity | Wind   | EnjoySport |
|---------|-------|-------------|----------|--------|------------|
| 1       | sunny | warm        | normal   | strong | yes        |
| 2       | sunny | warm        | high     | strong | yes        |
| 3       | rainy | cold        | high     | strong | <b>no</b>  |
| 4       | sunny | warm        | high     | strong | yes        |

To classify a new day, start at the root and follow the answers to a leaf.

The tree **is** the hypothesis. Unlike a weight vector, it can be read aloud as a set of rules — which is why trees are the standard choice when a decision has to be explained.

## Speaker notes

Interpretability is the selling point and should be named early. Ask what "explaining" a linear model with 200 weights would look like by comparison — this is Lecture 01's KDD criterion "ultimately understandable" returning, and the failures in its applications module show why it matters.



# Construction and Classification

## Construction (*recursive*)

If the data is pure enough, make a leaf.  
Otherwise: choose a **split criterion**,  
decompose  $D$  into  $D_1, \dots, D_s$ , and recurse  
on each part.

## Classification (*recursive*)

If the node is a leaf, return its class.  
Otherwise: follow the successor selected  
by the example's attribute value.

The whole learning problem reduces to one question: **which attribute to split on?**  
Everything else is bookkeeping.

The template is deliberately abstract — ID3, C4.5 and CART differ only in the split criterion and the stopping rule. Say that now so the next units read as variations, not as separate algorithms.

# When to Use Them

## Well suited

Attribute-value data; discrete or continuous features; missing values; a need for an explainable model; no assumption of linearity.

## Poorly suited

Smooth relationships (a tree approximates a diagonal boundary with a staircase); many equally relevant features; and — as we will see — stability.

# Mini-Exercise

Find a tree with a **single test** that classifies the four EnjoySport examples from the first slide correctly. Is it the only one?

Speaker notes

Expect Sky first. Ask for a second answer before revealing — someone finds Temperature. Humidity fails (example 2 is high and yes, example 3 is high and no), and Wind is constant. The tie survives the next module's criterion, which is the point: every learner needs a preference beyond the data.



# Mini-Exercise

Find a tree with a **single test** that classifies the four EnjoySport examples from the first slide correctly. Is it the only one?

Sky works (sunny  $\rightarrow$  yes, rainy  $\rightarrow$  no) — and so does Temperature (warm  $\rightarrow$  yes, cold  $\rightarrow$  no). Both separate the data perfectly; even the next module's split criterion scores them equally (0.811 bits each). The data cannot choose: the learner's **inductive bias**, not the examples, breaks the tie.

# Takeaway

A decision tree is a **readable hypothesis** built by recursive splitting. The only real decision is which attribute to split on next — and the data alone does not settle it.

## Speaker notes

The staircase point is worth a sketch on the board: trees split axis-parallel, so a diagonal boundary costs many splits. That limitation motivates ensembles at the end of the lecture.



# Module 2

---

## Choosing a Split

**Guiding Question:** *Every attribute could be split on. What makes one split better than another — and how do we measure it?*


# What Makes a Split Good?

A split is good if the resulting subsets are **purier** than what we started with — closer to containing a single class.

So we need a number that measures impurity: maximal when the classes are evenly mixed, zero when a node is pure.

Guiding question (P1). Let students propose a measure before showing entropy — "fraction of the majority class" is the usual proposal, and it is exactly misclassification error, whose weakness the visualization then shows.

# Three Measures



Stepped viz (P6): entropy, then Gini, then misclassification, then a concrete node at  $p = 0.8$ . Ask before the third click how misclassification will compare — its flatness is the reason it is rarely used as a split criterion, even though it is what we ultimately care about.



# Entropy and Gini

For a node with class probabilities  $P(A_1), \dots, P(A_k)$ :

$$H(\mathcal{A}) = - \sum_{i=1}^k P(A_i) \log_2 P(A_i) \quad G(\mathcal{A}) = 1 - \sum_{i=1}^k P(A_i)^2$$

**Entropy**  $H$  is measured in bits — the average number of yes/no questions needed to identify an example's class. **Gini impurity**  $G$  is the probability that two examples drawn at random from the node belong to different classes.

|             | Pure node | Two classes, 50/50 | $k$ equal classes |
|-------------|-----------|--------------------|-------------------|
| Entropy $H$ | 0         | 1 bit              | $\log_2 k$ bits   |
| Gini $G$    | 0         | 0.5                | $1 - 1/k$         |

## Speaker notes

The "questions needed" reading makes entropy concrete and is worth more than the formula. It also makes the next slide's subtraction obviously meaningful: information gain is *questions saved*. Gini is the curve from the previous slide's second step; CART uses it because it needs no logarithm and ranks splits almost exactly as entropy does.



# Conditional Entropy and Information Gain

After splitting on an attribute  $\mathcal{B}$  with outcomes  $B_1, \dots, B_s$ :

$$H(\mathcal{A} \mid \mathcal{B}) = \sum_{j=1}^s P(B_j) \cdot H(\mathcal{A} \mid B_j)$$

— the impurity of each part, weighted by how much data lands in it.

$$\text{Gain}(\mathcal{B}) = H(\mathcal{A}) - H(\mathcal{A} \mid \mathcal{B})$$

**Information gain** is the impurity removed by the split. Choose the attribute that maximises it.

The weighting by  $P(B_j)$  is what stops a split that isolates one example from looking good. Ask what would happen without it — a strong argument for reading formulas as decisions rather than as notation.

# Worked Example: The Root of PlayTennis

Quinlan's 14 weather days: 9 yes, 5 no, so  $H = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.940$  bits. Counts are (yes, no) per value.

| Attribute   | Split                                        | $H(\mathcal{A} \mid \mathcal{B})$ | Gain  |
|-------------|----------------------------------------------|-----------------------------------|-------|
| Outlook     | sunny (2, 3) · overcast (4, 0) · rain (3, 2) | 0.694                             | 0.247 |
| Humidity    | high (3, 4) · normal (6, 1)                  | 0.788                             | 0.152 |
| Wind        | weak (6, 2) · strong (3, 3)                  | 0.892                             | 0.048 |
| Temperature | hot (2, 2) · mild (4, 2) · cool (3, 1)       | 0.911                             | 0.029 |

Recomputed: sunny and rain branches each have  $H = 0.971$ , so  $H(A \mid \text{Outlook}) = (5/14 + 5/14) \cdot 0.971 = 0.694$ . Mitchell's textbook prints 0.246 and 0.151 because it rounds the intermediate entropies; exact values are 0.2467 and 0.1518. Ask before the fragment why Outlook wins with two mixed branches — the pure overcast branch is the answer. Gini agrees: the root has  $G = 0.459$  and Outlook again gives the largest reduction (0.116).

# Worked Example: The Root of PlayTennis

Quinlan's 14 weather days: 9 yes, 5 no, so  $H = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.940$  bits. Counts are (yes, no) per value.

| Attribute   | Split                                        | $H(\mathcal{A} \mid \mathcal{B})$ | Gain  |
|-------------|----------------------------------------------|-----------------------------------|-------|
| Outlook     | sunny (2, 3) · overcast (4, 0) · rain (3, 2) | 0.694                             | 0.247 |
| Humidity    | high (3, 4) · normal (6, 1)                  | 0.788                             | 0.152 |
| Wind        | weak (6, 2) · strong (3, 3)                  | 0.892                             | 0.048 |
| Temperature | hot (2, 2) · mild (4, 2) · cool (3, 1)       | 0.911                             | 0.029 |

Outlook wins the root although two of its branches are *more* mixed than the root (0.971 vs 0.940 bits): its overcast branch is **pure**, so 4 of the 14 days need no further question.

# Mini-Exercise

A node holds 4 yes and 4 no. A split yields [3 yes, 1 no] and [1 yes, 3 no]; a 3:1 node has entropy 0.811. What is the information gain?

The small number is the lesson: intuitively the split "looks decent", but it buys little. Comparing this with a perfect split ( $\text{gain} = 1$ ) calibrates expectations.



# Mini-Exercise

A node holds 4 yes and 4 no. A split yields [3 yes, 1 no] and [1 yes, 3 no]; a 3:1 node has entropy 0.811. What is the information gain?

Before:  $H = 1$  bit. After: both halves hold half the data and have  $H = 0.811$ , so  $H(\mathcal{A} \mid \mathcal{B}) = 0.811$ . Gain  $= 1 - 0.811 = 0.189$  bits — a split that "looks decent" removes less than a fifth of a bit.

# Takeaway

Impurity measures how mixed a node is; **information gain is the impurity a split removes**, weighted by the data it moves. Entropy and Gini agree closely; misclassification error is too flat to choose between splits.

# Module 3

---

## Building the Tree

*Guiding Question: ID3 never reconsiders a split. What does that greedy choice assume about the world?*

# ID3

**ID3**(examples  $D$ , attributes  $A$ , target)

1. If all examples share one class  $\rightarrow$  return a leaf with that class.
2. If  $A$  is empty  $\rightarrow$  return a leaf with the majority class.
3. Otherwise choose  $A^* \in A$  with the highest **information gain**.
4. For each value  $a$  of  $A^*$ : if  $D_a$  is empty, add a leaf with the majority class of  $D$ ; otherwise add the subtree **ID3**( $D_a$ ,  $A \setminus \{A^*\}$ , target).

Greedy and recursive: each split is chosen for immediate gain, and never reconsidered.

Speaker notes

Point at step 3 and name it: this is heuristic search from Lecture 04's principles unit. The hypothesis space of all trees is astronomically large; ID3 walks one path through it without backtracking.



# What ID3 Assumes

## Hypothesis space

**All** decision trees over the attributes — so the target function is certainly in there. No representational bias.

## Search bias

It prefers trees that place **high-gain attributes near the root**, and therefore **shorter trees**. It stops at the first tree that fits.

Two assumptions, not facts. **Short trees generalise** — Occam's razor as an algorithm. And **a useful attribute shows gain on its own**: for  $y = x_1 \text{ XOR } x_2$  on balanced data, each attribute alone has gain 0, yet together they decide  $y$  — greedy search cannot see it.

## Speaker notes

Force the interpretation (P2, P3): ask why a shorter tree should generalise better. There is an argument (fewer short trees exist, so fitting the data by accident is less likely) but it is not airtight, and students should see it as a defensible choice rather than a fact. The XOR case answers the guiding question: greedy search assumes that attributes which matter jointly also matter one at a time. Pruning returns to this (a weak split can enable a strong one beneath it).



# CART

## What differs

**Binary** splits only; **Gini** impurity instead of entropy; handles **regression** by splitting on variance reduction and predicting the mean in each leaf.

## What does not

The template is unchanged: choose a split, decompose, recurse. CART is a different set of choices inside the same skeleton.

Binary splits reduce a defect of multi-way splits: an attribute with many values (an ID number, in the extreme) achieves high information gain by fragmenting the data into pure singletons. C4.5 corrects for it with the **gain ratio**, the gain divided by the entropy of the split itself.

The many-valued-attribute defect is the reason C4.5 introduced gain ratio. It is a good concrete instance of a metric being gamed by the structure of the data — a theme students will meet in evaluation and in production systems.

# Mini-Exercise

Your dataset has a `customer_id` column, unique for every row. What information gain does a split on it achieve — and what does the resulting tree predict for a new customer?

## Speaker notes

The cleanest overfitting example in the course, and it arrives before pruning is introduced, which motivates the next module. If asked: the tree is a lookup table, one level deep; ID3 would send the unseen id to no branch at all, so an implementation has to fall back on the root's majority class.



# Mini-Exercise

Your dataset has a `customer_id` column, unique for every row. What information gain does a split on it achieve — and what does the resulting tree predict for a new customer?

The **maximum possible**: every branch holds one example, every leaf is pure, so the gain equals the root's entropy. And for a new customer — nothing: their id matches no branch. Perfect purity, zero generalisation; hence gain ratio, or binary splits.

# Takeaway

ID3 is a **greedy, non-backtracking search** whose bias is a preference for short trees. CART changes the split type and the impurity measure, not the idea. Both can be fooled by an attribute with too many values.

# Module 4


---

## Overfitting and Pruning

**Guiding Question:** *A tree can reach 100% training accuracy on any dataset without contradictory examples. Why is that a symptom rather than an achievement?*

# A Tree Grown to Purity Is Overfitted

Left alone, ID3 splits until every leaf is pure — memorising the training set, noise included.



Tree depth plays the role of polynomial degree: more splits, lower training error, and eventually rising test error.

Deliberate reuse of the visualization from the neural-network lecture (P6 and the spiral): the same phenomenon, a third model family. Students should now recognise the shape before the caption says it.

# Two Places to Intervene

## Stopping (pre-pruning)

Halt splitting **during** construction — when a node is small, when gain is below a threshold, or at a depth limit.

Cheap, but short-sighted: a weak split may enable a strong one beneath it.

## Pruning (post-pruning)

Grow the full tree, then **remove** subtrees that do not earn their keep.

More expensive, and generally more reliable.

The "weak split enabling a strong one" point is the standard argument for post-pruning, and it is the XOR structure again: neither attribute helps alone.

# Reduced Error Pruning

1. Split the data into a **training** and a **validation** set.
2. Grow the full tree on the training set.
3. For each inner node: consider replacing its subtree with a leaf carrying the majority class.
4. Keep the replacement if accuracy on the **validation** set does not fall.
5. Repeat until no beneficial replacement remains.

The decision is made on data the tree was **not** grown on. Pruning against the training set would remove nothing — every subtree helps there by construction.

## Speaker notes

This is the hold-out principle from Lecture 04, now used to make a structural decision rather than to report a number. Ask what happens if the validation set is small — pruning decisions become noisy, which is why cost-complexity pruning with cross-validation is preferred in practice.



# Mini-Exercise

A tree with 340 leaves scores 100% on training data, 78% on validation. Pruned to 41 leaves it scores 94% and 89%. Training accuracy fell by six points — did pruning fail?

Students hesitate here — accepting worse training performance feels like failure. Naming it as the point of the exercise reinforces that training error was never the objective.

# Mini-Exercise

A tree with 340 leaves scores 100% on training data, 78% on validation. Pruned to 41 leaves it scores 94% and 89%. Training accuracy fell by six points — did pruning fail?

No. The 299 removed leaves fitted noise — examples that are now misclassified, as intended. Unseen data gained **11 points**, and the tree became small enough to read. Training accuracy was never the objective.

# Takeaway

An unpruned tree memorises. **Pruning trades training accuracy for generalisation**, and the trade must be judged on data the tree never saw.

# Module 5

---

## Ensembles

**Guiding Question:** *Change a few examples and the tree changes completely. Can that instability be turned into an advantage?*

# Why One Learner Falls Short

Three reasons, after Dietterich (2000):

## Statistical

Too little data for the hypothesis space: many hypotheses fit equally well, and small changes in the data pick a very different one.

## Computational

The learner cannot search its hypothesis space exhaustively and settles for a heuristic path — a local optimum.

## Representational

The model family cannot express the target at all; a combination of models may.

**Decision trees are the textbook unstable learner** (the statistical reason): change a handful of examples and a different attribute wins at the root, producing a completely different tree.

The first reason is what "unstable learner" means: high variance. Instability is the property that makes ensembles work, so it must be understood as a fact about trees before bagging looks like anything other than a trick. Ask why a greedy root choice makes the whole tree fragile.

# The Idea

**Ensemble learning** builds several learners and combines their predictions — by majority vote for classification, by averaging for regression.

## Bagging

**Bootstrap aggregating** — train on resampled datasets.

## Boosting

Train weak learners in sequence, each focusing on the previous one's mistakes.  
*Named here; not developed in this course.*

## Random forest

Bagging plus randomised attribute selection at every split.

Ensembles exploit instability: **many different models, each better than chance and wrong in partly independent ways, are right together more often than any one of them.**

## Speaker notes

The takeaway is counterintuitive enough to state plainly, and both conditions matter: a committee of worse-than-chance voters gets *worse* with size, and a committee of identical voters does not change. Partial independence is what the next two units engineer deliberately. Boosting (AdaBoost, gradient boosting) is left to the literature; mention that it reduces bias rather than variance.



# Bootstrap Aggregating

**The problem.** A group of predictors helps only if they differ — but we have one algorithm and one dataset, which yield one model.

**The solution.** Manufacture different datasets: draw  $n$  examples from  $D$  **with replacement**, train a predictor on each sample, and combine them by vote or average.

The learner never changes. Only the data it sees does.

## Speaker notes

Bootstrapping was introduced in Lecture 04 as an error-estimation scheme; here the same resampling produces model diversity. Same tool, different job — worth naming as the spiral it is.



# What Resampling Actually Does

Each bootstrap sample contains about **63.2%** of the distinct examples of  $D$ . The probability that a particular example is *never* drawn in  $n$  tries is

$$\left(1 - \frac{1}{n}\right)^n \xrightarrow{n \rightarrow \infty} e^{-1} \approx 0.368$$

## The gain

Each model sees a different two-thirds of the data, so their errors are partly independent. **Averaging reduces variance** — and therefore overfitting.

## The left-out third

The ~36.8% not drawn is the *out-of-bag* sample, usable as a free validation set for that model.

The 63.2% is a nice concrete consequence of a limit students can verify. The out-of-bag idea is practically important: it gives an honest error estimate without a separate hold-out.

# Mini-Exercise

You bag 100 trees, each 70% accurate, and take a majority vote. What accuracy does the vote reach if all 100 trees are identical — and if their errors are independent?

Exact:  $P(\text{Bin}(100, 0.7) \geq 51) = 0.99998$ . Real bootstrap trees sit between the two extremes, far closer to the identical case, which is why bagging helps less than the independent number suggests. With 45%-accurate independent trees the same vote falls to 13% (a tie counted as an error) and to 0 as the committee grows — "better than chance" is not decoration. Bagging reduces mainly variance: averaging does not fix an error every tree shares. The identical case motivates random forests directly: bootstrap samples alone may not decorrelate trees enough, especially when one attribute dominates.

# Mini-Exercise

You bag 100 trees, each 70% accurate, and take a majority vote. What accuracy does the vote reach if all 100 trees are identical — and if their errors are independent?

Identical: **70%** — a hundred copies vote unanimously. Independent: the vote is wrong only if 50 or more trees err at once, which happens with probability  $2.2 \times 10^{-5}$  — about **99.998%**. Both conditions matter: each tree better than chance ( $> 50\%$ ), and errors that are not shared.

# Takeaway

Bagging manufactures diversity by **resampling the data**, and averages away variance. It helps exactly to the extent that the models' errors are independent — and not at all when they are not.

# Why Bagging Trees Is Not Enough

## Trees stay similar

If one attribute has much higher gain than the rest, it wins at the root of **every** bootstrap sample. The trees correlate, and averaging them buys little.

## And it costs

Interpretability is gone — you can no longer read the model as a set of rules — while the computation of 100 near-identical trees is largely wasted.

This answers the identical-trees case of the bagging mini-exercise. Resampling alone does not guarantee diversity; the dominant attribute defeats it.

# The Fix: Randomise the Split Too

## Random forest (Breiman 2001)

For each of  $T$  trees:

1. Draw a bootstrap sample from  $D$ .
2. Grow a tree — but **at every node**, choose the split from a random subset  $A' \subset A$  of  $m_A$  attributes rather than from all of them.
3. Do not prune.

Combine by majority vote (classification) or average (regression).

The dominant attribute is simply **absent** from most nodes' candidate sets, so other attributes get to define the tree. Diversity is engineered rather than hoped for.

Step 2 is the entire contribution, and it is one line. Typical  $m_A$  is  $\sqrt{|A|}$  for classification. Step 3 is counterintuitive — individual overfitting is tolerated because averaging handles the variance.

# What You Get and What You Give Up

## Gained

Strong accuracy with almost no tuning; out-of-bag error for free; a usable attribute-importance measure.

## Lost

The readable single tree. A forest is a **committee**, and no member speaks for it.

Attribute importance from a forest is not a causal statement. It reports which attributes the trees found useful for splitting — under this data, this sampling, this correlation structure.

## Speaker notes

The warning matters because attribute importance is routinely over-read in practice, including in published work. It connects back to the "Correlation Is Not Causation" slide in Lecture 01.



# Mini-Exercise

A pruned tree reaches 86% accuracy, a random forest 91%. Which do you deploy for an internal churn model — and which for automated loan refusals that must be justified to applicants?

## Speaker notes

Close the lecture on this (P2, and Lecture 01's KDD criterion "ultimately understandable", whose failures its applications module showed): model selection is a decision about consequences, not only about a metric. Expect disagreement on the loan case and let it run — that is the exercise working.



# Mini-Exercise

A pruned tree reaches 86% accuracy, a random forest 91%. Which do you deploy for an internal churn model — and which for automated loan refusals that must be justified to applicants?

Churn: the **forest** — five points matter and nobody needs the model explained. Loans: the **tree**, or a forest inside a documented, auditable process; a refusal that cannot be explained may be neither acceptable nor lawful. The best model fits the **decision**, not the leaderboard.

# Takeaway

Random forests **engineer diversity** by randomising the attributes available at each split, turning unstable trees into a stable committee — at the cost of the one thing a single tree was best at: being readable.

# The Churn Shop: A Tree You Can Read Aloud

## ONE TREE

- Contract **monthly**?
- No **app login** for 30 days?
- Then: **likely to churn** — call them.

The retention team can check every step against what they know.

## A FOREST OF 500 TREES

Usually more accurate, and stable when a few customers change.

Its answer is a vote. Asked *why* this customer, it has no sentence to offer.

The same data, two answers to Lecture 01's criteria: the tree is **understandable** and fragile, the forest is accurate and unreadable. Say which one you chose, and why.

## Speaker notes

The running example, now as a model. The tree is illustrative, not fitted: the point is its form — a path of questions a person can read out and dispute ("monthly contracts churn more — we knew that; the app login is new"). A forest built on the same customers would typically predict churn better, but its answer is a majority over hundreds of such paths.

Link to the decision the shop actually faces: if the retention team must explain the call to the customer, or check the model against their own knowledge, readability is worth some accuracy. If the list is only used to schedule calls, it may not be.



