

# Data Science

---

## Lecture 07: Data Exploration and Preprocessing

Prof. Dr. Michael Granitzer

Converted from LNKB-tex part-en-ds-chp3-data.tex (two source units re-encoded from ISO-8859). Uses Lecture 02's quantiles and covariance throughout, and its random variables and joint distribution in the unit on data as random variables.

# Where We Left Off

Three lectures fitted models to tables nobody questioned. Today steps back to KDD steps 1–2 from Lecture 01: what the columns are, whether they can be trusted, and what a plot shows that a summary hides.

## Lecture 01

data is never raw



## Lecture 02

summaries can agree while the data differs  
(Anscombe)



## Today

scale, quality, plots, preprocessing


From Lecture 08 on there is no label to check a model against — looking first stops being good practice and becomes the main check we have.

## Speaker notes

Name the order openly. Modelling came first because judging data needs the language (Lecture 02) and the models (Lectures 04–06); now the course runs the loop backwards, as the second roadmap slide of Lecture 01 showed. Attribute types were used informally in Lectures 04–06 (nominal attributes in naive Bayes and trees); today they get their names and their consequences.



# Argument Map



The scale of measurement licenses a method: a mean of postcodes computes and means nothing. A table is a sample from a joint distribution, and "i.i.d." is a claim about how it was collected. Why values are missing decides how they may be filled. Plots show what summaries hide. Preprocessing is modelling: every aggregation, sample, scaling and reduction trades information for tractability. The takeaway is the lecture's arc chip.

# Module 1

---

## What Kind of Data Is This?

**Guiding Question:** *The file loads and the numbers compute. Why is that not yet a reason to trust the average?*

# Four Attribute Types

## Nominal

*categorical,  
qualitative*

Labels only: eye  
colour, ZIP code, ID.

Meaningful:  $=$ ,  $\neq$

## Ordinal

*categorical,  
qualitative*

Ordered labels:  
grades, rankings,  
"small/medium/large".

Meaningful:  $<$ ,  $>$

## Interval

*numeric, quantitative*

Differences are  
meaningful, the zero  
point is arbitrary:  
temperature in  $^{\circ}\text{C}$ ,  
calendar dates.

Meaningful:  $+$ ,  $-$

## Ratio

*numeric, quantitative*

A true zero, so ratios  
are meaningful:  
length, mass, counts,  
duration.

Meaningful:  $\times$ ,  $\div$

The four types form a hierarchy: each level permits everything the previous one does, plus more. The type is a property of the *measurement*, not of the storage format — which is why a ZIP code stored as an integer is still nominal.

# Why the Type Decides the Method

## The classic error

Averaging a nominal attribute. The mean of ZIP codes 10115 and 81667 is a number, and it is meaningless.

## The subtle one

20 °C is not "twice as warm" as 10 °C — interval scales have no true zero, so ratios do not transfer.

The attribute type determines which **statistics**, which **distance measures** and which **plots** are permissible. Every method in this course depends on it.

# Mini-Exercise

A customer table records satisfaction (one to five stars) and year of birth. Which scale is each — and whose **mean** would you not report?

Solution on the same slide, click-gated. The two columns come back when a table is read as a sample from a joint distribution; payment method (nominal: only counting) and income (ratio: a true zero) complete the set if there is time.

# Mini-Exercise

A customer table records satisfaction (one to five stars) and year of birth. Which scale is each — and whose **mean** would you not report?

Satisfaction is **ordinal**: its mean assumes equal gaps between the stars — report the median.  
Year of birth is **interval**: the mean year is fine, but ratios are not ("born twice as late" means nothing); differences such as age are fine.

# Takeaway

An attribute's **scale of measurement** licenses the operations you may perform on it. Choosing a statistic before checking the scale is the most common silent error in data analysis.

# A Table Is a Sample From a Joint Distribution

| customer | payment | satisfaction | born | income (k€) | returned? |
|----------|---------|--------------|------|-------------|-----------|
| 1        | card    | 4            | 1987 | 52          | yes       |
| 2        | invoice | 2            | 1995 | 31          | no        |
| 3        | card    | 5            | 1979 | 78          | no        |
| ...      | ...     | ...          | ...  | ...         | ...       |

## Each column is a random variable

$X_{\text{payment}}, X_{\text{satisf.}}, \dots$  — a function from customers to values, exactly as a random variable was defined.

## Each row is one draw

The whole table is  $n$  draws from one joint distribution  $p(x_1, \dots, x_d)$  that nobody has written down.

Satisfaction and year of birth are the columns of the attribute-type mini-exercise — deliberately. There they were classified by scale; here they become random variables. Say the sentence that makes the rest of the course possible: *a dataset is not the distribution, it is a sample from it*. Everything we compute from the table is therefore an estimate, and an estimate comes with an error.

# The Scale Suggests the Model

| attribute    | scale                   | a first model                                              |
|--------------|-------------------------|------------------------------------------------------------|
| returned?    | nominal, two values     | <b>Bernoulli</b> ( $p$ ) — one parameter                   |
| payment      | nominal, $k$ values     | <b>categorical</b> ( $p_1, \dots, p_k$ ), $\sum_j p_j = 1$ |
| satisfaction | ordinal                 | categorical — the order is not used, or an ordinal model   |
| born         | interval                | normal, on a restricted range                              |
| income       | ratio, positive, skewed | log-normal, or normal on log income                        |

The scale limits the model as it limits the statistics: a normal distribution for ZIP codes is as meaningless as their mean.

## Speaker notes

Not a catalogue to memorise — a habit: before choosing a distribution, check what the scale permits. The categorical distribution is the Bernoulli with more than two values; it is the distribution naive Bayes uses for every nominal attribute. Satisfaction is the honest hard case: treating it as categorical throws the order away, treating it as numeric invents equal gaps.



# What "i.i.d." Claims

The standard reading of a table rests on one assumption: the rows are **independent and identically distributed**.

## Independent

One customer's row tells us nothing about another's.

## Identically distributed

Every row comes from the **same** joint distribution.

Which of these datasets break it: daily sales of one shop · two visits by the same customer · pixels of one image · a survey sent to a random sample?

The i.i.d. assumption is what turns "the table" into "a sample", and it is what lets a likelihood factorise into a product over rows — the step every fitting method takes. Link back to the shapes of data: record data is the case where i.i.d. is at least plausible; graphs, grids and time series are the cases where it is not.

# What "i.i.d." Claims

The standard reading of a table rests on one assumption: the rows are **independent and identically distributed**.

## Independent

One customer's row tells us nothing about another's.

## Identically distributed

Every row comes from the **same** joint distribution.

Which of these datasets break it: daily sales of one shop · two visits by the same customer · pixels of one image · a survey sent to a random sample?

Only the random survey is plausibly i.i.d. Time series, repeated measurements and spatial data all violate it — the arrangement *is* information.

# Takeaway

A dataset is  $n$  **draws from an unknown joint distribution**: columns are random variables, rows are draws, and the scale of each column suggests its model. Reading a table this way rests on an assumption — **i.i.d.** — that should be checked, not presumed.

# Four Shapes of Data

## Record data

Objects × attributes  
— the familiar table.  
Rows are *treated as* independent.

## Graph data

Objects carry **links** to other objects: web pages, social networks, citations.

## Fields / grids

Objects sit in a **spatial** arrangement: images, sensor lattices, simulations.

## Spatio-temporal

Position **and** time: monthly land temperatures, traffic movement, trajectories.

The shape decides which methods even apply. A table whose rows are treated as independent can be shuffled without loss; a grid or a time series cannot — the arrangement *is* information. Independence of rows is an assumption, not a property of the file.

Ask what is lost by flattening each non-record type into a table. That question is the bridge to network science and to time-series methods, and it recurs whenever someone loads a CSV.

# Three Data Quality Problems

## Noise

Random distortion of measured values. Cannot be removed, only modelled or smoothed.

## Outliers

Values far from the rest.  
Either an **error** to remove or the **signal** you were looking for — fraud, faults, discoveries.

## Missing values

Absent entries. Handling depends on **why** they are missing, which the data itself does not tell you.

"Outlier" is not a property of a data point; it is a judgement about a model of what is normal. Deleting outliers deletes evidence against that model.

# Mini-Exercise

A weather sensor misses 8% of its temperature readings, and it fails mostly **when it is very cold**. Is deleting the rows with missing values harmless?

## Speaker notes

This is Principle 3 in its most practical form: a preprocessing decision is a modelling assumption, and here an untestable one. Students should leave able to ask "why is this missing?" before "how do I impute it?". If time allows, ask what else could have told us about the mechanism (a thermometer next to it, the sensor's data sheet) — domain knowledge, not the data.



# Mini-Exercise

A weather sensor misses 8% of its temperature readings, and it fails mostly **when it is very cold**. Is deleting the rows with missing values harmless?

No. Deletion removes the coldest days, so every statistic afterwards — the mean temperature first — is biased upwards. Had the gaps been random transmission dropouts, deletion would be roughly harmless. The file looks the same in both cases; only the **mechanism** differs.

# Mini-Exercise

A weather sensor misses 8% of its temperature readings, and it fails mostly **when it is very cold**. Is deleting the rows with missing values harmless?

No. Deletion removes the coldest days, so every statistic afterwards — the mean temperature first — is biased upwards. Had the gaps been random transmission dropouts, deletion would be roughly harmless. The file looks the same in both cases; only the **mechanism** differs.

The shape of the data decides which methods apply; the **reason** for a quality problem decides how it may be handled. Neither is visible in the file.

# The Churn Shop's Customer Table

| customer | contract | satisfaction | tenure (months) | postcode | churned? |
|----------|----------|--------------|-----------------|----------|----------|
| 1        | monthly  | 4            | 3               | 94032    | yes      |
| 2        | yearly   | —            | 27              | 80331    | yes      |
| 3        | yearly   | 5            | 41              | 94032    | no       |
| ...      | ...      | ...          | ...             | ...      | ...      |

The running example as a table: the same subscription business whose app users churn less (Lecture 01), whose gap was tested (Lecture 02) and predicted (Lectures 04 and 06). Every earlier lecture used this table without asking what its columns are.

The missing ratings are the lecture's point in one number. Dropping the rows with a missing rating leaves 60 churners among 915 customers ( $60/915 \approx 6.6\%$ ): the missingness depends on the outcome itself, so the cleaned table is no longer a sample of the customers. Ask what a plot would have shown first — the share of missing ratings split by churned yes/no.

# The Churn Shop's Customer Table

| customer | contract | satisfaction | tenure (months) | postcode | churned? |
|----------|----------|--------------|-----------------|----------|----------|
| 1        | monthly  | 4            | 3               | 94032    | yes      |
| 2        | yearly   | —            | 27              | 80331    | yes      |
| 3        | yearly   | 5            | 41              | 94032    | no       |
| ...      | ...      | ...          | ...             | ...      | ...      |

### SCALES

Contract, postcode, churned: nominal.  
Satisfaction: ordinal. Tenure: ratio. The mean postcode computes — and means nothing.

### WHY IS SATISFACTION MISSING?

Customers who left stop answering: 40 of 100 churners gave no rating, 45 of 900 stayers. Drop those rows, and churn falls from **10%** to **6.6%**.

# Module 2

---

## Seeing the Data

**Guiding Question:** *Summary statistics answer the questions you thought to ask. What shows you the ones you did not?*

# Visualization Is Everywhere

Newspaper tables, subway maps, weather charts, stock graphs, a 3-D reconstruction of an injured knee, demographic bar charts, social-network layouts.

**What do these have in common?** All of them present data in a form the eye processes faster than the equivalent numbers.

Question-first (P1). Let students name the common property before it is stated; "you can see it at a glance" is the answer, and the next slide gives the reason.



# Why It Works

## Bandwidth

Vision is the human sense with the highest data rate — pattern detection runs in parallel and pre-attentively.

## Interaction

A picture is a natural surface for exploring: zoom, filter, select, compare.

## Discovery

Visualization surfaces the **unexpected** — the structure nobody thought to compute a statistic for.

Summary statistics answer questions you already asked. A plot lets the data ask you one — which is exactly what Anscombe's quartet demonstrated.

Visualization is not decoration for results; it is an **instrument of analysis**, and the only one that can show you what you did not think to measure.

Deliberate callback to L01 and L02 (spiral): the same argument, now with the tools to act on it. This unit is the hinge between "data does not speak for itself" and the concrete plot types that follow.



# The Boxplot

A visual summary of a distribution built from quartiles:

## The box

Spans  $Q_1$  to  $Q_3$  — the interquartile range, the middle half of the data.

## The line

The **median** inside the box. Its position shows skew.

## The whiskers

(Tukey) Reach to the most extreme data points inside  $[Q_1 - 1.5 \cdot \text{IQR}, Q_3 + 1.5 \cdot \text{IQR}]$ . Points outside are drawn singly as **outliers**.

The  $1.5 \times \text{IQR}$  rule is a **convention**, not a test. It flags candidates for inspection, not errors for deletion.

Boxplots are built entirely from the quantiles the students defined in Lecture 02 — say so. Their strength is comparison: ten boxplots side by side are readable where ten histograms are not.

# The Histogram

## What it shows

The **shape** of a distribution: modality, skew, gaps and clusters that a boxplot hides entirely.

## What it hides

Everything depends on the **bin width**. Too wide and structure vanishes; too narrow and noise looks like structure.

A bimodal distribution and a unimodal one can produce the same boxplot. Always look at one histogram before trusting a five-number summary.

# Mini-Exercise

Two datasets have the same minimum,  $Q_1$ , median,  $Q_3$  and maximum — so the same boxplot.  
Must their histograms look alike?

Speaker notes

If there is a minute left, let one student sketch the two-cluster version on the board. Some tools draw the whiskers to the minimum and maximum instead of the Tukey fences; always check which convention a plot uses.



# Mini-Exercise

Two datasets have the same minimum,  $Q_1$ , median,  $Q_3$  and maximum — so the same boxplot. Must their histograms look alike?

No. The five numbers fix five points; the mass between them can be arranged freely — spread evenly inside the box, or piled into two clusters at its edges. Same quartiles, opposite stories: Anscombe's quartet again.

# Mini-Exercise

Two datasets have the same minimum,  $Q_1$ , median,  $Q_3$  and maximum — so the same boxplot. Must their histograms look alike?

No. The five numbers fix five points; the mass between them can be arranged freely — spread evenly inside the box, or piled into two clusters at its edges. Same quartiles, opposite stories: Anscombe's quartet again.

A boxplot compares distributions compactly; a histogram shows their shape. **Neither alone is sufficient**, and the boxplot's outlier rule is a convention you may overrule.

# The Scatterplot

An X–Y plot of two dimensions against each other — the standard instrument for inspecting a **relationship** rather than a distribution.

With more than two dimensions there are four options:

## Subsetting

Let the user choose which dimensions to display.

## Reduction

Project to fewer dimensions (see preprocessing, and projections later).

## Embedding

Map extra dimensions to other channels: colour, size, shape.

## Multiple displays

Superimpose, or juxtapose in a **scatterplot matrix**.

The four options are exhaustive and worth memorising — every high-dimensional visualization technique is one of them or a combination. Colour and size are channels with far lower precision than position; use them for secondary variables.

# What a Scatterplot Reveals

Linear association, non-linear structure, clusters, outliers and heteroscedasticity — most of which no single correlation coefficient reports.

**Pearson correlation**  $r = \frac{\text{cov}(X, Y)}{s_X s_Y} \in [-1, 1]$ : the covariance measured in units of the two standard deviations. It measures **linear** association only.

A correlation of 0 is consistent with a perfect parabola. The scatterplot shows it; the coefficient does not.

# Tufte's Principle

**Maximise the data-ink ratio:** every drop of ink on the page should carry information. Delete anything that does not — heavy gridlines, boxes, redundant tick marks, decorative fills.

**Range frames** apply the principle to the axes: draw each axis only over the interval the data actually spans, so the frame itself reports the minimum and maximum.

The exercise: take a finished chart and ask of each element, *what would I lose by deleting this?*

A scatterplot shows the structure a coefficient summarises away. Beyond two dimensions you must **subset, reduce, embed or juxtapose** — and every mark you draw should earn its ink.

Tufte's rule is a good design heuristic and an over-applied dogma; note that some redundancy aids reading under time pressure. The point students should keep is intentionality: each element defended, not each element removed.

# When a Line Is Justified

Connecting points with a line asserts an **order** between them — and often an interpolation.

**Data requirement:** one or more quantitative value attributes on a comparable scale, plus one **ordered** key attribute (usually time).

Connecting nominal categories with a line claims an order that does not exist. The line is a statement, not a decoration.

This is the attribute-types unit paying off immediately: the permissible plot follows from the scale of the key attribute. Bar charts for nominal, lines for ordered.

# Aspect Ratio Is Not Cosmetic

Readers compare trends by the **orientation** of the line segments, so the shape of the plotting area changes what they conclude. The same series looks flat in a wide frame and dramatic in a tall one.

**Banking to 45°** chooses the aspect ratio so that the segments' average absolute orientation is 45°: differences in orientation are judged best near  $\pm 45^\circ$  (Cleveland et al. 1988).

## Additional channels

Colour and line type distinguish series — but only for a handful of them.

# Multiple Series

## Overlaid

Best for comparing a few series directly against each other.

## Stacked

Each series is drawn on top of the previous one. Shows the total well — and makes individual series **hard to read**, since each sits on a moving baseline.

## Sorted / small multiples

One panel per series, ordered by some property. Scales to many series where overlaying fails.

A line asserts order and interpolation; the **aspect ratio** decides what the reader sees and stacking trades the readability of the parts for the readability of the whole.

The stacked-plot critique is worth demonstrating: ask students to read the third band's trend off a stacked chart — they cannot, because the baseline moves. Small multiples are the usual better answer.

# Module 3

---

## Preprocessing

**Guiding Question:** *Lecture 01 claimed that preprocessing is modelling. Which choices made here change what a model can learn?*

# Aggregation

Combine several objects or attributes into one — customer purchases summarised by store, by month, by price band.

## Gains

Fewer objects to process; **less variation**, hence more stable estimates; the natural operation of a data warehouse.

## Costs

Detail is gone, irreversibly. Effects that exist only within a group become invisible.

Speaker notes

Aggregation is where Simpson's paradox lives: a trend can reverse when groups are pooled. Worth one sentence — students will meet it in dashboards long before they meet it in a statistics course.



# Sampling

Work with a subset of the objects: random sampling with or without replacement, or **stratified** sampling that preserves group proportions.

Sampling always loses information. The question is whether it loses anything that matters — a sufficiently large sample retains the overall structure.

Sample a sine wave once per period and every sample has the same value: the reconstruction is a flat line — not merely rougher, but **wrong**.

The sine example is aliasing: too sparse a sample does not blur a pattern, it produces a different one. The same holds for customer data sampled once a month when the behaviour of interest is weekly.

# Scaling and Encoding

## Scaling numeric attributes

Min–max:  $x' = \frac{x - \min}{\max - \min} \in [0, 1]$

Standardisation:  $z = \frac{x - \bar{x}}{s}$  — mean 0, standard deviation 1

## Encoding categorical attributes

**Nominal** → one-hot: one 0/1 column per value, so no false order.

**Ordinal** → ranks 1, 2, 3, ..., which keep the order.

Two customers differ by 1,000 € in income and 30 years in age. Unscaled, their distance is  $\approx 1,000$ : age is invisible. Standardised ( $s = 10,000$  € and 15 years) the differences are 0.1 and 2 — now age dominates.

## Speaker notes

Scaling is a modelling decision in the sense of this lecture: it decides which attribute a distance listens to. It matters for every distance-based method (clustering, nearest neighbours) and for PCA, which would otherwise find "the direction of income". Encoding nominal values as 1, 2, 3 invents an order and a distance between categories; one-hot avoids it at the cost of more columns. Exact numbers:  $\sqrt{(1000^2 + 30^2)} \approx 1000.4$  and  $\sqrt{(0.1^2 + 2^2)} \approx 2.0$ .



# Dimensionality Reduction

**The curse of dimensionality.** As the number of attributes grows, the data becomes sparse: distances between points concentrate, neighbourhoods lose meaning, and the volume to be covered grows exponentially.

## The problem

Many features are irrelevant, noisy or correlated.

## The cost

More computation, more parameters, more variance — and weaker distance-based methods.

## The response

**Feature selection** keeps a subset; **feature extraction** builds new combined dimensions.

Every preprocessing step **trades information for tractability**. The trade is usually worth it — but it is a modelling decision, and it should be recorded as one.

Point forward: projections (PCA, MDS) are the formal treatment of feature extraction, and clustering suffers directly from the curse — and from unscaled attributes. This unit is the motivation both later lectures rely on.

## Image Credits & References

Cleveland, William S. and McGill, Marylyn E. and McGill, Robert (1988). The Shape Parameter of a Two-Variable Graph. *Journal of the American Statistical Association*.