

Data Science

Lecture 08: Cluster Analysis

Prof. Dr. Michael Granitzer

Converted from LNKB-tex part-en-ds-chp6-clustering.tex. Builds on Lecture 07 (attribute types decide permissible distances) and Lecture 04 (the evaluation module reuses precision and recall).



Where We Left Off

Lecture 07 ended with a representation: typed, cleaned, scaled attributes. Today we ask that data for groups, with no label to check against — so the only definition of "group" is the one we hand the algorithm.

Lectures 04–06

labels to learn from



Today

no labels



The assumption

a distance, and a shape of cluster


Precision and recall from Lecture 04 return in Module 5 — but only where labels happen to exist.

Speaker notes

Say what changes: in Lectures 04–06 an error on unseen data could tell a good model from a bad one. Here every algorithm returns groups, and the check has to come from somewhere else. The churn shop's customers are the example: a segmentation is only as meaningful as the distance behind it.



Argument Map



Speaker notes

With no target, the distance defines similarity, so scaling changes the clusters before any algorithm runs. Three families, three definitions of "cluster": a hierarchy, with the linkage as a second choice; compactness around a centre (k-means: always k clusters, always an answer); density (DBSCAN: arbitrary shapes, and the right to say "noise"). Validation checks consistency with the distance you chose, not truth. That is the takeaway.



Module 1

Learning Without Labels

Guiding Question: *With no target concept, what makes one grouping of the data better than another?*

Learning Without a Target

Supervised

A **target concept** is given: each example carries a label, and the objective is to reproduce it on new data.

Unsupervised

No target concept. The data carries no labels, and the objective must come from the structure itself.

Cluster analysis is the unsupervised classification of a set of objects into groups such that objects **within** a group are similar and objects in **different** groups are dissimilar.

The definition contains the whole difficulty: "similar" is not defined. Ask what a cluster of customers *is* — the answer depends on a distance the analyst chooses, which is the epistemic-modelling point (P3) in its sharpest form.

The Stages of a Cluster Analysis

Represent

Choose the attributes and their scaling.

Measure

Choose a distance or similarity.

Group

Run an algorithm that optimises some criterion.

Interpret

Decide whether the groups mean anything.

Every clustering algorithm **always returns clusters**, including on data that has none. The output is never evidence that structure exists.

Clustering has no ground truth. The **representation and the distance encode what "similar" means** — the algorithm only optimises the criterion you handed it.

Stage 4 is the one that gets skipped. Return to it whenever a clustering is scored: internal measures score the criterion, not the meaning.

What a Distance Must Satisfy

A **metric** d satisfies, for all $\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3$:

1. $d(\mathbf{x}_1, \mathbf{x}_2) \geq 0$ — non-negativity
2. $d(\mathbf{x}_1, \mathbf{x}_2) = 0 \Leftrightarrow \mathbf{x}_1 = \mathbf{x}_2$ — identity
3. $d(\mathbf{x}_1, \mathbf{x}_2) = d(\mathbf{x}_2, \mathbf{x}_1)$ — symmetry
4. $d(\mathbf{x}_1, \mathbf{x}_3) \leq d(\mathbf{x}_1, \mathbf{x}_2) + d(\mathbf{x}_2, \mathbf{x}_3)$ — triangle inequality

The pairwise distances of n objects form the **distance matrix**, which is what most clustering algorithms actually consume.

Not every useful similarity is a metric — cosine similarity and many domain-specific scores violate the triangle inequality, and some algorithms depend on it. Worth naming so students check before assuming.

Numeric Attributes: the Minkowski Family

$$d(\mathbf{x}_1, \mathbf{x}_2) = \left(\sum_{i=1}^p |x_{1i} - x_{2i}|^r \right)^{1/r}$$

$$r = 1$$

Manhattan — sum of absolute differences. Robust to single large deviations.

$$r = 2$$

Euclidean — straight-line distance. The default.

$$r \rightarrow \infty$$

Chebyshev — the largest single-attribute difference dominates.

Minkowski distances are **not scale-invariant**. An attribute measured in metres and one in millimetres do not contribute equally — standardise first, or the unit of measurement chooses your clusters.

Binary Attributes

With f_{11} the count of attributes present in both objects, f_{00} absent in both, and f_{01}, f_{10} present in exactly one:

Simple Matching Coefficient

$$\text{SMC} = \frac{f_{11} + f_{00}}{f_{11} + f_{00} + f_{01} + f_{10}}$$

Counts **shared absences** as agreement.

Jaccard Coefficient

$$J = \frac{f_{11}}{f_{11} + f_{01} + f_{10}}$$

Ignores shared absences.

For sparse data the difference is decisive: two shopping baskets agree on thousands of products **neither** contains. SMC would call them nearly identical; Jaccard would not.

Mini-Exercise

Two documents over a 10,000-word vocabulary: A contains 30 words, B 35, and they share 20. SMC comes out at 0.9975. What is Jaccard — and which of the two is the honest similarity?

The distance measure **is** the definition of similarity. Scale sensitivity and the treatment of shared absences change the clustering before any algorithm runs.

Mini-Exercise

Two documents over a 10,000-word vocabulary: A contains 30 words, B 35, and they share 20. SMC comes out at 0.9975. What is Jaccard — and which of the two is the honest similarity?

$J = \frac{20}{20+10+15} = \frac{20}{45} \approx 0.44$. Jaccard is honest: SMC's 99.75% comes almost entirely from the 9,955 words **neither** document uses.

The distance measure **is** the definition of similarity. Scale sensitivity and the treatment of shared absences change the clustering before any algorithm runs.

The Churn Shop: Segmenting Customers

	age (years)	spend (€/year)
A	25	1,000
B	65	1,050
C	26	1,400

Euclidean distances, recompute on the board: raw, $d(A,B) = \sqrt{40^2 + 50^2} \approx 64$ and $d(A,C) = \sqrt{1^2 + 400^2} \approx 400$. Scaled (divide age by 15 and spend by 1,000, roughly their standard deviations in a customer base): $d(A,B) = \sqrt{2.67^2 + 0.05^2} \approx 2.67$ and $d(A,C) = \sqrt{0.067^2 + 0.4^2} \approx 0.41$.

This is the scaling step from Lecture 07's preprocessing, now with a consequence. Neither answer is "correct": the scaled one says age and spend should count about equally, which is itself an assumption about what makes two customers similar.

The Churn Shop: Segmenting Customers

	age (years)	spend (€/year)
A	25	1,000
B	65	1,050
C	26	1,400

RAW UNITS

$d(A, B) \approx 64$, $d(A, C) \approx 400$: **A groups with B**, the 65-year-old. Spend decides.

SCALED: AGE $\div$ 15, SPEND $\div$ 1,000

$d(A, B) \approx 2.7$, $d(A, C) \approx 0.4$: **A groups with C**. Age counts.

The Churn Shop: Segmenting Customers

	age (years)	spend (€/year)
A	25	1,000
B	65	1,050
C	26	1,400

RAW UNITS

$d(A, B) \approx 64, d(A, C) \approx 400$: **A groups with B**, the 65-year-old. Spend decides.

SCALED: AGE ÷ 15, SPEND ÷ 1,000

$d(A, B) \approx 2.7, d(A, C) \approx 0.4$: **A groups with C**. Age counts.

Same customers, opposite segments: the units were a choice.



Three Families

Hierarchical

Build a **tree** of nested groupings, by merging (agglomerative) or splitting (divisive). No k required in advance.

Partitional / exemplar-based

Split the data into k groups, each summarised by a representative, optimising a criterion. Requires k .

Density-based

Clusters are **dense regions** separated by sparse ones. Finds arbitrary shapes and marks noise; requires density parameters instead of k .

No family is best. They differ in what they assume a cluster *is* — nested, compact-around-a-centre, or dense — and that assumption is the inductive bias.

Choosing an algorithm family means choosing a **definition of cluster**. The question is never "which is best" but "which shape of structure am I claiming exists".

Foreshadow the failure examples: k-means on concentric rings, DBSCAN on varying density. Each family's failure case is exactly its assumption being violated.



Module 2

Hierarchical Clustering

Guiding Question: Can we avoid committing to a number of clusters — and what does that cost?

The Algorithm

HAC

Start with every object as its own cluster: $\mathcal{C} = \{\{x_1\}, \dots, \{x_n\}\}$.

while $|\mathcal{C}| > 1$: find the closest pair $\{C, C'\} = \arg \min_{C_i \neq C_j} d_{\mathcal{C}}(C_i, C_j)$ replace them by their union $C \cup C'$

The merge order is recorded as a **dendrogram** — a tree whose height shows the distance at which each merge happened. Cutting it at a chosen height yields a clustering.

No k is needed to run the algorithm — only to **read** the result.

The dendrogram is the real output: a whole family of clusterings, one per cut height. That is HAC's distinguishing advantage over k-means, and the reason it survives despite its cost.

The Linkage Criterion

The object distance d is given; d_c — the distance between **clusters** — must still be defined, and it changes everything.

Single link

distance of the **closest** pair.

Finds elongated shapes prone to chaining.

Complete link

distance of the **farthest** pair.

Prefers compact, roughly equal-sized clusters.

Group average

mean distance over all cross-pairs.

A compromise between the two extremes.

Ward

merge that increases **within-cluster variance** least.

Compact, similarly-sized clusters; the common default.

Same data, same distance, four different dendrograms. Have students predict which criterion produces long thin clusters before the chaining slide shows it.

Mini-Exercise

Four points on a line at 0, 1, 5, 6. The first merges give $\{0, 1\}$ and $\{5, 6\}$ at height 1. At what height do these two clusters merge under **single link**, and under **complete link**?

Mini-Exercise

Four points on a line at 0, 1, 5, 6. The first merges give $\{0, 1\}$ and $\{5, 6\}$ at height 1. At what height do these two clusters merge under **single link**, and under **complete link**?

Single link: $d(1, 5) = 4$, the closest cross-pair. Complete link: $d(0, 6) = 6$, the farthest. Same clusters, different heights: the linkage rescales the dendrogram and so moves where a cut lands.

Single Link's Characteristic Failure

Single link merges on the **closest pair only**, so a thin bridge of points between two dense groups joins them: each point is close to the next, and the chain propagates.

The same property, useful

Elongated and curved clusters are found: single link follows the shape of the data instead of imposing a compact one.

The same property, harmful

A few noise points between two well-separated groups merge them into one long cluster nobody wanted.

One property seen from two sides, not a bug: single link is **transitive** — near-of-near counts as near. This is the **chaining effect**.

Ask which real domains have genuinely chain-shaped clusters — road networks, river systems, gene sequences. There chaining is the desired behaviour. Choose single link when clusters may be elongated; avoid it when the data is noisy.

Merges Are Irrevocable

HAC is **greedy**: it takes the locally closest pair and never reconsiders. Two objects merged early can never be separated later.

The consequence

The dendrogram is **nested by construction**, whether or not the data is. An early mistake is inherited by every later level.

The comparison

ID3's greedy splitting, in the opposite direction: one grows a tree top-down, the other bottom-up, and neither backtracks.

Centroid linkage can even produce **inversions** — a later merge at a *lower* height. Single, complete, average and Ward linkage are monotonic.

The parallel to ID3 is worth drawing explicitly — students have just met greedy non-backtracking search in the decision-tree lecture. Read a dendrogram as one possible history of merges, not as the structure of the data.

What HAC Gives You

Strengths

No k required; a **whole family** of clusterings in one dendrogram; runs on a distance matrix alone, so objects need no vector representation.

Weaknesses

$O(n^2)$ memory for the distance matrix and $O(n^2 \log n)$ time or worse, so it does not scale. Greedy, irrevocable merges. Strong dependence on the linkage.

HAC returns a **dendrogram**, not a clustering, and the **linkage** is a modelling decision as consequential as the distance. It trades **scalability for insight**: greedy and $O(n^2)$, but it shows the whole hierarchy without a k .

Speaker notes

Use HAC when n is moderate, when you want to inspect the structure rather than commit to a k , or when the objects only admit pairwise dissimilarities. The "distance matrix alone" property is the underrated one: strings, graphs and sequences can be clustered without a vector embedding. k -means needs means, so it needs a vector space; k -medoids does not — it too works from a distance matrix, but it still asks for k .



Module 3

k-Means

Guiding Question: *k-means always returns k clusters. When is that an answer, and when is it an artefact?*

The Idea

Represent each cluster by one **exemplar** — a centroid — and assign every object to the nearest one. The quality criterion is the total squared distance to the assigned exemplar:

$$\text{SSE} = \sum_{j=1}^k \sum_{\mathbf{x} \in C_j} d(\mathbf{x}, \boldsymbol{\mu}_j)^2$$

Minimising SSE over all assignments is **NP-hard**. k-means is a heuristic that finds a local optimum — which is why its result depends on where it starts.

Name the criterion before the algorithm: k-means is not "the clustering algorithm", it is a greedy optimiser of one particular objective. Different objective, different notion of cluster.


The Algorithm

k-means (Lloyd)

1. Choose k initial centroids.
2. **Assign** each object to the nearest centroid.
3. **Move** each centroid to the mean of its assigned objects.
4. Repeat 2–3 until the assignment stops changing.



Move each centroid to the mean of its points (round 3)



Stepped viz (P6): odd clicks assign, even clicks move. The centroids start deliberately badly. Ask before each assignment step which points will change colour — students track the boundary moving, which is the mechanism.

Properties

Fast

$O(nkp)$ per iteration, and few iterations in practice — it never needs the full distance matrix.

Converges

SSE never increases, and there are finitely many assignments, so the procedure terminates.

To a local optimum

The result depends on initialisation. Standard practice: restart several times and keep the best SSE.

k-medoids uses an actual data point as the exemplar instead of the mean. It is slower but works with any distance matrix — including where a "mean" is undefined.

Mini-Exercise

k-means moves each centroid to the **mean** of its objects. Why the mean — and what would change with the **median**?

Speaker notes

Solution on the same slide, click-gated. The median half is the conceptual payoff and connects to the mean/median robustness discussion from descriptive statistics — same trade, again. If someone asks why k-means must stop: neither step can increase SSE, and the assignment changes only when SSE strictly drops, so no assignment repeats among finitely many.



Mini-Exercise

k-means moves each centroid to the **mean** of its objects. Why the mean — and what would change with the **median**?

The mean is exactly the point minimising the **sum of squared** distances to a set, so the move step solves its half of the SSE objective optimally. The median minimises the sum of **absolute** distances: you would be optimising a different, more outlier-robust criterion (k-medians).

Takeaway

k-means alternates **assign** and **move**, minimising SSE to a local optimum. The mean is not a convenience — it is the minimiser of the criterion, which is why changing the criterion changes the update.

Four Failure Modes

Non-globular shapes

Concentric rings, crescents, elongated bands. SSE rewards **compact, spherical** groups, so k-means cuts such shapes across the middle.

Differing sizes

A large sparse cluster next to a small dense one gets split, because splitting the large one lowers SSE more than respecting the boundary.

Differing densities

The criterion counts squared distance, not density — so a diffuse cluster is torn apart to reduce it.

Outliers

The mean is not robust: a single far point drags its centroid away from the group it is supposed to represent.

Every failure is the same failure: SSE encodes "compact and spherical, of similar spread". When that assumption holds, k-means is excellent; when it does not, the algorithm is not wrong — it is answering a different question.

What to Do About It

Change the representation

Scale or transform features so the clusters *become* globular.

Change the algorithm

Density-based methods for arbitrary shapes; k-medoids for outlier robustness.

Change k

Over-clustering and then merging can recover a non-globular shape from several spherical pieces.

k-means fails exactly where its criterion stops matching the data. **Diagnose by plotting** — the failures are obvious visually and invisible in the SSE.

The last sentence links back to the visual-analytics units: SSE keeps falling as k grows regardless of whether the clusters are meaningful, so the number cannot detect the failure that the plot shows at a glance.

Module 4

Density-Based Clustering

Guiding Question: *Some points belong to no group at all. Which algorithms are even allowed to say so?*

A Different Definition of Cluster

A cluster is a **dense region** of points, separated from other dense regions by sparser ones. Density is defined by two parameters:

ϵ

the radius of a neighbourhood

minPts

how many points that neighbourhood must contain to count as dense

Core point

Has at least *minPts* points within ϵ — it is inside a dense region.

Border point

Within ϵ of a core point, but not itself dense.

Noise point

Neither — and DBSCAN **leaves it unassigned.**

Speaker notes

The third category is what distinguishes DBSCAN from everything before it: it is the first algorithm in the course allowed to say "this point belongs to no cluster". Ask why k-means cannot do that.



The Algorithm

1. Label every point as core, border or noise.
2. Connect core points that lie within ε of each other — each connected group becomes a cluster.
3. Attach each border point to a cluster of an adjacent core point.
4. Leave noise points unassigned.

Strengths

Finds **arbitrarily shaped** clusters;
identifies noise explicitly; k is not required.

Weaknesses

Two parameters that interact and are hard to set; it struggles when clusters have **very different densities**, since one ε must serve all.

Mini-Exercise

DBSCAN returns one giant cluster plus a few noise points. You halve ε and get 40 tiny clusters and mostly noise. What went wrong in each run?

DBSCAN replaces "compact around a centre" with "**dense and connected**", which buys arbitrary shapes and an explicit noise category — at the price of two parameters and an assumption of uniform density.

If there is time: a systematic choice plots each point's sorted distance to its k-th nearest neighbour and takes the elbow, where the curve rises sharply. It previews the elbow heuristic that returns in the evaluation module for choosing k.

Mini-Exercise

DBSCAN returns one giant cluster plus a few noise points. You halve ε and get 40 tiny clusters and mostly noise. What went wrong in each run?

First run: ε too large — every dense region reaches the next, so all connect into one.

Second: too small — few points reach *minPts* neighbours, so nearly everything is noise. ε sets the density scale.

DBSCAN replaces "compact around a centre" with "**dense and connected**", which buys arbitrary shapes and an explicit noise category — at the price of two parameters and an assumption of uniform density.

Module 5

Evaluating a Clustering

Guiding Question: *There is no ground truth. How do you know whether the clusters mean anything?*

Three Kinds of Validation

External

Compare against **known labels**. Measures agreement with a ground truth we normally do not have.

Internal

Score the clustering using **the data alone** — compactness and separation. No labels needed.

Relative

Compare **several clusterings** of the same data to choose among them, e.g. across values of k .

Internal measures score the same geometric properties the algorithms optimise. A clustering that scores well on compactness may simply reflect that k-means minimised compactness — not that the groups mean anything.

The warning is the central caution of the module: the evaluation and the objective can be the same quantity in disguise. External validation escapes it, which is why labelled benchmark data matters even for unsupervised methods.

The Prior Question

Before asking *how good* a clustering is, ask whether the data has **cluster tendency** at all: uniformly random points produce clusters under every algorithm, and internal measures will happily score them.

Choose the validation type by what you have: **labels** → external, **only data** → internal (with suspicion), **several candidates** → relative. And check that structure exists before scoring it.

Reusing Precision and Recall

Given known classes, treat each **cluster** as a prediction of a **class**:

Recall of class i in cluster j

$$\text{rec}(i, j) = \frac{n_{ij}}{n_i}$$

the share of class i that landed in cluster j

Precision of class i in cluster j

$$\text{prec}(i, j) = \frac{n_{ij}}{n_j}$$

the share of cluster j that belongs to class i

$$F_1(i, j) = \frac{2 \cdot \text{prec}(i, j) \cdot \text{rec}(i, j)}{\text{prec}(i, j) + \text{rec}(i, j)}$$

Each class is matched to the cluster maximising its F_1 , and the overall score averages these, weighted by class size.

Speaker notes

Deliberate reuse from Lecture 04 (spiral): the same two conditional probabilities, now with clusters standing in for predicted labels. The extra wrinkle is the matching step — clusters carry no names, so they must be assigned.



Why the Matching Step Matters

A clustering has **no label for its clusters**. Any evaluation must first decide which cluster corresponds to which class — and different matching rules give different scores for the same clustering.

Splitting a class

Two pure clusters covering one class score high precision, low recall each. The clustering is arguably *right* but the score punishes it.

Merging two classes

One cluster containing both scores high recall, low precision. F_1 punishes this too — as it should.

Purity — The Simplest Matching

Give every cluster the label of its **majority class** and count the hits:

$$\text{purity} = \frac{1}{n} \sum_j \max_i n_{ij}$$

Purity only looks at precision. Splitting every class into many small clusters never lowers it — with $k = n$ singletons, purity is 1. It needs the recall side, which F_1 supplies.

Purity is the one number practitioners quote most often, so students should know its blind spot: it is the external analogue of SSE falling with k .

Mini-Exercise

100 objects, classes A and B (50 each). The clustering returns three pure clusters: [50 A], [25 B], [25 B]. What is class B's F_1 — and is the clustering wrong?

The interpretive point (P2): external validation assumes the labels are the truth. When clustering is used for discovery, disagreeing with the labels may be the finding. For the weighted overall F1: A scores 1, B 0.67, so 0.83.

Mini-Exercise

100 objects, classes A and B (50 each). The clustering returns three pure clusters: [50 A], [25 B], [25 B]. What is class B's F_1 — and is the clustering wrong?

Best cluster for B: precision 1, recall 0.5, so $F_1 = \frac{2 \cdot 0.5}{1.5} \approx 0.67$ (purity is still 1). Not necessarily wrong: B was split, and if B has two sub-types, the clustering found structure the labels do not record.

Mini-Exercise

100 objects, classes A and B (50 each). The clustering returns three pure clusters: [50 A], [25 B], [25 B]. What is class B's F_1 — and is the clustering wrong?

Best cluster for B: precision 1, recall 0.5, so $F_1 = \frac{2 \cdot 0.5}{1.5} \approx 0.67$ (purity is still 1). Not necessarily wrong: B was split, and if B has two sub-types, the clustering found structure the labels do not record.

External validation reuses precision and recall, plus a **matching** of clusters to classes. A low score can mean a bad clustering — or a label set coarser than the structure in the data.

Comparing Two Matrices

Proximity matrix

The pairwise **similarities** between objects: high for close pairs.

Incidence matrix

The **ideal** structure implied by the clustering: 1 if two objects share a cluster, 0 otherwise.

A good clustering makes the two agree: pairs in the same cluster should be similar, pairs in different clusters not. Their **correlation** measures that agreement in one number, and good means **close to +1**.

With a **distance** matrix instead, the sign flips: good means strongly **negative** correlation (same cluster $\leftrightarrow$ small distance).

The trick is worth naming: an arbitrary clustering is turned into a matrix so that two structures can be compared numerically. The same device appears in graph comparison and in kernel methods.

Reading It

High correlation

The similarity structure and the cluster structure agree — the partition follows the geometry.

What it does not tell you

Whether the geometry is meaningful. Correlation with a **self-imposed** ideal structure is not evidence of real groups.

Visual variant: sort the proximity matrix by cluster and display it. Genuine clusters appear as high-similarity blocks along the diagonal — and the absence of blocks is as informative as their presence.

Internal validation asks whether the partition **agrees with the distances**. It rewards consistency with the chosen distance measure, not correctness.

Two Quantities

Cohesion

How tight a cluster is — the within-cluster sum of distances (SSE is one form of it).

Lower is better.

Separation

How far apart clusters are — the between-cluster distance. **Higher is better.**

Every internal measure combines these two, and they trade against each other: splitting a cluster always improves cohesion and usually harms separation.

The Silhouette Coefficient

For a single object, with a its mean distance to its **own** cluster and b its mean distance to the **nearest other** cluster:

$$s = \frac{b - a}{\max(a, b)} \in [-1, 1]$$

Value	Reading
$s \approx 1$	well inside its cluster, far from the others
$s \approx 0$	on the boundary between two clusters
$s < 0$	closer to another cluster than to its own — probably misassigned

Example: $a = 1$, $b = 4$ gives $s = 3/4 = 0.75$; swap them and $s = -0.75$.

Speaker notes

The per-object nature is what makes the silhouette useful in practice: it localises the problem to specific points and specific clusters, where SSE gives one global number that cannot be diagnosed.



Reading the Silhouette

Average the silhouette over all objects for one score — but **plot the distribution** too. A good average can hide a cluster whose members are all near zero.

Unlike SSE, the average silhouette does not simply improve with k : it peaks where clusters are both tight and well separated, so it can select k .

Cohesion and separation are the two axes of every internal measure. The **silhouette** combines them per object, which is why its distribution is more informative than its mean.

Worst Separation over Worst Cohesion

$$D = \frac{\min_{i \neq j} \delta(C_i, C_j)}{\max_k \Delta(C_k)}$$

where δ is the distance between two clusters and Δ the diameter of a cluster. **Higher is better.**

The numerator is the **closest** pair of clusters; the denominator is the **widest** cluster. The index is therefore governed entirely by the two worst cases.

Contrast with the silhouette, which averages: Dunn is a minimax criterion. Ask which is preferable when one bad cluster among twenty matters — and when it does not.

Consequences

The virtue

A single badly separated pair or one sprawling cluster **cannot** be hidden by good behaviour elsewhere.

The cost

Extremely sensitive to outliers and noise: one distant point inflates a diameter and collapses the index.

Dunn scores the **worst case**, so it is strict and fragile. Use it when no bad cluster is acceptable; prefer averaging measures when a few outliers are tolerable.

The Problem

SSE **always** decreases as k grows — with $k = n$ every point is its own cluster and SSE is zero. So the criterion cannot choose k by itself.

The same structure as training error against model capacity: the objective improves monotonically with complexity, so **the objective cannot select the complexity**.

Deliberate parallel to overfitting (P7 spiral, third appearance after polynomials and trees). Students who see it will not be surprised that clustering needs an external criterion too.

The Elbow

Plot SSE against k and look for the point where the curve **stops falling steeply** — where additional clusters buy little.

The reasoning

Up to the elbow, each new cluster splits a genuinely heterogeneous group. Past it, new clusters merely subdivide homogeneous ones.

The catch

The elbow is often **not sharp**, and different observers read different values. It is a heuristic, not a criterion.

Alternatives

Silhouette over k

Pick the k maximising the average silhouette — a genuine maximum rather than a bend.

Gap statistic

Compare SSE against the SSE expected on **random data with no clusters**.

Domain constraints

Often decisive: how many segments can the business actually act on?

SSE alone cannot select k : it always improves with k . The elbow is a **heuristic reading** of that trade-off; criteria with a genuine optimum over k (silhouette, Dunn, gap statistic) formalise it, and the domain often overrules them all.

End the lecture on the last card: the number of clusters is very often decided by what can be done with them, not by the geometry. That is the epistemic point the whole module has been building toward.



