

Data Science

Lecture 09: Association Rule Mining

Prof. Dr. Michael Granitzer

Where We Left Off

Lecture 08 grouped *objects* by similarity. Today groups *items* by co-occurrence — and every measure is a probability from Lecture 02 in disguise.

Support

a joint probability



Confidence

a conditional one




Lift

a comparison with independence


The base-rate trap from Lecture 02 returns: a rule can hold 80% of the time and say nothing.

Argument Map

Lecture 09 · argument map



Course arc
assume → check



Module 1

Itemsets and Their Search Space

Guiding Question: *Counting how often items occur together is easy. Why does it become hard with ten thousand products?*

The Question

Given a set of transactions — shopping baskets, clicked pages, co-occurring symptoms — which **items appear together** more often than chance, and which **imply** which?

Support

$$P(A \cap B) \geq \sigma_{\min}$$

How often do these items occur in the same transaction?

Confidence

$$P(B \mid A) \geq \gamma$$

Given A is in the transaction, how often is B too?

Both are quantities the students already know: support is a **joint** probability, confidence a **conditional** one, both estimated by counting transactions.

Formalisation

Let I be a set of **items** and X a set of **transactions**, each $x \subseteq I$.

- An **itemset** is any $A \subseteq I$; a **k -itemset** has $|A| = k$.
- Its **support** is $\sigma(A) = \frac{|\{x \in X : A \subseteq x\}|}{|X|}$.
- A is **frequent** if $\sigma(A) \geq \sigma_{\min}$.

$\sigma(A \cup B)$ counts baskets that contain **all** items of A **and** of B . A union of *itemsets* is an intersection of *events*: $\sigma(A \cup B) = P(A \cap B)$, not "A or B".

The Search Space

All itemsets over I form a **lattice** ordered by inclusion — with $2^{|I|}$ elements.

A supermarket with 10,000 products has 2^{10000} candidate itemsets. Enumeration is impossible; the entire chapter is about avoiding it.

Support and confidence are joint and conditional probabilities estimated by counting. The difficulty is not the measure but the **size of the search space**.

Module 2

The A-Priori Algorithm

Guiding Question: *What single structural fact about support lets us discard most of an exponential space without looking at the data?*

The Property That Makes It Possible

A-priori property (anti-monotonicity). Every subset of a frequent itemset is frequent — equivalently: if an itemset is **infrequent**, **every superset of it is infrequent too**.

Adding an item can only ever *reduce* the number of transactions containing the set. So an infrequent itemset prunes its entire superset subtree from the search.

The Algorithm

A-Priori

L^1 = all frequent 1-itemsets.

for $k = 2$; $L^{k-1} \neq \emptyset$; $k++$: **generate** candidates $C_k = \{a \cup b \mid a \in L^{k-1}, b \in L^1, b \not\subseteq a\}$
prune any candidate having an infrequent subset **count** the support of the survivors in one pass over the transactions $L^k = \{c \in C_k \mid \sigma(c) \geq \sigma_{\min}\}$

Generate

build k -candidates from frequent $(k-1)$ -itemsets

Prune

discard candidates with any infrequent subset — no data access needed

Count

one database pass per level

Cost and Improvements

What it costs

One pass over the transactions **per level k** , plus candidate generation. The bottleneck is I/O and the number of candidates.

How it is improved

Hashing candidates into buckets, transaction reduction (dropping transactions containing no frequent itemset), partitioning and sampling.

A-Priori turns one structural fact — **infrequent sets have no frequent supersets** — into a level-wise search that prunes most of an exponential space without touching the data.

Mini-Exercise

$\{\text{bread, milk}\}$ has support 0.18, and $\sigma_{\min} = 0.22$.

What do you know about $\{\text{bread, milk, eggs}\}$ — and how many database passes did it take to know it?

Mini-Exercise

$\{\text{bread, milk}\}$ has support 0.18, and $\sigma_{\min} = 0.22$.

What do you know about $\{\text{bread, milk, eggs}\}$ — and how many database passes did it take to know it?

It is **infrequent**: adding eggs cannot raise the support above 0.18. And it took **no pass at all** — the property alone decides. That is where A-Priori's savings come from.

Module 3

From Itemsets to Rules

Guiding Question: *A rule holds 80% of the time — and tells you nothing. How is that possible, and what should you measure instead?*

Generating Rules

An **association rule** $A \Rightarrow B$ (with $A \cap B = \emptyset$) has

$$\text{support} = \sigma(A \cup B), \quad \text{confidence} = \frac{\sigma(A \cup B)}{\sigma(A)}$$

Every frequent itemset F is split in all possible ways into A and $B = F \setminus A$; rules meeting the confidence threshold are kept.

Support comes from the itemset and is already known. Only confidence must be computed — so rule generation needs **no further database passes**.

Confidence Is Not Enough

Suppose 80% of all transactions contain milk. A rule $\text{bread} \Rightarrow \text{milk}$ with **confidence 0.80** has learned nothing: milk appears in 80% of baskets regardless of bread.

$$\text{lift}(A \Rightarrow B) = \frac{P(B \mid A)}{P(B)} = \frac{\sigma(A \cup B)}{\sigma(A)\sigma(B)}$$

Lift compares the joint probability against the product of the marginals — the **definition of independence** from Lecture 02, written as a ratio. It is symmetric: $\text{lift}(A \Rightarrow B) = \text{lift}(B \Rightarrow A)$.

Reading Lift

lift	meaning
> 1	positive association: A and B co-occur more often than independence predicts
$= 1$	independent: the rule carries no information
< 1	negative association: A and B co-occur less often than independence predicts

Lift is a ratio, not a test. On a rule with low support it rests on a handful of baskets, and a lift of 1.5 can be chance. Before acting, check it for **significance** — the hypothesis tests of Lecture 02 apply.

Mini-Exercise

A retailer reports nappies $\Rightarrow$ beer with support 0.02 and confidence 0.60; beer is in 50% of all baskets.

What is the lift — and why is the rule still a weak basis for action?

Mini-Exercise

A retailer reports nappies $\Rightarrow$ beer with support 0.02 and confidence 0.60; beer is in 50% of all baskets.

What is the lift — and why is the rule still a weak basis for action?

$\text{lift} = 0.60 / 0.50 = 1.2$: beer is 20% more likely with nappies. But the rule covers only 2% of baskets, so the effect may be chance — test it, and check that it persists.

Mini-Exercise

A retailer reports nappies $\Rightarrow$ beer with support 0.02 and confidence 0.60; beer is in 50% of all baskets.

What is the lift — and why is the rule still a weak basis for action?

$\text{lift} = 0.60 / 0.50 = 1.2$: beer is 20% more likely with nappies. But the rule covers only 2% of baskets, so the effect may be chance — test it, and check that it persists.

Confidence alone rewards **popular** consequents. **Lift** corrects for the base rate by comparing against independence — the same correction Bayes' theorem makes for priors.
Low support calls for a significance check.

The Churn Shop's Baskets

The shop delivers grocery boxes. Of **100** boxes, **80** contain milk, **20** contain cereal, **16** contain both.

The rule $\text{cereal} \Rightarrow \text{milk}$ holds with confidence 80%. Is it worth a "customers who order cereal also order milk" banner?

The Churn Shop's Baskets

The shop delivers grocery boxes. Of **100** boxes, **80** contain milk, **20** contain cereal, **16** contain both.

The rule $\text{cereal} \Rightarrow \text{milk}$ holds with confidence 80%. Is it worth a "customers who order cereal also order milk" banner?

$$\text{conf} = \frac{16}{20} = 0.8 \quad \text{lift} = \frac{0.8}{0.8} = 1.0$$

No: milk is in 80% of **all** boxes. Cereal buyers take milk exactly as often as everyone else — the rule restates the base rate.