

Exercise Sheet 2

Graph Theory

5942UE Network Science

18 March 2026

Prof. Dr. Michael Granitzer Dr. Kanishka Ghosh Dastidar Dr. Jelena Mitrovic

Python Setup

Python Setup 1/2: Install Conda

We will use Python with [NetworkX](#).

Install **Miniconda** from:

<https://docs.conda.io/en/latest/miniconda.html>

Choose your system:

- **Windows:** download the installer, run it, then open **Anaconda Prompt**
- **Mac:** download the .pkg installer and install normally
- **Linux:** download the .sh installer and run:

```
bash Miniconda3-latest-Linux-x86_64.sh
```

After installation, restart your terminal.

Python Setup 2/2: Create Environment

Check Conda:

```
conda --version
```

Create and activate the course environment:

```
conda create -n network-science python=3.11
conda activate network-science
conda install -c conda-forge networkx matplotlib numpy scipy pandas jupyter
```

Start Jupyter:

```
jupyter notebook
```

Test in a notebook:

```
import networkx as nx
import matplotlib.pyplot as plt

G = nx.karate_club_graph()
nx.draw(G, with_labels=True)
plt.show()
```

If a graph appears, the setup works.



2.A Properties of Networks

Exercise 2.A.1: The Five Modeling Questions

A hospital wants to study how infections spread among patients and staff.

Answer the five modeling questions from the lecture for this system:

1. **Nodes:** What should each node represent? Is there more than one natural choice?
2. **Edges:** What should an edge represent? (Hint: an edge should encode a transmission-relevant relationship.)
3. **Direction:** Is the relation directed or undirected? Does it depend on how you define the edge?
4. **Weight:** Should edges have weights? If so, what should the weight encode?
5. **Time:** Does this network change over time? Does the order of events matter for your analysis?

Finally: what does your modeling choice **reveal**, and what does it **hide**?

Exercise 2.A.2: Two Meanings of the Same Network

The ARPANET had two distinct kinds of nodes: **universities/labs** and **routers**. Depending on which you model as nodes, you get different graphs of the same physical system.

- **Model R** (router-level): nodes are routers; edges are physical cables.
 - **Model U** (institution-level): nodes are institutions; an edge connects two institutions if at least one router-level path exists between them.
1. For a path $A \rightarrow B$ in Model R, what does it mean in the physical system?
 2. For the same path in Model U, what does it represent? How do the two models differ in what they express?
 3. If your question is "*Can MIT reach Stanford in at most 3 hops?*", which model do you use and why?
 4. If your question is "*Which physical cable, if cut, would disconnect the ARPANET?*", which model do you use and why?
 5. What general principle do these questions illustrate about modeling choices?

Exercise 2.A.3: Bipartite vs Projected Networks

Consider a university course enrollment system. Two researchers model it differently:


- **Model A:** A bipartite graph $G_A = (S, C, E)$ where S = students, C = courses, and $(s, c) \in E$ means student s is enrolled in course c .
 - **Model B:** A projected graph $G_B = (S, E')$ where two students are connected if they share at least one course.
1. Draw a small example: 3 students (Alice, Bob, Carol) and 3 courses (Math, CS, Physics). Alice takes Math and CS; Bob takes CS and Physics; Carol takes Math.
 2. Draw both G_A and G_B for this example.
 3. What does G_A let you analyse that G_B does not? Give a concrete example question.
 4. What does G_B let you analyse that G_A does not? Give a concrete example question.
 5. What information is irreversibly lost in the projection from G_A to G_B ?

2.B Representing Graphs

Exercise 2.B.1: From Picture to Formal Notation

Consider an undirected graph with five nodes A, B, C, D, E and edges $A-B, A-C, B-C, B-D, D-E$.

1. Write the formal set notation: $V = \dots, E = \dots$
2. Draw the adjacency matrix (nodes in alphabetical order).
3. Write the adjacency list as a Python dictionary.
4. Compute $|V|$ and $|E|$.
5. What is the degree of each node? Verify the handshaking lemma:
$$\sum_{v \in V} \deg(v) = 2|E|.$$
6. Now suppose $B-D$ becomes directed ($B \rightarrow D$). What changes?



Exercise 2.B.2: Graph Representations in Python

Task (10 min): Work with the graph from Exercise B.1 using [NetworkX](#), the standard Python package for the creation, manipulation, and study of complex networks.

```
import networkx as nx

G = nx.Graph()
G.add_edges_from([("A", "B"), ("A", "C"), ("B", "C"), ("B", "D"), ("D", "E")])
```

1. Print the adjacency matrix using `nx.to_numpy_array`.
2. Print the degree of each node. Does it match your hand calculation?
3. Compute and print the **density** of the graph. What does density = 1 mean? What does density = 0 mean?
4. Add a directed version: `DG = nx.DiGraph()`. Print in-degree and out-degree of *B* and *D*.
5. **Bonus:** Plot both graphs side by side with `matplotlib`.

Exercise 2.B.3: Reading a Network Diagram

Task (10 min): Use Python and NetworkX to build the Zachary karate club graph — one of the most studied social networks — and answer the following programmatically:

```
import networkx as nx
import matplotlib.pyplot as plt

G = nx.karate_club_graph()
```

1. How many nodes and edges does the network have?
2. What is the average degree? What does this mean for the "average member"?
3. Are there any isolated nodes (degree 0)? What would an isolated node mean socially?
4. Plot the network. Can you visually identify the two factions that the club split into?

Hint: `G.degree()` returns degree for all nodes. `nx.draw_spring` gives a reasonable layout.

2.C Paths and Cycles

Exercise 2.C.1: Walks, Paths, and Cycles

Consider the undirected graph with edges $A-B$, $B-C$, $C-D$, $D-A$, $A-C$.

Classify each sequence as a **walk**, **path**, **cycle**, or **none** (edge doesn't exist). Justify each answer.

Sequence	Classification	Reason
A, B, C, A	?	?
A, C, D, A	?	?
A, B, C, D, A	?	?
A, B, A, C	?	?
A, B, C, D, C	?	?
A, E	?	?

Then: what is the **diameter** of this graph?



Exercise 2.C.2: The Königsberg Bridge Problem

The Königsberg bridges problem (Euler, 1736) asks: can you walk through the city crossing each of the seven bridges exactly once and return to your starting point?

The city has four land masses: **North bank (N)**, **South bank (S)**, **Island (I)**, **East bank (E)**. Bridges: N–I ($\times 2$), S–I ($\times 2$), N–E ($\times 1$), S–E ($\times 1$), I–E ($\times 1$).

1. Model this as a multigraph $G = (V, E)$. Write down V and E .
2. Write the degree of each node.
3. Euler proved that an **Eulerian circuit** exists iff **every vertex has even degree**. Does Königsberg satisfy this condition?
4. An **Eulerian path** exists iff **exactly two vertices have odd degree**. Does Königsberg have an Eulerian path?
5. What does this tell us?


Exercise 2.C.3: Shortest Paths with NetworkX

Task (10 min): The ARPANET had 20 nodes in its early topology. We will work with a small synthetic version.

```
import networkx as nx
import matplotlib.pyplot as plt

arpa = nx.Graph()
arpa.add_edges_from([
    ("UCLA", "SRI"), ("UCLA", "UCSB"), ("UCLA", "UTAH"),
    ("SRI", "UTAH"), ("SRI", "BBN"),
    ("UCSB", "SRI"), ("UTAH", "MIT"),
    ("BBN", "HARVARD"), ("BBN", "MIT"), ("MIT", "HARVARD"),
])
```

1. Plot the network using `nx.draw_spring()`.
2. Find the shortest path from "UCLA" to "HARVARD". Print path and length.
3. Compute eccentricities, **diameter**, and **radius**.
4. Which node is the center? What does this mean structurally?
5. Remove "BBN" and recompute the diameter. What changed and why?




2.D Breadth-First Search (BFS)

Exercise 2.D.1: BFS by Hand — Layer Tracing

Run BFS starting from node 1 on the following undirected graph:

Edges: 1–2, 1–3, 2–4, 2–5, 3–6, 4–7, 5–7, 6–8

1. Fill in the BFS discovery table with layer and parent for each node.
2. Draw the BFS tree.
3. What is $d(1, 7)$ and $d(1, 8)$? Is there more than one shortest path to 7?
4. If we add weights $w(2-4) = 10$, $w(5-7) = 1$, does BFS still give shortest paths?



Exercise 2.D.2: Implementing BFS and Computing Distances

Task (15 min): Implement BFS from scratch and use it to answer structural questions about the Zachary karate club network.

```
from collections import deque
import networkx as nx

def bfs_distances(graph, source):
    """Return a dict {node: shortest-path distance from source}."""
    dist = {source: 0}
    queue = deque([source])
    while queue:
        node = queue.popleft()
        # --- your code here ---
    return dist

G = nx.karate_club_graph()
adj = {n: list(G.neighbors(n)) for n in G.nodes()}
```


1. Complete `bfs_distances`. Verify it against `nx.single_source_shortest_path_length`.
2. Use it to find the **eccentricity** of node 0 (the instructor "Mr. Hi").
3. What fraction of the network is reachable within 2 hops from node 0?
4. Repeat for node 33 (the administrator "Officer"). Compare eccentricities and reachable fractions. What does the difference tell you?

2.E Connectivity and Components

Exercise 2.E.1: Components and Connectivity by Hand

Consider a directed graph with nodes 1, 2, 3, 4, 5, 6 and directed edges: $1 \rightarrow 2$, $2 \rightarrow 3$, $3 \rightarrow 1$, $4 \rightarrow 3$, $5 \rightarrow 6$, $6 \rightarrow 5$.

1. List all **weakly connected components** (ignoring direction).
2. List all **strongly connected components** (following directed paths).
3. Node 4 is a "source" node. Explain what this means for information flow. Can node 4 receive information from nodes 1, 2, or 3?
4. If you removed edge $3 \rightarrow 1$, how would the SCCs change?



Exercise 2.E.2: Giant Components in Practice

Task (10 min): Use NetworkX to investigate what happens to connectivity when we progressively remove nodes from the karate club network.

```
import networkx as nx
import matplotlib.pyplot as plt
import random

G = nx.karate_club_graph()
```

1. What fraction of nodes belong to the giant component of the original network?
2. **Random failure:** repeatedly remove a randomly chosen node and record the size of the largest component after each removal. Plot the result.
3. **Targeted attack:** repeatedly remove the node with the **highest current degree**. Plot on the same axes.
4. Which strategy destroys the giant component faster? Why?

Hint: `max(nx.connected_components(G), key=len)` gives the largest component.