

# Network Science

---

## Lecture 02: Graph Theory

Prof. Dr. Michael Granitzer Dr. Kanishka Ghosh Dastidar

Dr. Jelena Mitrovic

# From Networks to Graphs

Last time we saw that networks are everywhere: social, informational, economic, technical. Today we build the formal language to describe them precisely.

## Today

- what properties do networks have?
- how do we formalize networks as graphs?
- paths and distance
- breadth-first search
- connectivity and components

# Properties of Networks

---

**Guiding Question:** What properties does a network have? What do we need to describe it?

# What Does a Node Mean?

- In a social network: a **person**
- In the Web: a **page**
- In a financial network: an **institution**
- In a transport network: a **city** or **station**

A node is whatever entity you choose to model. That choice is never neutral — it determines what you can see and what you miss.

# What Does an Edge Mean?

- Friendship: **symmetric** — if A is friends with B, B is friends with A
- Following on Twitter/X: **asymmetric** — A follows B does not mean B follows A
- Email: **directed** — sender and receiver are distinct roles
- Co-authorship: **symmetric** — both contributed to the same paper

An edge encodes a relationship. But relationships differ:

- symmetric vs. directed
- strong vs. weak
- binary vs. weighted.

# What About Time?

- Networks **grow**: new people join, new links form
- Networks **shrink**: people leave, links break
- Networks **change**: relationships strengthen or weaken
- The **order** of events matters: who connected to whom first?

Most real networks are dynamic. A static graph is a snapshot — useful, but incomplete.

# From Properties to Formal Language

We now have a list of properties we need to capture:

| Property       | Formal concept                 |
|----------------|--------------------------------|
| Entities       | Nodes (vertices)               |
| Relationships  | Edges (links)                  |
| Directionality | Directed vs. undirected graphs |
| Strength       | Weighted vs. unweighted graphs |
| Structure      | Adjacency matrix, degree       |
| Time           | Temporal / dynamic graphs      |

# Modeling Is a Choice

Before we analyse a graph, we decide how to turn a situation into one.

## From reality to graph model



Example: an affiliation setting can become a bipartite graph, or a projection, depending on the question.

Graph theory gives us a **formal language**, but the first step is still **epistemic**: choose the entities, choose the relations, and decide what abstraction is good enough for the question.

# Takeaway

Describing a network precisely requires answering:

- What are the nodes? What are the edges?
- Are they directed? Weighted?
- Changing over time?

Graph theory gives us a formal language to answer each of these questions.

# Quick Check

Pick a real-world network you interact with (e.g. your university's course system, a messaging group, a transport network). Write down:

1. What are the **nodes**?
2. What are the **edges**?
3. Is the relation **directed** or **undirected**?
4. Should edges be **weighted**? If so, what does the weight represent?
5. Does the network **change over time**?

# Quick Check — Sample Answer

## Example: University course enrollment system

1. **Nodes:** Students and courses (two types → bipartite)
2. **Edges:** "Student X is enrolled in course Y"
3. **Direction:** Directed — enrollment is an action from student to course
4. **Weights:** Could represent credit hours, or grade achieved
5. **Time:** Yes — students enroll and drop courses each semester

This is already a bipartite graph with temporal dynamics.

# Graphs as Models

---

**Guiding Question:** When is a graph as model good enough — and what gets lost?

# Why Graph Theory

We identified properties we need to describe. Now we need precise definitions. Graph theory gives us that language and lets us move from intuition to analysis.

# Formal Definition


A graph  $G = (V, E)$  consists of:

- a set of **vertices** (nodes)  $V$
- a set of **edges** (links)  $E \subseteq V \times V$

This compact notation already encodes a modeling choice: what counts as a node, what counts as a relationship.

# Undirected Graphs

**Undirected Graph.** An undirected graph treats relationships as symmetric. An edge  $u, v$  connects two vertices without orientation:  $u, v = v, u$ . Source and target are not distinguished.



Edges have no arrows, meaning the relation is symmetric (e.g. friendship).

# Directed Graphs

**Directed Graph.** An edge  $(u, v) \in E$  has a strict orientation from  $u$  to  $v$ .  $(u, v)$  does not imply  $(v, u)$ . Directions matter whenever the relation is asymmetric (e.g. following, citing).



Edges have arrows indicating direction. The same set of nodes can yield very different models depending on this choice.

# Weighted Graph

**Weighted Graph.** A graph equipped with a weight function  $w : E \rightarrow \mathbb{R}$ . The weight  $w(e)$  lets edges carry strength, cost, length, probability, or capacity.



Each edge carries a numerical value. The same topological structure can mean vastly different things depending on what the weights encode.

# Bipartite Graphs

**Bipartite Graph.** A graph  $G = (U, V, E)$  with two disjoint node sets ( $U \cap V = \emptyset$ ) where all edges connect a node from  $U$  to a node from  $V$ . Edges never connect two nodes within the same set.



Modelling of different kinds of nodes and limited edges. Useful when tracking relations between two inherently different sets of entities (e.g. students and courses, users and products).

# Sparse, Dense, and Random Graphs

Not all graphs differ qualitatively. They also differ quantitatively.

**Sparse graph.** A graph with relatively few edges compared to the maximum possible number.

**Dense graph.** A graph with many edges, often close to the maximum possible number.

**Random graph.** A graph generated by a probabilistic rule, e.g. each possible edge appears independently with probability  $p$ .

- many real social and information networks are **sparse**
- dense graphs are often easier to reason about theoretically, but rarer at scale
- random graphs are important as a **baseline model** for comparison

# Graph Representations

How would you store (write down) a graph on paper?





# Graph Representations

How would you store (write down) a graph on paper?

## Edge list

(u1, v1)  
(u1, v2)  
(u2, v2)  
(u3, v1)  
(u3, v2)

Compact and easy to exchange,  
but neighbor lookup is slower.

## Adjacency list

u1: v1, v2  
u2: v2  
u3: v1, v2  
v1: u1, u3  
v2: u1, u2, u3

Best default for traversal:  
neighbors are stored directly.

## Adjacency matrix

|    | v1 | v2 | v3 |
|----|----|----|----|
| v1 | 0  | 1  | 1  |
| v2 | 0  | 0  | 1  |
| v3 | 1  | 1  | 1  |

Rows and columns are node labels.  
Fast lookup, but space-heavy.

*Common representations are edge lists, adjacency lists, and adjacency matrices. They encode the same graph, but each supports different algorithms and scales differently.*

# Why Representations Matter


Different representations make different operations easier.

- edge lists are compact, easy to exchange, and natural for tabular data
- adjacency lists are good for traversal because neighbors are stored directly
- adjacency matrices are good for algebraic methods and constant-time edge lookup

| Representation   | Strength                                  | Limitation                                   |
|------------------|-------------------------------------------|----------------------------------------------|
| Edge list        | Lean, simple, good for import/export      | Slow to test whether a specific edge exists  |
| Adjacency list   | Best default for BFS/DFS on sparse graphs | Less convenient for matrix-based methods     |
| Adjacency matrix | Fast edge lookup, linear algebra friendly | Uses $O( V ^2)$ space even for sparse graphs |

# Real-World Example: ARPANET

The ARPANET already shows how quickly a technical system becomes a graph with non-trivial structure (Easley & Kleinberg, 2010).



## Adjacency List (1970):

```
SRI:      [UCSB, UCLA, Stanford, UTAH]
UCSB:     [SRI, UCLA]
UCLA:     [SRI, UCSB, RAND, Stanford]
RAND:     [UCLA, BBN, SDC]
SDC:      [RAND, UTAH]
Stanford: [SRI, UCLA]
Utah:     [SRI, SDC, MIT]
Case:     [Carnegie, Lincoln]
Carnegie: [Case, Harvard]
Harvard:  [Carnegie, BBN]
BBN:      [RAND, MIT, Harvard]
MIT:      [UTAH, BBN, Lincoln]
Lincoln:  [Case, MIT]
```

***ARPANET topology (1969–1977).*** Nodes are universities and research labs, edges are data links. One of the earliest computer networks.

# Neighbourhood and Degree

**Neighbourhood.** The neighbourhood  $N(v)$  of a node  $v$  is the set of all nodes connected to  $v$  by an edge.

**Degree.** The degree  $\deg(v)$  is the number of edges incident to  $v$ . In simple graphs,  $\deg(v) = |N(v)|$ .



# Beyond Simple Graphs

The basic definition  $G = (V, E)$  is only the start. In practice, we often need richer graph models.

| Graph type     | What is added?                      | Typical use                                |
|----------------|-------------------------------------|--------------------------------------------|
| Labelled graph | names or labels on nodes/edges      | named entities, relation types             |
| Typed graph    | explicit node and edge classes      | users, posts, organizations                |
| Property graph | attributes on nodes and edges       | metadata, timestamps, weights              |
| Multigraph     | multiple edges between same nodes   | repeated interactions, transport links     |
| Temporal graph | time attached to edges or snapshots | evolving communication or contact networks |

**Key point:** the graph formalism is flexible. The "right" graph depends on what information must be preserved for the question.

# Teaser: Properties We Will Use Later

Once we have a graph, we can ask structural questions beyond "is there an edge?"

- **degree distribution:** how are node degrees distributed across the whole graph?
- **centrality:** which nodes are structurally important or influential?
- **components:** where is the graph connected or fragmented?
- **path lengths:** how many steps separate typical pairs of nodes?

These are only teaser concepts for now. We will return to them in later lectures with formal definitions and applications.

# Takeaway

Graph theory is useful because it gives a precise language for the properties we identified, but every graph still reflects a modeling choice that must be interpreted.

# Quick Check

Consider a simple 5-node undirected graph: A-B, B-C, C-D, D-A, A-C. Here is an edge list and a function to convert it to an adjacency dictionary, but two key lines are missing. **What goes in the blanks?**

```
edges = [("A","B"), ("B","C"), ("C","D"), ("D","A"), ("A","C")]
```

```
def edges_to_adj(edges):  
    adj = {}  
    for u, v in edges:  
        if u not in adj: adj[u] = []  
        if v not in adj: adj[v] = []  
  
        adj[u].append(____)  
        adj[____].append(____)  
    return adj
```

# Quick Check — Solution

```
edges = [("A","B"), ("B","C"), ("C","D"), ("D","A"), ("A","C")]

def edges_to_adj(edges):
    adj = {}
    for u, v in edges:
        if u not in adj: adj[u] = []
        if v not in adj: adj[v] = []

        adj[u].append(v)
        # Undirected graph implies symmetric edges
        adj[v].append(u)
    return adj
```

**Key insight:** Undirected graphs require symmetric entries. If a graph is directed, the second append line is omitted. Why should we not omit the second append line in undirected cases?

# Paths and Distance

---

**Guiding Question:** What does it mean for one node to "reach" another?

# Paths

A path is a sequence of vertices where each consecutive pair is connected by an edge.






***A path highlighted in a graph.** The colored edges trace a route from a start node to an end node. The path length equals the number of edges traversed.*

# Walk, Path, Cycle

**Walk.** A walk of length  $k$  is a sequence  $v_0, v_1, \dots, v_k$  where  $(v_{i-1}, v_i) \in E$  for all  $i$ . Vertices may be repeated.

**Path.** A path is a walk where all vertices are distinct:  $v_i \neq v_j$  for  $i \neq j$ . Also called *simple path*.

**Cycle.** A cycle is a walk  $v_0, v_1, \dots, v_k$  with  $v_0 = v_k$  and all intermediate vertices distinct.



**Directed vs. undirected.** In an undirected graph, an edge can be traversed in either direction. In a directed graph, each step must follow the edge direction.

# Path Length, Shortest Path, Diameter

**Path length.** The length of a path  $v_0, \dots, v_k$  is  $k$  (number of edges). In a weighted graph:  $\sum_{i=1}^k w(v_{i-1}, v_i)$ .

**Shortest path.**  $\text{dist}(u, v) = \min\{k \mid \text{path of length } k \text{ from } u \text{ to } v\}$ . If no path exists,  $\text{dist}(u, v) = \infty$ .

**Diameter.**  $\text{diam}(G) = \max_{u,v \in V} \text{dist}(u, v)$  — the longest shortest path in the graph.

## Path Length

Number of edges in the path




Path  $A \rightarrow B \rightarrow C \rightarrow D$  has length 3  
(3 edges traversed)

Weighted: length = sum of edge weights

## Diameter

Longest shortest path in the graph



$\text{diam}(G) = \text{dist}(A, D) = 3$

The farthest pair determines the diameter


In directed graphs,  $\text{dist}(u, v) \neq \text{dist}(v, u)$  in general

# Why Paths Matter

- they formalize reachability
- they give us a notion of distance
- they underpin routing, diffusion, and traversal algorithms

# A Classical Motivation

The Königsberg bridges problem is one of the classic entry points into graph theory.



***The Seven Bridges of Königsberg (1736).*** Euler asked whether one could walk through the city crossing each bridge exactly once. He proved it impossible — and in doing so, founded graph theory.

# Eulerian Path and Circuit

**Eulerian Path.** A walk traversing every edge exactly once. Exists iff exactly 0 or 2 vertices have an odd degree.

**Eulerian Circuit.** An Eulerian path that starts and ends on the same vertex. Exists iff *every* vertex has an even degree.

## Eulerian Path

Exactly two odd-degree vertices



Odd degrees: 2 endpoints

## Eulerian Circuit

All vertices have even degree



Closed walk possible

# Eulerian Path and Circuit


**Eulerian Path.** A walk traversing every edge exactly once. Exists iff exactly 0 or 2 vertices have an odd degree.

**Intuitive Proof:** Think of rooms with doors (=edges). If a room has an even number of doors, you can enter and leave it an equal number of times. If it has an odd number of doors, you must start or end in that room.

**Eulerian Circuit.** An Eulerian path that starts and ends on the same vertex. Exists iff *every* vertex has an even degree.

## Eulerian Path

Exactly two odd-degree vertices



Odd degrees: 2 endpoints

## Eulerian Circuit

All vertices have even degree



Closed walk possible

# Eulerian Path and Circuit

**Eulerian Path.** A walk traversing every edge exactly once. Exists iff exactly 0 or 2 vertices have an odd degree.

**Intuitive Proof:** Think of rooms with doors (=edges). If a room has an even number of doors, you can enter and leave it an equal number of times. If it has an odd number of doors, you must start or end in that room.

**Eulerian Circuit.** An Eulerian path that starts and ends on the same vertex. Exists iff *every* vertex has an even degree.

**Intuitive Proof:** If there are no rooms with an odd number of doors, you can only start and end at the same room.

## Eulerian Path

Exactly two odd-degree vertices



Odd degrees: 2 endpoints

## Eulerian Circuit

All vertices have even degree



Closed walk possible

# The Königsberg Solution



Connecting back to Königsberg: all four land masses had odd degrees (3, 3, 3, 5), violating both conditions.

# Takeaway

Paths turn connectivity into something analyzable: reachability, distance, redundancy, and traversal constraints.

# Quick Check

Consider a 5-node graph: `adj = {1: [2, 3], 2: [1, 4], 3: [1, 4], 4: [2, 3, 5], 5: [4]}`

Here is a recursive function to find *any* path between two nodes, but the crucial recursive step is missing. **What goes in the blank?**

```
def find_path(graph, start, end, path=[]):
    path = path + [start]
    if start == end: return path
    for node in graph.get(start, []):
        if node not in path:
            # How do we continue the search?
            result = _____
            if result is not None: return result
    return None
```

# Quick Check — Solution

The algorithm is called depth first search (DFS) - we traverse every edge as soon as we encounter it.

```
def find_path(graph, start, end, path=[]):  
    path = path + [start]  
    if start == end: return path  
    for node in graph.get(start, []):  
        if node not in path:  
            # Recursive call on the neighbor  
            result = find_path(graph, node, end, path)  
            if result is not None: return result  
    return None
```

**Key insight:** DFS explores deep first and finds *any* path. How would we change the approach to guarantee finding the *shortest* path?

# Graph Search Algorithms

---

**Guiding Question:** How can we systematically find a node — or the shortest route to it — in a large graph?

# Breadth-First Search

Breadth-first search explores a graph layer by layer, guaranteeing shortest paths in unweighted graphs.

## Algorithm: BFS

*(Input: Graph  $G = (V, E)$ , source node  $s$ )*

1. Initialize a **queue**  $Q$  and add  $s$
2. Mark  $s$  as **visited**
3. While  $Q$  is not empty:
  - Extract current node  $u$  from the front of  $Q$
  - For each neighbor  $v$  of  $u$ : if  $v$  is **not visited**, mark  $v$  as visited and add  $v$  to the back of  $Q$

## Key Properties:

- Uses a **FIFO** (First-In, First-Out) queue.
- Discovers nodes strictly layer by layer.
- Runs in  $O(|V| + |E|)$  time.
- Useful for shortest paths, components, and reachability.


# Visual Intuition



**BFS layer-by-layer expansion.** Starting from "You" (layer 0), BFS discovers all direct neighbors (layer 1), then all not-yet-visited neighbors of those (layer 2), and so on.



# BFS on the ARPANET



BFS applied to the ARPANET topology starting from UCLA. Node color encodes BFS layer (distance).

**Discussion:** The BFS visualization shows some nodes with multiple incoming edges from the previous layer. Is this possible in a strict BFS spanning tree?

# BFS on the ARPANET



BFS applied to the ARPANET topology starting from UCLA. Node color encodes BFS layer (distance).

**Discussion:** The BFS visualization shows some nodes with multiple incoming edges from the previous layer. Is this possible in a strict BFS spanning tree?

**Answer:** No. A true BFS tree has exactly one parent per node. Drawing *all* shortest-path edges creates a "level graph", not a spanning tree!

# Complexity

- Queue-based BFS runs in  $O(|V| + |E|)$  — each vertex and each edge is visited at most once
- This makes BFS the standard exact method for shortest paths in **unweighted** graphs

**Directed graphs:** BFS works identically — it follows edges in their given direction.

**Weighted graphs:** BFS does **not** give shortest paths when edges have different weights. For weighted shortest paths, use Dijkstra's algorithm.

# Depth-First Search

Depth-first search traverses as deep as possible along each branch before backtracking.

## Algorithm: DFS

(Input: Graph  $G = (V, E)$ , source node  $s$ )

1. Initialize a **stack**  $S$  and add  $s$
2. While  $S$  is not empty:
  - Pop current node  $u$  from the top of  $S$
  - If  $u$  is **not visited**:
    - Mark  $u$  as **visited**
    - Push all neighbors of  $u$  onto the stack  $S$

## Key Properties:

- Uses a **LIFO** (Last-In, First-Out) stack.
- Explores deep before wide.
- Runs in  $O(|V| + |E|)$  time.
- Standard for finding cycles and topological sorting.

# DFS Visual Intuition





**DFS deep expansion.** Starting from "You", DFS picks one neighbor and immediately proceeds to that neighbor's neighbors. It only returns to explore other branches (backtracking) when it hits a "dead end" where all neighbors of the current node are already visited.

# DFS on the ARPANET

## DFS on ARPANET — starting from UCLA

Discovery order assumes neighbors are processed in adjacency-list order.



DFS applied to the ARPANET. Node numbers indicate the discovery order if neighbors are processed in the listed adjacency-list order. Notice how the search can wander far across the network before exploring nodes physically close to the start.

**Discussion:** Why is DFS usually a poor choice for finding the shortest path between two nodes (e.g., California to Massachusetts)?

# DFS on the ARPANET

## DFS on ARPANET — starting from UCLA

Discovery order assumes neighbors are processed in adjacency-list order.



DFS applied to the ARPANET. Node numbers indicate the discovery order if neighbors are processed in the listed adjacency-list order. Notice how the search can wander far across the network before exploring nodes physically close to the start.

**Discussion:** Why is DFS usually a poor choice for finding the shortest path between two nodes (e.g., California to Massachusetts)?

# BFS vs. DFS Transition

| Property       | Breadth-First Search (BFS) | Depth-First Search (DFS) |
|----------------|----------------------------|--------------------------|
| Data Structure | Queue (FIFO)               | Stack (LIFO)             |
| Expansion      | Layer by layer (Ripple)    | Deep and narrow (Snake)  |
| Shortest Path  | Guaranteed (unweighted)    | No guarantee             |
| Complexity     | $O( V  +  E )$             | $O( V  +  E )$           |

# Dijkstra's Algorithm

BFS explores layer by layer. For **weighted graphs with non-negative additive edge costs**, we use Dijkstra's algorithm.

**Algorithm: Dijkstra** (Input: Graph  $G = (V, E)$ , source  $s$ , additive weights  $w \geq 0$ )

1. Set distance to  $s$  as 0 and all others as  $\infty$ .
2. Initialize a **priority queue**  $PQ$  with all nodes.
3. While  $PQ$  is not empty:
  - Extract node  $u$  with the minimum distance.
  - For each neighbor  $v$  of  $u$ :
    - If path through  $u$  is shorter: update  $v$ 's distance.

Dijkstra's Algorithm (Step C)



**Priority Queue**  
(Extracting C)

C : d=6

B : d=7

D : d= $\infty$

**BFS vs. Dijkstra.** BFS explores identical layers. Dijkstra uses a priority queue to always expand the currently cheapest unvisited node. Complexity:  $O((|V| + |E|) \log |V|)$ .

**Priority Queue (Min-Heap).** A data structure that allows inserting elements and efficiently extracting the element with the *minimum* value. In Dijkstra, it acts as a smart queue: instead of First-In-First-Out, the node with the shortest accumulated distance is always extracted next.

**Important restriction.** Dijkstra assumes path costs are **additive** and every edge weight is **non-negative**. With negative weights, a path that looked final can later become cheaper, so the greedy rule breaks.

# Quick Check

Consider our 5-node graph: `adj = {1: [2, 3], 2: [1, 4], 3: [1, 4], 4: [2, 3, 5], 5: [4]}`

Here is a BFS function to calculate the shortest path *distance* from start to all reachable nodes.  
**What goes in the two blanks?**

```
from collections import deque

def bfs(graph, start):
    distances = {start: 0}
    queue = deque([start])
    while queue:
        current = queue.popleft()
        for neighbor in graph.get(current, []):
            if neighbor not in distances:
                distances[neighbor] = _____
                queue._____(neighbor)
    return distances
```

# Quick Check — Solution

```
from collections import deque

def bfs(graph, start):
    distances = {start: 0}
    queue = deque([start])
    while queue:
        current = queue.popleft()
        for neighbor in graph.get(current, []):
            if neighbor not in distances:
                distances[neighbor] = distances[current] + 1
                queue.append(neighbor)
    return distances
```

**Key insight:** Because the queue is FIFO, we always finish layer  $k$  before starting layer  $k + 1$ . The first time we see a node, we have found its shortest path.

# Summary: Graph Search

| Algorithm       | Structure    | Strategy       | Guarantees                                            |
|-----------------|--------------|----------------|-------------------------------------------------------|
| <b>BFS</b>      | Queue (FIFO) | Layer-by-layer | Shortest path (unweighted)                            |
| <b>DFS</b>      | Stack (LIFO) | Deepest-first  | Reachability, Cycles                                  |
| <b>Dijkstra</b> | Priority Q   | Cheapest-first | Shortest path (weighted, non-negative additive costs) |

**Key Insight:** BFS matters because a simple local rule produces exact shortest-path information in unweighted graphs. Dijkstra extends this logic to weighted networks using a priority queue.

# Connectivity

---

**Guiding Question:** When is a network still one system — and when has it split apart?

# Connectedness

**Connected (undirected).** Let  $G = (V, E)$  be a graph.  $G$  is *connected* if for every pair  $u, v \in V$  there exists a path from  $u$  to  $v$ .


## Connected Graph

A path exists between every pair of nodes



## Disconnected Graph

No path exists between the two clusters



# Connected Components

Connected components are maximal connected subgraphs.

**Connected component.** A *connected component* of  $G$  is a subgraph  $H \subseteq G$  such that:

1.  $H$  is **connected** (every pair of vertices in  $H$  has a path within  $H$ ), and
2.  $H$  is **maximal** — no vertex outside  $H$  has a path to any vertex in  $H$ .



Each colored cluster is a maximal connected subgraph — meaning it cannot be extended without losing separation.

Typical examples:

- separate friendship groups if no acquaintance links the groups
- islands in a road network if no bridge or ferry connects them
- disconnected subnetworks in technical infrastructure after a link failure

# Connectedness in Directed Graphs

How does connectivity change when edges have direction?

# Connectedness in Directed Graphs

How does connectivity change when edges have direction?

**Strongly connected.** For every pair  $u, v \in V$  there is a directed path from  $u$  to  $v$  **and** from  $v$  to  $u$ .

**Weakly connected.** The graph is connected when all directed edges are replaced with undirected edges.

**Weakly Connected**  
(directed, NOT strongly connected)



A can reach D, but D cannot reach A

**Strongly Connected**  
(directed, MUTUALLY reachable)



Every node can reach every other node

# Real-World Component Patterns

Many real networks contain one large giant component and several much smaller disconnected parts. Below is a collaboration network with one giant component plus several small peripheral groups.



Source: Easley & Kleinberg 2010

Giant components are quite common:

- **Co-authorship networks:** one giant component plus a few isolated teams
- **Infrastructure networks:** failures can split one system into several components
- **Web graph:** one weakly connected macro-

structure, but many  
different strongly  
connected regions

# Directed Components: The Web Bow Tie

The Web is the classic example that weak and strong connectivity are not the same.



Source: Easley & Kleinberg 2010, p. Ch. 13


Broder et al. (2000), *Graph structure in the Web*, [Computer Networks 33\(1-6\):309-320](#), showed that the Web decomposes into:

- **SCC:** pages mutually reachable by directed paths
- **IN:** pages that can reach the SCC but are not reachable from it
- **OUT:** pages reachable from the SCC but unable to link back
- **Tendrils / tubes:** side regions attached to IN or OUT

This is why directed networks need more than the undirected idea of "one big connected piece."

# Giant Components

**Giant component.** A connected component that is significantly larger than all others, usually containing a macroscopic fraction of all vertices. Appears quite often in real world networks.



Source: Easley & Kleinberg 2010

# Takeaway

Connected components and giant components tell us whether interaction, diffusion, or coordination can happen across the whole graph or only inside local regions.

# Quick Check

We can use BFS to find all connected components by repeatedly starting search from unvisited nodes. **What goes in the blanks?**

```
def find_components(graph):
    visited = set()
    components = []
    for node in graph:
        if node not in visited:
            # bfs returns {node: distance} for all reachable nodes
            reachable = bfs(graph, node)
            component = set(reachable.keys())

            # Update global visited set and store result
            visited |= _____
            components.append(_____)
    return components
```

## Discussion Points:

- How does BFS naturally identify a "maximal" component?
- For directed graphs, does this find SCCs or WCCs?

# Quick Check — Solution

**Strategy:** Start from an unvisited node, run BFS to find all reachable nodes. Repeat until all nodes visited.

```
def find_components(graph):
    visited = set()
    components = []
    for node in graph:
        if node not in visited:
            reachable = bfs(graph, node)
            component = set(reachable.keys())

            # Solution:
            visited |= component
            components.append(component)
    return components

# Test with disconnected graph
graph = {
    "A": ["B"], "B": ["A", "C"], "C": ["B"], # component 1
    "X": ["Y"], "Y": ["X", "Z"], "Z": ["Y"] # component 2
}
print(find_components(graph))
# [{'A', 'B', 'C'}, {'X', 'Y', 'Z'}]
```

# Wrap-Up

- **Networks have properties** that demand formal description
- **Graphs** provide a precise model: nodes, edges, weights, direction, time
- **Graph variants** can be labelled, typed, attributed, sparse, dense, or random
- **Paths** give us reachability and distance
- **BFS** is a central traversal algorithm with guaranteed shortest paths
- **Dijkstra's** extends BFS to weighted graphs with non-negative additive costs
- **Connectivity** helps us detect structure and giant components

## Reading

Chapter 2 of Easley & Kleinberg (2010) covers graph theory foundations, paths, BFS, and connectivity in detail.

# Frequently Asked Questions (FAQ)

# Why would we use DFS in practice?

- to test **reachability**: is there any path at all from a start node to a target, even if we do not care about the shortest one?
- to detect **cycles**: if DFS encounters a node that is already on the current recursion stack, we have found a loop
- to build **topological orderings** in DAG-like settings: finishing times from DFS give an ordering that respects dependencies
- to explore one branch deeply before trying alternatives

DFS is not the shortest-path tool. It is the structural exploration tool.

# Why is DFS useful for connected component detection?

Because one DFS from a start node visits **exactly the nodes reachable from that node**. In an undirected graph, that reachable set is one full connected component. Repeating DFS from each still-unvisited node therefore partitions the graph into components.

DFS is also a common practical choice because the implementation is simple: we only need a visited set plus recursion or an explicit stack. BFS works just as well for component detection, but it is not generally true that DFS always uses less memory; that depends on the graph shape.

# Why do negative weights break Dijkstra?

Dijkstra assumes that once the currently cheapest node is removed from the priority queue, its distance is final. That is only true when every extra edge can only **increase** total cost.

With a negative edge, a later detour can make an already-settled node cheaper after all:

$A \rightarrow B = 2$ ,  $A \rightarrow C = 5$ ,  $C \rightarrow B = -10$



Negative edges can lower a distance after a node was settled

At first, Dijkstra would settle B with cost 2. But the path  $A \rightarrow C \rightarrow B$  has total cost -5, so the earlier decision was wrong.

## Image Credits & References

Easley, David and Kleinberg, Jon (2010). Networks, Crowds, and Markets: Reasoning about a Highly Connected World. *Cambridge University Press*. <https://www.cs.cornell.edu/home/kleinber/networks-book/>