

Network Science

Lecture 07: Small-World Networks

Prof. Dr. Michael Granitzer

Dr. Kanishka Ghosh Dastidar

Dr. Jelena Mitrovic

Milgram's Experiment — The Question That Started It All

In 1967, the social psychologist Stanley Milgram mailed 296 letters to random people in Nebraska and Kansas. Each letter named a **target person** — a stockbroker in Boston — and asked the recipient to forward it to someone they knew *on a first-name basis* who might be "closer" to the target.

Milgram's Experiment — The Question That Started It All

In 1967, the social psychologist Stanley Milgram mailed 296 letters to random people in Nebraska and Kansas. Each letter named a **target person** — a stockbroker in Boston — and asked the recipient to forward it to someone they knew *on a first-name basis* who might be "closer" to the target.


The letters that arrived took a **median of about 6 steps**.

The phrase "six degrees of separation" entered popular culture — but the experiment raised more questions than it answered.

The Chain

Milgram's small-world experiment (1967)

Can a letter reach a stranger in Boston through a chain of personal acquaintances?



Questions this experiment raises

- ① Why do short paths exist at all in a network of 200 million people?
- ② How can a network be locally clustered AND have short paths to strangers?
- ③ If short paths exist, can ordinary people find them using only local knowledge?
- ④ Only ~25% of chains arrived. Why did most fail?


Each intermediary sees only their own contacts and must guess who is "closer" to the target. About 25% of chains completed; the rest died when someone refused to forward. The experiment demonstrates both that short paths *exist* and that finding them is *hard*.

Four Questions from One Experiment

1. **Existence.** Why do short paths exist at all in a network of 200 million people?
2. **Coexistence.** How can a network be *locally clustered* (your friends know each other) and *globally short* (six steps to a stranger in another state)?
3. **Navigability.** If short paths exist, can ordinary people find them using only local knowledge — no map, no search engine?
4. **Failure.** Only ~25% of chains arrived. Does this mean short paths are rare, or that *finding* them is the bottleneck?

Every concept in today's lecture answers one of these.

Milgram's small-world experiment (1967)
Can a letter reach a stranger in Boston through a chain of personal acquaintances?



A New Kind of Gap — Existence vs. Findability

L03–L04 had a **computational** gap (NP-hard ideals). L05 had a **causal** gap. L06 had a **structural** gap (completeness required).

L07 has a **navigational** gap: short paths may *exist* without anyone being able to *find* them using local information.

Lecture	Gap type	Ideal	Reality	Strategy
L03–L04	Computational	Exact optimization	NP-hard	Polynomial heuristics
L05	Causal	Mechanism identification	Snapshots insufficient	Longitudinal data
L06	Structural	Complete signed graph	Sparse, noisy data	Approximate balance
L07	Navigational	Short navigable paths	Local information only	Geometric link distribution

Takeaway

The small-world phenomenon is really two phenomena: the *existence* of short paths (explained by a few random long-range links) and the *findability* of those paths (explained only by a specific geometric distribution of links). Milgram's experiment revealed both — and the gap between them.

The Small-World Phenomenon

Guiding Question: Why are distances in large networks often surprisingly short?

Logarithmic Distances

Small-world property. A network exhibits the small-world property if the average shortest-path distance $\bar{d}$ grows at most logarithmically with the number of nodes:

$$\bar{d} \propto \log |V|$$

In a random graph $G(n, p)$ with average degree k , the typical distance is approximately $\bar{d} \approx \frac{\log n}{\log k}$.

Logarithmic Distances


Small-world property. A network exhibits the small-world property if the average shortest-path distance $\bar{d}$ grows at most logarithmically with the number of nodes:

$$\bar{d} \propto \log |V|$$

In a random graph $G(n, p)$ with average degree k , the typical distance is approximately $\bar{d} \approx \frac{\log n}{\log k}$.

For a social network of $n = 300$ million people with average degree $k \approx 100$: $\bar{d} \approx \frac{\log(3 \times 10^8)}{\log 100} \approx \frac{19.5}{4.6} \approx 4.2$. Milgram's 6 is in the right ballpark.

Empirical Examples



Source: Easley & Kleinberg 2010, p. Ch. 20




Figure 2.12: Ron Graham's hand-drawn picture of a part of the mathematics collaboration graph, centered on Paul Erdős [189]. (Image from <http://www.oakland.edu/enp/cgraph.jpg>)

Source: Easley & Kleinberg 2010

Left: Microsoft Instant Messenger network — 180 million users, median distance ~ 7 . **Right:** Mathematical collaboration network — Erdős numbers rarely exceed 5. Both confirm logarithmic scaling across orders of magnitude.

Takeaway

Logarithmic distances are a robust empirical fact across social, communication, and collaboration networks. In any network with moderate average degree, even billions of nodes remain only a handful of hops apart.

Quick Check

The Facebook social graph has approximately 3 billion users and an average degree of about 300.

1. Estimate the average shortest-path distance using $\bar{d} \approx \log n / \log k$.
2. Facebook reported an empirical average distance of 3.57 in 2016. Is this consistent with your estimate?
3. Why might the empirical value be *lower* than the logarithmic estimate?

Quick Check — Solution

$$\bar{d} \approx \frac{\log(3 \times 10^9)}{\log 300} \approx \frac{21.8}{5.7} \approx 3.8$$

Yes — 3.57 is remarkably close to the logarithmic estimate.

The $\log n / \log k$ formula assumes homogeneous degree (random graph). Facebook has heavy-tailed degree — celebrities and organizations with millions of friends act as hubs, creating shortcuts that reduce $\bar{d}$ below the random-graph prediction.

Quick Check — Solution

1. $\bar{d} \approx \frac{\log(3 \times 10^9)}{\log 300} \approx \frac{21.8}{5.7} \approx 3.8$

Yes — 3.57 is remarkably close to the logarithmic estimate.

The $\log n / \log k$ formula assumes homogeneous degree (random graph). Facebook has heavy-tailed degree — celebrities and organizations with millions of friends act as hubs, creating shortcuts that reduce $\bar{d}$ below the random-graph prediction.

Quick Check — Solution

1. $\bar{d} \approx \frac{\log(3 \times 10^9)}{\log 300} \approx \frac{21.8}{5.7} \approx 3.8$

2. Yes — 3.57 is remarkably close to the logarithmic estimate.

The $\log n / \log k$ formula assumes homogeneous degree (random graph). Facebook has heavy-tailed degree — celebrities and organizations with millions of friends act as hubs, creating shortcuts that reduce $\bar{d}$ below the random-graph prediction.

Quick Check — Solution

1. $\bar{d} \approx \frac{\log(3 \times 10^9)}{\log 300} \approx \frac{21.8}{5.7} \approx 3.8$

2. Yes — 3.57 is remarkably close to the logarithmic estimate.

3. The $\log n / \log k$ formula assumes homogeneous degree (random graph). Facebook has heavy-tailed degree — celebrities and organizations with millions of friends act as hubs, creating shortcuts that reduce $\bar{d}$ below the random-graph prediction.

The Watts–Strogatz Model

Guiding Question: How can a network be both highly clustered *and* have short paths?

The Paradox

Real social networks have two seemingly contradictory properties:

Property	Regular lattice	Random graph	Real social network
Clustering C	High (~ 0.5)	Low ($\sim k/n$)	High ($\sim 0.1-0.5$)
Avg. path length L	Long ($\sim n/2k$)	Short ($\sim \log n / \log k$)	Short ($\sim \log n$)

The Paradox

Real social networks have two seemingly contradictory properties:

Property	Regular lattice	Random graph	Real social network
Clustering C	High (~ 0.5)	Low ($\sim k/n$)	High ($\sim 0.1-0.5$)
Avg. path length L	Long ($\sim n/2k$)	Short ($\sim \log n / \log k$)	Short ($\sim \log n$)

A regular lattice is clustered but large. A random graph is small but unclustered. **Real networks are both** — high C and low L simultaneously.

The Model

Watts–Strogatz model (1998).

1. Start with a **ring lattice**: n nodes arranged in a circle, each connected to its k nearest neighbors.
2. For each edge, with probability p , **rewire** one endpoint to a uniformly random node (avoiding self-loops and duplicates).

The parameter $p \in [0, 1]$ controls the interpolation:

- $p = 0$: regular lattice (high C , high L)
- $p = 1$: random graph (low C , low L)
- $0 < p \ll 1$: **small-world regime** (high C , low L)

The Characteristic Curves




(a) *Pure exponential growth produces a small world*




(b) *Triadic closure reduces the growth rate*

Source: Easley & Kleinberg 2010, p. Ch. 20



Source: Easley & Kleinberg 2010, p. Ch. 20



Source: Easley & Kleinberg 2010, p. Ch. 20

$C(p)/C(0)$ stays near 1 while $L(p)/L(0)$ drops sharply. The small-world regime (shaded) spans roughly $0.001 \leq p \leq 0.1$ — a few percent of edges rewired is enough to make the world small.

Takeaway

The Watts-Strogatz model resolves the clustering/path-length paradox: a tiny fraction of random long-range shortcuts collapses distances without destroying local structure. The small-world regime exists because $L(p)$ falls much faster than $C(p)$.

Quick Check

A ring lattice has $n = 1000$ nodes and $k = 10$ neighbors per node.

1. Estimate the average path length L at $p = 0$.
2. After rewiring 1% of edges ($p = 0.01$), how many long-range shortcuts are created?
3. Qualitatively, what happens to L and C ?

Quick Check — Solution

$L(0) \approx n/(2k) = 1000/20 = 50$ steps.

2. $nk/2 = 5000$ edges total; 1% = .

3. L collapses — 50 random shortcuts in a 1000-node lattice bring L close to $\log n / \log k \approx 3$. C barely changes — only 50 of 5000 edges were rewired, so the vast majority of local triangles survive.

Quick Check — Solution

1. $L(0) \approx n/(2k) = 1000/20 = 50$ steps.
2. $nk/2 = 5000$ edges total; 1% = .
3. L collapses — 50 random shortcuts in a 1000-node lattice bring L close to $\log n / \log k \approx 3$. C barely changes — only 50 of 5000 edges were rewired, so the vast majority of local triangles survive.

Quick Check — Solution

1. $L(0) \approx n/(2k) = 1000/20 = 50$ steps.
2. $nk/2 = 5000$ edges total; 1% = **50 long-range shortcuts**.
3. L collapses — 50 random shortcuts in a 1000-node lattice bring L close to $\log n / \log k \approx 3$. C barely changes — only 50 of 5000 edges were rewired, so the vast majority of local triangles survive.

Quick Check — Solution

1. $L(0) \approx n/(2k) = 1000/20 = 50$ steps.
2. $nk/2 = 5000$ edges total; 1% = **50 long-range shortcuts**.
3. L collapses — 50 random shortcuts in a 1000-node lattice bring L close to $\log n / \log k \approx 3$. C barely changes — only 50 of 5000 edges were rewired, so the vast majority of local triangles survive. **This is the small-world regime.**

Kleinberg's Decentralized Search

Guiding Question: Short paths exist — but can local actors *find* them without a map?

The Problem

In Milgram's experiment, each person forwarded the letter based on **local information only**:

- they knew their own contacts
- they knew the target's name, city, and profession
- they did *not* know the global network structure

The Problem

In Milgram's experiment, each person forwarded the letter based on **local information only**:

- they knew their own contacts
- they knew the target's name, city, and profession
- they did *not* know the global network structure

Decentralized search: at each step, the current holder forwards to the neighbor *closest* to the target (in some distance metric). No global routing table, no BFS.

The Problem

In Milgram's experiment, each person forwarded the letter based on **local information only**:

- they knew their own contacts
- they knew the target's name, city, and profession
- they did *not* know the global network structure

Decentralized search: at each step, the current holder forwards to the neighbor *closest* to the target (in some distance metric). No global routing table, no BFS.

Watts-Strogatz does not guarantee this works. Short paths exist, but the greedy algorithm may not find them — it depends on *where* the long-range links point.

Kleinberg's Grid Model (2000)

Setup. Place n nodes on a d -dimensional grid. Each node has:

- **local contacts:** all grid neighbors within lattice distance p
- **one long-range link:** drawn with probability proportional to $d(v, u)^{-r}$, where d is grid distance and $r \geq 0$ is the distance exponent.

Greedy routing. At each step, forward to the neighbor (local or long-range) closest to the target in grid distance.

The parameter r controls the *geometry* of long-range links:


- $r = 0$: uniform random (long-range links go anywhere)
- $r = d$: inverse-power law matched to dimension (links span all scales equally)
- $r \gg d$: long-range links are effectively local

The Three Regimes

Kleinberg's grid model — the geometry of navigability

Each node has local grid neighbors + one long-range link drawn with $P(\text{link to } u) \propto d(v,u)^{-r}$

$r = 0$ (uniform random)




short paths exist,
but search is $\Omega(n^{2/3})$

$r = 2 = d$ (inverse-square)



short paths exist
AND search is $O(\log^2 n)$

$r \gg d$ (too clustered)



short paths exist,
but "shortcuts" are too local

Kleinberg's Theorem

Theorem (Kleinberg, 2000).

In the d -dimensional grid model with long-range exponent r :

- If $r = d$: greedy routing finds the target in $O(\log^2 n)$ steps.
- If $r \neq d$: **no** decentralized algorithm (using only local information) can find the target in fewer than $\Omega(n^c)$ steps for some constant $c > 0$.

Kleinberg's Theorem


Theorem (Kleinberg, 2000).

In the d -dimensional grid model with long-range exponent r :


- If $r = d$: greedy routing finds the target in $O(\log^2 n)$ steps.
- If $r \neq d$: **no** decentralized algorithm (using only local information) can find the target in fewer than $\Omega(n^c)$ steps for some constant $c > 0$.

The exponent $r = d$ is **unique** — it is the only value at which decentralized search is polylogarithmic. Any other value, and search degrades to polynomial.

Simulation Evidence



Source: Easley & Kleinberg 2010, p. Ch. 20



Source: Easley & Kleinberg 2010, p. Ch. 20

Left: delivery time vs. exponent r — a sharp minimum at $r = d = 2$. **Right:** the proof intuition — at $r = d$, each "ring" of nodes at distance 2^i from the target receives roughly equal link probability, creating a hierarchy of scales the algorithm descends one level at a time.

Takeaway

Short paths exist in almost any small-world network. But navigability — the ability of local actors to *find* those paths — requires a specific geometric distribution of long-range links ($r = d$). Watts-Strogatz gives existence; Kleinberg gives findability.

Quick Check

A social network is well-modeled by a 2D grid with random long-range links.

1. What exponent r makes greedy search efficient?
2. If the network's long-range links are uniformly random ($r = 0$), what happens to (a) the existence of short paths and (b) greedy routing performance?
3. Why is the Watts-Strogatz model insufficient to explain Milgram's result?

Quick Check — Solution

$r = 2 = d$ for a 2D grid.

- (a) Short paths — uniform random shortcuts reduce average path length.
- (b) Greedy routing — it takes $\Omega(n^{2/3})$ steps because the random shortcuts provide no gradient toward the target.

3. Watts-Strogatz uses **uniform** rewiring ($r = 0$). This gives short paths but not short paths. Milgram's participants actually found short paths using local information — that requires $r = d$, a much stronger structural condition.

Quick Check — Solution

1. $r = 2 = d$ for a 2D grid.

- (a) Short paths — uniform random shortcuts reduce average path length.
- (b) Greedy routing — it takes $\Omega(n^{2/3})$ steps because the random shortcuts provide no gradient toward the target.

3. Watts-Strogatz uses **uniform** rewiring ($r = 0$). This gives short paths but not short paths. Milgram's participants actually found short paths using local information — that requires $r = d$, a much stronger structural condition.

Quick Check — Solution

1. $r = 2 = d$ for a 2D grid.

2.

- (a) Short paths — uniform random shortcuts reduce average path length.
- (b) Greedy routing — it takes $\Omega(n^{2/3})$ steps because the random shortcuts provide no gradient toward the target.

3. Watts-Strogatz uses **uniform** rewiring ($r = 0$). This gives short paths but not short paths.

Milgram's participants actually found short paths using local information — that requires $r = d$, a much stronger structural condition.

Quick Check — Solution

1. $r = 2 = d$ for a 2D grid.

2.

- (a) Short paths **still exist** — uniform random shortcuts reduce average path length.
- (b) Greedy routing — it takes $\Omega(n^{2/3})$ steps because the random shortcuts provide no gradient toward the target.

3. Watts-Strogatz uses **uniform** rewiring ($r = 0$). This gives short paths but not short paths.

Milgram's participants actually found short paths using local information — that requires $r = d$, a much stronger structural condition.

Quick Check — Solution

1. $r = 2 = d$ for a 2D grid.

2.

- (a) Short paths **still exist** — uniform random shortcuts reduce average path length.
- (b) Greedy routing **fails** — it takes $\Omega(n^{2/3})$ steps because the random shortcuts provide no gradient toward the target.

3. Watts-Strogatz uses **uniform** rewiring ($r = 0$). This gives short paths but not short paths.
Milgram's participants actually found short paths using local information — that requires $r = d$, a much stronger structural condition.

Quick Check — Solution

1. $r = 2 = d$ for a 2D grid.

2.

- (a) Short paths **still exist** — uniform random shortcuts reduce average path length.
- (b) Greedy routing **fails** — it takes $\Omega(n^{2/3})$ steps because the random shortcuts provide no gradient toward the target.

3. Watts-Strogatz uses **uniform** rewiring ($r = 0$). This gives short paths but not *navigable* short paths. Milgram's participants actually found short paths using local information — that requires $r = d$, a much stronger structural condition.

Beyond the Grid — Where Else Does This Show Up?

Kleinberg's grid is a toy model. But the *principle* — mix **short-range** links (for local refinement) with **long-range** links at every scale (for jumping) — generalises far beyond social networks.

Beyond the Grid — Where Else Does This Show Up?

Kleinberg's grid is a toy model. But the *principle* — mix **short-range** links (for local refinement) with **long-range** links at every scale (for jumping) — generalises far beyond social networks.

The grid gave us a **known, low-dimensional** metric. What happens when the "distance" is something we have to *learn* from data — high-dimensional embeddings of images, documents, or molecules?

Beyond the Grid — Where Else Does This Show Up?

Kleinberg's grid is a toy model. But the *principle* — mix **short-range** links (for local refinement) with **long-range** links at every scale (for jumping) — generalises far beyond social networks.

The grid gave us a **known, low-dimensional** metric. What happens when the "distance" is something we have to *learn* from data — high-dimensional embeddings of images, documents, or molecules?

The same small-world recipe powers **modern vector search**: the index graphs used inside Pinecone, Weaviate, Milvus, FAISS, and every retrieval-augmented LLM pipeline are Kleinberg's grid in disguise.

From Kleinberg's Grid to HNSW




From Kleinberg's Grid to Modern Vector Search (HNSW)

Same principle: long-range + short-range links enable greedy routing

Kleinberg 2000 — Geometric Grid

known 2-D coordinates · $r = d = 2 \cdot O(\log^2 n)$




known metric · distance = $|u - v|$ in the grid

generalise
metric
→ high-dim
vectors

HNSW — Hierarchical Navigable Small World

learned vector space · layered graph · $O(\log n)$



top layer = long-range jumps

bottom layer = local refinement

HNSW in Three Ideas

- 1. Layers as scales.** Each node is assigned a maximum level $l \sim \text{Geom}(1 / \ln M)$. Layer k contains only nodes with level $\geq k$ — so higher layers are exponentially sparser.
- 2. Greedy per layer.** Search starts at the top-layer entry point, greedily moves to the neighbour closest to the query, then descends one layer and repeats — refining locally.
- 3. Long-range by construction.** Top-layer edges span the whole space (like Kleinberg's $r = d$ links); bottom-layer edges are local. The geometric distribution is *baked into the index*, not drawn from a random process.

HNSW in Three Ideas

- 1. Layers as scales.** Each node is assigned a maximum level $l \sim \text{Geom}(1 / \ln M)$. Layer k contains only nodes with level $\geq k$ — so higher layers are exponentially sparser.
- 2. Greedy per layer.** Search starts at the top-layer entry point, greedily moves to the neighbour closest to the query, then descends one layer and repeats — refining locally.
- 3. Long-range by construction.** Top-layer edges span the whole space (like Kleinberg's $r = d$ links); bottom-layer edges are local. The geometric distribution is *baked into the index*, not drawn from a random process.

Complexity. $O(\log n)$ query time in practice; billions of vectors routinely indexed in production.

Why This Matters for RAG and LLM Retrieval

Every retrieval-augmented generation (RAG) pipeline has the same structure:

1. **Embed** each document chunk into $\mathbb{R}^d$ (typically $d = 384-4096$).
2. **Index** the vectors in an HNSW (or IVF, or DiskANN) graph.
3. At query time, **greedy-search** from the entry point for the top- k nearest neighbours.

Why This Matters for RAG and LLM Retrieval

Every retrieval-augmented generation (RAG) pipeline has the same structure:

1. **Embed** each document chunk into $\mathbb{R}^d$ (typically $d = 384-4096$).
2. **Index** the vectors in an HNSW (or IVF, or DiskANN) graph.
3. At query time, **greedy-search** from the entry point for the top- k nearest neighbours.

A 2026-era LLM stack searching 100M+ embeddings in <10 ms is doing *Kleinberg-style decentralized search* — on a learned metric, in a graph someone built to guarantee the $r = d$ property.

Why This Matters for RAG and LLM Retrieval

Every retrieval-augmented generation (RAG) pipeline has the same structure:

1. **Embed** each document chunk into $\mathbb{R}^d$ (typically $d = 384\text{--}4096$).
2. **Index** the vectors in an HNSW (or IVF, or DiskANN) graph.
3. At query time, **greedy-search** from the entry point for the top- k nearest neighbours.

A 2026-era LLM stack searching 100M+ embeddings in <10 ms is doing *Kleinberg-style decentralized search* — on a learned metric, in a graph someone built to guarantee the $r = d$ property.

We will come back to **how** these embeddings are constructed in **Lecture 09** (node and document representations, spectral methods, DeepWalk, node2vec, GNNs). The *geometry* of the embedding space determines whether navigation works — which closes the loop to Kleinberg.

Empirical Anchor: The Global Email Experiment

Guiding Question: Does Milgram's result replicate at scale — and what does chain failure tell us?

The Study

Dodds, Muhamad & Watts (2003), *An Experimental Study of Search in Global Social Networks*, [Science 301, 827–829](#).

They replicated Milgram's experiment using **email** — a global medium with no geographic constraint.

Parameter	Value
Participants recruited	~60 000 from 166 countries
Target persons	18 in 13 countries
Chains started	24 163
Chains completed	384 (~1.6%)
Median chain length (completed)	4–7 steps

What They Found

- 1. Completed chains are short.** Median 4–7 steps, consistent with Milgram and with $\log n / \log k$ scaling.
- 2. Most chains die early.** The completion rate of ~1.6% is not because paths don't exist — it's because intermediaries *choose not to forward*. The bottleneck is **motivation**, not topology.

What They Found

- 1. Completed chains are short.** Median 4–7 steps, consistent with Milgram and with $\log n / \log k$ scaling.
- 2. Most chains die early.** The completion rate of ~1.6% is not because paths don't exist — it's because intermediaries *choose not to forward*. The bottleneck is **motivation**, not topology.
- 3. Successful chains use geography and profession as search dimensions.** Forwarders who brought the letter closer in *geographic* or *professional* distance were more effective — consistent with Kleinberg's model where search exploits multiple distance dimensions simultaneously.

What the Study Could Not See

- 1. Chain death $\neq$ path absence.** A chain that dies after 3 steps does not mean no short path existed — it means someone chose not to forward. The study measures *willingness to search*, not *path existence*. The true completion rate under full cooperation is unknown.
- 2. Selection bias in completions.** The 384 completed chains are a biased sample — they may systematically over-represent easy targets (well-known, geographically close, professionally distinctive). The median chain length of 4–7 steps may underestimate the true difficulty of reaching a random target.
- 3. Email $\neq$ acquaintance.** Milgram's experiment required *personal acquaintance*; the email study allowed forwarding to anyone whose email address you knew. The social network traversed by email may be larger and weaker-tied than the face-to-face network Milgram assumed.
- 4. No ground-truth network.** Without observing the full network, the study cannot measure whether the successful chains actually followed near-optimal paths or just happened to complete.


Takeaway

Dodds et al. (2003) confirmed Milgram at scale: short paths exist and people can sometimes find them. But the 1.6% completion rate reveals that *motivation*, not topology, is the primary bottleneck. The navigational gap is real — short paths exist but are rarely found in practice.

The Web as a Directed Small World


Guiding Question: What happens to navigability when links are directed?

From Social to Information Networks



Source: Easley & Kleinberg 2010


Source: Easley & Kleinberg 2010



Source: Easley & Kleinberg 2010

Hyperlinks are **directed** — page A can link to page B without B linking back. This means reachability is asymmetric: you can follow links *from A to B* but not necessarily in reverse.

The Bow-Tie Structure



Source: Easley & Kleinberg 2010, p. Ch. 13

Navigability in Directed Networks

The small-world property and navigability both change in directed networks:

- **Short paths exist** within the SCC — the strongly connected core is a small world (~16 clicks between random page pairs within SCC).
- **Short paths do not exist** from OUT back to SCC, or between tendrils.
- **Greedy search is harder** because you can only follow outgoing links — you cannot "pull" information from a page that doesn't link to you.

Navigability in Directed Networks

The small-world property and navigability both change in directed networks:

- **Short paths exist** within the SCC — the strongly connected core is a small world (~16 clicks between random page pairs within SCC).
- **Short paths do not exist** from OUT back to SCC, or between tendrils.
- **Greedy search is harder** because you can only follow outgoing links — you cannot "pull" information from a page that doesn't link to you.

Search engines (Google, 1998) solve the navigability problem by building a *global index* — they observe the entire graph and compute optimal paths. This is the opposite of Kleinberg's decentralized search: it works precisely because it is *centralized*.

Takeaway

The Web's directed links create asymmetric reachability. Only the Strongly Connected Component (~28%) is a "small world." Search engines solve the navigability problem through centralized indexing — the opposite of Milgram-style decentralized routing.

Quick Check

A directed network has 1000 nodes. The SCC contains 300 nodes. The IN component has 200 nodes, and the OUT component has 250 nodes. The rest are tendrils.

1. From a random node in IN, can you reach a random node in OUT? Explain your path.
2. From a random node in OUT, can you reach a node in IN?
3. What fraction of all possible source-target pairs have a directed path between them?

Quick Check — Solution

IN $\rightarrow$ SCC $\rightarrow$ OUT. By definition, IN nodes can reach SCC, and SCC nodes can reach OUT. The path goes through the SCC as a relay.

OUT nodes cannot reach back to the SCC (that is what defines them as OUT, not SCC). So OUT $\rightarrow$ IN has no directed path.

3. Reachable pairs include: SCC $\rightarrow$ SCC (300×299), IN $\rightarrow$ SCC (200×300), SCC $\rightarrow$ OUT (300×250), IN $\rightarrow$ OUT via SCC (200×250). That totals $\sim 240,000$ pairs out of $1000 \times 999 \approx 10^6$ — roughly 24%. Most pairs are mutually reachable.

Quick Check — Solution

1. Yes: $IN \rightarrow SCC \rightarrow OUT$. By definition, IN nodes can reach SCC, and SCC nodes can reach OUT. The path goes through the SCC as a relay.

OUT nodes cannot reach back to the SCC (that is what defines them as OUT, not SCC). So $OUT \rightarrow IN$ has no directed path.

3. Reachable pairs include: $SCC \rightarrow SCC$ (300×299), $IN \rightarrow SCC$ (200×300), $SCC \rightarrow OUT$ (300×250), $IN \rightarrow OUT$ via SCC (200×250). That totals $\sim 240,000$ pairs out of $1000 \times 999 \approx 10^6$ — roughly 24%. Most pairs are mutually reachable.

Quick Check — Solution

- 1. Yes:** $IN \rightarrow SCC \rightarrow OUT$. By definition, IN nodes can reach SCC, and SCC nodes can reach OUT. The path goes through the SCC as a relay.
- 2. No.** OUT nodes cannot reach back to the SCC (that is what defines them as OUT, not SCC). So $OUT \rightarrow IN$ has no directed path.
- 3. Reachable pairs include:** $SCC \rightarrow SCC$ (300×299), $IN \rightarrow SCC$ (200×300), $SCC \rightarrow OUT$ (300×250), $IN \rightarrow OUT$ via SCC (200×250). That totals $\sim 240,000$ pairs out of $1000 \times 999 \approx 10^6$ — roughly 24%. Most pairs are mutually reachable.

Quick Check — Solution

- 1. Yes:** $IN \rightarrow SCC \rightarrow OUT$. By definition, IN nodes can reach SCC, and SCC nodes can reach OUT. The path goes through the SCC as a relay.
- 2. No.** OUT nodes cannot reach back to the SCC (that is what defines them as OUT, not SCC). So $OUT \rightarrow IN$ has no directed path.
- 3. Reachable pairs include:** $SCC \rightarrow SCC$ (300×299), $IN \rightarrow SCC$ (200×300), $SCC \rightarrow OUT$ (300×250), $IN \rightarrow OUT$ via SCC (200×250). That totals $\sim 240,000$ pairs out of $1000 \times 999 \approx 10^6$ — roughly 24%. Most pairs are *not* mutually reachable.

Wrap-Up — The Five Gaps

Lecture	Gap type	Ideal	Reality	Strategy
L03– L04	Computational	Exact optimization	NP-hard	Polynomial heuristics
L05	Causal	Mechanism identification	Snapshots insufficient	Longitudinal data
L06	Structural	Complete signed graph	Sparse, noisy data	Approximate balance
L07	Navigational	Short navigable paths	Local info only; directed links	Geometric link distribution; centralized search

The small-world phenomenon says short paths *exist*. Watts-Strogatz explains *why*. Kleinberg explains *when* they are *findable*. Milgram and Dodds et al. tested *whether* real people find them — and the answer is: sometimes, with effort, and only when the network's geometry cooperates.

Looking Ahead: Lecture 08



How do things *spread* on networks — diseases, information, behaviors? The structure we have built (ties, communities, balance, small-world shortcuts) now becomes the substrate for dynamic processes: contagion, diffusion, and cascading failure.

Reading

Chapter 20 of Easley & Kleinberg (2010) covers the small-world phenomenon and Watts-Strogatz. Kleinberg (2000), *The Small-World Phenomenon: An Algorithmic Perspective*, is the source for the decentralized search theorem. Dodds, Muhamad & Watts (2003) is the email experiment. Broder et al. (2000) is the bow-tie structure of the Web.

Image Credits & References

Easley, David and Kleinberg, Jon (2010). Networks, Crowds, and Markets: Reasoning about a Highly Connected World. *Cambridge University Press*. <https://www.cs.cornell.edu/home/kleinber/networks-book/>