

# Probabilistic Machine Learning

---

## Lecture 02: Probability and Statistics

Prof. Dr. Michael Granitzer

## Speaker notes

Two sessions. Session 1 covers Modules 1–2 (the language and conditioning), session 2 covers Modules 3–4 (quantities, then data, statistics and a parameter). The module slides mark the boundary.


2026-09-24: Module 4 rebuilt around data. Attribute types come from the Data Science data lesson; a new unit reads a table as a sample from a joint; Bernoulli/binomial moved here from Module 3 as the first model of data; the quantiles/Anscombe unit is dropped; a new Beta–Bernoulli unit conditions a parameter on data and closes with the frequentist/Bayesian comparison.

Reuses the data-science probability units. Trimmed for a bachelor audience: the counting unit and the axiomatic-foundations examples are dropped, since this course needs the calculus of probability rather than its construction. Combinatorics can be reinstated by adding `units/math/stochastic/counting.md` to Module 1, and the extra Laplace worked examples by adding `probability-measure-definition-examples.md` after it.



# Where We Left Off

Lecture 01 ended on an ideal — model the joint, condition on the data, read off answers — and on the shortcuts most methods take. Today builds the language all of them are written in:



It also left a claim unpaid: that **epistemic** uncertainty is the kind more data removes. Module 4 turns that slogan into a rate.


## Speaker notes

Pick up the thread deliberately (the roadmap slide of L01 instructs the lecturer to return to it at the start of every lecture). Say which move each module equips, and say that the modules are met in the order the *mathematics* needs, not the order the three moves are performed in — conditioning is defined for events long before there is a joint distribution to condition.

Module 4 runs all three moves once, for the simplest possible model: yes/no data, one parameter, a Beta prior. That is the whole ideal of Lecture 01 on a single number, small enough to draw. Lecture 03 then takes the first shortcut (keep only the peak), and Lecture 07 returns to the full treatment for many parameters.



# Argument Map



The lecture on one page. Top row: the language and the one operation that matters (conditioning), with the trap it resolves. Middle row: from outcomes to quantities to the joint — the object Lecture 01's first move refers to. Bottom row: a table is a sample from a joint, statistics computed from it are estimates with an error (SE), and conditioning a parameter on the data gives a posterior — the three moves of Lecture 01, carried out once for a single number. Point at the strip: this lecture builds the language every later shortcut is written in.

# Module 1

---

## A Language for Uncertainty

**Guiding Question:** *We need a language that can say "how strongly", not just "true or false". What must such a language satisfy — and who decides the numbers?*

# Example: Two Fruit Boxes

## Red box

2 apples, 6 oranges

## Blue box

3 apples, 1 orange

We pick a box, then a fruit from it.

- If we pick the **red box** **40%** of the time, how often do we get an apple?
- We got an **orange** — which box did it probably come from?

Deliberate overview slide (P4): the two questions are the whole chapter in miniature. Do not formalise anything yet — the point is that both questions are answerable by counting, and that the second one feels harder than the first.

# The Counts, in One Grid

A hundred picks with the red box chosen 40% of the time, as a table of counts — box across, fruit down:

|                  | $B = r$ | $B = b$ | row total |
|------------------|---------|---------|-----------|
| $F = a$ (apple)  | 10      | 45      | 55        |
| $F = o$ (orange) | 30      | 15      | 45        |
| column total     | 40      | 60      | 100       |

## Joint

a single cell over the total

$$P(B = r, F = a) = \frac{10}{100}$$

## Marginal

a row or column total over the total

$$P(F = a) = \frac{55}{100}$$

## Conditional

a cell over its own row or column

$$P(F = a \mid B = r) = \frac{10}{40}$$

The counts are what the boxes' contents would produce on average: of 40 red picks a quarter are apples, of 60 blue picks three quarters. Everything on this slide is division of counts. Have students compute one of each kind before naming them. The names will be redefined formally later; here they are labels for three ways of dividing the same table.

# The Two Rules

## Sum rule

$$P(X) = \sum_Y P(X, Y)$$

Add along a row or column to remove a variable.

## Product rule

$$P(X, Y) = P(Y | X)P(X)$$

A cell is its column total times its share of that column.

Almost all of probability theory is these two rules applied repeatedly.

This is the Bishop framing: two rules, everything else derived. Students who carry only these two out of the chapter are in good shape — the formal apparatus that follows exists to justify them and to extend them beyond finite tables.

# Bayes' Theorem, in Advance

Combining the two rules and the symmetry  $P(X, Y) = P(Y, X)$ :

$$P(Y \mid X) = \frac{P(X \mid Y)P(Y)}{P(X)}, \quad P(X) = \sum_Y P(X \mid Y)P(Y)$$

**Which box did the orange come from?**

## Speaker notes

Bishop's numbers: the answer moves from 0.4 to  $2/3$ , and the table on the previous slide confirms it (30 of 45 oranges came from red). Then the follow-up question: with 3 apples and 9 oranges in blue, both boxes are 75% oranges and the posterior is the prior, 0.4 — the fruit says nothing about the box. That is the deeper lesson: evidence updates belief only to the extent that the hypotheses predict different things.



# Bayes' Theorem, in Advance

Combining the two rules and the symmetry  $P(X, Y) = P(Y, X)$ :

$$P(Y \mid X) = \frac{P(X \mid Y)P(Y)}{P(X)}, \quad P(X) = \sum_Y P(X \mid Y)P(Y)$$

Which box did the orange come from?

$$P(B = r \mid F = o) = \frac{P(F = o \mid B = r)P(B = r)}{P(F = o)} = \frac{\frac{6}{8} \cdot 0.4}{\frac{6}{8} \cdot 0.4 + \frac{1}{4} \cdot 0.6} = \frac{2}{3}$$

Before the pick: 40% red. After seeing an orange: 67% red. The orange is **evidence** — because the red box is full of oranges and the blue one is not.

# Bayes' Theorem, in Advance

Combining the two rules and the symmetry  $P(X, Y) = P(Y, X)$ :

$$P(Y \mid X) = \frac{P(X \mid Y)P(Y)}{P(X)}, \quad P(X) = \sum_Y P(X \mid Y)P(Y)$$

Which box did the orange come from?

$$P(B = r \mid F = o) = \frac{P(F = o \mid B = r)P(B = r)}{P(F = o)} = \frac{\frac{6}{8} \cdot 0.4}{\frac{6}{8} \cdot 0.4 + \frac{1}{4} \cdot 0.6} = \frac{2}{3}$$

Before the pick: 40% red. After seeing an orange: 67% red. The orange is **evidence** — because the red box is full of oranges and the blue one is not.

And if the blue box held 3 apples and 9 oranges — the same orange share as the red one?

# Takeaway

Two rules — **sum** to remove a variable, **product** to split a joint — generate joint, marginal, conditional probabilities and Bayes' theorem. What follows makes them precise and extends them beyond counting.

Signpost explicitly: everything just used informally gets a definition in the coming modules. This unit is the map, not the territory.



# What Makes an Experiment "Random"?

**Definition (Random experiment).** A procedure that can — at least in theory — be repeated infinitely often, characterised by:

1. **Configuration** — a precisely specified system that can be reconstructed.
2. **Procedure** — an instruction for executing the experiment.
3. **Unpredictability** of the outcome.

**Coin flip:** a 1-Euro coin (configuration), placed on the thumb and snapped into the air (procedure), record heads or tails (outcome).

Speaker notes

Note what the definition does *not* say: nothing about the world being indeterministic. Ask students for a "natural random experiment" — rainfall, click-through, a queue — where nobody designed the configuration.



# Randomness as Missing Information

A random experiment is still causal. Its unpredictability reflects **missing information about the causal chain**, not an absence of causes.

Consequence: a random experiment can become a deterministic process once new information is available.

A coin flip filmed at 10,000 fps and measured precisely enough stops being random. What does that imply for a model that calls something "random"?

## Speaker notes

This is the epistemic-modeling principle in probability form (P3): calling something random is a statement about our knowledge, not only about the system. Good place to ask where that matters in data science — noise vs. unmodelled signal is the same distinction.



# Sample Space and Events

**Definition (Sample space, event space).** A set  $\Omega = \{\omega_1, \omega_2, \dots, \omega_n\}$  is the **sample space** of a random experiment if every outcome is associated with exactly one element  $\omega \in \Omega$ . The elements are called **outcomes**.

Let  $\Omega$  be finite. Every subset  $A \subseteq \Omega$  is an **event**; the event  $A$  **occurs** iff the observed outcome  $\omega$  is a member of  $A$ . The set of all events, the power set  $\mathcal{P}(\Omega)$ , is the **event space**.

An outcome is a single result; an event is a *question you can ask* about the result.


For infinite  $\Omega$  — the drug dosage below — the event space is a suitable family of subsets rather than the full power set; we return to this when continuous distributions appear.

Math first (P5). The load-bearing word is "exactly one" — outcomes must be mutually exclusive and exhaustive, which is a modelling decision, not a mathematical one. The outcome/event distinction is the one students blur.


# Important Event Types




$\emptyset$   
impossible event  
no outcome belongs to it




$\Omega$   
certain event  
every outcome belongs to it




$\bar{A} = \Omega \setminus A$   
complementary event  
 $A = \{1, 2\}$ ,  $\bar{A} = \{3, 4, 5, 6\}$



$|A| = 1$   
elementary event  
 $A = \{5\}$ : one outcome





$A \subseteq B$   
A entails B  
 $A = \{1, 2\}$ ,  $B = \{1, 2, 4, 5\}$



$A \cap B = \emptyset$   
incompatible events  
 $A = \{1, 2\}$ ,  $B = \{3, 6\}$

Six panels on one sample space, the six faces of a die, so that every event is a set of visible outcomes; read them as vocabulary, not results. "Entails" is worth dwelling on: set inclusion is logical implication between events, which is why the next unit can do algebra with pictures.


# Examples



An event is a set of outcomes you can shade — "heads", "at least one head" — and the sample space fixes which such questions can be asked.

Have students write down the event "exactly one head" ( $\{HT, TH\}$ ) and the event "the first flip is heads" ( $\{HH, HT\}$ ) — same sample space, different questions.

# The Event Space Grows Fast



$|\Omega| = n \rightarrow 2^n \text{ events}$

|    |           |
|----|-----------|
| 1  | 2         |
| 2  | 4         |
| 3  | 8         |
| 6  | 64        |
| 10 | 1,024     |
| 20 | 1,048,576 |

lines: one element added

For  $|\Omega| = n$  the event space has  $2^n$  elements — far too many to list, which is why probability is done by counting, not by enumeration.

The lattice diagram of the original is replaced by three inclusion levels; the point is the  $2^n$  growth, not the drawing. This motivates counting and, later, why we cannot always take the full power set for infinite  $\Omega$ .

# Not Every Sample Space Is Finite

Dosage of a drug:  $\Omega = (0, \infty)$



uncountably many outcomes — probability  
can only live on intervals, not on points

Race with 7 teams: all orderings



$7! = 5,040$  outcomes; shaded: the event "team 3 wins"

Choosing  $\Omega$  is a modelling decision: it fixes what counts as an outcome, and therefore which questions can be asked at all.


Speaker notes

Close the module on the modelling move (P3). Ask what is lost by modelling dosage as a positive real (measurement precision, dosing intervals) — the same question returns for every feature in every dataset later in the course.



# Venn Diagrams

Venn diagrams are the graphical representation of events:



- the **rectangle** is the sample space  $\Omega$
- a **circle** is an event  $A \subseteq \Omega$
- every **region** is a set — and every set operation is a region you can point at


Bridge from sets to pictures before any formula (P4). The payoff of the diagram: once probabilities are attached, probability laws become area arguments.



# Important Event Types

Before each step: which region will light up?

$$A \cap B = \emptyset$$



A and B are disjoint

## Speaker notes

Stepped visualization (P6, tier 2), one click per event type: the event  $A$ ; the complement  $\bar{A} = \Omega \setminus A$ ; the subset  $B \subset A$ ; disjoint events  $A \cap B = \emptyset$ . Predict-then-check on every click — students name the region first, the highlight confirms. Disjointness is the one to slow down on: it is a statement about events, not about probabilities.




# Algebra of Events

| Law          | Union form                                       | Intersection form                                |
|--------------|--------------------------------------------------|--------------------------------------------------|
| Commutative  | $A \cup B = B \cup A$                            | $A \cap B = B \cap A$                            |
| Associative  | $(A \cup B) \cup C = A \cup (B \cup C)$          | $(A \cap B) \cap C = A \cap (B \cap C)$          |
| Distributive | $(A \cup B) \cap C = (A \cap C) \cup (B \cap C)$ | $(A \cap B) \cup C = (A \cup C) \cap (B \cup C)$ |

Events form a set algebra — the same laws you know, but now every identity is a picture.

# Set Operations as Regions

$$A \cup B \cup C$$



Union of three events

Second stepped pass with three sets:  $A \cup B$ ,  $A \cap B$ ,  $A \cap B \cap C$ ,  $A \cup B \cup C$ . Same predict-then-check rhythm.



# Mini-Exercise

Are  $(A \cup B) \cap C$  and  $(A \cap B) \cup C$  the same region?


Speaker notes

Two minutes: students picture the regions in their head — which circle is the outer operation? — and commit to yes or no before the click. The figure is the answer; the algebra confirms it. If there is time, ask for a case where the two *are* equal (exactly when  $A \cap B \subseteq C \subseteq A \cup B$ ).




# Mini-Exercise

Are  $(A \cup B) \cap C$  and  $(A \cap B) \cup C$  the same region?



inside C only



all of C, plus the lens  $A \cap B$

**No.** The first lies entirely inside  $C$ ; the second is all of  $C$  plus the lens  $A \cap B$ . The distributive laws expand them into different expressions.

# An Intuition: Probability as Area

Draw a Venn diagram on the ground, with each area proportional to the number of elementary events it contains. Throw a stone into the rectangle at random.

The probability of  $A$  is how likely the stone lands inside  $A$ , assuming

1. we hit the rectangle  $\Omega$  for certain, and
2. every part of  $\Omega$  is equally likely to be hit.

Probability is then a **measure of the state**  $A \subseteq \Omega$  of a random experiment — literally, how much of the sample space the event occupies.

Both assumptions are doing real work.  
Assumption 2 is the one we will have to give up.

Intuition first (P4), and it connects directly to the Venn unit: events were regions, now regions have size. Ask what happens to the picture if some parts of  $\Omega$  are "heavier" than others — that is precisely what the axioms will allow.



# Two Ways to Read a Probability

## Frequentist

The probability of an elementary event  $a$  is the fraction of repetitions of the experiment that yield  $a$ .

*Requires* the experiment to be repeatable.

## Bayesian

The probability of  $a$  is the degree of certainty that  $a$  will happen — even without conducting the experiment.

*Requires* a stated state of knowledge.

Both interpretations yield the same mathematics. They are mental models for what the number *means*, not different theories.

Speaker notes

Interpretation over recitation (P2): ask which reading makes sense for "the probability that this patient has the disease" (no repetition available) versus "the probability that this machine part fails within a year". The distinction returns whenever we call something a prior.



# The Classical (Laplace) Definition

**Definition (Classical / Laplace probability).** If every elementary event in  $\Omega$  is assigned the same probability, then

$$P(A) = \frac{|A|}{|\Omega|} = \frac{\text{number of cases favourable for } A}{\text{number of possible outcomes}}$$

An experiment whose configuration and procedure make outcomes equiprobable — by definition or by construction — is a **Laplace experiment**; its probabilities need not be estimated, they are implied by the setup.

This is the definition that counting serves. Emphasise that it is only usable under the Laplace assumption; when that fails, probabilities must come from data.

# Why That Definition Is Not Enough

**The circularity.** "Laplace probability" defines probability using *equiprobability* — that is, using another notion of probability. Strictly speaking it is a circular definition.

The frequentist repair — defining probability as the limit of relative frequencies — fails too: such a limit exists only in mathematical settings, never in an experiment that can actually be repeated.

**So what do we do?** Stop trying to *define* what probability is, and instead specify what it must *do*.

Speaker notes

This slide is the pivot of the module and answers the guiding question. Let the tension sit for a moment before the axioms arrive: students should feel why an axiomatic move is the honest way out, not a formal ritual.



# The Kolmogorov Axioms

**Definition (Probability measure, Kolmogorov 1933).** Let  $\Omega$  be a sample space and  $\mathcal{P}(\Omega)$  its event space. A function  $P : \mathcal{P}(\Omega) \rightarrow \mathbb{R}$  is a **probability measure** if:

## Axiom I

$$P(A) \geq 0$$

non-negative

## Axiom II

$$P(\Omega) = 1$$

normalised

## Axiom III

$$A \cap B = \emptyset \Rightarrow$$

$$P(A \cup B) = P(A) + P(B)$$

additive

Nothing is said about *which* probabilities the events get — only about how any assignment must behave.

**Beyond finite sample spaces.** When  $\Omega$  is infinite, the event space is not the full power set but a  $\sigma$ -**algebra** of measurable sets, and Axiom III is required for **countably many** disjoint events. Everything we compute in this course is unaffected; the construction belongs to measure theory.

## Speaker notes

Math first (P5). Ask: which axiom does the area picture make obvious? (All three — areas are non-negative, the whole rectangle is everything, disjoint regions have additive areas.) The axioms are the area intuition, minus the requirement that the areas be equal.



# Probability Space

**Definition.** Let  $\Omega$  be a sample space,  $\mathcal{P}(\Omega)$  an event space, and  $P$  a probability measure. The triple  $(\Omega, \mathcal{P}(\Omega), P)$  — also written  $(\Omega, P)$  — is a **probability space**.

Every probabilistic statement in this course is made inside some probability space. Naming it is the discipline that keeps modelling assumptions visible.

Whenever a later theorem opens with "let  $(\Omega, \mathcal{P}(\Omega), P)$  be a probability space", students should hear "we have agreed on outcomes, questions, and a measure".



# What the Axioms Already Give Us

$$P(A) + P(\overline{A}) = 1$$

*(Axioms II, III)*

$$P(\emptyset) = 0$$

*(from the above with  $A = \Omega$ )*

$$A \subseteq B \Rightarrow P(A) \leq P(B)$$

monotonicity *(Axioms I, III)*

$$P(A \cup B) = P(A) + P(B) - P(A \cap B)$$

inclusion–exclusion *(Axiom III)*

For mutually exclusive  $A_1, \dots, A_p$ :  $P(A_1 \cup \dots \cup A_p) = P(A_1) + \dots + P(A_p)$ .

Speaker notes

2x2 grid of consequences. Have students justify inclusion–exclusion on the Venn picture — the overlap was counted twice. Note that "A entails B implies A is no more likely than B" is monotonicity: the logic of events transfers to numbers.



# Takeaway

Probability is not defined by counting or by repetition — it is any assignment of numbers to events that is **non-negative, normalised, and additive**. The Laplace formula is one such assignment, not the definition.

Stabilise the epistemic move: an axiomatic definition says what the object must do, leaving where the numbers come from to the modeller. That freedom is what lets the same theory carry frequentist and Bayesian readings.

# Module 2

---

## Conditioning Is Inference

**Guiding Question:** *A 99%-accurate test says you are sick. Why is that not a 99% chance of being sick — and what exactly does the evidence do to your belief?*

# Motivating Example

You roll a die twice. **After the first roll you see a 3.** Knowing this, what is the probability that both rolls sum to 8?

## Speaker notes

Intuition first (P4): conditioning is not a new kind of probability, it is the old probability computed in a **smaller sample space**. Let students do the count before the definition appears.



# Motivating Example

You roll a die twice. **After the first roll you see a 3.** Knowing this, what is the probability that both rolls sum to 8?

- Before the first roll:  $\Omega = \{(1, 1), (1, 2), \dots, (6, 6)\}$ ,  $|\Omega| = 36$
- After seeing the 3, only  $\Omega_{red} = \{(3, 1), \dots, (3, 6)\}$  remains
- Outcomes like  $(2, 4)$  have become **impossible**
- The reduced space now carries all the probability:  $P(\Omega_{red}) = 1$
- "Sum is 8" inside it:  $E = \{(3, 5)\}$

$$P = \frac{1}{6}$$

# Definition

**Definition (Conditional probability).**


$$P(A \mid B) = \frac{P(A \cap B)}{P(B)}, \quad \text{if } P(B) > 0$$

$P(A \mid B)$  is the probability of  $A$  **under the condition**  $B$ .

Back to the stone: *given that the stone landed in area  $B$ , how likely is it that it also landed in  $A$ ?*  
We renormalise by the area of  $B$ , because  $B$  is now the whole world.

Math first (P5) after the dice intuition. The division by  $P(B)$  is exactly the renormalisation students performed implicitly when they used the six remaining outcomes.

# Conditioning, Drawn



**Before the last click:** what must the two shaded cells be divided by?

## Speaker notes

Stepped visualization (P6): the A columns; B's share inside each column; keep only B; renormalise. The stretch factor *is* the  $1/P(B)$  in the definition, so the picture performs the formula rather than illustrating it.

Then use the sliders. Two settings are worth doing live: drag  $P(B|A)$  and  $P(B|\neg A)$  to the same value — the bar's split stops moving and  $P(A|B) = P(A)$ , which is independence, met before it has a name. Then set  $P(B|\neg A)$  near zero: B becomes nearly conclusive evidence for A. Ask what has to be true of the two conditionals for evidence to be worth anything.



# Example: Two Coins

Two coins are flipped. The event of interest is  $B = \{(h, h)\}$ .

**(a) "The first coin is heads"**

$$F = \{(h, h), (h, t)\}$$

$$P(B \mid F) = \frac{P(B \cap F)}{P(F)} = \frac{1/4}{1/2} = \frac{1}{2}$$

**(b) "At least one coin is heads"**

$$A = \{(h, h), (h, t), (t, h)\}$$

$$P(B \mid A) = \frac{P(B \cap A)}{P(A)} = \frac{1/4}{3/4} = \frac{1}{3}$$

Same event  $B$ , two different pieces of information, two different answers. **Why is (b) smaller?**

## Speaker notes

Force interpretation (P2). The information in (a) is stronger — it names *which* coin — so it eliminates more of the sample space. This is the classic "boy or girl" puzzle in disguise; expect disagreement, which is the point.



# Total Probability

**Theorem (Total probability).** Let  $B_1, \dots, B_k$  be mutually exclusive with  $\Omega = B_1 \cup \dots \cup B_k$  and  $P(B_i) > 0$ . Then for any event  $A$ :

$$P(A) = \sum_{i=1}^k P(B_i) \cdot P(A \mid B_i)$$

**Proof.** Decompose  $A$  along the partition and use additivity:

$$P(A) = P((B_1 \cup \dots \cup B_k) \cap A) = \sum_{i=1}^k P(B_i \cap A) = \sum_{i=1}^k P(B_i) \cdot P(A \mid B_i)$$

Two axioms do all the work: completeness of the partition, then exclusiveness for additivity. This theorem is the denominator of Bayes' rule — say so now, so that Bayes looks assembled rather than conjured.

# Rules That Follow

## Product rule

$$\begin{aligned} P(A \cap B) &= P(B) \cdot P(A \mid B) \\ &= P(A) \cdot P(B \mid A) \end{aligned}$$

## Reversing the condition

$$P(B \mid A) = \frac{P(B) \cdot P(A \mid B)}{P(A)}$$

## Complement, same condition

$$P(\overline{A} \mid B) = 1 - P(A \mid B)$$

## But in general

$$P(A \mid \overline{B}) \neq 1 - P(A \mid B)$$

Held fixed at  $B$ , the map  $A \mapsto P(A \mid B)$  satisfies the Kolmogorov axioms — conditioning yields a **new probability measure**  $P_B$ .

2x2 grid. The top-right box is already Bayes' theorem — flag it, then move on; the Bayes unit gives it its name and its interpretation. The bottom two boxes are the most common exam mistake: you may complement the *event*, not the *condition*.

# Chain Rule

$$P(E_1 \cap E_2 \cap \dots \cap E_n) = P(E_1) \cdot P(E_2 \mid E_1) \cdot P(E_3 \mid E_1 \cap E_2) \cdots P(E_n \mid E_1 \cap \dots \cap E_{n-1})$$

By successive application of the product rule.

Any joint probability can be written as a chain of conditionals. Choosing a *good* order — and dropping conditions that do not matter — is exactly what graphical models will do later.

Do not belabour the proof; the payoff is the remark. The chain rule plus conditional independence is the entire idea behind belief networks.

# Takeaway

Conditioning **shrinks the sample space and renormalises**. Everything else — product rule, total probability, Bayes, the chain rule — is bookkeeping on that single move.

# Bayes' Theorem — An Intuition First

A test for a rare disease is **99% accurate**: positive for 99% of the sick, negative for 99% of the healthy. About **1 in 1,000 people** has the disease.

**Your test comes back positive. How likely is it that you are actually sick?**

Intuition-first entry (P4). Collect guesses by show of hands *before* clicking: most say 99%. The surprise is the motivation for the whole module. Do not explain yet — the next slide lets them count it themselves.

# Bayes' Theorem — An Intuition First

A test for a rare disease is **99% accurate**: positive for 99% of the sick, negative for 99% of the healthy. About **1 in 1,000 people** has the disease.

**Your test comes back positive. How likely is it that you are actually sick?**

Only about **9%** — the answer most people miss by a factor of ten.

# Why? Count People, Not Percentages

Restricting to the positives is exactly the conditioning step

$$P(\text{sick} \mid \text{positive}) = 1/11 \approx 9\%$$



Stepped visualization (P6 tier 2), one click per step: the population; the one sick person; the ~10 false alarms among the 999 healthy; restricting to the 11 positives. Ask before the last click: "of these 11 flagged people, how many are sick?" Everything here is counting and restriction of the sample space — no new machinery. This *is* conditioning, drawn.

# What Just Happened

Restricting attention to the people who tested positive is exactly the move from the conditional-probability definition:

$$P(A \mid B) = \frac{P(A \cap B)}{P(B)}$$

The 11 positives are the new sample space  $B$ ; the single sick person among them is  $A \cap B$ . The rarity of the disease survives the evidence because there were so many more healthy people to misclassify.

## The trap

"99% accurate" describes

$P(\text{positive} \mid \text{sick})$ .

You wanted  $P(\text{sick} \mid \text{positive})$ .


These are different numbers.

Name the confusion explicitly — swapping the two conditionals is the single most common probabilistic error, in exams and in the newspapers. Bayes' theorem is the machine that converts one into the other.

# Bayes' Theorem — Two Events

**Theorem (Bayes, 2 events).** Let  $(\Omega, \mathcal{P}(\Omega), P)$  be a probability space and  $A, B \in \mathcal{P}(\Omega)$  two events with  $P(B) > 0$ . Then:

$$P(A | B) = \frac{P(B | A) \cdot P(A)}{P(B)}$$



Which of the four was "99% accurate", and which one did the intuition ignore?

Math-first-then-meaning (P5), legitimate here because the intuition already exists. Stepped equation walk-through: each click highlights one part and its brace — posterior, likelihood, prior, evidence — while the rest dims. Map the disease test onto the symbols as you go: A = sick, B = positive test.  $P(A) = 1/1000$  is the term everyone forgets. Reference: (Sheldon & others 2002).

# Proof — Two Events

From the definition of conditional probability, twice:

$$P(A \mid B) = \frac{P(A \cap B)}{P(B)} = \frac{\frac{P(A \cap B)}{P(A)} \cdot P(A)}{P(B)} = \frac{P(B \mid A) \cdot P(A)}{P(B)}$$

## Speaker notes

One line: expand, multiply by  $P(A)/P(A)$ , recognise the likelihood. Worth asking which step needed  $P(B) > 0$  and which needed  $P(A) > 0$ .



# Bayes' Theorem — Multiple Hypotheses

**Theorem (Bayes, multiple events).** Let  $A_i \in \mathcal{P}(\Omega)$  be mutually exclusive with  $\Omega = A_1 \cup \dots \cup A_k$  and  $P(A_i) > 0$ , and let  $B$  be an event with  $P(B) > 0$ . Then

$$P(A_i | B) = \frac{P(B | A_i) \cdot P(A_i)}{P(B)} = \frac{P(B | A_i) \cdot P(A_i)}{\sum_{j=1}^k P(B | A_j) \cdot P(A_j)}$$

The denominator is the **total probability** of the evidence: the same observation, explained by every competing hypothesis, weighted by how much we believed each one beforehand.

## Speaker notes

The hypotheses partition  $\Omega$  — "sick vs. healthy" was the  $k = 2$  case. This is the form that reappears as a classifier (naive Bayes) and in belief networks: score every hypothesis by likelihood  $\times$  prior, then normalise.



# Proof — Multiple Hypotheses

As before, with the theorem of total probability supplying the denominator:

$$P(B) = \sum_{j=1}^k P(B \mid A_j) \cdot P(A_j)$$

# Mini-Exercise

Same test, but now screening a **high-risk group** where 1 in 20 people is sick — the test is unchanged.

Estimate  $P(\text{sick} \mid \text{positive})$  before computing it.

# Mini-Exercise — Solution

Solution fragmented per the animation policy. The punchline is the comparison with 9%: identical likelihoods, different prior, completely different clinical meaning.

# Mini-Exercise — Solution

Out of 1,000: 50 sick → 49.5 true positives; 950 healthy → 9.5 false alarms.

$$P(\text{sick} \mid \text{positive}) = \frac{49.5}{49.5 + 9.5} = \frac{49.5}{59} \approx 84\%$$

**The test did not change. The prior did.** This is why screening decisions depend on who is screened.

# Takeaway

Bayes' theorem is the arithmetic of changing your mind: **posterior**  $\propto$  **likelihood**  $\times$  **prior**.  
Evidence never speaks alone — it is always weighed against how plausible the hypothesis was to begin with.

Stabilise both halves: the base-rate trap (9%) and the prior's power (84%). Both return in Bayesian decision making and graphical models.

# Back to Lecture 01: the Model Says "Fraud"

A fraud model is **99% accurate** in both directions. One transaction in **1,000** is fraudulent. Per 100,000 transactions:

|            | flagged      | not flagged | total   |
|------------|--------------|-------------|---------|
| fraud      | 99           | 1           | 100     |
| legitimate | 999          | 98,901      | 99,900  |
| total      | <b>1,098</b> | 98,902      | 100,000 |

The disease test was not a medical curiosity: it is every classifier for a rare class. Let students fill the table before the click; it is the same counting as the natural-frequency picture, with the numbers of Lecture 01's opening example.

Two consequences to say out loud. First, "accuracy" is the wrong number to ask for when classes are rare — Lecture 06 opens with exactly that model. Second, this is why the opening slide of Lecture 01 insisted on a number rather than a label: a calibrated score of 0.09 would have told the analyst what the word "fraud" did not.

# Back to Lecture 01: the Model Says "Fraud"

A fraud model is **99% accurate** in both directions. One transaction in **1,000** is fraudulent. Per 100,000 transactions:

|            | flagged      | not flagged | total   |
|------------|--------------|-------------|---------|
| fraud      | 99           | 1           | 100     |
| legitimate | 999          | 98,901      | 99,900  |
| total      | <b>1,098</b> | 98,902      | 100,000 |

$$P(\text{fraud} \mid \text{flagged}) = \frac{99}{1,098} \approx 9\%$$

Nine in ten blocked cards belong to honest customers — from a "99% accurate" model. The label "fraud" hid the **prior**; a score would at least have carried it.

# Independence

**Definition (Statistical independence).** Two events  $A$  and  $B$  are **independent** iff

$$P(A \cap B) = P(A) \cdot P(B)$$

If additionally  $0 < P(B) < 1$ , this is equivalent to

$$P(A \mid B) = P(A) \quad \text{and} \quad P(A \mid \overline{B}) = P(A)$$

Independence means: **learning  $B$  changes nothing about  $A$ .**

The conditional form is the meaningful one; the product form is the one that is easy to check and stays defined when  $P(B) = 0$ . Emphasise the reading: knowing B tells you nothing new.

# The Classic Confusion

Statistical independence does not mean the events are disjoint.

Disjoint events with positive probability are about as **dependent** as events can be: if  $A$  occurs,  $B$  certainly did not.

$$A \cap B = \emptyset \Rightarrow P(A \mid B) = 0 \neq P(A)$$


## Speaker notes

Ask students to state, in one sentence, why disjointness is the opposite of independence. The Venn picture misleads here — non-overlapping circles look "unrelated", which is exactly wrong.



# Independence Is About Areas, Not Overlap

$$0.30 \cdot 0.40 = 0.12 = P(A \cap B) \rightarrow \text{independent}$$



$$P(A) = 0.30$$

$$P(B) = 0.40$$

$$P(A \cap B) = 0.12$$


## Speaker notes

Stepped viz (P6): band A, band B, their intersection, then the test.  $0.3 \cdot 0.4 = 0.12$  — independent. Ask before the last click whether they expect equality. The lesson: independence is a numerical coincidence of areas, not a visual property of the drawing.



# The Same Picture, Dependent

$$0.30 \cdot 0.40 = 0.12 \neq 0.20 = P(A \cap C) \rightarrow \text{dependent}$$



Here  $P(A) \cdot P(C) = 0.12$  but  $P(A \cap C) = 0.20$ : the intersection is **larger** than independence would predict —  $A$  and  $C$  occur together more often than chance.

## Speaker notes

Same marginals as the previous slide, different joint. This is the payoff of the pair: the marginals never tell you whether events are independent. Later, that is exactly why joint distributions are needed.



# More Than Two Events

Events  $A_1, \dots, A_n$  are (mutually) independent iff for **every** subset of indices  $i_1, \dots, i_l$ :

$$P(A_{i_1} \cap \dots \cap A_{i_l}) = P(A_{i_1}) \cdots P(A_{i_l})$$

Pairwise independence is **not** enough — every subset must factorise.

Give the standard counterexample if time allows (two fair coins and the event "the results agree" are pairwise independent but not mutually independent). The distinction matters for naive Bayes, which assumes far more than pairwise.

# Takeaway

Independence is a statement about **numbers**, not pictures:  $P(A \cap B)$  equals the product of the marginals. Disjoint is not independent, and identical marginals can hide either.

# Module 3

---

## From Outcomes to the Joint

**Guiding Question:** *We rarely care which outcome occurred, only what it was worth — and never about one quantity alone. How do we get from a sample space to the joint distribution the three moves start from?*

Second session starts here. Open by recalling the three moves: this module builds the object that "model the joint" refers to.



# Motivation: We Rarely Care About the Outcome Itself

Toss three fair coins:  $\Omega = \{H, T\} \times \{H, T\} \times \{H, T\}$ , eight outcomes.

Usually the interesting quantity is not *which* sequence occurred, but the **number of heads**  $Y$ :

|                          |                                   |                                   |                          |
|--------------------------|-----------------------------------|-----------------------------------|--------------------------|
| $P(Y = 0) = \frac{1}{8}$ | $P(Y = 1) = \frac{3}{8}$          | $P(Y = 2) = \frac{3}{8}$          | $P(Y = 3) = \frac{1}{8}$ |
| $(T, T, T)$              | $(H, T, T), (T, H, T), (T, T, H)$ | $(H, H, T), (T, H, H), (H, T, H)$ | $(H, H, H)$              |

$Y$  compresses eight outcomes into four numbers — and the probabilities come along automatically.

Intuition first (P4). Ask what was lost in the compression (which coin showed heads) and what was gained (a quantity we can average, compare, plot). That trade is the entire point of a random variable.

# Definition

**Definition (Random variable).** A random variable  $X$  maps outcomes to a measurable range  $E$ :  $X : \Omega \rightarrow E$ . Its probabilities are inherited from the underlying probability space:

$$P(X \in S) = P(\{\omega \in \Omega \mid X(\omega) \in S\})$$

A random variable is neither random nor a variable — it is a **function** on the sample space. The randomness lives in  $\omega$ , not in  $X$ .

Speaker notes

The naming trap is worth one minute; it removes a persistent confusion. The formula says: to get the probability of a value, collect all outcomes that map to it and use the measure we already have.




# The Function, Drawn

Three arrows arrive at 1 and at 2 — which is the whole shape

$\Omega$  (outcomes)

$Y$  (number of heads)



Stepped visualization (P6), three coins: the eight outcomes; the arrows, one per outcome; the four values with their counts; the resulting mass function. Two things to say out loud while stepping. First, *every* outcome has exactly one arrow leaving it — that is what makes  $Y$  a function. Second, three arrows arrive at 1 and three at 2, and that is the entire reason those values are more probable. The probability was never invented; it was inherited.

# Examples

## Two coins

$$\Omega = \{HH, HT, TH, TT\}$$

$X$  = number of heads

## People

$$\Omega = \{\text{person}_1, \text{person}_2, \dots\}$$

$X$  = height of the person

## Two dice

$$\Omega = \{\langle n_1, n_2 \rangle\}$$

$$X(\omega) = n_1 + n_2$$

The middle example is the one that connects to data science: a dataset column *is* a random variable — a function from records to values. Say it explicitly; it makes the next chapters feel continuous with this one.

# Discrete and Continuous

## Discrete random variables

Countable range, usually  $E \subseteq \mathbb{N}$ .

- number of children of a person
- sum of two dice

## Continuous random variables

Real-valued range,  $E = \mathbb{R}$ .

- height of a person
- horsepower of a car

The distinction decides *how* we may describe the probabilities — with a mass function or with a density.

Keep this light: the deep measure-theoretic reasons are out of scope. What students need is the practical consequence — sums for discrete, integrals for continuous.

# A Random Variable With Infinitely Many Values

Flip a biased coin ( $P(H) = p$ ) until the first head. Let  $X$  be the number of flips needed:

$$P(X = 1) = p, \quad P(X = 2) = (1 - p)p, \quad P(X = 3) = (1 - p)^2 p$$

$$P(X = n) = (1 - p)^{n-1} p$$

Countably infinite range, and the probabilities still sum to 1.

Speaker notes

This is the geometric distribution, met before it is named — fine, because it is constructed from the chain rule the students already have. Ask them to verify the sum is 1 (geometric series) if the group is comfortable with it.



# Takeaway

A random variable is a **function from outcomes to numbers** that carries the probability measure with it. It lets us stop talking about sample spaces and start talking about quantities — which is what the rest of the course does.

# Probability Mass Function

**Definition (PMF).** For a discrete random variable  $X$ , the **probability mass function** is

$$p(a) = P(X = a)$$

$p(a) > 0$  for at most countably  
many  $a$


$$\sum_i p(a_i) = 1$$

$p$  is itself a probability measure

The PMF is just the probability measure, re-indexed by values instead of outcomes. Everything from the axioms still applies — which is why the three properties are consequences, not new assumptions.

# Example: Sum of Two Dice

Sum of two dice — probability mass



The triangular shape is counting, drawn: six of the 36 outcomes give 7, only one gives 2.

Connect the picture to the Laplace assumption: 36 equally likely outcomes, counted by sum. Ask why the shape is symmetric and peaked — more ways to make a middling sum. This is the first hint of why sums of random quantities concentrate.

# Expectation

Definition (Expected value).

$$E[X] = \sum_i a_i \cdot p(a_i)$$

- Frequentist reading: the **average outcome** if the experiment were repeated infinitely often
- The first moment:  $\mu = E[X^1]$ ; the first *central* moment  $E[X - E[X]]$  is always 0

For a real function  $g$ :

$$E[g(X)] = \sum_i g(a_i)p(a_i)$$

and for constants  $a, b$ :

$$E[aX + b] = aE[X] + b$$

## Speaker notes

The linearity identity is the one used constantly later (in regression, in evaluation). Note that  $E[g(X)] \neq g(E[X])$  in general — ask for a counterexample with  $g(x) = x^2$ .



# Variance

## Definition (Variance).

$$\text{Var}(X) = E[(X - \mu)^2] = E[X^2] - (E[X])^2$$

- a measure of **dispersion** around the mean — larger variance, larger typical deviation
- the second central moment
- $\text{Var}(aX + b) = a^2 \text{Var}(X)$
- $\sigma = \sqrt{\text{Var}(X)}$  is the **standard deviation**, in the same unit as  $X$

## Speaker notes

Two things to stabilise: the computational identity  $E[X^2] - E[X]^2$ , and why we report  $\sigma$  rather than the variance (units). The  $b$  disappearing under variance is worth a question: shifting data does not change its spread.



# Takeaway

A discrete distribution is a PMF; expectation and variance compress it into **where it sits** and **how much it spreads**.

# From Mass to Density

**Definition (PDF).** A continuous random variable  $X$  has a **probability density function**  $f$  with

$$P(b \leq X \leq c) = \int_b^c f(a)da, \quad \int_{-\infty}^{\infty} f(a)da = 1$$

$f(a)$  is **not** a probability — it is probability *per unit interval*. Consequently  $f(a) > 1$  is perfectly possible, and  $P(X = a) = 0$  for every single value  $a$ .


## Speaker notes

The density-is-not-probability point is the single biggest source of confusion in this unit. A useful question: what is the probability that a person is exactly 180.000... cm tall? Zero — probability lives on intervals now.



# Probability Is Area Under the Curve

Normal  $N(\mu = 0, \sigma = 1)$  — density



The shaded region is  $P(-1 \leq X \leq 1)$  — the same "probability as area" intuition from the axioms of probability, now with a curve instead of a Venn diagram.

Deliberate callback (P7 spiral): area was the intuition for the axioms, and it still is. Ask what the total area must be, and why that forces the normalisation constant in the density formula.

# Expectation and Variance — Sums Become Integrals

## Discrete

$$E[X] = \sum_i a_i p(a_i)$$

$$E[g(X)] = \sum_i g(a_i) p(a_i)$$

## Continuous

$$E[X] = \int_{-\infty}^{\infty} a f(a) da$$

$$E[g(X)] = \int_{-\infty}^{\infty} g(a) f(a) da$$

The rules are unchanged:  $E[aX + b] = aE[X] + b$ , and  $\text{Var}(X) = E[(X - \mu)^2]$ .

Emphasise the continuity of concepts across the discrete/continuous divide: only the summation symbol changed. Students who see this stop treating the two cases as separate theories.

# Two Basic Continuous Distributions

**Uniform on  $[a, b]$**

$$f(x) = \frac{1}{b-a} \text{ for } x \in [a, b]$$

Every interval of equal length is equally likely — the Laplace assumption, continuous version.

**Normal  $N(\mu, \sigma^2)$**


$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

$E[X] = \mu$ ,  $\text{Var}(X) = \sigma^2$ ;  $N(0, 1)$  is the **standard** normal.

The uniform is worth naming as the continuous Laplace assumption — it makes the earlier discussion feel like one idea rather than two.

# What $\sigma$ Does to the Shape

Normal  $N(\mu = 0, \sigma = 2.4)$  — density



The peak falls as the curve widens — the area is always 1.

Predict-then-check: ask what must happen to the height when the width grows. The constraint "total area = 1" is doing the work, and it is the same normalisation axiom that defined a probability measure.

# Why the Normal Is Everywhere

The **central limit theorem**: under mild conditions, the average of  $n$  independent observations is approximately normal for large  $n$  — whatever the original distribution was.

This is why the normal is the default when the true distribution is unknown — and why a sum of many independent coin flips, the binomial, looks bell-shaped.

State the CLT, do not prove it. If the binomial sweep has already been shown, call back to it — students already saw it happen; if it comes later, it will confirm the theorem. This also sets up the sampling distribution of a mean, which every error bar computed from a sample rests on.

# Other Distributions You Will Meet

- **Exponential** — time until an event occurs (e.g. until the next earthquake)
- **Gamma** — time until  $n$  events have occurred
- **Beta** — values in  $[0, 1]$ , e.g. an unknown probability
- **Dirichlet** — distributions over probabilities

Continuous distributions replace mass with **density** and sums with **integrals**. Probability remains area, expectation remains a weighted average, and the normal is the default shape because of the central limit theorem.

## Speaker notes

The last two (Beta, Dirichlet) are named for recognition only — they return wherever a probability is itself unknown, as priors over probabilities. That is a legitimate forward *pointer*, not a forward dependency: nothing here uses them.



# More Than One Random Variable

## Univariate

one random variable

## Bivariate

two random variables

## Multivariate

three or more

A **joint distribution** describes the outcomes of several random variables *at the same time* — which is what every dataset with more than one column is.

Make the dataset connection explicit: a table with columns age, income, and churn is a sample from a trivariate joint distribution. Everything in the rest of the course operates on joints.

# Joint Mass and Density

**Discrete** — joint PMF

$$p_{X,Y}(a,b) = P(X = a, Y = b)$$

**Continuous** — joint PDF

$$P((X,Y) \in R) = \iint_R f_{X,Y}(a,b) da db$$

**Notational convention from here on.** We stop distinguishing discrete from continuous variables and simply write  $p(a,b,c,\dots)$  for the probability that  $\{X_1, X_2, X_3, \dots\}$  is in state  $a, b, c, \dots$

Flag the convention as a deliberate simplification, and say why it is safe: the rules (marginalise, condition, factorise) are identical; only the mechanics of summing versus integrating differ.

# Marginalisation

**Marginalisation** eliminates a random variable by summing over all its states. For  $p(x_1, \dots, x_n)$ , a **marginal** over a subset is

$$p(x_1) = \sum_{x_2} \cdots \sum_{x_n} p(x_1, x_2, \dots, x_n)$$

"Forget about a variable, but keep its influence correctly accounted for."

Marginalisation is the total-probability theorem, applied to random variables — say so, since students met the theorem with conditional probability. That continuity is worth more than the formula.

# Example: Boys and Girls

15% of families have no children, 20% have one, 35% have two, 30% have three. Each child is independently a boy or a girl with probability  $\frac{1}{2}$ .

Let  $B$  = number of boys,  $G$  = number of girls.

| $p(B, G)$ | $G = 0$ | $G = 1$ | $G = 2$ | $G = 3$ | $p(B)$ |
|-----------|---------|---------|---------|---------|--------|
| $B = 0$   | .15     | .10     | .0875   | .0375   | .375   |
| $B = 1$   | .10     | .175    | .1125   | 0       | .3875  |
| $B = 2$   | .0875   | .1125   | 0       | 0       | .2     |
| $B = 3$   | .0375   | 0       | 0       | 0       | .0375  |
| $p(G)$    | .375    | .3875   | .2      | .0375   | 1      |

The margins of the table *are* the marginals — which is where the name comes from.

Speaker notes

Have students read one cell aloud as a sentence ("11.25% of families have one boy and two girls"). Then ask whether B and G are independent — they are not, since the family size couples them. Good preparation for covariance.



# Covariance

$$\text{cov}(A, B) = E[(A - E[A])(B - E[B])]$$

$$\text{cov} > 0$$

$A$  large when  $B$  large —  
positively related

$$\text{cov} < 0$$

$A$  large when  $B$  small —  
negatively related

$$\text{cov}(A, A) = \text{Var}(A)$$

variance is self-covariance

Independence implies zero covariance. **The converse is false** — covariance only detects *monotone linear* association.


## Speaker notes

The warning is the load-bearing statement. A classic counterexample:  $X$  uniform on  $[-1,1]$  and  $Y = X^2$  are dependent but have zero covariance. This is why later methods that rely on correlation can miss real structure.



# Covariance Sees Only One Kind of Structure

Slide  $\rho$ , then switch the relation. Predict the covariance first.



**Predict first:** with  $Y = X^2 - 1$ , every point is fixed by its  $x$ . The covariance?

## Speaker notes

Stepped, then manipulable (P6 tiers 2 and 3). The quadrant signs make the definition visible: covariance averages the sign of  $(x - \bar{x})(y - \bar{y})$ , so a cloud that is symmetric about the vertical line has the two signs in balance and averages to zero — however tightly the points sit on a curve.

Collect the prediction before switching. Almost everyone says "large": the dependence is total. The reading afterwards is  $\approx 0$ . That is the counterexample from the warning above, on screen, and it is the reason "uncorrelated" and "independent" must stay different words.



# Takeaway

The joint distribution holds everything; **marginalisation** sums a variable away, **conditioning** restricts, and **covariance** summarises linear co-variation. Zero covariance is weaker than independence.

# Module 4

---

## Data, Statistics and a Parameter

**Guiding Question:** *A model has a parameter nobody told us, and all we have is a table.  
What does the table say about the parameter — and how sure may we be?*

## Speaker notes

Three steps, in this order: what data is (attribute types, and a table as a sample from a joint), what we compute from it (statistics, estimators and their error), and what it says about a parameter (the Beta–Bernoulli posterior). The module ends where Lecture 03 begins: the peak of that posterior is maximum likelihood.

If time is short, drop the descriptive units (means, median, variance) — the sampling unit recalls what it needs.



# Four Attribute Types

## Nominal

*categorical,  
qualitative*

Labels only: eye  
colour, ZIP code, ID.

Meaningful:  $=$ ,  $\neq$

## Ordinal

*categorical,  
qualitative*

Ordered labels:  
grades, rankings,  
"small/medium/large".

Meaningful:  $<$ ,  $>$

## Interval

*numeric, quantitative*

Differences are  
meaningful, the zero  
point is arbitrary:  
temperature in  $^{\circ}\text{C}$ ,  
calendar dates.

Meaningful:  $+$ ,  $-$

## Ratio

*numeric, quantitative*

A true zero, so ratios  
are meaningful:  
length, mass, counts,  
duration.

Meaningful:  $\times$ ,  $\div$

The four types form a hierarchy: each level permits everything the previous one does, plus more. The type is a property of the *measurement*, not of the storage format — which is why a ZIP code stored as an integer is still nominal.

# Why the Type Decides the Method

## The classic error

Averaging a nominal attribute. The mean of ZIP codes 10115 and 81667 is a number, and it is meaningless.

## The subtle one

20 °C is not "twice as warm" as 10 °C — interval scales have no true zero, so ratios do not transfer.

The attribute type determines which **statistics**, which **distance measures** and which **plots** are permissible. Every method in this course depends on it.

# Mini-Exercise

A customer table records satisfaction (one to five stars) and year of birth. Which scale is each — and whose **mean** would you not report?

Solution on the same slide, click-gated. The two columns come back when a table is read as a sample from a joint distribution; payment method (nominal: only counting) and income (ratio: a true zero) complete the set if there is time.

# Mini-Exercise

A customer table records satisfaction (one to five stars) and year of birth. Which scale is each — and whose **mean** would you not report?

Satisfaction is **ordinal**: its mean assumes equal gaps between the stars — report the median. Year of birth is **interval**: the mean year is fine, but ratios are not ("born twice as late" means nothing); differences such as age are fine.

# Takeaway

An attribute's **scale of measurement** licenses the operations you may perform on it. Choosing a statistic before checking the scale is the most common silent error in data analysis.

# A Table Is a Sample From a Joint Distribution

| customer | payment | satisfaction | born | income (k€) | returned? |
|----------|---------|--------------|------|-------------|-----------|
| 1        | card    | 4            | 1987 | 52          | yes       |
| 2        | invoice | 2            | 1995 | 31          | no        |
| 3        | card    | 5            | 1979 | 78          | no        |
| ...      | ...     | ...          | ...  | ...         | ...       |

## Each column is a random variable

$X_{\text{payment}}, X_{\text{satisf.}}, \dots$  — a function from customers to values, exactly as a random variable was defined.

## Each row is one draw

The whole table is  $n$  draws from one joint distribution  $p(x_1, \dots, x_d)$  that nobody has written down.

Satisfaction and year of birth are the columns of the attribute-type mini-exercise — deliberately. There they were classified by scale; here they become random variables. Say the sentence that makes the rest of the course possible: *a dataset is not the distribution, it is a sample from it*. Everything we compute from the table is therefore an estimate, and an estimate comes with an error.

# The Scale Suggests the Model

| attribute    | scale                   | a first model                                              |
|--------------|-------------------------|------------------------------------------------------------|
| returned?    | nominal, two values     | <b>Bernoulli</b> ( $p$ ) — one parameter                   |
| payment      | nominal, $k$ values     | <b>categorical</b> ( $p_1, \dots, p_k$ ), $\sum_j p_j = 1$ |
| satisfaction | ordinal                 | categorical — the order is not used, or an ordinal model   |
| born         | interval                | normal, on a restricted range                              |
| income       | ratio, positive, skewed | log-normal, or normal on log income                        |

The scale limits the model as it limits the statistics: a normal distribution for ZIP codes is as meaningless as their mean.

Not a catalogue to memorise — a habit: before choosing a distribution, check what the scale permits. The categorical distribution is the Bernoulli with more than two values; it is the distribution naive Bayes uses for every nominal attribute. Satisfaction is the honest hard case: treating it as categorical throws the order away, treating it as numeric invents equal gaps.

# What "i.i.d." Claims

The standard reading of a table rests on one assumption: the rows are **independent and identically distributed**.

## Independent

One customer's row tells us nothing about another's.

## Identically distributed

Every row comes from the **same** joint distribution.

Which of these datasets break it: daily sales of one shop · two visits by the same customer · pixels of one image · a survey sent to a random sample?

The i.i.d. assumption is what turns "the table" into "a sample", and it is what lets a likelihood factorise into a product over rows — the step every fitting method takes. Link back to the shapes of data: record data is the case where i.i.d. is at least plausible; graphs, grids and time series are the cases where it is not.

# What "i.i.d." Claims

The standard reading of a table rests on one assumption: the rows are **independent and identically distributed**.

## Independent

One customer's row tells us nothing about another's.

## Identically distributed

Every row comes from the **same** joint distribution.

Which of these datasets break it: daily sales of one shop · two visits by the same customer · pixels of one image · a survey sent to a random sample?

Only the random survey is plausibly i.i.d. Time series, repeated measurements and spatial data all violate it — the arrangement *is* information.

# Takeaway

A dataset is  $n$  **draws from an unknown joint distribution**: columns are random variables, rows are draws, and the scale of each column suggests its model. Reading a table this way rests on an assumption — **i.i.d.** — that should be checked, not presumed.

# Bernoulli and Binomial

**Bernoulli**( $p$ ) — one trial

$X = 1$  success,  $X = 0$  failure

$$E[X] = p, \quad \text{Var}(X) = p(1 - p)$$

**Binomial**  $B(n, p)$  —  $n$  independent trials

$$p(k) = \binom{n}{k} p^k (1 - p)^{n-k}$$


$$E[X] = np, \quad \text{Var}(X) = np(1 - p)$$

$\binom{n}{k} = \frac{n!}{k!(n-k)!}$  counts the sequences with  $k$  successes; the PMF multiplies it by the probability of one such sequence — combinatorics and independence in a single line.

Point at where each ingredient comes from: the binomial coefficient counts the ways to place  $k$  successes among  $n$  trials, the product  $p^k(1-p)^{(n-k)}$  comes from independence. Nothing here is new machinery, which is exactly the message.

# How the Shape Depends on the Parameters

Binomial  $B(n = 60, p = 0.5)$  — probability mass



**Predict before each click:** as  $n$  grows, does the distribution get wider or narrower — and what happens to its shape?

## Speaker notes

Predict-then-check (P6). Two things happen at once: the spread in absolute terms grows ( $\sqrt{n}$ ) while the shape becomes symmetric and bell-like. The second observation is the central limit theorem at work — a sum of many independent coin flips. Name it as an observation; the theorem itself is stated with the normal distribution.



# Mini-Exercise

A student guesses all **10** questions of a quiz, each with **4** options.

How many correct answers should they expect — and is 5 or more likely?

## Speaker notes

Two minutes: the expectation is head arithmetic, the second half only asks for a judgement. If someone wants the exact tail, it is  $P(X \geq 5) \approx 0.078$  — a sum over six PMF values, not one. The modelling point matters more than the number: the binomial needs independent guesses with a constant success probability; ask when that would fail in a real exam.



# Mini-Exercise

A student guesses all **10** questions of a quiz, each with **4** options.

How many correct answers should they expect — and is 5 or more likely?

$B(10, 0.25)$ : on average  $10 \cdot 0.25 = 2.5$  correct. Five or more is twice the expectation — unlikely (about 8%).

# Takeaway

A yes/no attribute is a **Bernoulli** variable with one parameter  $p$ ; counting the yeses in  $n$  independent trials gives a **binomial**. Both are counting in probabilistic clothing — and  $p$  is the number data will have to tell us.

# We Never Know the Distribution

Up to now we assumed a distribution was given. In practice it is **unknown**, and we must estimate its properties from data.

## Population

The entire group — all humans on earth, all transactions, all documents.

Its properties  $(\mu, \sigma^2)$  are the truth we cannot observe.

## Sample

The observations we actually have.

Its properties  $(\bar{x}, s^2)$  are what we compute; estimates of the population's  $(\bar{x}, \hat{\sigma}^2)$  are built from them.

**Descriptive statistics** summarise a sample, usually while assuming something about the underlying distribution.

Speaker notes

This slide is the hinge of the whole module: everything before it was deduction from a model, everything after is inference from data. The population/sample distinction is exactly the gap that statistics manages.



# What We Summarise

## Central tendency

1st order moment  
sample mean, median

## Dispersion

2nd central moment  
variance, standard deviation,  
interquartile range

## Shape

3rd central moment  
skewness — or, without moments,  
quantiles

Univariate statistics describe one random variable; **multivariate** statistics describe  $X_1, \dots, X_d$  together.

The moment ordering gives a clean mental index: where is it, how spread, how lopsided. Each has a robust and a non-robust representative: the median against the mean for location, the interquartile range against the standard deviation for spread.

# Means

**Sample mean**

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

**Weighted mean**

$$\bar{x}_w = \frac{\sum_i w_i x_i}{\sum_i w_i}$$

**Geometric mean** — for rates and ratios

$$\left( \prod_i x_i \right)^{1/n}$$

**Harmonic mean** — for averaging rates

$$\frac{n}{\sum_i 1/x_i}$$

**Trimmed mean:** drop the  $k$  smallest and largest values, then average — a cheap way to buy robustness.

## Speaker notes

The sample mean estimates  $E[X]$ . Ask when the arithmetic mean is simply wrong: average speed over equal distances (harmonic), average growth rate (geometric). Choosing the mean is a modelling decision, not a formality.



# Median

The **median** is the middle value of the sorted sample: for odd  $n$  the  $\frac{n+1}{2}$ -th value, for even  $n$  the average of the two middle values.

**Robust** against outliers — one extreme value moves the median by at most one position


**Costlier** to compute: requires sorting, or an interval approximation over binned data

For binned data, the median is interpolated within its interval from the cumulative counts, the interval's lower border and width.

The mean-versus-median contrast is the first robustness discussion of the course, and it returns in clustering (k-means vs. k-medoids) and in evaluation. Ask which one a salary report should quote, and why the answer is political.

# One Value, Two Very Different Reactions

Drag the light point — which marker follows it?



**Commit first:** drag the top salary to the far right. Which marker follows?



## Speaker notes

Manipulable widget (P6 tier 3) with the prediction taken first, as the policy requires. The mean is drawn as the fulcrum the beam balances on, so the drag is visibly a lever: one value far out has to be balanced by everything else.

Drag it all the way and stop. The mean is now above six of the seven salaries — a "typical" value that almost nobody earns. The median has not moved at all — it is still 44 k€. Then ask the political question: which number goes in the recruitment advert, and which one in the union's press release?



# Mini-Exercise

Salaries in a small firm (in k€): 38, 41, 42, 44, 45, 47, 260.

Compute the mean and the median. Which number should appear in a recruitment ad, and which one describes a typical employee?

# Mini-Exercise — Solution

The punchline: neither number is wrong, but only one is honest as a summary of "typical". This is the first appearance of the idea that a statistic encodes a question.

# Mini-Exercise — Solution

Mean  $\approx 73.9$ , median = 44.

The single large value pulls the mean above **every** salary except one. The median describes the typical employee; the mean describes the payroll divided by headcount. Both are correct — they answer different questions.

# Dispersion

**Dispersion** describes how stretched or squeezed the data is around its centre.

**Sample variance** — the variance *of the observed data*

$$s^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

**Standard deviation**  $s = \sqrt{s^2}$ , in the unit of the data.

Computable in **one pass** over the data,  $O(n)$ , using

$$\sum x_i^2 \text{ and } \sum x_i$$

— which matters when the data does not fit in memory.

The one-pass remark is the data-engineering link: streaming aggregations rely on it. Mention the numerical-stability caveat if the group is advanced.

# Why the Denominator Changes

The **population variance**  $\sigma^2$  is the variance over the whole population. If the true mean  $\mu$  is unknown, we must estimate it — and using  $\bar{x}$  in its place makes the deviations systematically too small.

$$\hat{\sigma}^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

Dividing by  $n - 1$  compensates: one degree of freedom was spent estimating the mean from the same data.

Naming is not uniform: many textbooks, and pandas, call the  $n - 1$  version  $s^2$ ; NumPy's `var` divides by  $n$  unless told otherwise.

Do not derive the correction; give the intuition. A good probe: with  $n = 1$ , the sample variance is 0 (the point equals its own mean) — clearly not an estimate of population spread, and the corrected formula refuses to answer at all.

# Takeaway

$s^2$  describes **the sample**;  $\hat{\sigma}^2$  estimates **the population**. The  $n - 1$  is not a convention — it is the price of having estimated the mean from the same data.

# You Computed a Number. Which Number Did You Want?

Seven salaries gave a sample mean of  $\bar{x} = 73.9$  k€.

What is the probability that the population mean  $\mu$  is exactly 73.9 k€?

## Speaker notes

Open on the gap, not on a definition (P1, P4). Most students will answer with a number or hedge; the point is that the question is malformed. Everything in this unit exists to replace " $\mu = 73.9$ " with " $\mu = 73.9 \pm \text{something}$ ", and to say where the something comes from.



# You Computed a Number. Which Number Did You Want?

Seven salaries gave a sample mean of  $\bar{x} = 73.9$  k€.

What is the probability that the population mean  $\mu$  is exactly 73.9 k€?

Zero.  $\bar{x}$  is not  $\mu$  — it is one **estimate** of it, computed from the seven people who happened to be in the sample.

The honest question is not *what is  $\mu$*  but **how far from  $\mu$  can this number be?**

# Estimator and Estimate


**Definition.** An **estimator** is a function of the sample —  $\hat{\theta} = g(X_1, \dots, X_n)$  — used to approximate an unknown quantity  $\theta$  of the population. The number it takes on for one particular sample is an **estimate**.

The sample is random, so the estimator is a **function of random variables**:  $\hat{\theta}$  is itself a random variable, with a distribution of its own.

$\bar{x}$ ,  $s^2$  and  $\hat{\sigma}^2$  from descriptive statistics are all estimators. We used them without saying what kind of object they are.

Math first (P5) — this is a second pass over objects the students already compute. The load-bearing sentence is the middle one: a statistic inherits randomness from the sample. Everything else follows from taking that seriously.

# An Estimator Has a Distribution



Stepped visualization (P6 tier 2 with tier-3 controls). One click per step: draw one sample; record its mean; fifty samples; a thousand, with the normal curve and the standard error. Both pictures share the x-axis, so "the means are more tightly packed than the population" is something students see rather than something we assert. Collect predictions before the third click — the common guess is that the spread stays the same.

# Three Things the Picture Shows

- The means pile up **around**  $\mu$  — the estimator is not systematically off
- The pile is **narrower than the population** — averaging cancels noise
- The pile is **bell-shaped**, even though the population was not

**Definition (Sampling distribution).** The distribution of an estimator  $\hat{\theta}$  over all possible samples of size  $n$ .  $\hat{\theta}$  is **unbiased** if  $E[\hat{\theta}] = \theta$ .

## Speaker notes

The three observations are, in order: unbiasedness, the standard error, and the central limit theorem — named on the next slides. Naming them here would be faster and worse. Ask which of the three would still hold for the *median* (the first two do; the third does too, but that needs a different argument).



# Use the Widget: Change the Population

Set the population to **skewed** and then to **bimodal**, keeping  $n = 5$ .

Which of the three observations survive a population that is not symmetric?

Tier-3 exploration, with the prediction committed first as the policy requires. The bimodal population is the persuasive one: two separate humps, and the means of five draws still form a single bell. Push n to 2 to show where it breaks down — the claim is asymptotic, not magic.

# Use the Widget: Change the Population

Set the population to **skewed** and then to **bimodal**, keeping  $n = 5$ .

Which of the three observations survive a population that is not symmetric?

All three. The shape of the sampling distribution stops depending on the shape of the population as soon as  $n$  is more than a handful — which is why one formula can cover cases we never inspect.

# Standard Error

**Definition (Standard error).** The standard deviation of an estimator's sampling distribution. For the sample mean of independent draws:

$$\text{SE}(\bar{X}) = \frac{\sigma}{\sqrt{n}}$$

It is measured in the unit of the data and answers "how far off is a typical estimate", where  $\sigma$  answers "how far apart are typical *observations*".

$\sigma$  and SE are different quantities. More data does nothing to  $\sigma$  — the world stays as variable as it was — and shrinks SE.

## Speaker notes

The warning is the most valuable line in the unit:  $\sigma$  is the variability of the world (aleatoric uncertainty), SE is our ignorance about the estimate (epistemic uncertainty) — the two words return at the end of the unit. Ask which one a weather forecast's " $\pm 3$  °C" refers to; the answer is usually both, mixed.



# Precision Costs Data

| Multiply the data by... | ...and the standard error is multiplied by |
|-------------------------|--------------------------------------------|
| 2                       | 0.71                                       |
| <b>4</b>                | <b><math>\frac{1}{2}</math></b>            |
| 25                      | $\frac{1}{5}$                              |
| <b>100</b>              | <b><math>\frac{1}{10}</math></b>           |

The  $\sqrt{n}$  is why the last digit of an estimate is so expensive: halving the error costs four times the data, and a tenfold improvement costs a hundredfold.

Speaker notes

The economics of data collection in one line, and a sanity check students can apply to any claim about "more data". Ask what this implies for an A/B test that wants to detect a 0.1% difference. It also foreshadows why we later care about better *models* rather than only bigger *samples*.



# Two Theorems We Have Already Seen Happen

## Law of large numbers

$\bar{X}_n \rightarrow \mu$  as  $n \rightarrow \infty$ .

The estimate converges on the truth — the pile keeps narrowing.

## Central limit theorem

$\bar{X}_n$  is approximately  $N(\mu, \sigma^2/n)$  for large  $n$ .

And we know the *shape* of what is left.

The first says the error goes away; the second says how much of it is left after  $n$  observations. Only the second lets us put a number on our ignorance.

Deliberate spiral (P7): the CLT was stated in the continuous-distributions unit and was visible in the binomial sweep. Here it acquires a job. Emphasise the division of labour between the two theorems — students routinely quote the LLN where they need the CLT.

# From a Number to an Interval

If  $\bar{X} \approx N(\mu, \sigma^2/n)$ , then for 95% of samples  $\bar{X}$  lands within 1.96 SE of  $\mu$  — and therefore  $\mu$  lands within 1.96 SE of  $\bar{X}$ .

$$\bar{x} \pm 1.96 \cdot \frac{\sigma}{\sqrt{n}}$$

the estimate — our best single guess


how confident we insist on being

how noisy the estimate is

Math-first-then-meaning (P5), legitimate because the picture came first. Walk the three parts and ask what each one would have to do to make the interval wider.  $\sigma$  is treated as known here; the correction for an estimated  $\sigma$  (the t distribution) is a refinement, not a different idea.

# What Does "95% Confident" Mean?

Change the level or n — which number follows, and which does not?



Stepped, then manipulable. One interval; twenty; sixty with the coverage tally; then the invitation to change the level and n. Ask before the third click how many of sixty intervals should miss. The red rows are the whole argument: the misses look exactly like the hits from the inside.

# The Sentence Almost Everyone Gets Wrong

## WHAT IT DOES NOT MEAN

"There is a 95% probability that  $\mu$  lies in **this** interval."

$\mu$  is a fixed number. It is either in this interval or it is not — there is no probability left to have.

## WHAT IT MEANS

"95% of the intervals produced **by this procedure** contain  $\mu$ ."

The probability is a property of the recipe, evaluated before the data is seen.

This is the distinction the picture just made — the rows are the procedure, one row is your data. Expect resistance: the wrong reading is the useful-sounding one. Note also that a Bayesian credible interval *does* license the first sentence, at the price of requiring a prior — a forward pointer, not a dependency.

# Where This Leaves Us

Two kinds of uncertainty meet in one formula. We can now measure one of them.

- $\sigma$  — how variable the world is. **Aleatoric**. Unchanged by data.
- $SE = \sigma / \sqrt{n}$  — how badly we know our own estimate.  
**Epistemic**. Shrinks with data.

"More data helps" was a slogan.

It is now a rate:  $1/\sqrt{n}$ .

If the course has already named the two kinds of uncertainty, close that loop here; if not, the two bullets introduce the words. This is the slide to point back to whenever a later lecture claims a model is uncertain "because it has not seen enough" — that claim has a magnitude, and this is where it comes from.

# Mini-Exercise

A shop timed  $n = 100$  customers at a counter:  $\bar{x} = 180$  s, and past records give  $\sigma = 60$  s. What is the standard error — and how many customers would halve it?

## Speaker notes

Solution on the same slide, click-gated. The second part is the one that transfers: the  $\sqrt{n}$  law is the answer to nearly every "how much data do we need" question a student will be asked. The interval is  $180 \pm 1.96 \cdot 6 \approx [168, 192]$  s. If time allows, ask what would change if  $\sigma$  had to be estimated too.



# Mini-Exercise

A shop timed  $n = 100$  customers at a counter:  $\bar{x} = 180$  s, and past records give  $\sigma = 60$  s. What is the standard error — and how many customers would halve it?

$SE = 60/\sqrt{100} = 6$  s, so the 95% interval is about  $180 \pm 12$  s. Halving the SE needs twice the  $\sqrt{n}$  — **four times the data**,  $n = 400$ .  $\sigma = 60$  s does not change: customers stay as variable as they were; only our knowledge of their average improves.

# Takeaway

A statistic is a **random variable**. Its distribution over repeated samples is the sampling distribution; its spread is the **standard error**,  $\sigma/\sqrt{n}$ . A confidence interval reports that spread — and its percentage describes the **procedure**, not the one interval in front of you.

## Speaker notes

Stabilise three things: statistic-as-random-variable, the  $1/\sqrt{n}$  rate, and the procedure reading of confidence. All three return whenever a measured number — an accuracy, an error rate — is reported with an error bar, and a Bayesian treatment gives the same question a different, more directly stated answer.



# An Unknown Number, Ten Observations

A new feature went live. Of the first **10 users**, **3 came back**.

What is  $p$ , the probability that a user comes back — and how sure are we?

## Speaker notes

Pick up the tension the confidence-interval unit left: the useful-sounding sentence was forbidden, because  $\mu$  (here  $p$ ) was not random. This unit takes the other road: if we are uncertain about  $p$ , describe that uncertainty with a distribution — exactly as we describe any other unknown.

The frequentist numbers:  $SE = \sqrt{(0.3 \cdot 0.7 / 10)} \approx 0.145$ , so  $0.3 \pm 1.96 \cdot 0.145$ . Note in passing that the normal approximation is poor at  $n = 10$ ; that is one more reason to want something else.



# An Unknown Number, Ten Observations

A new feature went live. Of the first **10 users**, **3 came back**.

What is  $p$ , the probability that a user comes back — and how sure are we?

## WHAT THE LAST UNIT CAN SAY

$\hat{p} = 0.3$ , standard error  $\approx 0.14$ , a 95% interval of roughly  $[0.02, 0.58]$  — about the *procedure*, not about  $p$ .


## WHAT WE WOULD LIKE TO SAY

" $p$  is probably between 0.1 and 0.6, most likely around 0.3" — a statement **about**  $p$ .

# Conditioning on the Data — for a Parameter

Treat the unknown  $p$  as a random variable  $\theta$  and apply **Bayes' theorem**, with  $\theta$  in place of the hypothesis and the data  $D$  as the evidence:

$$p(\theta \mid D) = \frac{p(D \mid \theta) \cdot p(\theta)}{p(D)}$$



For i.i.d. yes/no data with  $h$  yeses in  $n$  trials:  $p(D \mid \theta) = \theta^h (1 - \theta)^{n-h}$  — read as a function of  $\theta$ , with the data fixed.

The same four parts as the disease test, one level up: the hypothesis is now a number, so the prior and posterior are densities over  $[0, 1]$  instead of two probabilities. The likelihood is the Bernoulli model from the data module, multiplied over rows — the i.i.d. assumption doing its job. Say explicitly that the likelihood is *not* a distribution over  $\theta$ ; Lecture 03 builds on that distinction.

# A Prior That Behaves Like Counts

**Beta prior.**  $p(\theta) \propto \theta^{a-1}(1 - \theta)^{b-1}$  on  $[0, 1]$ , written  $\text{Beta}(a, b)$ .  $\text{Beta}(1, 1)$  is flat: every  $p$  equally plausible.

Multiply by the likelihood  $\theta^h(1 - \theta)^{n-h}$ :

$$p(\theta \mid D) \propto \theta^{a+h-1}(1 - \theta)^{b+n-h-1} \implies \theta \mid D \sim \text{Beta}(a + h, b + n - h)$$


The update is **adding counts**:  $a$  and  $b$  act like yeses and noes seen before the data. A prior whose posterior stays in the same family is called **conjugate**.

## Speaker notes

No integral needed — the exponents simply add. That is the whole reason for the Beta: it makes the posterior computable by hand. The pseudo-count reading is the one to stabilise, because it makes the next two slides predictable: a prior worth 2 observations is swamped by 1,000 real ones.



# Watching the Posterior Form



**Predict before each click:** which way does the curve move after a ✗ — and does it get narrower every time?

## Speaker notes

One observation per click, flat prior. Stop after the first ✖ and after the first ✔ and ask what the curve now says. Two things to see: each observation moves the peak toward what was observed, and the curve narrows with every observation regardless of direction — the narrowing is the epistemic uncertainty shrinking, as promised in Lecture 01. After ten: Beta(4, 8).



# What the Posterior Gives You

After 3 of 10, with a flat prior:  $\theta \mid D \sim \text{Beta}(4, 8)$ .

| question                    | answer       | from                                   |
|-----------------------------|--------------|----------------------------------------|
| most plausible $p$          | $0.30 = h/n$ | the peak (mode)                        |
| best guess on average       | 0.33         | the posterior mean $\frac{a+h}{a+b+n}$ |
| where is $p$ , 95%?         | [0.11, 0.61] | a <b>credible interval</b>             |
| after 1,000 users, 300 back | [0.27, 0.33] | same model, more data                  |

The peak under a flat prior is exactly  $h/n$ . **Lecture 03** keeps only this peak and calls it maximum likelihood — the first shortcut of the course.

Every row is read off the same distribution — Lecture 01's third move. The credible interval licenses the sentence the confidence interval forbade: "p lies in [0.11, 0.61] with probability 0.95", given the prior. The last row is the  $1/\sqrt{n}$  law again: the posterior standard deviation falls from 0.13 to 0.015.

# How Much Does the Prior Matter?

Two analysts: one with a flat prior  $\text{Beta}(1, 1)$ , one convinced that  $p$  is near a half,  $\text{Beta}(10, 10)$ .

| data         | flat prior | "near a half" prior |
|--------------|------------|---------------------|
| 3 of 10      | mean 0.33  | mean 0.43           |
| 300 of 1,000 | mean 0.30  | mean 0.30           |

With little data the prior matters; with enough data the analysts **agree**. A prior is a stated assumption — it can be argued with, and it washes out.

## Speaker notes

The usual objection to Bayesian methods is that the prior is subjective. The answer has two halves, both on the slide: the prior is written down, so it can be criticised like any other modelling choice, and its influence is worth  $a+b$  pseudo-observations — 20 here — which data overwhelms. Where data is scarce the prior matters, and then it *should*: it is the only information there is.



# Mini-Exercise

Prior  $\text{Beta}(2, 2)$ . Of 10 users, **8** come back.

Is the posterior mean above or below 0.8 — and why?

## Speaker notes

One question, answerable from the pseudo-count reading alone; the numbers are for checking. The peak is  $9/12 = 0.75$ , also below 0.8. The transferable point: a prior acts like regularisation, shrinking an estimate towards a default — exactly how Lecture 03 reads weight decay.



# Mini-Exercise

Prior  $\text{Beta}(2, 2)$ . Of 10 users, **8** come back.

Is the posterior mean above or below 0.8 — and why?

**Below.** The posterior is  $\text{Beta}(2 + 8, 2 + 2) = \text{Beta}(10, 4)$ , mean  $10/14 \approx 0.71$ . The prior counts two yeses and two noes, so it pulls towards 0.5 — and the pull fades as data grows.

# Two Ways to Answer "Where Is $p$ ?"

## FREQUENTIST

$p$  is fixed; the data is random

The confidence interval is a **procedure** that catches  $p$  in 95% of samples. About this one interval, no probability statement is allowed.

*Needs:* nothing but the sampling model.

## BAYESIAN

the data is fixed;  $p$  is uncertain

Condition on the sample and get a distribution **over**  $p$ : " $p$  lies in this interval with probability 0.95" becomes a true sentence.

*Needs:* a **prior** — what we believed about  $p$  before the data.

The axioms are the same; the question is not. The rest of the course mostly takes the second road — and Lecture 03 starts with its cheapest version: keep only the peak.

## Speaker notes

This is the place where "both interpretations yield the same mathematics" (the axioms unit) stops being the whole story: the mathematics agrees, but the two schools assign probabilities to different things, and so they answer different questions. Do not stage a war; state the trade. The frequentist interval needs no prior and makes a weaker statement; the Bayesian interval makes the statement everyone wanted and pays with a prior.

With a flat prior and plenty of data the two intervals nearly coincide numerically (compare  $[0.27, 0.33]$  with  $0.3 \pm 1.96 \cdot 0.0145$ ) — worth one sentence, because students will notice and suspect a trick.



# Takeaway

An unknown parameter is just another unknown: give it a **prior**, condition on the data, and the **posterior** says what the data justifies — as a peak, a mean and an interval. For yes/no data with a Beta prior the update is **adding counts**, and the prior's influence fades as data grows.

## Image Credits & References

Sheldon, Ross and others (2002). A first course in probability. *Pearson Education India*.