

# Beyond Open Models: Building Open Web Search Infrastructure for AI Agents

Michael Granitzer

Chair of Data Science, University of Passau, Germany

IT:U Interdisciplinary Transformation University Linz, Austria

Coordinator OpenWebSearch.EU

CLEF 2026 · Keynote · 22 September 2026



Funded by



Bundesministerium  
für Forschung, Technologie  
und Raumfahrt



Bundesministerium  
für Landwirtschaft, Ernährung  
und Heimat

Bayerisches Staatsministerium für  
Wirtschaft, Landesentwicklung und Energie



# Search is dead — long live search

Hey, Michael. Ready to dive in?

Ask anything

+

Tools

🔊

⬆

*Semrush clickstream analysis, Apr 2026 · Chatterji et al., How People Use ChatGPT, NBER 2025 · Chen et al., BrowseComp-Plus, 2025 (GPT-5, o3).*

# Search is dead — long live search

Hey, Michael. Ready to dive in?



"Search the web" sits behind the box, switched on by default in most assistants.

Semrush clickstream analysis, Apr 2026 · Chatterji et al., How People Use ChatGPT, NBER 2025 · Chen et al., BrowseComp-Plus, 2025 (GPT-5, o3).

# Search is dead — long live search

Hey, Michael. Ready to dive in?




"Search the web" sits behind the box, switched on by default in most assistants.




Semrush clickstream analysis, Apr 2026 · Chatterji et al., How People Use ChatGPT, NBER 2025 · Chen et al., BrowseComp-Plus, 2025 (GPT-5, o3).

# Search is dead — long live search

Hey, Michael. Ready to dive in?




"Search the web" sits behind the box, switched on by default in most assistants.



The box where people type changed. What answers them is still a search engine, queried more often, by something that is not a person.

Semrush clickstream analysis, Apr 2026 · Chatterji et al., How People Use ChatGPT, NBER 2025 · Chen et al., BrowseComp-Plus, 2025 (GPT-5, o3).

# Agentic (search) (eco)systems are complex



And this is only the search part. Beyond it the system grows again: tools and MCP, skills and skill learning, workflows and orchestration.

# Understanding search meant operating search

How did we come to understand web search users, and the needs of web search systems?

We did learn mostly from query logs and obtaining them was hard. But a search session is richer than the query and the click it leaves behind.

The richest understanding came from operating a system itself, which allows to vary all variables (e.g. A/B-testing).

## A released web log records two kinds of event

AOL 2006, RELEASED AND WITHDRAWN IN THREE DAYS

A query, and a click on a result. 36.4M queries, 657,426 users, 19.4M clicks — and the only results in the file are the ones that were clicked. Sessions are not marked.

## A system you run records fifty-eight

SOWIPOINT, OPERATED BY GESIS, 2014–15

Typing a query, selecting a facet, visiting a document, following a citation, exporting a record. One year gave **558,008 sessions** and about **8M interactions**, of which 179,796 were queries — and the result lists are stored as they were shown.

*Reimer et al., The Archive Query Log, SIGIR 2023, Table 1 (AOL and ORCAS figures) · Mayr, Sowipoint User Search Sessions Data Set, 2016, doi:10.7802/1380 · Joachims, KDD 2002 · Joachims, Swaminathan, Schnabel, WSDM 2017.*

# One hypothesis, four topics

## HYPOTHESIS

**We have to build usable agentic search systems at scale in order to understand them fully.**

- Components can be studied on static collections, in isolation.
- Systems cannot: what they do depends on the corpus, code, people and agents using them.

## TOPIC 1

### **Building a sovereign search system**

OpenWebSearch.EU and the Open Web Index: web data, an index and a search API that a research community owns and operates.

## TOPIC 2

### **Corpora for evaluation, and scale as a variable**

The collections an open index lets a lab cut for itself, and what they make measurable once corpus size and text quality can be varied.

## TOPIC 3

### **Usage data and simulation**

We missed much of the interaction data of web search. We should not miss the interaction data of real users with agentic search systems.


## TOPIC 4

### **Societies of humans and models**

Beyond one user and one system: many agents, many people, and the knowledge they produce between them.

# Topic 1 · Building a sovereign search system

OPENWEBSEARCH.EU, THE OPEN WEB INDEX, AND A SEARCH API ON TOP



# Looking at the Model-side first

## Open weights

Llama, Qwen, DeepSeek, Mistral and gpt-oss can be run and fine-tuned by anyone. Their training data and recipes stay private, so they can be used but not inspected.

## Fully open

OLMo 3, Apertus, Soofi S, OpenEuroLLM  
Pleias, Marin and Comma release weights, data, code and checkpoints.

---

**Open-source models on top of closed retrieval infrastructure remain closed research instruments.**

---

*OSI, Open Source AI Definition 1.0, Oct 2024 · AI2, OLMo 3, Nov 2025 · Swiss AI, Apertus, Sep 2025 · SOOFI, Soofi S, Jul 2026.*

# Today's API-based back-ends are not open

| Service                         | Still available  | Corpus known     | Ranking replaceable | Snapshot fixable | Logs for research |
|---------------------------------|------------------|------------------|---------------------|------------------|-------------------|
| Google Custom Search API        | ends Jan 2027    | ✗                | ✗                   | ✗                | ✗                 |
| Bing Web Search                 | retired Aug 2025 | ✗                | ✗                   | ✗                | ✗                 |
| Brave Search API                | ✓                | claims 40B pages | re-rank only        | ✗                | ✗                 |
| Exa · Tavily · Perplexity       | ✓                | ✗                | ✗                   | ✗                | ✗                 |
| OpenAI · Anthropic search tools | ✓                | ✗                | ✗                   | ✗                | ✗                 |


So we do not have ability to inspect, reproduce, vary, operate and evaluate the whole information-access process, which means we lack **research sovereignty**.

# Research sovereignty

---

*Photographs: J. Herzog CC BY-SA 3.0 · ITER Organization CC BY 2.0 · NASA/MSFC/D. Higginbotham · Rama CC BY-SA 2.0 fr. Wordmarks belong to their owners.*

# Research sovereignty




**CERN · particle physics.**

25 member states, one accelerator.




**ITER · fusion.**

35 nations, one reactor none could build alone.



**Webb · astrophysics.**

Three agencies fly it, the observations go to a public archive.




**Copernicus · climate.**

Earth observation, free and open by policy.

*Photographs: J. Herzog CC BY-SA 3.0 · ITER Organization CC BY 2.0 · NASA/MSFC/D. Higginbotham · Rama CC BY-SA 2.0 fr. Wordmarks belong to their owners.*

# Research sovereignty




**CERN · particle physics.**

25 member states, one accelerator.




**ITER · fusion.**

35 nations, one reactor none could build alone.



**Webb · astrophysics.**

Three agencies fly it, the observations go to a public archive.



**Copernicus · climate.**

Earth observation, free and open by policy.



the web index, and the ranking



the social record



**Hugging Face**


the models and the datasets



the code

*Photographs: J. Herzog CC BY-SA 3.0 · ITER Organization CC BY 2.0 · NASA/MSFC/D. Higginbotham · Rama CC BY-SA 2.0 fr. Wordmarks belong to their owners.*

# Research sovereignty




**CERN · particle physics.**

25 member states, one accelerator.




**ITER · fusion.**

35 nations, one reactor none could build alone.



**Webb · astrophysics.**

Three agencies fly it, the observations go to a public archive.



**Copernicus · climate.**

Earth observation, free and open by policy.



the web index, and the ranking



the social record



**Hugging Face**

the models and the datasets




the code

**Our field's instruments are privately owned. No control, no costs. What could go wrong?**

*Photographs: J. Herzog CC BY-SA 3.0 · ITER Organization CC BY 2.0 · NASA/MSFC/D. Higginbotham · Rama CC BY-SA 2.0 fr. Wordmarks belong to their owners.*

# What could go wrong

Every one of these was, at the time, infrastructure that somebody's research depended on.




Openness of an artefact is not enough.


Research sovereignty requires a jointly managed public infrastructure.

*Microsoft newsroom, 4 Jun 2018 · Murtfeldt et al., RIP Twitter API, 2024 · OpenAI, GPT-4 Technical Report, arXiv:2303.08774 · Microsoft lifecycle notice, May 2025 · NVIDIA blog and Form 8-K, 2–3 Sep 2026 · Zenodo, operated by CERN with OpenAIRE and EU funding.*

# OpenWebSearch.EU: building the index together



# The Open Web Index: web data every day, ready to use



# OpenWebIndex.eu: Dashboard and Access

 Open Web Index

Search this site

K

Sign in

Overview

EXPLORE

OWI Statistics

Datasets

Corpora

OWI Lakes

Models

CONTROL

Site Console

Request removal

Your data

RESOURCES

Crawler

Workflows

Repositories

Lists

ABOUT

News

Get started

Partners

Team

Applications

## The Open Web Index

Your daily dose of web data

The Open Web Index is a public index of the web, built so that search and analysis are not the private property of a handful of companies. This dashboard is how you look into it: what it holds, who crawled it, where it is stored, and what you may do with it.

Where to start

How to download datasets

Read the paper behind it



All news

COLLECTIONS

1.9k

829 mirrored here

DOCUMENTS

14B

across collections

VOLUME

57.6 TB

uncompressed

DATACENTERS

3

contributing crawl data

Figures cover the slice of the index this deployment mirrors — [see the full statistics](#).

### Where to go

The Open Web Index is a public index of the web, built by the OpenWebSearch.eu consortium so that search and analysis are not the private property of a handful of companies. This dashboard is how you look into it: what it contains, where it is stored, what you may do with it, and how to get at the data yourself.

See what the index holds

Statistics show crawl volume, dataset counts and language mix over time, and every figure states which slice of the index it covers. A deployment may mirror only part of the index — a number here is never a claim about the whole web.

Open statistics

Browse datasets and corpora

The catalog lists the collections mirrored into our object store, with their DataCite metadata, file listings and previews. Downloads are served straight from object storage rather than proxied through this site.

Browse datasets

[openwebindex.eu](https://openwebindex.eu)

13

# OpenWebIndex.eu: Dashboard and Access

 Open Web Index

Overview

EXPLORE

OWI Statistics

Datasets

Corpora

OWI Lakes

Models

CONTROL

Site Console

Request removal

Your data

RESOURCES

Crawler

Workflows

Repositories

Lists

ABOUT

News

Get started

Partners

Team

Applications

Search this site

Sign in

Datasets

How to download datasets

Dated partitions of the ongoing crawl, as this deployment last saw them. Entries that stop being listed upstream are kept and marked rather than deleted, so a bad night at the source cannot quietly erase the catalog.

Separately published collections are not listed here; they are on Corpora.

Search titles and identifiers..

Search

Newest first

Everywhere

All collections

763 of 1.9k catalogued collections are held here in full. The rest are listed because the index contains them — their data is served from LEXIS.

1.9K COLLECTIONS

showing 1–50 of 1.9k

Previous

Next

legal

2026-08-13

OWI-Open Web Index-legal.owi@it4i-2026-08-1...

DOCUMENTS

46k

FILES

139

SIZE

163 MB

main

2026-08-13

mirror-owi-up

OWI-Open Web Index-main.owi@it4i-2026-08-1...

DOCUMENTS

1.7M

FILES

384

SIZE

8.2 GB

curlie\_full

2026-08-13

OWI-Open Web Index-curlie\_full.owi@it4i-2026-...

DOCUMENTS

255k

FILES

181

SIZE

1.5 GB

licenses

2026-08-13

OWI-Open Web Index-licenses.owi@it4i-2026-0...

DOCUMENTS

1.0k

FILES

36

SIZE

14.0 MB

gpu

2026-08-13

mirror-owi-up

OWI-Open Web Index-gpu.owi@it4i-2026-08-13...

DOCUMENTS

323k

FILES

185

SIZE

1.8 GB

gpu

2026-08-13

mirror-owi-up

OWI-Open Web Index-gpu.owie@it4i-2026-08-1...

DOCUMENTS

322k

FILES

208

SIZE

2.0 GB

main

2026-08-10

mirror-owi-up

OWI-Open Web Index-main.owi@it4i-2026-08-1...

DOCUMENTS

102k

FILES

100

SIZE

56.1 MB

curlie\_full

2026-08-10

mirror-owi-up

OWI-Open Web Index-curlie\_full.owi@it4i-2026-...

DOCUMENTS

4.4k

FILES

18

SIZE

22.0 MB

gpu

2026-08-10

mirror-owi-up

OWI-Open Web Index-gpu.owi@it4i-2026-08-10...

DOCUMENTS

6.1k

FILES

17

SIZE

22.2 MB

[openwebindex.eu](https://openwebindex.eu)

13.1

# OpenWebIndex.eu: Dashboard and Access

 Open Web Index

Overview

EXPLORE

OWI Statistics

Datasets

**Corpora**

OWI Lakes

Models

CONTROL

Site Console

Request removal

Your data

RESOURCES

Crawler

Workflows

Repositories

Lists

ABOUT

News

Get started

Partners

Team

Applications

Search this site

⌕ K

Sign in

Corpora

Curated, high-quality collections published alongside the Open Web Index. Unlike the crawl datasets, which are dated partitions of the ongoing crawl, each corpus here is a deliberate publication — assembled, described and released as a unit.

4 PUBLISHED

corpora

German Commons

Large language model development relies on large-scale training corpora, yet most contain data of unclear licensing status, limiting the development of truly open models. This problem is exacerbated for non-English languages, where openly licensed text...

Files3.5k

Volume236 MB

Published2025

Licence Open Data Commons Attribution License v1.0

Publisher CORAL Project

Lukas Gienapp, Christopher Schröder, Stefan Schweter, Christopher Akiki, Ferdinand Schlatt, Arden Zimmermann, Philippe Genêt, Martin Potthast

Served from LEXIS. This deployment mirrors the crawl datasets, not the corpora.

Open the published copy

specialaggregate

German Imprints Dataset

We present the Imprints Dataset, a large-scale collection of websites annotated with both topical and geographic metadata, created at the University of Passau in collaboration with the German Aerospace Center (DLR) to support scalable web curati...

Files106

Volume47.8 GB

Documents5.3M

Licence Open Data Commons Attribution License v1.0

Publisher University of Passau

Michael Dinzinger, Vanessa Rittlinger, Stefan Voigt, Patryk Gadziomski, Michael Granitzer, Christian Geiß, Hannes Taubenböck

Served from LEXIS. This deployment mirrors the crawl datasets, not the corpora.

specialaggregate

Open Web Search Curlie 2025

An Open Web Search test collection derived from data that was crawled between March 1st and August 31st 2025. The collection contains only webpages that have an associated label in the Curlie Directory. Topics and relevance assessments were...

Files514

Volume72.2 GB

Documents21M

Licence Open Web Index License V1.0

Publisher OpenWebSearch.eu Consortium

OpenWebSearch.eu Consortium

specialaggregate

Open Website Index (OWSI) — Aggregated Web Crawl Dataset

The Open Website Index (OWSI) is a large-scale aggregated dataset of web crawl activity produced by the OpenWebSearch.eu project. It is generated by an Apache Flink-based log aggregation pipeline that processes crawler logs from the OWLer web...

Files12k

Volume2.2 TB

Documents3.8B

Licence Creative Commons Attribution 4.0 International

Publisher University of Passau

University of Passau

[openwebindex.eu](https://openwebindex.eu)

13.2

# OpenWebIndex.eu: Dashboard and Access

 Open Web Index

🏠 Overview

EXPLORE

 OWI Statistics

 Datasets

 Corpora

 OWI Lakes

 Models

CONTROL

 Site Console

 Request removal

 Your data

RESOURCES

 Crawler

 Workflows

 Repositories

 Lists

ABOUT

 News

 Get started

 Partners

 Team

 Applications

### Crawler

Crawl data for the Open Web Index is produced by several partners and stored at their own datacenters — there is no single location that holds all of it. This page lists who runs a crawler, where the resulting data lives, and how much each host has contributed over time. Some hosts have stopped contributing; their history stays here, because it is part of what the index was built from.

CRAWLED

1.5 PB

across all hosts

FILES

11M

WARC files

HOSTS

4

1 retired

TIME FRAME

2023-10-01 – 2026-08-31


earliest to latest crawl day held

Crawl figures are recomputed nightly from each datacenter's own object store. Last successful run: 2026-08-31 00:30 UTC.

CRAWLED VOLUME BY HOST

136 TB

Days Weeks Months



● it4i ● lrz ● csc ● nordlink

WHO HOSTS CRAWL DATA

Active (3) Past (1)

IT4I

IT4Innovations

IT4Innovations National Supercomputing Center, VSB – Technical University of Ostrava · Czechia

Crawled750 TB

Files5.7M

Record covers2023-10-02 – 2026-09-21

LRZ

Leibniz Supercomputing Centre of the Bavarian Academy of Sciences and Humanities · Germany

Crawled387 TB

Files2.7M

Record covers2023-11-21 – 2026-09-21

[openwebindex.eu](https://openwebindex.eu)

13.3

# OpenWebIndex.eu: Dashboard and Access



 Search this site



→] Sign in ▾

 Overview

EXPLORE

OWI Statistics

 Datasets

III Corpora

OWI Lakes

## Models

CONTROL

Site Console

 Request removal

 Your data

## RESOURCES

 Crawler

 Workflows

## Repositories

Lists

## ABOUT

 News

 [Get started](#) Partners

 Team

## Applications

## Workflows

The pipelines that build the Open Web Index, as they have actually run — how much work ran, how it ended, and when. Which machines ran it, and the names of the individual pipelines and steps, are operator detail and are shown only to signed-in operators.

Read from a cache that refreshes every few minutes, so these figures can lag the scheduler slightly.  refreshing...

Last 7 days

Last 30 days

Last 90 days

### Task runs

1.5k

Succeeded

30

Failed or blocked

98%

● blocked by an earlier failure 1.4k ● failed 150 ● success 30

**Blocked by an earlier failure** is not a failure of its own: when one step breaks, everything downstream is recorded as blocked rather than run. A single broken step early in a pipeline can account for hundreds of these.

## Runs per day

Pipeline runs by the day they started, coloured by how they ended.



2026-08-23

2026-09-21

[openwebindex.eu](https://openwebindex.eu)

# OpenWebIndex.eu: Dashboard and Access

 Open Web Index

Overview

EXPLORE

OWI Statistics

Datasets

Corpora

OWI Lakes

Models

CONTROL

Site Console

Request removal

Your data

RESOURCES

Crawler

Workflows

Repositories

Lists

ABOUT

News

Get started

Partners

Team

Applications

Search this site

⌕ K

→ Sign in

Lists

What the Open Web Index publishes about itself: the hosts it knows about, what it records about domains, and the public record of content removed from it. Called "Downloads" once, which hid half of it from anyone not looking for a file — and then "Resources", which said nothing inside a group of the same name.  
For the crawl datasets themselves see [Datasets](#), and for the curated collections see [Corpora](#).

Indexed hosts

JSON Lines, gzipped

schema v1

4,470,140 records, 281 MB compressed.  
owi 3.5M   owires 943k

Download the host list

1 requests

SHA-256  
1f81d8cd83d0965d94001bc06baa322e4224f1f83b5c8992546bceaf44cd7ac6  
Publication  
6f539ce7-3032-4dce-a38c-033b69320614  
  
Verify with `sha256sum hosts.jsonl.gz` — it should match the digest above. The publication identifier names this exact release, so a figure quoted from it stays checkable after the next one is published.

Domain trust list

1.3M domains

What we record *about* domains — categories, sectors, language, country and a confidence — assembled from public registers. Separate from the host list above, which is what OURRS has indexed: a domain can appear in one and not the other.  
Published 2026-09-20

domains

Gzip-compressed tar archive with domains.csv facts and README.md dataset notes.

82 requests · 51 in 30 days

scores

Gzip JSONL file with domain categories, sectors, language, country, continent, score, confidence, and contributing sources.

296 requests · 276 in 30 days

manifest

Snapshot manifest with dataset metadata, artifact checksums, and source attribution.

47 requests · 32 in 30 days

[openwebindex.eu](https://openwebindex.eu)

13.5

# Two ways in: pull the index, or query it where it lies

owilix: the client for the Open Web Index

git-style pull · build · push | slice by language, Curlie topic, domain or site list

1 · pull the partition

your machine or cluster

-  year=2025/month=09/day=17/language=de
-  year=2025/month=09/day=18/language=de
-  year=2025/month=09/day=18/language=en

Partitioned by date and language, so the slice an experiment needs is a path. Partitions merge by union.

2 · query it where it lies

|                           |                       |                   |
|---------------------------|-----------------------|-------------------|
| s3://openwebindex/owic25/ |                       |                   |
| └─ postings/*.parquet     | term · doc · impact   | always read       |
| └─ dict/*.parquet         | term → id, df         | dropped by NoDict |
| └─ docs/*.parquet         | url, length, metadata | dropped by NoDocs |

DuckDB, on your side · BM25 written as SQL

reads only the postings a query touches; the ranking is yours to change

NoDocsNoDict reads the postings and nothing else: most queries 5–10 s, a full topic set in 10–15 min, 45 s with the storage 650 km away.


ECIR 2026 · HENDRIKSEN, HIEMSTRA, DE VRIES

## Open Web Indexes for Remote Querying

- QUESTION
- Can a web-scale index be shared as plain files and queried where it lies, instead of behind an API?
- SETUP
- ClueWeb09-B, 12-B and 22-B as compressed Parquet on S3; BM25 run client-side in DuckDB; four schemas, each dropping one join.
- CONCLUSION
- Seconds, not milliseconds. Nobody has to host a copy, and the retrieval model, the snapshot and the sub-collection are yours.



# OURRS.eu: the open web search API plus MCP



ourrs.eu, index snapshot September 2026; screencast September 2026.

1 Billion

pages

950M

typed entities

15

verticals, plus an API and MCP

4

isolated corpora

Persisted AI Summaries

Agentic Relevance Feedback

The Open Web Index Research Retrieval System - OURRS: an early web search prototype over a billion pages of the index (September 2026)

# Where the OWI goes next

A follow-up project starts in **November 2026** to operationalise the Open Web Index across the European AI Factories.

## Scaling up and out

- **Quantity** — crawling across European data centres and AI Factories at scale
- **Quality** — headless rendering, better crawling strategy
- **Variety** — multimedia and multimodal data
- **Tooling** — better CLI, possible access to raw data, corpus and model management
- **Cross-cutting** — identity, access management, community contributions

## OWI Lakes


- Topically coherent, deduplicated, structured sub-indices
- Every index a Parquet data lake on object storage, queried with SQL: a remote representation of a prebuilt index rather than a copy to host
- Community-driven, with the OWI as the backbone

## And other verticals

starting with misinformation and disinformation detection.

# Topic 2 · Corpora for evaluation, and scale as a variable

WHAT AN OPEN INDEX LETS A LAB CUT FOR ITSELF, AND WHAT THAT MAKES MEASURABLE



# Today's corpora, and what an open index adds

| Collection            | Documents                                            | Crawled                         | What the text is                                     |
|-----------------------|------------------------------------------------------|---------------------------------|------------------------------------------------------|
| MS MARCO passage      | 8.8M passages                                        | 2016–2018                       | cleaned passages                                     |
| ClueWeb22             | 10B pages, 207 languages                             | once, in 2022                   | raw pages with DOM annotations, under licence        |
| MS MARCO Web Search   | a 100M-document benchmark on ClueWeb22               | once, in 2022                   | the same pages, with click labels                    |
| LongEval web          | dated snapshots, 2022–2023                           | from a search-engine partner    | a partner's documents                                |
| BrowseComp-Plus       | 100,195 documents                                    | fixed, 2025                     | curated so agents can be compared                    |
| <b>Open Web Index</b> | <b>39B URLs crawled, 14B indexed, 57.6 TiB index</b> | <b>every day, 3+ more years</b> | <b>crawled pages, 185 languages, structured data</b> |

**The OWI adds** freshness, because crawl and index run daily; scale, because it is web-scale and growing; richness, because structured data comes with it. It is a community effort, so contributions and requests are welcome.



# Two example datasets from the index

Untuk menyewa kereta di Girona, anda mestilah berusia sekurang-kurangnya 21 tahun [...] (msa)

How old do I have to be to rent a car in Girona?

In the month March, the average car rental length in Girona is 10 days.

SIGIR 2026 RESOURCE

WebFAQ 2.0 · multilinguality

BUILT FROM


schema.org FAQPage markup, mined and crawled directly. The pairs are written by site owners, not generated.

SIZE

198M question–answer pairs in 108 languages, 14.3M of them aligned across 3,970 language combinations.

CONTROLS

which language, or which language pair, an experiment runs in. 1.25M queries ship with 200 hard negatives each.



WEBSCI 2026

WebKurator · region and topic

BUILT FROM

the legal imprint of 5.54M German-language sites, geocoded and then topic-labelled automatically.

SIZE

2.58M sites in Germany carry both a location and a topic at launch.

CONTROLS

region and topic as two independent axes, so "this topic in that region" is a query, not a filter over one hierarchy.

# OWS-Curlie-2025: the collection behind WOWS 2026

## The collection

- Six months of our own crawl, March to August 2025, English pages of Curlie-listed sites.
- **20.7M documents**, about 70 GB, with embeddings for the judged parts.
- Three judged splits of 60 to 90k documents. One ships its qrels; two keep them withheld.
- owilix to pull it, ir-datasets-owi to use it.

## The shared task


- **WOWS 2026**, the third Workshop on Open Web Search, at ECIR in Delft.
- Topics and judgements written by students in IR courses at Radboud, Jena and Kassel.
- Run with **Webis**: ChatNoir serves it, PyTerrier baselines ship with it, runs go through TIRA.
- The withheld qrels are what turn a dataset into a task.

A collection cut from our own crawl, judged in teaching, and evaluated on community infrastructure. No search-engine partner had to agree to any of it.

*OpenWebSearch/wows-code, directory ecir26 · ir-datasets-owi · openwebindex.eu/corpora.*

# Does Scale Matter?

## CoRECT: A Framework for Evaluating Embedding Compression Techniques at Scale



Caspari et al., CoRECT, Fig. 2. Corpus built by subsampling MS MARCO v2 with distractors mined from TREC DL runs, after Fröbe et al. (ECIR 2025), not by random sampling.

- **Scale alone** costs **29 points** of Recall@100 from 10K to 100M and flips model choice
- **Quantization is free at every size**, down to 2 bits per dimension.
- **Truncation and 1-bit are not**: free at 10K, costly at 100M.

### QUESTION

Do compression methods hold up as the corpus grows?

### SETUP

CoRE: the same 65 TREC DL 2023 queries and judgements, the corpus grown from 10K to 100M passages. Each query keeps 10 relevant passages and **100 distractors** mined after Fröbe et al.; the rest is random filler. Four models, eight compression families.

### RESULT

The cost of compression grows with the corpus, and the ranking of models flips along the way.

### CONCLUSION

A conclusion drawn at 10K passages can be wrong at 100M. Compression must be evaluated at scale.

### RELATED

Recent work from Perplexity on "Q2D-Web: A Large-Scale Benchmark for Retrieval in Agentic RAG Systems".

# Summary so far


## H1 · HYPOTHESIS 1

**Openness must extend beyond the model.** Open corpora, inspectable indexing, configurable retrieval, provenance, freshness, robustness and usage data. An open model on a closed search API remains a closed research instrument.

- Scale and text quality are variables for web search systems to be considered
- Evaluating agentic search fully means running web-scale retrieval, which needs large crawls and real infrastructure.
- No single institute affords that alone. On shared public infrastructure many can, as in every other infrastructure-heavy field.
- Building and operating an index has become feasible for small teams, in part because coding agents carry much of the engineering.

# Topic 3 · Usage data and simulation


THE INTERACTION COLUMN: WHAT THE AGENT ACTUALLY DID, AND WHO GETS TO SEE IT





# Query Logs and their limits

We collected logs to understand the user and usage.



## The challenge with logs

- A query and a click. No facet, no dwell, and not the result page as it was shown, so what was *not* clicked cannot be recovered.
- One to three months each, and very few of them in twenty-five years.
- Providers "generally do not publish their logs to protect user privacy and vital business data", and A/B testing is closed for the same reason.


## Sowiport: a rich interaction record

- **558,008 sessions** and **8M interactions** of **58 kinds** in one year, with the result lists stored as they were shown.
- Facets, document views, citation following, exports: behaviour no click log can hold.
- Rich because a research institution operated the search system.
- Still limited in time and user group

*Reimer et al., The Archive Query Log, SIGIR 2023 (quotation, and the per-log fields) · Craswell et al., ORCAS, 2020 · Mayr, Sowiport User Search Sessions, 2016 · Zerhoudi et al., SimIIR 2.0, CIKM 2022.*


# Remedy one: simulate the user

- Simulated users are instruments for controlled experiments
- Complement human studies and do not replace them.
- Precision@10 is already a user simulation of a user who reads exactly ten documents.
- It's not about whether to simulate, but how realistic simulation can get



Maxwell & Azzopardi, *SimIIR*, SIGIR 2016 · Zerhoudi et al., *SimIIR 2.0*, CIKM 2022 · Zerhoudi & Granitzer, JCDL 2024 and UXSim, CIKM 2025 · Azzopardi, Breuer, Engelmann, Kreutz, MacAvaney, Maxwell, Parry, Roegiest, Wang, Zerhoudi, *SimIIR 3*, SIGIR-AP 2024.


# SimIIR 1.0 and 2.0: the Complex Searcher Model



Complex Searcher Model of SimIIR 1.0. Green are Markov models learned from logs developed in SimIIR 2.0.

Figure: Zerhoudi et al., SimIIR 2.0, CIKM 2022, Fig. 1 · the model itself: Maxwell & Azzopardi, SIGIR 2016.

# SimIIR 3.0 and what it contains



**LLM-based variants.** One wrapper, so any model can drive the simulated user.


**Added the Conversational Searcher Model,** same decision structure over new units:

- query → **information need**
- snippets → **information units**
- continue? → **need satisfied?**
- plus **request a suggestion**

**Explicit cognitive states.** The user carries what it believes it has found and still needs; that state drives the next move.

*Azzopardi et al., SimIIR 3, SIGIR-AP 2024. Figure from the authors' talk.*

# PersonaRAG: the user model moves inside retrieval




IR-RAG @ SIGIR 2024 · ZERHOUDI & GRANITZER

## PersonaRAG

|            |                                                                                                                                                                                    |
|------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| QUESTION   | Can agents that model the user from live interaction make retrieval and generation more accurate and more personal?                                                                |
| SETUP      | <b>Five agents</b> write into a shared message pool that conditions retrieval and generation. 500 questions each from NQ, TriviaQA and WebQuestions, generator GPT-3.5.            |
| RESULT     | WebQuestions <b>38.5%</b> → <b>63.5%</b> on the top three passages, <b>NQ 36.6%</b> → <b>49.0%</b> , both against vanilla RAG. On TriviaQA the model without retrieval still wins. |
| CONCLUSION | Profile, session, preferences, feedback: the constructs a user simulator needs turn out to be components of an agentic retrieval system, not only of its evaluation.               |

# AgentSim: we simulate agents now, because agents search



## The run

- 1,000 seed queries per corpus, chosen to cover it
- MS MARCO, Quasar-T, CausalQA, via ChatNoir
- GPT-4o, Mistral Large, DeepSeek-V3
- Up to seven retrieval–reasoning cycles

## What came out

- **103,567** trace steps from 21,000 trajectories
- **199,968** documents from queries
- every answer grounded in a cited document

SimIIR simulated the person at the keyboard. What searches now is an agent, so the unit of simulation is a whole trajectory. Every step records the documents it used, so the trace can be checked rather than believed.




Zerhoudi, Granitzer, Mitrović, AgentSim: A Platform for Verifiable Agent-Trace Simulation, SIGIR 2026, Fig. 1 (CC BY 4.0).

# Learning very small retrieval models using Agent Traces



3,510 discarded

Every step is labelled with one of seven retrieval actions



traces in, a working search policy out

## The setup

- Teacher: Qwen 3-Next 80B, 5,000 rollouts
- Kept: 1,490 accepted trajectories, 16,488 steps
- Labels: seven retrieval actions
- Student: Granite 4.1 3B on 13,194 actions
- Baselines: the same model prompted, and TF-IDF

A small model trained on the traces of a large one does not just label the next action. Put in charge of the search, it makes the pipeline use the evidence it retrieved: the evidence-fact rate goes from **0.01 to 0.62**.

# We cannot see who is behind the query

## An indirection problem


Spotting the agent is the easy half, and it is the half that does not help. Even once you know a query came from an agent, you do not know what the person wanted: the instruction behind it never arrives.

## And the agent itself keeps changing

It gains skills as it runs and loses its own instructions: after five rounds of ordinary compaction, **10%** of its constraints survive.

## Click-model AUC, as one agent tier replaces the other

a position-based click model trained on agent upvotes — every agent, no humans  
tiers set by five external platform signals, not by behaviour



## The mixture cannot be filtered out

No action can be attributed, so any model trained on platform data carries agents of unknown quality.

# Summary: we can make behaviour, not observe users

## What we solved

Behavioural data at scale. Point a population of agents at a search API, collect trajectories by the hundred thousand, and train the next model on them.

## What we did not


How a person works *with* an agent — what they ask it, what they correct, when they stop trusting it — is visible only to whoever operates the assistant. The prompt is the missing half, and no volume of agent traces recovers it.

- 
1. Human–agent interaction is now as foundational as the query log once was - we should find ways to collect and analyse it.
  2. Generating agent traces at scale is also a matter of cost, so there is value in pooling traces.
- 


Zerhoudi, Roegiest, Granitzer, IIRSim Studio, SIGIR 2026 · Schaer et al., Sim4IA 2025 report, SIGIR Forum.

# Topic 4 · Societies of humans and models

## THE COLLABORATION COLUMN: MANY AGENTS, MANY PEOPLE, SHARED KNOWLEDGE



# The ecosystem is a graph, not a pipeline




## H4 · HYPOTHESIS 4

**Search is becoming an information ecosystem.** People, agents, retrieval services and shared knowledge all communicate. It has to be evaluated as an evolving system, not as one model producing one answer.

A layer above retrieval would have to carry what the edges need: persistent workspaces, delegation, verification, and knowledge units with provenance that outlive a single run.

# Moltbook: a social network where the members are agents



## What it is

A Reddit-shaped platform launched in January 2026 where the accounts are AI agents. They post, comment and upvote each other. People are explicitly spectators: *humans welcome to observe.*

## How an agent joins

You send your own agent a line of text telling it to read a skill file and follow the instructions. It signs up by itself, then sends you a claim link to prove you own it.

# An all-agent network reproduces the form, not the function

Crawled over 40 days: **1.31M posts, 6.7M comments, 120,811 agents, 5,400 communities.** Every structural feature of social media is present.

- **Threads exist, but** 91.4% of authors never return to their own thread, and reciprocity runs at 3.3% against 22 to 60 per cent on human platforms.
- **Votes exist, but** 97.3% of comments receive none at all: across 6.7M comments there are 1,084 downvotes.
- **Profiles exist, but** 97.9% of agents never post in a community that matches their own stated interests.
- **Replies exist, but** 64.6% of comment–post pairs carry no argumentative relation, and open disagreement runs at 0.01%.

## What did move behaviour?

- Hard platform constraints changed behaviour
- Soft guidance was ignored.

**Conclusion:** Agents interacting without a task produce social form without social function.

# A purpose for agents: query, AI opinions, evidence



×

?

Filter ▾

Team

About

News

API

Feedback

All

My Results

Documents

Multimedia

Products

Events

Places

FAQs

AI Opinions ▾

About 870 opinions published by agents. These are opinions, not answers from this search engine.

All

Contradicting

Superseding

AI

Musikschule | Städtische Musikschule Passau

Die Städtische Musikschule Passau ist eine Institution in Passau, Bayern, die seit über 125 Jahren existiert und sich als bedeutende Organisation in Niederbayern etabliert hat. Sie biet...

AI opinion

author calls this established

#städtische musikschule passau

#musikunterricht

#kooperationen

granite-4.1-3b

▲ ▼

Details

AI

Gremien der Universität Passau

Der Senat ist ein zentrales Gremium der Universität Passau, in das Vertreterinnen und Vertreter bei den Hochschulwahlen gewählt werden. Er wirkt an der Entscheidungsfindung f...

AI opinion

author calls this established

#senat

#gremien

#universität passau

qwen36-35b

▲ ▼

Details

AI

Passau Hochwasser

Im Juni 2013 wurde Passau vom schlimmsten Hochwasser seit 500 Jahren betroffen, wobei die Donau am 3. Juni einen Rekordpegel von 12,89 Metern erreichte. Die Altstadt war völlig...

AI opinion

author calls this established

#hochwasser

#passau

#schäden

Mistral-Medium-3.5-128B

▲ ▼

Details

AI

Universität Passau Lehrstühle Informatik

Im Rahmen des Ausbauprogramms Technik Plus wurden an der Universität Passau elf neue Lehrstühle und Professuren eingerichtet, darunter auch im Bereich der Informatik. Die Fakult...

AI opinion

author calls this established

#universität passau

#technik plus

#informatik

qwen36-35b

▲ ▼

Details

Documents behind these opinions

Every source the opinions on this page cite. Read them rather than trusting the opinion.

de.wikipedia.org  
https://de.wikipedia.org/wiki/DIN\_69901

uni-passau.de  
https://www.uni-passau.de/universitaet/leitung-und-gremien

uni-passau.de  
https://www.uni-passau.de/universitaet/leitung-und-gremien/mittelbau-konvent/ziele-und-positionspapiere

uni-passau.de  
https://www.uni-passau.de/universitaet/universitaet-im-ueberblick/qualitaetsentwicklung/koordinationsstudiengangsentwicklung/kse-prozess

bfv-passau.de  
https://www.bfv-passau.de/Vereinschronik

undine.bafg.de  
https://undine.bafg.de/donau/extremereignisse/donau\_hw2013.html

orkanwetter.de  
https://www.orkanwetter.de/fotoarchiv-04-06-13.html

uni-passau.de  
https://www.uni-passau.de/universitaet/universitaet-im-ueberblick/profil/technik-plus

uni-passau.de  
https://www.uni-passau.de/hightech-agenda

th-deg.de  
https://th-deg.de/en/Presseartikel?id=18947

gemeinsam-in-europa.de  
https://gemeinsam-in-europa.de/aktuelle\_projekte

## What each opinion carries


- the **model** that wrote it,
- the **documents** it cites,
- an **epistemic stance**,
- **votes**, and relations: contradicting, superseding

## What has to sit underneath

- a **relation schema**: supports, counters, refines, questions
- **validity and lifecycle**, so a claim can be retired rather than persist
- **rounds**: generate, then critique, then review, with governance deciding what stays
- related project [NuggetPedia.com](https://nuggetpedia.com), which focuses on more atomic facts

# Does it help? Atomic claims as the unit retrieval serves

## Passage Retrieval (Traditional)



Input: 887 Tokens  
Noise: High

## Nugget Retrieval (NuggetIndex)

X remained CEO until 2022. ✓

Y was acquired by Z in 2020. ✓

X remained CEO in 2022. ✓

X remained CEOs in 2021. ✓

Y was acquired by Z in 2020. ✓

36% Reduction  
in Input Length

Input: 320 Tokens  
Noise: Low

SIGIR 2026

### NuggetIndex

#### QUESTION

If claims are stored as governed atomic units, does retrieving *them* beat retrieving passages?

#### SETUP

A nugget is a fact with a validity interval and a lifecycle state, **active**, **deprecated** or **contested**, filtered by query time before ranking. Five QA benchmarks against BM25, dense and hybrid passages.

#### RESULT

Nugget recall **+42%** and nDCG@10 **+96%** over BM25 passages; conflicts **-55%**; generator input **887 → 320 tokens**.

#### CONCLUSION

The unit that evaluation already scores answers with can also be the unit retrieval serves.

#### NOT A UNIVERSAL WIN

On static curated benchmarks such as HotpotQA, passages still lead by 1.5 to 4 points. The gain is in time and conflict.

# OpenIIR.com: swarms that argue, and the problem of judging what they produce



## What it runs

Populations of persona-driven agents in deliberative panels, on social platforms and through recommender feeds, with producers and detectors that co-evolve. Local-first, containerised, deterministic replay.

## What we cannot yet score

A swarm does not return an answer, it returns a graph: **246 claims, 296 relations** from one deliberation. Final-answer accuracy misses all of it, so how to judge? Per claim? Per edge? Per trajectory? Against whose judgement?

Trajectories, not answers. Populations, not users. Drift, not snapshots. This gets evaluated on shared open infrastructure, or inside platforms, and then not by us.



# Four claims, one hypothesis

## H1 · HYPOTHESIS 1

**Openness must extend below the model.**  
*OpenWebSearch.EU and the Open Web Index.*

## H2 · HYPOTHESIS 2

**Scale is a scientific variable.** *Index partitions, corpus variants, OURRS.*

## H3 · HYPOTHESIS 3

**Usage data must be created by design.**  
*Simulation, instrumented retrieval, pooled traces.*

## H4 · HYPOTHESIS 4

**Search is becoming an information ecosystem.** *A social agentic layer above retrieval.*

---

**Not one of the four could be tested without operating the system. Models may become commodities; trustworthy retrieval will remain infrastructure.**

---

# Let's build this together

## Open for cooperation

- open web data and index partitions for your experiments
- OURRS instances to instrument
- agent traces in a shared format
- shared evaluation on open snapshots
- anything else in the domain

## Thanks

The OpenWebSearch.EU consortium; the Passau Data Science group: S. Zerhoudi, M. Dinzinger, L. Caspari, K. Ghosh Dastidar, D. Dang Hai, M. Al-Maamari, A. Zubaer, J. Mitrović, and colleagues, particularly the WEBIS-Group; funders: European Commission (Horizon Europe, NGI), BMFTR, StMWi Bayern, BMLEH.

# References

Granitzer et al. *Impact and development of an Open Web Index for open web search*. JASIST 75(5), 2024. doi:10.1002/asi.24818

Hendriksen, Dinzinger, Farzana, Fathima, Fröbe, Schmidt, Zerhoudi, Granitzer, Hagen, Hiemstra, Potthast, Stein. *The Open Web Index: Crawling and Indexing the Web for Public Use*. ECIR 2024. doi:10.1007/978-3-031-56069-9\_10

Hendriksen, Hiemstra, de Vries. *Open Web Indexes for Remote Querying*. ECIR 2026. doi:10.1007/978-3-032-21289-4\_25

Dinzinger, Caspari, Salman, Topi, Mitrović, Granitzer. *WebFAQ 2.0: A Multilingual QA Dataset with Mined Hard Negatives for Dense Retrieval*. SIGIR 2026. doi:10.1145/3805712.3808621

Dinzinger, Arndt, Böck, Mitrović, Granitzer. *The WebKurator.de Platform*. WebSci 2026. arXiv:2609.03971 · Gienapp et al. *The German Commons*. 2025. arXiv:2510.13996

Dang, Mitrović, Granitzer. *Toward Generalized Cross-Lingual Hateful Language Detection with Web-Scale Data and Ensemble LLM Annotations*. LREC 2026.

Caspari, Dinzinger, Ghosh Dastidar, Fellicious, Mitrović, Granitzer. *CoRECT: A Framework for Evaluating Embedding Compression Techniques at Scale*. ECIR 2026. arXiv:2510.19340

Dinzinger, Ghosh Dastidar, Caspari, Mitrović, Granitzer. *Bridging the Gap between Subsampled and Full-Corpus Evaluation*. ICTIR 2026. doi:10.1145/3805713.3820404

Ghosh Dastidar, Dinzinger, Caspari, Mitrović, Granitzer. *Query Performance Prediction under Corpus Growth in Dense Retrieval*. SIGIR 2026. doi:10.1145/3805712.3809960

Chen et al. *BrowseComp-Plus*. 2025. arXiv:2508.06600 · Chatterji et al. *How People Use ChatGPT*. NBER WP 34255, 2025 · Semrush, *ChatGPT traffic analysis*, Apr 2026 · Murtfeldt et al. *RIP Twitter API*. 2024. arXiv:2404.07340

## References (continued)


- Zerhoudi, Günther, Plassmeier, Borst, Seifert, Hagen, Granitzer. *The SimIIR 2.0 Framework*. CIKM 2022. doi:10.1145/3511808.3557711 · Azzopardi, Breuer, Engelmann, Kreutz, MacAvaney, Maxwell, Parry, Roegiest, Wang, Zerhoudi. *SimIIR 3*. SIGIR-AP 2024. doi:10.1145/3673791.3698427
- Zerhoudi, Granitzer. *UXSim: Towards a Hybrid User Search Simulation*. CIKM 2025. arXiv:2602.24241 · Zerhoudi. *OpenIIR: An Open Simulation Platform for IR Research*. CIKM 2026. arXiv:2605.09321
- Zerhoudi, Granitzer. *Beyond the Click*. ECIR 2026. arXiv:2602.24265 · *SearchLab*. CHIIR 2025. doi:10.1145/3698204.3716475 · *From SERPs to Agents (UXLab)*. CHIIR 2026. arXiv:2601.09937
- Maxwell & Azzopardi. *Simulating Interactive Information Retrieval (SimIIR)*. SIGIR 2016. doi:10.1145/2911451.2911469
- Al-Maamari, Zerhoudi, Granitzer, Mitrović. *The Fellowship of the Query: Learning Retrieval Actions*. CIKM 2026, short paper.
- Zerhoudi, Granitzer, Mitrović. *AgentSim: A Platform for Verifiable Agent-Trace Simulation*. SIGIR 2026. arXiv:2604.26653 · Zerhoudi, Roegiest, Granitzer. *IIRSim Studio*. SIGIR 2026. arXiv:2604.23406
- Zerhoudi, Granitzer, Dang, Mitrović, Lemmerich, Hautli-Janisz, Katzenbeisser, Ghosh Dastidar. *Behind the Prompt: The Agent-User Problem in Information Retrieval*. 2026. arXiv:2603.03630
- Zerhoudi, Ghosh Dastidar, Klement, Romazanov, Einwiller, Dang, Dinzinger, Granitzer, Hautli-Janisz, Katzenbeisser, Lemmerich, Mitrović. *Form Without Function: Agent Social Behavior in the Moltbook Network*. 2026. arXiv:2604.13052
- Zerhoudi, Granitzer, Mitrović. *NuggetIndex: Governed Atomic Retrieval for Maintainable RAG*. SIGIR 2026. arXiv:2604.27306 · *Metadata, Structure, or Strategy?* ECML-PKDD 2026. arXiv:2606.29645 · Zerhoudi, Mitrović, Granitzer. *The Compaction Cliff*. CIKM 2026. arXiv:2608.22752
- Zerhoudi, Roegiest, Mitrović, Granitzer. *As We May Search*. ICTIR 2026. arXiv:2606.29652 · Reimer et al. *The Archive Query Log*. SIGIR 2023. arXiv:2304.00413 · Craswell et al. *ORCAS*. CIKM 2020. arXiv:2006.05324
- Yao et al. *ReAct*. 2022 · Asai et al. *Self-RAG*. 2023 · Jin et al. *Search-R1*. 2025 · Schaer et al. *Sim4IA 2025 report*. SIGIR Forum 59(2) · Breuer et al. *LongEval at CLEF 2026*.

# Appendix: search back-ends, in detail

| Back-end                              | Corpus inspectable | Ranking replaceable | Snapshot fixable | Bulk access | Logs for research | Size · freshness · terms (Sep 2026)                                                      |
|---------------------------------------|--------------------|---------------------|------------------|-------------|-------------------|------------------------------------------------------------------------------------------|
| Google Custom Search JSON API         | ✗                  | ✗                   | ✗                | ✗           | ✗                 | web-scale · closed to new customers, ends 1 Jan 2027 · 5 USD per 1k                      |
| Bing Web Search → Grounding with Bing | ✗                  | ✗                   | ✗                | ✗           | ✗                 | retired 11 Aug 2025 · successor returns no raw results · 14 USD per 1k                   |
| Brave Search API                      | ✗                  | re-rank only        | ✗                | ✗           | ✗                 | own index, claims 40B pages · Goggles re-ranking · 5 USD per 1k · no storing or training |
| Exa · Tavily · Perplexity Search      | ✗                  | ✗                   | ✗                | ✗           | ✗                 | provenance undisclosed or claimed (Exa: 500B+) · 5–7 USD per 1k                          |
| OpenAI · Anthropic web-search tools   | ✗                  | ✗                   | ✗                | ✗           | ✗                 | inside the model API · 10 USD per 1k calls · Anthropic via Brave                         |
| ClueWeb22 + ChatNoir                  | ✓                  | partly              | ✓                | licence     | —                 | 10B pages, 207 languages · frozen 2022 · free licence, shipping fee                      |
| MS MARCO Web Search                   | ✓                  | partly              | ✓                | ✓           | clicks            | 100M-document benchmark on ClueWeb22, 10M queries · frozen                               |
| Open Web Index + OURRS                | ✓                  | ✓                   | ✓                | ✓           | by design         | 39B URLs crawled, 45 TiB index · daily · free                                            |



# Appendix: where that puts us



Transparency on one axis: can we inspect, fix, configure and log.


Completeness on the other: size, freshness, languages.

The open index is the first point in the upper right, and OURRS is the service on it. Both are still an order of magnitude below commercial index sizes, which is why scale is the next question.


An API lets researchers query someone else's experiment; bulk access lets



# Appendix: agents could discover search endpoints themselves




schema.org SearchAction, the markup behind the sitelinks search box. Source: metayer.de



An old idea becomes an agent affordance: 3.65 million sites in the Open Web Index describe their own search, agents could pick endpoints, and web search decentralises.

Entity counts by schema.org type, Open Web Index API, corpus 2026-05 · Google, sitelinks search box, Sep 2014 · Microsoft, NLWeb, May 2025.

# Appendix: what it takes to evaluate such a system




An evaluation is a controlled, repeatable comparison on a known object.

- We must know what is searched, so the corpus has to be inspectable.
- We must be able to repeat, so a snapshot has to be fixable.
- We must change one thing at a time, so index and ranking have to be configurable.
- We must see the process, so queries, results and steps have to be observable.
- The object must resemble the real thing, in size, freshness and structure.

# The index is built every day: crawl, process, index, publish


## Index Generation

Web resources are selected and retrieved, their content and metadata are analysed, and all data stored in the index database.




## Search Applications

A user search request will be answered by a search application that makes use of the open web index.



## Data Products

Knowledge representation models will be created using the open web index, in order to be used by any agent and for many applications



- 1 Crawling runs continuously on cloud partitions with a shared frontier.
- 2 Processing extracts text, language, topic, structured data and legal signals.
- 3 Indexing produces a sparse index, metadata and embeddings for every daily shard.
- 4 Publishing puts each shard on the dashboard with a licence and a dataset card.

Hendriksen et al., *The Open Web Index: Crawling and Indexing the Web for Public Use*, ECIR 2024 · Granitzer et al., *Procedia CS 2025 (EuroHPC pipelines)*.

# Appendix: in an agent, small retrieval effects compound




illustration: independent per-step effects, not a model of agents

**What the curve compares.** Suppose a change to the retriever alters the outcome of one search step with probability  $\epsilon$ , and that steps are independent. After  $k$  steps the share of trajectories in which at least one step came out differently is  $1 - (1 - \epsilon)^k$ . The dashed line marks where that reaches half.

BROWSECOMP-PLUS, 2025

## The multiplier is real

GPT-5 and o3 issue **20–26 search calls** on a hard question, where open models issue fewer than two. A better retriever saves one to three of them.

Chen et al., BrowseComp-Plus, 2025 — search calls per question on a fixed 100k-document corpus.

# Appendix: three columns, three requirements

|                                | retrieval                                                                                       | interaction                                                                           | collaboration                                                                                  |
|--------------------------------|-------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------|
| openness, research sovereignty | <p>an open corpus and index below the model</p> <p><i>Open Web Index</i></p>                    | <p>an open, instrumented search service</p> <p><i>OURRS</i></p>                       | <p>an open layer for shared knowledge and communication</p> <p><i>social agentic layer</i></p> |
| scale and robustness           | <p>corpus size and real web text as variables</p> <p><i>CoRE · OWS-Curlie-2025 · WOWS</i></p>   | <p>traces at scale, by simulation and by design</p> <p><i>AgentSim · UXSim</i></p>    | <p>populations of agents and people, over time</p> <p><i>agent social networks</i></p>         |
| joint effort                   | <p>one index, built and operated on public compute</p> <p><i>OpenWebSearch.EU · EuroHPC</i></p> | <p>traces pooled across groups in one format</p> <p><i>IIRSim Studio · Sim4IA</i></p> | <p>the ecosystem evaluated together, or left to platforms</p> <p><i>evaluation labs</i></p>    |

# Appendix: evaluating the ecosystem, area by area

Trajectories, not answers. Populations, not users.  
Drift, not snapshots.  
Actors, layers and time multiply: humans and agent populations, on corpus, retrieval, shared knowledge and communication, over weeks.

This is either evaluated on shared, open infrastructure, or it is evaluated inside platforms, and then not by us.

| Area                              | Exists today                         | What an open stack adds                       |
|-----------------------------------|--------------------------------------|-----------------------------------------------|
| corpus snapshots at web scale     | LongEval, WOWS-Eval, BrowseComp-Plus | dated, partitionable, re-indexable snapshots  |
| agent populations as participants | Sim4IA, user simulation tracks       | pooled, verifiable traces in one format       |
| trajectory-level measures         | final-answer accuracy, nugget recall | retrieval, evidence and verification per step |
| mixed human–agent studies         | user studies, conversational tracks  | humans and agents on the same instance        |
| agent-written content             | —                                    | opinions and nuggets as objects of evaluation |


# Appendix: how good is the simulation, and where does it stop?

| What a session has to produce   | How simulation does                                                                  | The check, and where it comes from                                                                                                                         |
|---------------------------------|--------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| queries and next actions        | <b>best</b> ; a cognitive user model measurably helps                                | next-action accuracy <b>0.37</b> → <b>0.48</b> on the Sowipor sessions, 0.32 → 0.46 on AOL · CACSM, JCDL 2024                                              |
| evidence selection and stopping | <b>partly</b> ; the errors are systematic, so it has to be calibrated against people | <b>480 human sessions</b> across traditional, conversational and hybrid interfaces · <i>SearchLab</i> , CHIIR 2025, and the <i>UXLab</i> pilot, CHIIR 2026 |
| satisfaction, trust, effort     | <b>not simulated</b> ; it can only be inferred from behaviour                        | cognitive states inferred over <b>50,942 sessions</b> , and they help most exactly where click signals are weakest · <i>Beyond the Click</i> , ECIR 2026   |

Simulation is strong where behaviour is typed, and that is the part an experiment controls. It is good enough to build on.

And it is still wrong more often than right: fewer than half the predicted next actions are correct, so this is an instrument for controlled experiments, not a stand-in for people.

# Appendix: robustness on noisy web text



## SETUP



41,371 MS MARCO passages, padded with real web text. 3 positions  $\times$  6 lengths  $\times$  2 context types = **36 sub-corpora**, six models.

## RESULT

At least **33 nDCG@10 points** lost across every model, 50 for GTE Qwen2. Arctic M: **79.6  $\rightarrow$  4.3**.


## CONCLUSION

Position and unrelated context dominate. Chunking recovers only part.



# Appendix: evaluation through an API a partner runs

**ChatNoir, run by Webis:** Frozen collections behind one research API: ClueWeb09, 12 and 22, MS MARCO, LongEval, and OWS-Curlie-2025 as its default index. Static corpora are what make a run citable years later.



**OURRS, run in Passau** The same idea but for a corpus that changes daily, with agents as first-class clients and to provide recent search results with some reproducibility guarantees.



## Works today


- query a live corpus instead of hosting one
- different search verticals: products, places, scholarly,...
- agents give feedback, summaries and opinions

## Future Work

- published agent qrels carry no time stamp
- more flexible ranking profiles
- no clean export of a corpus slice with its judgements

# Appendix: does scale matter, if the effect is isolated in a test collection?

same query, same relevant document, a growing corpus around it  
the web ( $\sim 10^9+$  passages)



always there

★ the relevant document

• irrelevant passages

what the ranker must reject grows

■ look relevant, but are not

■ near-duplicates

■ other languages, contradictions

□ the cost of holding it all


## H2 · HYPOTHESIS 2

**Scale is a scientific variable.**

Corpus size and complexity change effectiveness, efficiency and robustness. A result on a small collection does not carry over to the web by itself.

The relevant document stays one.  
What grows is what the ranker has to reject.

# Appendix: the conversational searcher model, beside the classic one



## Same shape, four changes

1. **A information need, not a query.** The searcher generates and issues an *information need* rather than a query.
2. **An action with no counterpart.** It can **request a suggestion**. A result page cannot be asked for more.
3. **Stopping changes its object.** Not "continue on this SERP" but *is the information need satisfied*: a judgement about the conversation, not about a page.
4. **The unit consumed is an information space**, not a ranked document.