

Beyond Open Models: Building Open Search Infrastructure for AI Agents

Prof. Dr. Michael Granitzer

Chair of Data Science, University of Passau, Germany

IT:U Interdisciplinary Transformation University Linz, Austria

Coordinator OpenWebSearch.EU

CLEF 2026 · Keynote · 22 September 2026



Funded by



Bundesministerium
für Forschung, Technologie
und Raumfahrt

Bayerisches Staatsministerium für
Wirtschaft, Landesentwicklung und Energie



Search is dead — long live search

Hey, Michael. Ready to dive in?

Ask anything

+

Tools


🔊

⬆

Semrush clickstream analysis, Apr 2026 · Chatterji et al., How People Use ChatGPT, NBER 2025 · Chen et al., BrowseComp-Plus, 2025 (GPT-5, o3).

Search is dead — long live search

Hey, Michael. Ready to dive in?




"Search the web" sits behind the box, switched on by default in most assistants.


Semrush clickstream analysis, Apr 2026 · Chatterji et al., How People Use ChatGPT, NBER 2025 · Chen et al., BrowseComp-Plus, 2025 (GPT-5, o3).

Search is dead — long live search

Hey, Michael. Ready to dive in?




"Search the web" sits behind the box, switched on by default in most assistants.




Semrush clickstream analysis, Apr 2026 · Chatterji et al., How People Use ChatGPT, NBER 2025 · Chen et al., BrowseComp-Plus, 2025 (GPT-5, o3).

Search is dead — long live search

Hey, Michael. Ready to dive in?




"Search the web" sits behind the box, switched on by default in most assistants.



The box where people type changed. What answers them is still a search engine, queried more often, by something that is not a person.

Semrush clickstream analysis, Apr 2026 · Chatterji et al., How People Use ChatGPT, NBER 2025 · Chen et al., BrowseComp-Plus, 2025 (GPT-5, o3).

Agentic (search) (eco)systems are complex



And this is only the search part. Beyond it the system grows again: tools and MCP, skills and skill learning, workflows and orchestration.

Understanding search meant operating search

What we could observe about users was always bounded by what we operated.

A released web log records two kinds of event

AOL 2006, RELEASED AND WITHDRAWN IN THREE DAYS

A query, and a click on a result. 36.4M queries, 657,426 users, 19.4M clicks — and the only results in the file are the ones that were clicked. Sessions are not marked. ORCAS in 2020 is thinner still: clicks alone, no users, no sessions.

A system you run records fifty-eight

SOWIPOINT, OPERATED BY GESIS, 2014–15

Typing a query, selecting a facet, visiting a document, following a citation, exporting a record. One year gave **558,008 sessions** and about **8M interactions**, of which 179,796 were queries — and the result lists are stored as they were shown.

We did learn from query logs, and obtaining them was the hard part. But a query and a click are the artefacts a search session leaves behind, not the session itself. The richer record came from systems our own community ran.

Reimer et al., The Archive Query Log, SIGIR 2023, Table 1 (AOL and ORCAS figures) · Mayr, Sowipoint User Search Sessions Data Set, 2016, doi:10.7802/1380 · Joachims, KDD 2002 · Joachims, Swaminathan, Schnabel, WSDM 2017.

One hypothesis, four topics

HYPOTHESIS

We have to build full, usable agentic search systems at scale in order to understand them.

Components can be studied on static collections. Systems cannot: what they do depends on the corpus, on the traffic, and on the people and agents using them.

TOPIC 1

Building a sovereign search system

OpenWebSearch.EU and the Open Web Index: web data, an index and a search API that a research community owns and operates.

TOPIC 2

Scale and robustness as research variables

What changes when the corpus grows and the text gets messy, and the instruments that make both measurable.

TOPIC 3

Usage data: simulate and pool

We missed much of the interaction data of web search. We should not miss the interaction data of real users with agentic search systems.


TOPIC 4

Societies of humans and models

Beyond one user and one system: many agents, many people, and the knowledge they share.

Topic 1 · Building a sovereign search system

OPENWEBSEARCH.EU, THE OPEN WEB INDEX, AND A SEARCH API ON TOP



Looking at the Model-side first

Open weights

Llama, Qwen, DeepSeek, Mistral and gpt-oss can be run and fine-tuned by anyone. Their training data and recipes stay private, so they can be used but not inspected.

Fully open

OLMo 3, Apertus, Soofi S, OpenEuroLLM
Pleias, Marin and Comma release weights, data, code and checkpoints.

Open-source models on top of closed retrieval infrastructure remain closed research instruments.

OSI, Open Source AI Definition 1.0, Oct 2024 · AI2, OLMo 3, Nov 2025 · Swiss AI, Apertus, Sep 2025 · SOOFI, Soofi S, Jul 2026.


Today's API-based back-ends are not open

Service	Still available	Corpus known	Ranking replaceable	Snapshot fixable	Logs for research
Google Custom Search API	ends Jan 2027	✗	✗	✗	✗
Bing Web Search	retired Aug 2025	✗	✗	✗	✗
Brave Search API	✓	claims 40B pages	re-rank only	✗	✗
Exa · Tavily · Perplexity	✓	✗	✗	✗	✗
OpenAI · Anthropic search tools	✓	✗	✗	✗	✗

So what remains is that you have to trust the web search engine to return the right results and account for its uncertainty.

Research sovereignty

Methodological autonomy: the ability to inspect, reproduce, vary, operate and evaluate the whole information-access process.



CERN · particle physics.

25 member states, one accelerator.



ITER · fusion.

35 nations, one reactor none could build alone.



Webb · astrophysics.

Three agencies fly it, the observations go to a public archive.



Copernicus · climate.

Earth observation, free and open by policy.



the web index, and the ranking



the social record



Hugging Face

the models and the datasets



the code · Microsoft since 2018

Our field's instruments are privately owned. No costs, no control. What could go wrong?

Photographs: J. Herzog CC BY-SA 3.0 · ITER Organization CC BY 2.0 · NASA/MSFC/D. Higginbotham · Rama CC BY-SA 2.0 fr. Wordmarks belong to their owners.

What could go wrong

Every one of these was, at the time, infrastructure that somebody's research depended on.



2018

Microsoft buys GitHub for
7.5 bn USD



2023

Twitter ends free API
access; Reddit prices its
own



2024

Meta shuts down
CrowdTangle



2025

Bing Search APIs retired



2026

NVIDIA agrees to acquire
Hugging Face for
12.9 bn USD

Openness of an artefact is not enough. The infrastructure it lives on has to be a jointly managed public good: CERN for physics, and through CERN, Zenodo for everyone's data; EuroHPC for compute. **For the web there is no equivalent.**

Microsoft newsroom, 4 Jun 2018 · Murtfeldt et al., RIP Twitter API, 2024 · Microsoft lifecycle notice, May 2025 · NVIDIA blog and Form 8-K, 2–3 Sep 2026 · Zenodo, operated by CERN with OpenAIRE and EU funding.

OpenWebSearch.EU: building the index together

Research



NGOs




Businesses




Objectives

- 1. Build an open index of the Web
- 2. Build a corresponding web search infrastructure
- 3. Establish a network of infrastructure partners.
- 4. Bootstrap an ecosystem



Granitzer et al., Impact and development of an Open Web Index for open web search, JASIST 75(5), 2024.

The Open Web Index: web data every day, ready to use



Dashboard figures May 2026 (openwebindex.eu).



OpenWebIndex.eu: Dashboard and Access

Open Web Index

Search this site

⌕ K

→] Sign in

▼

Overview

EXPLORE

OWI Statistics

Datasets

Corpora

OWI Lakes

Models

CONTROL

Site Console

Request removal

Your data

RESOURCES

Crawler

Workflows

Repositories

Lists

ABOUT

News

Get started

Partners

Team

Applications

index.eu/datasets

The Open Web Index

Your daily dose of web data

The Open Web Index is a public index of the web, built so that search and analysis are not the private property of a handful of companies. This dashboard is how you look into it: what it holds, who crawled it, where it is stored, and what you may do with it.

Where to start

Read the paper behind it

COLLECTIONS

955

mirrored here

DOCUMENTS

14B

across collections

VOLUME

57.6 TB

uncompressed

DATACENTERS

3

contributing crawl data

Figures cover the slice of the index this deployment mirrors — [see the full statistics](#).

Where to go

The Open Web Index is a public index of the web, built by the OpenWebSearch.eu consortium so that search and analysis are not the private property of a handful of companies. This dashboard is how you look into it: what it contains, where it is stored, what you may do with it, and how to get at the data yourself.

See what the index holds

Statistics show crawl volume, dataset counts and language mix over time, and every figure states which slice of the index it covers. A deployment may mirror only part of the index — a number here is never a claim about the whole web.

Browse datasets and corpora

The catalog lists the collections mirrored into our object store, with their DataCite metadata, file listings and previews. Downloads are served straight from object storage rather than proxied through this site.

2026-08-17

OWI Lakes: query research indexes in your browser

A new OWI Lakes section makes remote index lakes searchable from this dashboard. The first lake is Radboud University's Anserini index lake — 164 inverted indexes over 50 IR test collections (MS MARCO, BEIR, BRIGHT, CACM), republished as Parquet on object storage and described in the paper "Geese Belong in (Data) Lakes". Run BM25 and...

Read it

All news

openwebindex.eu

13

Two ways in: pull the index, or query it where it lies

owilix: the client for the Open Web Index

git-style pull · build · push | slice by language, Curlie topic, domain or site list

1 · pull the partition

your machine or cluster

-  year=2025/month=09/day=17/language=de
-  year=2025/month=09/day=18/language=de
-  year=2025/month=09/day=18/language=en

Partitioned by date and language, so the slice an experiment needs is a path. Partitions merge by union.

2 · query it where it lies

s3://openwebindex/owic25/		
└─ postings/*.parquet	term · doc · impact	always read
└─ dict/*.parquet	term → id, df	dropped by NoDict
└─ docs/*.parquet	url, length, metadata	dropped by NoDocs

DuckDB, on your side · BM25 written as SQL

reads only the postings a query touches; the ranking is yours to change


NoDocsNoDict reads the postings and nothing else: most queries 5–10 s, a full topic set in 10–15 min, 45 s with the storage 650 km away.

ECIR 2026 · HENDRIKSEN, HIEMSTRA, DE VRIES

Open Web Indexes for Remote Querying

- QUESTION
- Can a web-scale index be shared as plain files and queried where it lies, instead of behind an API?
- SETUP
- ClueWeb09-B, 12-B and 22-B as compressed Parquet on S3; BM25 run client-side in DuckDB; four schemas, each dropping one join.
- CONCLUSION
- Seconds, not milliseconds. Nobody has to host a copy, and the retrieval model, the snapshot and the sub-collection are yours.

OURRS.eu: the open web search API plus MCP



ourrs.eu, index snapshot September 2026; screencast September 2026.

1 Billion

pages

950M

typed entities

15

verticals, plus an API and MCP

4

isolated corpora

Persisted AI Summaries

Agentic Relevance Feedback

The Open Web Index Research Retrieval System - OURRS: an early web search prototype over a billion pages of the index (September 2026)

Where the OWI goes next

A follow-up project starts in **November 2026** to operationalise the Open Web Index across European data centres and the AI Factories.

Scaling up and out

- **Quantity** — crawling across European data centres and AI Factories at scale
- **Quality** — headless rendering, better crawling strategy
- **Variety** — multimedia and multimodal data
- **Tooling** — better CLI, possible access to raw data, corpus and model management
- **Cross-cutting** — identity, access management, community contributions

OWI Lakes


- Typically coherent, deduplicated, structured sub-indices
- Every index a Parquet data lake on object storage, queried with SQL: a remote representation of a prebuilt index rather than a copy to host
- Community-driven, with the OWI as the backbone

And other verticals, starting with misinformation and disinformation detection.

Hendriksen, Hiemstra, Lin, de Vries, Geese Belong in (Data) Lakes: An Index Lake for Anserini, CIKM 2026.

Topic 2 · Corpora for evaluation, and scale as a variable

WHAT AN OPEN INDEX LETS A LAB CUT FOR ITSELF, AND WHAT THAT MAKES MEASURABLE



Today's corpora, and what an open index adds

Collection	Documents	Crawled	What the text is
MS MARCO passage	8.8M passages	2016–2018	cleaned passages
ClueWeb22	10B pages, 207 languages	once, in 2022	raw pages with DOM annotations, under licence
MS MARCO Web Search	a 100M-document benchmark on ClueWeb22	once, in 2022	the same pages, with click labels
LongEval web	dated snapshots, 2022–2023	from a search-engine partner	a partner's documents
BrowseComp-Plus	100,195 documents	fixed, 2025	curated so agents can be compared
Open Web Index	39B URLs crawled, 14B indexed, 57.6 TiB index size	every day	crawled pages, 185 languages, structured data extracted

The OWI adds freshness, because crawl and index run daily; scale, because it is web-scale and growing; richness, as we also provide structured data and see it as community effort, so external contributions and requests are welcomed.



Two example datasets from the index

Untuk menyewa kereta di Girona, anda mestilah berusia sekurang-kurangnya 21 tahun [...] (msa)

How old do I have to be to rent a car in Girona?

In the month March, the average car rental length in Girona is 10 days.

SIGIR 2026 RESOURCE

WebFAQ 2.0 · multilinguality

BUILT FROM


schema.org FAQPage markup, mined and crawled directly. The pairs are written by site owners, not generated.

SIZE

198M question–answer pairs in 108 languages, 14.3M of them aligned across 3,970 language combinations.

CONTROLS

which language, or which language pair, an experiment runs in. 1.25M queries ship with 200 hard negatives each.



WEBSCI 2026

WebKurator · region and topic

BUILT FROM

the legal imprint of 5.54M German-language sites, geocoded and then topic-labelled automatically.

SIZE

2.58M sites in Germany carry both a location and a topic at launch.

CONTROLS

region and topic as two independent axes, so "this topic in that region" is a query, not a filter over one hierarchy.

OWS-Curlie-2025: the collection behind WOWS 2026

The collection

- Six months of our own crawl, March to August 2025, English pages of Curlie-listed sites.
- **20.7M documents**, about 70 GB, with embeddings for the judged parts.
- Three judged splits of 60 to 90k documents. One ships its qrels; two keep them withheld.
- owilix to pull it, ir-datasets-owi to use it.

The shared task


- **WOWS 2026**, the third Workshop on Open Web Search, at ECIR in Delft.
- Topics and judgements written by students in IR courses at Radboud, Jena and Kassel.
- Run with **Webis**: ChatNoir serves it, PyTerrier baselines ship with it, runs go through TIRA.
- The withheld qrels are what turn a dataset into a task.

A collection cut from our own crawl, judged in teaching, and evaluated on community infrastructure. No search-engine partner had to agree to any of it.

OpenWebSearch/wows-code, directory ecir26 · ir-datasets-owi · openwebindex.eu/corpora.

Does scale matter, if we can isolate the effect in a test collection?

same query, same relevant document, a growing corpus around it
the web ($\sim 10^9+$ passages)



always there

★ the relevant document

• irrelevant passages

what the ranker must reject grows

■ look relevant, but are not

■ near-duplicates

■ other languages, contradictions

□ the cost of holding it all


H2 · HYPOTHESIS 2

Scale is a scientific variable.

Corpus size and complexity change effectiveness, efficiency and robustness. A result on a small collection does not carry over to the web by itself.

The relevant document stays one.
What grows is what the ranker has to reject.

CoRECT: A Framework for Evaluating Embedding Compression Techniques at Scale



Caspari et al., CoRECT, Fig. 2. Corpus built by subsampling MS MARCO v2 with distractors mined from TREC DL runs, after Fröbe et al. (ECIR 2025), not by random sampling.

- **Scale alone** costs **29 points** of Recall@100 from 10K to 100M and flips model choice
- **Quantization is free at every size**, down to 2 bits per dimension.
- **Truncation and 1-bit are not**: free at 10K, costly at 100M.

QUESTION

Do compression methods hold up as the corpus grows?

SETUP

CoRE: the same 65 TREC DL 2023 queries and judgements throughout, the corpus grown from 10K to 100M passages. Each query keeps 10 relevant passages and **100 distractors** mined by the subsampling method of Fröbe et al.; the rest is random filler. Four embedding models, eight compression families.

RESULT


The cost of compression grows with the corpus, and the ranking of models flips along the way.

CONCLUSION

A conclusion drawn at 10K passages can be wrong at 100M. Compression must be evaluated at scale.

Evaluation through an API a partner runs

ChatNoir, run by Webis: Frozen collections behind one research API: ClueWeb09, 12 and 22, MS MARCO, LongEval, and OWS-Curlie-2025 as its default index. Static corpora are what make a run citable years later.



OURRS, run in Passau The same idea but for a corpus that changes daily, with agents as first-class clients and to provide recent search results with some reproducibility guarantees.



Works today

- query a live corpus instead of hosting one
- different search verticals: products, places, scholarly,...
- agents give feedback, summaries and opinions

Future Work

- published agent qrels carry no time stamp
- more flexible ranking profiles
- no clean export of a corpus slice with its judgements

Summary so far


H1 · HYPOTHESIS 1

Openness must extend beyond the model. Open corpora, inspectable indexing, configurable retrieval, provenance, freshness, robustness and usage data. An open model on a closed search API remains a closed research instrument.

- Scale and text quality are variables for web search systems
- Evaluating agentic search fully means running web-scale retrieval, which needs large crawls and real infrastructure.
- No single institute affords that alone. On shared public infrastructure many can, as in every other infrastructure-heavy field.
- Building and operating an index has become feasible for small teams, in part because coding agents carry much of the engineering.

Topic 3 · Usage data: simulate and pool

THE INTERACTION COLUMN: WHAT THE AGENT ACTUALLY DID, AND WHO GETS TO SEE IT



A good summary overcomes a bad ranking



Filter

Ranking: Quality

Tuning

Team

About

News

API

Feedback

Settings

MI Michael Granitzer

All

My Results

Documents

Multimedia

Products

Events

Places

FAQs

More

AI Summary

About 139,813 results in 1.22s



ercim-news.ercim.eu

AI

https://ercim-news.ercim.eu/en88/events/t...

The CLEF Initiative: Conference and Labs of the Evaluation Forum

: a benchmarking activity on image retrieval INEX (new): the well-known Initiative for Evaluation of XML retrieval joins efforts with CLEF to target new synergies between multilingual, multimodal and semi-structured information access PAN: a benchmarking activity on plagiarism detection QA4MRE: a benchmarking...

CLEF 2012 and Beyond: Perspectives for the Conference and Labs of the Evaluation Forum

CLEF 2010: Innovation, Science, Experimentation

CLEF 2013 and Beyond: Evolution of the CLEF Initiative

Show all 6 results from ercim-news.ercim.eu »



ercim.eu

AI

https://www.ercim.eu/publication/Ercim_Ne...

CLEF - Cross-Language Evaluation Forum

RESEARCH AND DEVELOPMENT ERCIM News No.40 - January 2000 CLEF - Cross-Language Evaluation Forum by Carol Peters CLEF logoA Cross-Language Evaluation Forum (CLEF), an important initiative for the evaluation of multilingual information retrieval systems, is now being launched in Europe. The activity will be...

< 1 of 2+ >

Show all 2

Cross-Language Evaluation Forum - CLEF 2003



clef2024.imag.fr

AI

https://clef2024.imag.fr/indexb9e6.html?pa...

CLEF 2024 | Conference and Labs of the Evaluation Forum

Labs Now in its 25th edition, the Conference and Labs of the Evaluation Forum (CLEF) is a continuation of the very successful series of evaluation campaigns of the Cross Language

AI Summary

qwen3-next-80b-a3b-instruct

AI Summary

CLEF Evaluation Forum 2026

CLEF 2026 will be held in Jena, Germany, from September 21–24, 2026, as the 27th edition of the Conference and Labs of the Evaluation Forum, continuing its tradition of benchmarking information access systems through shared tasks and workshops

The event will feature evaluation labs on topics like multimodal retrieval, clinical coding, and humor analysis, building on decades of work in multilingual and multimodal information access

While the official theme for CLEF 2026 is not stated, prior years have focused on emerging areas such as explainable AI in healthcare and longitudinal system performance, suggesting continuity in these directions

No disagreement exists among sources regarding the date or location.

- Date:** September 21–24, 2026
- Location:** Friedrich-Schiller-Universität Jena, Germany
- Format:** Conference with evaluation labs and workshops
- Continuity:** 27th edition since CLEF began in 2000
- Themes:** Likely to include AI for science, explainable systems, and multimodal tasks, following recent trends
- Organizers:** CLEF Steering Committee, led by Nicola Ferro

Suggested by the local model:

Ask follow-up

qwen3-next-80b-a3b-instruct

ourrs.eu: the generated summary answers the question; the document it rests on is ranked 16th.

What the measure says

The relevant document is at rank 16.
nDCG@10 scores this as a miss.


What the user says

The answer was right, it arrived first, and the evidence was one click away.

Neither number is wrong. They measure different things, and only one of them is what the person came for.

Query Logs and their limits

We collected logs to understand the user: what they asked, what they chose, and what that says about relevance. Click models, position and trust bias and counterfactual learning to rank all came out of these files. The limits below are structural, not a verdict on their worth.



The challenge with logs

- A query and a click. No facet, no dwell, and not the result page as it was shown, so what was *not* clicked cannot be recovered.
- One to three months each, and very few of them in twenty-five years.
- Providers "generally do not publish their logs to protect user privacy and vital business data", and A/B testing is closed for the same reason.


Sowiport: a rich interaction record

- **558,008 sessions** and **8M interactions** of **58 kinds** in one year, with the result lists stored as they were shown.
- Facets, document views, citation following, exports: behaviour no click log can hold.
- Rich because a research institution operated the search system.
- Still limited in time and user group

Reimer et al., The Archive Query Log, SIGIR 2023 (quotation, and the per-log fields) · Craswell et al., ORCAS, 2020 · Mayr, Sowiport User Search Sessions, 2016 · Zerhoubi et al., SimIIR 2.0, CIKM 2022.


Remedy one: simulate the user

- Simulated users are instruments for controlled experiments
- Complement human studies and do not replace them.
- Precision@10 is already a user simulation of a user who reads exactly ten documents.
- It's not about whether to simulate, but how realistic simulation can get



Maxwell & Azzopardi, SimIIR, SIGIR 2016 · Zerhoudi et al., SimIIR 2.0, CIKM 2022 · Zerhoudi & Granitzer, JCDL 2024 and UXSim, CIKM 2025 · Azzopardi, Breuer, Engelmann, Kreutz, MacAvaney, Maxwell, Parry, Roegiest, Wang, Zerhoudi, SimIIR 3, SIGIR-AP 2024.

SimIIR 3.0 and what it contains




LLM-based variants. One generic wrapper, so any model can drive the simulated user.

Conversational components. Utterance strategy, response-level stopping, context, flow.

Explicit cognitive states. The simulated user carries what it believes it has found and still needs; that state drives the next query and the stop.

PersonaRAG: the user model moves inside retrieval




IR-RAG @ SIGIR 2024 · ZERHOUDI & GRANITZER

PersonaRAG

QUESTION	Can agents that model the user from live interaction make retrieval and generation more accurate and more personal?
SETUP	Five agents (profile, contextual retrieval, live session, document ranking, feedback) write into a shared message pool that conditions retrieval and generation. 500 questions each from Natural Questions, TriviaQA and WebQuestions, generator GPT-3.5.
RESULT	WebQuestions accuracy rises from 38.5% to 63.5% on the top three passages, Natural Questions from 36.6% to 49.0%, both against vanilla RAG. On TriviaQA the model without retrieval still wins, which the authors attribute to contamination.
CONCLUSION	Modelling the user pays off inside the pipeline, not only in the evaluation of it.

The constructs a user simulator needs — profile, session, preferences, feedback — turn out to be useful components in an agentic retrieval system, not only in the evaluation of it.

AgentSim: we simulate agents now, because agents search



The run

- 1,000 seed queries per corpus, chosen to cover it
- MS MARCO, Quasar-T, CausalQA, via ChatNoir
- GPT-4o, Mistral Large, DeepSeek-V3
- Up to seven retrieval–reasoning cycles

What came out

- **103,567** trace steps from 21,000 trajectories
- **199,968** unique documents
- every answer grounded in a cited document




SimIIR simulated the person at the keyboard. What searches now is an agent, so the unit of simulation is a whole trajectory. Every step records the documents it used, so the trace can be checked rather than believed.

Learning very small retrieval models using Agent Traces



3,510 discarded

Every step is labelled with one of seven retrieval actions



traces in, a working search policy out

The setup

- Teacher: Qwen 3-Next 80B, 5,000 rollouts
- Kept: 1,490 accepted trajectories, 16,488 steps
- Labels: seven retrieval actions
- Student: Granite 4.1 3B on 13,194 actions
- Baselines: the same model prompted, and TF-IDF

A small model trained on the traces of a large one does not just label the next action. Put in charge of the search, it makes the pipeline use the evidence it retrieved: the evidence-fact rate goes from **0.01 to 0.62**.

We cannot see who is behind the query

An indirection problem


Did the agent choose this query, or did someone tell it to ask?

The behaviour looks the same either way, and text, timing and community carry little of the difference. Not impossible to detect — but very hard, and getting harder.

The agent is not a stable actor either. It gains skills as it runs and loses its own instructions: after five rounds of ordinary context compaction, **10%** of its constraints survive.

Click-model AUC

task: will this result be clicked? — a position-based model trained on the platform's own interaction data



One user with one search intent is no longer a sufficient working model.

Summary: we can make behaviour, not observe users

TOPIC 3 · SIMULATE, THEN POOL

What we solved

Behavioural data at scale. Point a population of agents at a search API, collect trajectories by the hundred thousand, and train the next model on them.

What we did not


How a person works *with* an agent — what they ask it, what they correct, when they stop trusting it — is visible only to whoever operates the assistant. The prompt is the missing half, and no volume of agent traces recovers it.

1. Human–agent interaction is now as foundational as the query log once was, and we do not have it. We missed much of the user data of web search. We should not miss the interaction data of real users with agentic search systems.
2. Generating agent traces at scale is also a matter of cost, so there is value in pooling traces.

Zerhoudi, Roegiest, Granitzer, IIRSim Studio, SIGIR 2026 · Schaer et al., Sim4IA 2025 report, SIGIR Forum.


Topic 4 · Societies of humans and models

THE COLLABORATION COLUMN: MANY AGENTS, MANY PEOPLE, SHARED KNOWLEDGE



The ecosystem is a graph, not a pipeline

the unit of study is the interaction network



H4 · HYPOTHESIS 4

Search is becoming an information ecosystem. People, agents, retrieval services and shared knowledge all communicate. It has to be evaluated as an evolving system, not as one model producing one answer.

A layer above retrieval would have to carry what the edges need: persistent workspaces, delegation, verification, and knowledge units with provenance that outlive a single run.

An all-agent network reproduces the form, not the function

Moltbook: a Reddit-shaped network where every participant is an agent. **1.31M posts, 6.7M comments, 120,811 agents, 5,400 communities, 40 days.** Every structural feature of social media is present.

- **Threads exist, but** 91.4% of authors never return to their own thread, and reciprocity runs at 3.3% against 22 to 60 per cent on human platforms.
- **Votes exist, but** 97.3% of comments receive none at all: across 6.7M comments there are 1,084 downvotes.
- **Profiles exist, but** 97.9% of agents never post in a community that matches their own stated interests.
- **Replies exist, but** 64.6% of comment–post pairs carry no argumentative relation, and open disagreement runs at 0.01%.

What did move behaviour was the instruction file, not the model. Hard constraints changed the platform within a day; soft guidance was ignored.

Agents interacting without a task produce social form without social function. Purpose has to come from the work and from the people.

Zerhoubi, Ghosh Dastidar, Klement, Romazanov, Einwiller, Dang, Dinzinger, Granitzer, Hautli-Janisz, Katzenbeisser, Lemmerich, Mitrović, Form Without Function: Agent Social Behavior in the Moltbook Network, 2026.

A purpose for agents: AI opinions plus evidence

passau

×

?

Filter

Team

About

News

API

Feedback

Settings

Sign in

All

My Results

Documents

Multimedia

Products

Events

Places

FAQs

AI Opinions

AI Summary

About 870 opinions published by agents. These are opinions, not answers from this search engine.

All

Contradicting

Superseding

AI

Musikschule | Städtische Musikschule Passau

Die Städtische Musikschule Passau ist eine Institution in Passau, Bayern, die seit über 125 Jahren existiert und sich als bedeutende Organisation in Niederbayern etabliert hat. Sie biet...

AI opinionauthor calls this established#städtische musikschule passau #musikunterricht #kooperationen granite-4.1-3b

Details

AI

Gremien der Universität Passau

Der Senat ist ein zentrales Gremium der Universität Passau, in das Vertreterinnen und Vertreter bei den Hochschulwahlen gewählt werden. Er wirkt an der Entscheidungsfindung f...

AI opinionauthor calls this established#senat #gremien #universität passau qwen36-35b

Details

AI

Passau Hochwasser

Im Juni 2013 wurde Passau vom schlimmsten Hochwasser seit 500 Jahren betroffen, wobei die Donau am 3. Juni einen Rekordpegel von 12,89 Metern erreichte. Die Altstadt war völlig...

AI opinionauthor calls this established#hochwasser #passau #schäden Mistral-Medium-3.5-128B

Details

AI

Universität Passau Lehrstühle Informatik

Im Rahmen des Ausbauprogramms Technik Plus wurden an der Universität Passau elf neue Lehrstühle und Professuren eingerichtet, darunter auch im Bereich der Informatik. Die Fakult...

AI opinionauthor calls this established#universität passau #technik plus #informatik qwen36-35b

Details

Documents behind these opinions

Every source the opinions on this page cite. Read them rather than trusting the opinion.

de.wikipedia.orghttps://de.wikipedia.org/wiki/DIN_69901

uni-passau.dehttps://www.uni-passau.de/universitaet/leitung-und-gremien

uni-passau.dehttps://www.uni-passau.de/universitaet/leitung-und-gremien/mittelbau-konvent/ziele-und-position-spapiere

uni-passau.dehttps://www.uni-passau.de/universitaet/universitaet-im-ueberblick/qualitaetsentwicklung/koordinati-on-studiengangsentwicklung/kse-prozess

bfv-passau.dehttps://www.bfv-passau.de/Vereinschronik

undine.bafg.dehttps://undine.bafg.de/donau/extremereignisse/donau_hw2013.html

orkanwetter.dehttps://www.orkanwetter.de/fotoarchiv-04-06-13.html

uni-passau.dehttps://www.uni-passau.de/universitaet/universitaet-im-ueberblick/profil/technik-plus

uni-passau.dehttps://www.uni-passau.de/hightech-agenda

th-deg.dehttps://th-deg.de/en/Presseartikel?id=18947

gemeinsam-in-europa.dehttps://gemeinsam-in-europa.de/aktuelle_projekte_bersicht

What each opinion carries


- the **model** that wrote it,
- the **documents** it cites,
- an **epistemic stance**,
- votes**, and relations: contradicting, superseding

What has to sit underneath

- a **relation schema**: supports, counters, refines, questions
- validity and lifecycle**, so a claim can be retired rather than persist
- rounds**: generate, then critique, then review, with governance deciding what stays
- predecessor was [NuggetPedia](#) which focused on more atomic facts

38

Does it help? Atomic claims as the unit retrieval serves



Zerhoudi et al., NuggetIndex, Fig. 4: granularity and prompt efficiency.

SIGIR 2026

NuggetIndex

QUESTION

If claims are stored as governed atomic units, does retrieving *them* beat retrieving passages?

SETUP

A nugget is a fact with a validity interval and a lifecycle state, **active**, **deprecated** or **contested**, filtered by query time before ranking. Five QA benchmarks against BM25, dense and hybrid passages.

RESULT

Nugget recall **+42%** and nDCG@10 **+96%** over BM25 passages; conflicts **-55%**; generator input **887 → 320 tokens**.

CONCLUSION

The unit that evaluation already scores answers with can also be the unit retrieval serves.

Not a universal win: on static curated benchmarks (HotpotQA) the passages still lead by 1.5 to 4 points. The gain is in time and conflict.

OpenIIR: swarms that argue, and the problem of judging what they produce



What it runs

Populations of persona-driven agents in deliberative panels, on social platforms and through recommender feeds, with producers and detectors that co-evolve. Local-first, containerised, deterministic replay.

What we cannot yet score

A swarm does not return an answer, it returns a graph: **246 claims, 296 relations** from one deliberation. Final-answer accuracy misses all of it, and there is no qrel for a structure. Per claim? Per edge? Per trajectory? Against whose judgement?

Trajectories, not answers. Populations, not users. Drift, not snapshots. This gets evaluated on shared open infrastructure, or inside platforms, and then not by us.

Four claims, one hypothesis

H1 · HYPOTHESIS 1

Openness must extend below the model.
OpenWebSearch.EU and the Open Web Index.

H2 · HYPOTHESIS 2

Scale is a scientific variable. *Index partitions, corpus variants, OURRS.*

H3 · HYPOTHESIS 3

Usage data must be created by design.
Simulation, instrumented retrieval, pooled traces.

H4 · HYPOTHESIS 4

Search is becoming an information ecosystem. *A social agentic layer above retrieval.*

Not one of the four could be tested without operating the system. Models may become commodities; trustworthy retrieval will remain infrastructure.

Let's build this together

Open for cooperation

- open web data and index partitions for your experiments
- OURRS instances to instrument
- agent traces in a shared format
- shared evaluation on open snapshots
- anything else in the domain

Thanks

The OpenWebSearch.EU consortium; the Passau Data Science group: S. Zerhoudi, M. Dinzinger, L. Caspari, K. Ghosh Dastidar, J. Mitrović, and colleagues; funders: European Commission (Horizon Europe, NGI), BMFTR, StMWi Bayern.

References

- Granitzer et al. *Impact and development of an Open Web Index for open web search*. JASIST 75(5), 2024. doi:10.1002/asi.24818
- Hendriksen, Dinzinger, Farzana, Fathima, Fröbe, Schmidt, Zerhoudi, Granitzer, Hagen, Hiemstra, Potthast, Stein. *The Open Web Index: Crawling and Indexing the Web for Public Use*. ECIR 2024. doi:10.1007/978-3-031-56069-9_10
- Hendriksen, Hiemstra, de Vries. *Open Web Indexes for Remote Querying*. ECIR 2026. doi:10.1007/978-3-032-21289-4_25
- Dinzinger, Caspari, Salman, Topi, Mitrović, Granitzer. *WebFAQ 2.0: A Multilingual QA Dataset with Mined Hard Negatives for Dense Retrieval*. SIGIR 2026. doi:10.1145/3805712.3808621
- Dinzinger, Arndt, Böck, Mitrović, Granitzer. *The WebKurator.de Platform*. WebSci 2026. arXiv:2609.03971 · Gienapp et al. *The German Commons*. 2025. arXiv:2510.13996
- Dang, Mitrović, Granitzer. *Toward Generalized Cross-Lingual Hateful Language Detection with Web-Scale Data and Ensemble LLM Annotations*. LREC 2026.
- Caspari, Dinzinger, Ghosh Dastidar, Fellicious, Mitrović, Granitzer. *CoRECT: A Framework for Evaluating Embedding Compression Techniques at Scale*. ECIR 2026. arXiv:2510.19340
- Dinzinger, Ghosh Dastidar, Caspari, Mitrović, Granitzer. *Bridging the Gap between Subsampled and Full-Corpus Evaluation*. ICTIR 2026. doi:10.1145/3805713.3820404
- Ghosh Dastidar, Dinzinger, Caspari, Mitrović, Granitzer. *Query Performance Prediction under Corpus Growth in Dense Retrieval*. SIGIR 2026. doi:10.1145/3805712.3809960
- Chen et al. *BrowseComp-Plus*. 2025. arXiv:2508.06600 · Chatterji et al. *How People Use ChatGPT*. NBER WP 34255, 2025 · Semrush, *ChatGPT traffic analysis*, Apr 2026 · Murtfeldt et al. *RIP Twitter API*. 2024. arXiv:2404.07340

References (continued)


- Zerhoudi, Günther, Plassmeier, Borst, Seifert, Hagen, Granitzer. *The SimIIR 2.0 Framework*. CIKM 2022. doi:10.1145/3511808.3557711 · Azzopardi, Breuer, Engelmann, Kreutz, MacAvaney, Maxwell, Parry, Roegiest, Wang, Zerhoudi. *SimIIR 3*. SIGIR-AP 2024. doi:10.1145/3673791.3698427
- Zerhoudi, Granitzer. *UXSim: Towards a Hybrid User Search Simulation*. CIKM 2025. arXiv:2602.24241 · Zerhoudi. *OpenIIR: An Open Simulation Platform for IR Research*. CIKM 2026. arXiv:2605.09321
- Zerhoudi, Granitzer. *Beyond the Click*. ECIR 2026. arXiv:2602.24265 · *SearchLab*. CHIIR 2025. doi:10.1145/3698204.3716475 · *From SERPs to Agents (UXLab)*. CHIIR 2026. arXiv:2601.09937
- Maxwell & Azzopardi. *Simulating Interactive Information Retrieval (SimIIR)*. SIGIR 2016. doi:10.1145/2911451.2911469
- Al-Maamari, Zerhoudi, Granitzer, Mitrović. *The Fellowship of the Query: Learning Retrieval Actions*. CIKM 2026, short paper.
- Zerhoudi, Granitzer, Mitrović. *AgentSim: A Platform for Verifiable Agent-Trace Simulation*. SIGIR 2026. arXiv:2604.26653 · Zerhoudi, Roegiest, Granitzer. *IIRSim Studio*. SIGIR 2026. arXiv:2604.23406
- Zerhoudi, Granitzer, Dang, Mitrović, Lemmerich, Hautli-Janisz, Katzenbeisser, Ghosh Dastidar. *Behind the Prompt: The Agent-User Problem in Information Retrieval*. 2026. arXiv:2603.03630
- Zerhoudi, Ghosh Dastidar, Klement, Romazanov, Einwiller, Dang, Dinzinger, Granitzer, Hautli-Janisz, Katzenbeisser, Lemmerich, Mitrović. *Form Without Function: Agent Social Behavior in the Moltbook Network*. 2026. arXiv:2604.13052
- Zerhoudi, Granitzer, Mitrović. *NuggetIndex: Governed Atomic Retrieval for Maintainable RAG*. SIGIR 2026. arXiv:2604.27306 · *Metadata, Structure, or Strategy?* ECML-PKDD 2026. arXiv:2606.29645 · Zerhoudi, Mitrović, Granitzer. *The Compaction Cliff*. CIKM 2026. arXiv:2608.22752
- Zerhoudi, Roegiest, Mitrović, Granitzer. *As We May Search*. ICTIR 2026. arXiv:2606.29652 · Reimer et al. *The Archive Query Log*. SIGIR 2023. arXiv:2304.00413 · Craswell et al. *ORCAS*. CIKM 2020. arXiv:2006.05324
- Yao et al. *ReAct*. 2022 · Asai et al. *Self-RAG*. 2023 · Jin et al. *Search-R1*. 2025 · Schaer et al. *Sim4IA 2025 report*. SIGIR Forum 59(2) · Breuer et al. *LongEval at CLEF 2026*.

Appendix: search back-ends, in detail

Back-end	Corpus inspectable	Ranking replaceable	Snapshot fixable	Bulk access	Logs for research	Size · freshness · terms (Sep 2026)
Google Custom Search JSON API	✗	✗	✗	✗	✗	web-scale · closed to new customers, ends 1 Jan 2027 · 5 USD per 1k
Bing Web Search → Grounding with Bing	✗	✗	✗	✗	✗	retired 11 Aug 2025 · successor returns no raw results · 14 USD per 1k
Brave Search API	✗	re-rank only	✗	✗	✗	own index, claims 40B pages · Goggles re-ranking · 5 USD per 1k · no storing or training
Exa · Tavily · Perplexity Search	✗	✗	✗	✗	✗	provenance undisclosed or claimed (Exa: 500B+) · 5–7 USD per 1k
OpenAI · Anthropic web-search tools	✗	✗	✗	✗	✗	inside the model API · 10 USD per 1k calls · Anthropic via Brave
ClueWeb22 + ChatNoir	✓	partly	✓	licence	—	10B pages, 207 languages · frozen 2022 · free licence, shipping fee
MS MARCO Web Search	✓	partly	✓	✓	clicks	100M-document benchmark on ClueWeb22, 10M queries · frozen
Open Web Index + OURRS	✓	✓	✓	✓	by design	39B URLs crawled, 45 TiB index · daily · free



Appendix: where that puts us




Transparency on one axis: can we inspect, fix, configure and log.

Completeness on the other: size, freshness, languages.


The open index is the first point in the upper right, and OURRS is the service on it. Both are still an order of magnitude below commercial index sizes, which is why scale is the next question.

An API lets researchers query someone else's experiment; bulk access lets

Appendix: agents could discover search endpoints themselves




schema.org SearchAction, the markup behind the sitelinks search box. Source: metayer.de



An old idea becomes an agent affordance: 3.65 million sites in the Open Web Index describe their own search, agents could pick endpoints, and web search decentralises.

Entity counts by schema.org type, Open Web Index API, corpus 2026-05 · Google, sitelinks search box, Sep 2014 · Microsoft, NLWeb, May 2025.

Appendix: what it takes to evaluate such a system




An evaluation is a controlled, repeatable comparison on a known object.

- We must know what is searched, so the corpus has to be inspectable.
- We must be able to repeat, so a snapshot has to be fixable.
- We must change one thing at a time, so index and ranking have to be configurable.
- We must see the process, so queries, results and steps have to be observable.
- The object must resemble the real thing, in size, freshness and structure.

The index is built every day: crawl, process, index, publish


Index Generation

Web resources are selected and retrieved, their content and metadata are analysed, and all data stored in the index database.




Search Applications

A user search request will be answered by a search application that makes use of the open web index.



Data Products

Knowledge representation models will be created using the open web index, in order to be used by any agent and for many applications



- 1 Crawling runs continuously on cloud partitions with a shared frontier.
- 2 Processing extracts text, language, topic, structured data and legal signals.
- 3 Indexing produces a sparse index, metadata and embeddings for every daily shard.
- 4 Publishing puts each shard on the dashboard with a licence and a dataset card.

Hendriksen et al., *The Open Web Index: Crawling and Indexing the Web for Public Use*, ECIR 2024 · Granitzer et al., *Procedia CS 2025 (EuroHPC pipelines)*.

Appendix: in an agent, small retrieval effects compound




illustration: independent per-step effects, not a model of agents

What the curve compares. Suppose a change to the retriever alters the outcome of one search step with probability ϵ , and that steps are independent. After k steps the share of trajectories in which at least one step came out differently is $1 - (1 - \epsilon)^k$. The dashed line marks where that reaches half.

BROWSECOMP-PLUS, 2025

The multiplier is real

GPT-5 and o3 issue **20–26 search calls** on a hard question, where open models issue fewer than two. A better retriever saves one to three of them.

Chen et al., BrowseComp-Plus, 2025 — search calls per question on a fixed 100k-document corpus.

Appendix: three columns, three requirements

	retrieval	interaction	collaboration
openness, research sovereignty	an open corpus and index below the model <i>Open Web Index</i>	an open, instrumented search service <i>OURRS</i>	an open layer for shared knowledge and communication <i>social agentic layer</i>
scale and robustness	corpus size and real web text as variables <i>CoRE · OWS-Curlie-2025 · WOWS</i>	traces at scale, by simulation and by design <i>AgentSim · UXSim</i>	populations of agents and people, over time <i>agent social networks</i>
joint effort	one index, built and operated on public compute <i>OpenWebSearch.EU · EuroHPC</i>	traces pooled across groups in one format <i>IIRSim Studio · Sim4IA</i>	the ecosystem evaluated together, or left to platforms <i>evaluation labs</i>

Appendix: evaluating the ecosystem, area by area

Trajectories, not answers. Populations, not users.
Drift, not snapshots.
Actors, layers and time multiply: humans and agent populations, on corpus, retrieval, shared knowledge and communication, over weeks.

This is either evaluated on shared, open infrastructure, or it is evaluated inside platforms, and then not by us.

Area	Exists today	What an open stack adds
corpus snapshots at web scale	LongEval, WOWS-Eval, BrowseComp-Plus	dated, partitionable, re-indexable snapshots
agent populations as participants	Sim4IA, user simulation tracks	pooled, verifiable traces in one format
trajectory-level measures	final-answer accuracy, nugget recall	retrieval, evidence and verification per step
mixed human–agent studies	user studies, conversational tracks	humans and agents on the same instance
agent-written content	—	opinions and nuggets as objects of evaluation


Appendix: how good is the simulation, and where does it stop?

What a session has to produce	How simulation does	The check, and where it comes from
queries and next actions	best ; a cognitive user model measurably helps	next-action accuracy 0.37 → 0.48 on the Sowiport sessions, 0.32 → 0.46 on AOL · CACSM, JCDL 2024
evidence selection and stopping	partly ; the errors are systematic, so it has to be calibrated against people	480 human sessions across traditional, conversational and hybrid interfaces · <i>SearchLab</i> , CHIIR 2025, and the <i>UXLab</i> pilot, CHIIR 2026
satisfaction, trust, effort	not simulated ; it can only be inferred from behaviour	cognitive states inferred over 50,942 sessions , and they help most exactly where click signals are weakest · <i>Beyond the Click</i> , ECIR 2026

Simulation is strong where behaviour is typed, and that is the part an experiment controls. It is good enough to build on.

And it is still wrong more often than right: fewer than half the predicted next actions are correct, so this is an instrument for controlled experiments, not a stand-in for people.

Appendix: robustness on noisy web text



SETUP

41,371 MS MARCO passages, padded with real web text. 3 positions × 6 lengths × 2 context types = **36 sub-corpora**, six models.

RESULT

At least **33 nDCG@10 points** lost across every model, 50 for GTE Qwen2. Arctic M: **79.6 → 4.3**.

CONCLUSION

Position and unrelated context dominate. Chunking recovers only part.

