


OpenWebSearch.eu — Large Scale Data for Statistical Scrapping

Prof. Dr. Michael Granitzer Chair of Data Science, University of Passau Coordinator
OpenWebSearch.eu



Web Search Today

A critical resource managed as an oligopoly

Two properties that don't fit

- A **critical infrastructure** for society, comparable to satellite navigation
- A **market oligopoly**: dominated by a small number of gatekeepers (Google, Microsoft, Baidu, ...)



But does it matter? Yes — on four dimensions

User Choice & Democracy

- Single gatekeeper for society's information
- Search rankings can shift voting preferences by **20–72%** (SEME)
- Monopolistic ranking = single point of failure for information quality

Untapped Richest Information Source

- 60 % of web pages carry structured data
- No open access to web-scale document collections
- Science, research and technology heavily depend on the Web as platform

Societal & Security Risks

- European OSINT depends on US-controlled infrastructure
- Filter bubbles, algorithmic bias
- Misinformation & Disinformation

AI & Innovation

- Web data is the fuel for LLMs, RAG, and agentic AI
- Closed indices block reproducible research & open innovation

Being able to navigate the Web is as fundamental as GPS or the power grid

Challenges

THE WEB: VOLUME, VELOCITY, VARIETY

Careful: domains, sites, pages, and indexed documents are different units

VOLUME

1.15 Billion

Netcraft site responses

272.6 M domains · Dec 2024 Netcraft survey

Top 0.1%–1% of sites

may account for roughly half of all pages
heavy-tail intuition from large web graph studies

~ 400 Billion

documents in Google's search index
2020 figure cited in US v. Google testimony

Sites ≠ domains ≠ pages; use exact counts cautiously

VELOCITY

+8.6 M / month

monthly net change in Netcraft site responses

Dec 2024 versus Nov 2024 Netcraft survey

~ 3.2 / second

same monthly net change, averaged across December

149 ZB (global)

all digital data created / copied worldwide in 2024
IDC / Statista global datasphere estimate

The web changes constantly; crawling is continuous, not one-off

VARIETY

150+ languages

observed across websites

W3Techs usage survey

249 CMS Platforms

CMS platforms identified in the HTTP Archive dataset
HTTP Archive Web Almanac 2024 (Wappalyzer-based)

HTML · PDF · JSON-LD · media · APIs

content types to parse and process

HTML · PDFs · JS apps · schemas · media · APIs

Sources: Netcraft Web Server Survey Dec 2024 · large web graph studies (heavy-tail intuition) ·
US v. Google testimony on ~400B documents (2020 figure) · HTTP Archive Web Almanac 2024 · W3Techs · IDC/Statista



- High infrastructure costs
- Broad range of technical skills required
 - Big Data Infrastructure
 - Natural Language Processing, Image Analysis
 - Web Technologies
- Growing legal uncertainties

Challenges

THE WEB: VOLUME, VELOCITY, VARIETY

Careful: domains, sites, pages, and indexed documents are different units

VOLUME

1.15 Billion

Netcraft site responses

272.6 M domains · Dec 2024 Netcraft survey

Top 0.1%–1% of sites

may account for roughly half of all pages
heavy-tail intuition from large web graph studies

~ 400 Billion

documents in Google's search index
2020 figure cited in US v. Google testimony

Sites ≠ domains ≠ pages; use exact counts cautiously

VELOCITY

+8.6 M / month

monthly net change in Netcraft site responses

Dec 2024 versus Nov 2024 Netcraft survey

~ 3.2 / second

same monthly net change, averaged across December

149 ZB (global)

all digital data created / copied worldwide in 2024
IDC / Statista global datasphere estimate

The web changes constantly; crawling is continuous, not one-off

VARIETY

150+ languages

observed across websites

W3Techs usage survey

249 CMS Platforms

CMS platforms identified in the HTTP Archive dataset
HTTP Archive Web Almanac 2024 (Wappalyzer-based)

HTML · PDF · JSON-LD · media · APIs

content types to parse and process

HTML · PDFs · JS apps · schemas · media · APIs

Sources: Netcraft Web Server Survey Dec 2024 · large web graph studies (heavy-tail intuition) ·
US v. Google testimony on ~400B documents (2020 figure) · HTTP Archive Web Almanac 2024 · W3Techs · IDC/Statista



- High infrastructure costs
- Broad range of technical skills required
 - Big Data Infrastructure
 - Natural Language Processing, Image Analysis
 - Web Technologies
- Growing legal uncertainties

High upfront costs to tap the Web as resources
— especially small innovators and researchers
are left behind

OpenWebSearch.eu — Our Mission

Building an Open Index of the Web and an open infrastructure to reduce upfront costs for researchers, innovators and companies in tapping the Web as a resource for search, web-analytics and AI.

Four Objectives

- 1. Open Technology Stack** Open-source pipelines for crawling, preprocessing, indexing
- 2. Resource Provision** Building a network of infrastructure providers (EuroHPC / AI Factories)
- 3. Added Value Services** Dashboard, tools, data access APIs
- 4. Bootstrapping the Ecosystem** Third-party calls, community building, applications and use-cases

14 Core Partners



Funded by **Horizon Europe** (NGI programme)

The Open Web Index (OWI)

A **Web Index** = a data structure for fast query-based access and ranking of web documents — the core of every web search engine

Our Proposition: A collaboratively created, open and transparent Web Index running on existing HPC computing centres / AI Factories in Europe



What is the Open Web Index?

The Open Web Index

The Open Web Index — a federated, daily-published web index built across European HPC sites and consumed as composable shards.

The Open Web Index

The Open Web Index — a federated, daily-published web index built across European HPC sites and consumed as composable shards.



Federated production


- each HPC site runs the same pipeline independently
- daily shards by date, language, topic

Three access verbs

- owi pull — download pre-built shards
- owi build — index your own data
- owi push — contribute shards back

The Open Web Index

The Open Web Index — a federated, daily-published web index built across European HPC sites and consumed as composable shards.



Federated production

- each HPC site runs the same pipeline independently
- daily shards by date, language, topic

Three access verbs

- owi pull — download pre-built shards
- owi build — index your own data
- owi push — contribute shards back

Three output types per daily run — the building blocks of a shard:

- index.ciff.gz — sparse inverted index (CIFF format; pre-configured tokenizer)
- metadata_*.parquet — document metadata & enrichments (topics, microdata, locations, language, etc.; extensible)
- Dense embeddings (GPU-enriched slice, Jina V3) for AI-based search and classification (on selected subsets)

OWI at a Glance — Statistics (Feb 2026)

Crawling

Pages crawled	~39 B total, ~12 B unique
Unique hosts	~500 M (extrapolated)
Crawling effort	1,927 machine-days
Avg. throughput	~20 M URLs / day

Index

Documents indexed	~10.7 B
Unique domains	~75 M
Public datasets	511 main / 1,511 total
Index size	45 TB (main) / 50 TB total
Embeddings	119 M docs · 884 GB

Top languages: English (38.9%), German (7.2%), Spanish (6.4%), French (5.7%), Russian (5.2%)

Multimedia Objects: Roughly 2-3 quality images and 0.5 Videos per Web-Document

Comparison: To meet commercial index size, OWI would need to scale by a factor of 20× — factor 10–50× for multimedia




Building the Open Web Index

Conceptual Workflow

Index Generation

Web resources are selected and retrieved, their content and metadata are analysed, and all data stored in the index database.




Search Applications

A user search request will be answered by a search application that makes use of the open web index.




Data Products

Knowledge representation models will be created using the open web index, in order to be used by any agent and for many applications



...

Daily Workflow Pipeline



Crawlig Infrastructure (2026)

- **25 VMs** across IT4I, LRZ, CSC, CERN
- **100 M URLs / day** crawling capacity
- **~1.1 PB** raw data · **~50 TB** index · **810+ public datasets**
- Federated storage: iRODS + S3, orchestrated via LEXIS

Limitations as of 2026

- TLR 5/6: it is not a product (yet)
- No 24/7 Operations
- Limited funding - minimum engineering / improvement

Accessing the Open Web Index

The Dashboard — openwebindex.eu

OWI

Dashboard

SERVICES

Websites

Search

Serci Search EXTERNAL

Crawling

EXPERIMENTAL

Deep Research

UTILITIES

Parse Preview

Ethics Readiness Check

Impressum

Privacy Statement

Terms of Use

Open Web Search Book

OpenWebSearch

Open Web Index

Overview OWI Statistics OWI Datasets OWI Corpora OWI Models

The Open Web Index

Your daily dose of web data!

Welcome to the Open Web Index (OWI)- your way to slice and dice the Petabytes in the Web to your needs. Our mission is to contribute to a fair, open, diverse and free Web by providing an open Index of the Web for you to use. This dashboard gives an overview over this [Open Web Index](#) and provides some prototype services on top of the OWI.



What can you find here?

Discover our comprehensive suite of tools and services designed to help you explore and analyze web data.



OWI Statistics
Get an overview of the data we offer



OWI Datasets
Browse through available datasets



OUR Search
Search through URLs and Titles



Documentation
Access tutorials and detailed documentation

Expect Bugs: While we try to provide as many valuable datasets and services as possible, please always remember that this is still a research project. So we cannot guarantee that everything works perfectly fine and we cannot take any responsibility if something is not working as intended. Bugs are to be expected, but if you encounter some, please get in touch with us so that we can improve.

Central entry point for the Open Web Index openwebindex.eu

- **OWI Statistics** — crawl volume, index size, language distribution, dataset timeline
- **OWI Datasets** — browse, filter and download index shards (by DC, date, format, language)
- **Search** — demo content search + SERCI integration

424 registered users · 17,000+ unique visitors · 70,000+ page views

Dashboard — Statistics & Datasets



OWI Statistics (04/26)

- **1,389 TB** crawled · **281** WARC datasets · **1,696** public datasets · **35.09 TB** index
- Language distribution chart (English 38%, German 8%, French 6%, ...)
- Crawl volume over time per data center + dataset collection distribution by type
- Datasets over time

Dashboard — Statistics & Datasets

Available Datasets

The data shows the datasets available for download (i.e. public owi datasets) or upon request (mostly project private warc datasets) by using our [LEXIS Platform](#) or our [OWI Command Line Tool](#). A dataset is a temporal slice, most often a single day, of crawled (warc), preprocessed and indexed data (owi).

All systems are online

All components are functioning properly. Data repositories and services are accessible.

Last Updated: 12/04/26, 15:36

SYSTEMS	DESCRIPTION	COMPONENTS			STATUS	
IT4I	IT4I	connected	Repository	Project	Public	ONLINE
LRZ	LRZ	connected	Repository	Project	Public	ONLINE

Filter datasets...

All Collections

All Datacenters

All Resources

Public

Sort by: Collection Date

OWI - Open Web Index

05/04/2026 → 05/04/2026

LEGAL

IT4I

PUBLIC

204.5 MB

253 files

All OWI data crawled between 2026-04-05 and 2026-04-05 at CSC2

owilix remote pull all/internalID=fcb68e42-3191-11f1-8cbe-3e9ca2b0f439

VIEW STATISTICS

View Dataset License

Open in LEXIS Portal

OWI - Open Web Index

05/04/2026 → 05/04/2026

MAIN

IT4I

PUBLIC

24.3 GB

1,065 files

All OWI data crawled between 2026-04-05 and 2026-04-05 at CSC2

owilix remote pull all/internalID=faba586e-3191-11f1-9f25-3e9ca2b0f439

VIEW STATISTICS

View Dataset License

Open in LEXIS Portal

OWI - Open Web Index

05/04/2026 → 05/04/2026

CURLIE_FULL

IT4I

PUBLIC

2.1 GB

365 files

All OWI data crawled between 2026-04-05 and 2026-04-05 at CSC2

owilix remote pull all/internalID=fabf3850-3191-11f1-af7e-3e9ca2b0f439

VIEW STATISTICS

View Dataset License

Open in LEXIS Portal

OWI - Open Web Index

05/04/2026 → 05/04/2026

LICENSES

IT4I

PUBLIC

63.6 MB

163 files

All OWI data crawled between 2026-04-05 and 2026-04-05 at CSC2

owilix remote pull all/internalID=fac99de2-3191-11f1-b11f-3e9ca2b0f439

VIEW STATISTICS

View Dataset License

Open in LEXIS Portal

OWI - Open Web Index

05/04/2026 → 05/04/2026

GPU

IT4I

PUBLIC

3.6 GB

374 files

All OWI data crawled between 2026-04-05 and 2026-04-05 at CSC2

owilix remote pull all/internalID=faf69d860-3191-11f1-9f25-3e9ca2b0f439

VIEW STATISTICS

View Dataset License

Open in LEXIS Portal

OWI - Open Web Index

03/04/2026 → 03/04/2026

GPU

IT4I

PUBLIC

40.8 MB

6 files

All OWI data crawled between 2026-04-03 and 2026-04-03 at CSC2

owilix remote pull all/internalID=2f8f678c-2fd5-11f1-8ad8-3e9ca2b0f439

VIEW STATISTICS

View Dataset License

Open in LEXIS Portal

OWI - Open Web Index

03/04/2026 → 03/04/2026

LICENSES

IT4I

PUBLIC

275.8 KB

2 files

All OWI data crawled between 2026-04-03 and 2026-04-03 at CSC2

VIEW STATISTICS

View Dataset License

Open in LEXIS Portal

OWI - Open Web Index

03/04/2026 → 03/04/2026

CURLIE_FULL

IT4I

PUBLIC

5.3 MB

8 files

All OWI data crawled between 2026-04-03 and 2026-04-03 at CSC2

VIEW STATISTICS

View Dataset License

Open in LEXIS Portal

OWI - Open Web Index

03/04/2026 → 03/04/2026

MAIN

IT4I

PUBLIC

108.4 MB

73 files

All OWI data crawled between 2026-04-03 and 2026-04-03 at CSC2

VIEW STATISTICS

View Dataset License

Open in LEXIS Portal

OWI Datasets Tab

- Status of repositories
- Filter by datacenter, format (WARC / Parquet / CIFF / Embeddings), language, date
- Dataset cards with download links — LEXIS Portal or owilix CLI
- Parquet file preview directly in the browser

18

Download Tool: owilix

owilix — The OWI Command-Line Tool

Git-like versioning for index datasets — pull, build, push:

```
owilix remote pull all:latest/collectionName=main
owilix local ls all:latest
```

Slice by: **domain**, **language**, **topic** (Curlie), **sitelist**, **structured data**, **outlinks**
Authentication via B2ACCESS / Keycloak (LEXIS SSO)

Remote Index — Query first, then Download

- Parquet partitions directly queryable via **DuckDB** or **Apache Spark**
- Full-text search
- SQL access: filter by language, domain, date — no full-shard download needed
- Enables search + filtering + download (but still in alpha phase)

OWILIX - Open Web Index CLI

version 5.3.1

python 3.11+

license Apache 2.0

A command-line tool for managing [OpenWebIndex](#) datasets across local and remote repositories.

Quick Start

One-Line Install (Recommended)

```
curl -fsSL https://openwebindex.net/owilix/install.sh | sh
```

<https://opencode.it4i.eu/openwebsearcheu-public/owi-cli/>

```
[→ ~ owilix remote search "michael granitzer" --limit 10
Searching remote index: terms=['michael', 'granitzer'] language=eng representation=main_content mode=OR limit=10
Found 10 results in 113.4s
Search results (10 results in 113.4s)
```

score	date	url
3.088	2025-12-30	https://easychair.org/smart-program/ICADL2024/person16.html
3.010	2026-02-19	https://www.wolframalpha.com/input/?i=michael
3.010	2026-01-19	https://www.michaelmoennich.com/
3.010	2025-12-08	https://michaelkoenig.de/
3.010	2026-01-05	http://www.michaelmoerk.dk/
3.010	2026-01-22	https://michaelkoenig.de/
3.010	2026-02-18	https://www.sktthemes.org/forums/users/dedicationpt/replies/
3.010	2026-02-26	https://www.abeldiaz3.com/tag/michael/
3.010	2026-02-20	https://discourse.psychopy.org/u/michael
3.010	2026-02-27	https://thedigitalgobag.com/author/michael/



The OpenWebSearch Book

Comprehensive Documentation at openwebindex.eu/book

Covers the full OWI ecosystem:

- Getting started with owilix
- Building vertical search engines with MOSAIC
- Data formats: CIFF, Parquet, embeddings
- Pipeline architecture and HPC deployment
- Legal & ethical guidelines for index users

Published as an open, continuously updated online resource — aligned with deliverables and community feedback.

Access Modes at a Glance

Tool	Mode	Best for
Dashboard	Browser	Explore, monitor
owilix	CLI	Download, slice, build
Remote index	DuckDB/SQL	Query + download
MosaicRAG	Conversational	RAG & chat over OWI

All software, data formats and access tools are **open source**



Using the Open Web Index — Applications

MosaicRAG — Conversational Search on the OWI

Modular RAG framework extending MOSAIC with LLM-based capabilities and dense retrieval (WEBIS / TU Graz)

- **Two retrieval paths:** inverted index (CIFF) + dense retrieval
- **4 reranking strategies:** TF-IDF, Embedding cross-encoder, Tournament-style (LLM pairwise), Group-style (LLM window)
- **LLM backends:** self-hosted on-premise models for summarization
- **Conversational search:** follow-up questions about retrieval results; dual-window UI (search results + chat)
- **Workflow:** define retrieval pipeline → issue query → converse about results

Key Features

- Relevance marking — LLM highlights relevant passages within documents
- Summarization — user-provided instructions applied per document
- Domain-specific indices: science, health, arts (8M+ docs)
- Fully open-source modular pipeline

Data Pipeline


OWI daily datasets → CIFF (inverted index) + Parquet (metadata) + dense embeddings (since Dec 2025) → MosaicRAG



Mosaic Search Applications · Earth Observation Search

Mosaic Vertical Search (TU Graz)


Modular framework for building vertical search engines on OWI slices — domain-specific indices for science, health, arts (8M+ docs)



Demo: mosaic.ows.eu · mosaic.ows.eu/demos

Earth Observation Search (DLR / Radboud)

RAG-based QA over scientific papers, datasets & web content; geo-contextualized multi-genre search on interactive maps; KG with 2M+ nodes (OpenAlex, PANGAEA, DLR Geoservice)



Demo: eo-rag-demo.2.rahtiapp.fi · ows-demo.2.rahtiapp.fi

Multilingual Scientific Index · Restaurant Recommender

Multilingual Scientific Index (CSC)

Enriches Finland's Research.fi portal with OWI web data — Finnish & English science content classified by type (publication, blog, event); embedded directly in national research portal

Energy-centric optimization of large-scale additive manufacturing deployment for sustainability through electrification						
Rahoitetun hankkeen kuvaus	The employment of additive manufacturing (AM), commonly called 3D printing, in the industrial sector has been growing for many years. As the technology matures, it is expected that its use will become widespread and substitute other established forms of production. This project is... Näytä enemmän					
Aloitusvuosi	2026					
Päättymisvuosi	2029					
Myönnetty rahoitus	Humberto Almeida Junior	Lappeenrannan–Lahden teknillinen yliopisto LUT	305 350 €			
Rooli Suomen Akatemian konsortiossa	Partneri					
Muut osapuolet	Partneri	Aalto-yliopisto (372725)	298 731 €			
	Johtaja	Lappeenrannan–Lahden teknillinen yliopisto LUT (372723)	395 919 €			
Rahoittaja	Suomen Akatemia					
Rahoitusmuoto	Suunnattu akatemiahanke					
Haku	2025 Tulevaisuuden kestävien energiaratkaisujen tutkimuksen haku 2025					
Muut tiedot						
Rahoituspäätöksen numero	372724					
Tieteenalat	Kone- ja valmistustekniikka					
Tutkimusalat	Kone- ja valmistustekniikka					

Myöntöön liittyvä verkkosivu

Rahoitusmyöntöön ei ole tarjolla linkejä tässä portaaliassa

Viittaukset muissa palveluissa

Viittauksia yhteensä: 20

Service	Percentage	Count
Other	40%	(8)
Article	30%	(6)
News	20%	(4)
Blog	10%	(2)
Event	0%	(0)


Hakutulokset 1 - 5

- <https://www.industryweek.com/...>
3 Ways Additive Is Shaping the Factory of the Future | IndustryWeek
- <https://www.industryweek.com/...>
These Industries are the Future of Additive Manufacturing | IndustryWeek
- <https://global-sci.com/...>
Global Science Press: A Review of Process Optimization for Additive Manufacturing Based on Machine Learning
- <https://www.engineering.com/...>
Customized Production Will Become Scalable with 3D Printing - Engineering.com
- <https://axial.acs.org/...>
Welcome the Newest Associate Editors of ACS Publications Journals: Q3 2018 | ACS Publications Chemistry Blog

Restaurant Recommender (A1 Slovenia)

Privacy-preserving mobile app; on-device personalization via query imputation; LLM-extracted restaurant features from OWI; 5-month pilot with 214 users — overall satisfaction 4.2 / 5

Restaurant Recommender



Or search for restaurants in 1 km radius

Select City ▼

Bascarsija grill food



Baščaršija

[Visit Restaurant Page](#)

Poshtna ulica 8, Maribor 2000 Slovenia 📞 +386 2 250 63 59

Cuisines: Barbecue; European; Grill; Slovenian; Eastern European Meals: Lunch; Dinner

Price: \$4 - \$13 Features: outdoor seating;

Okrepcevalnica Sarajevo

[Visit Restaurant Page](#)

Gospodsvetska cesta 43c, Maribor 2000 Slovenia 📞 +386 2 252 34 31

Cuisines: Barbecue; European Meals: Lunch; Dinner Price: \$2 - \$11 Features: outdoor seating;

Argumentation Search · CERN Search Applications

Argumentation Search (WEBIS)

args.me extended with OWI: pro & contra arguments on 31 controversial topics; argument quality scoring, key point matching, temporal trend analysis

CERN Search Applications

- **Institutional**: 25K+ pages crawled from CERN internal web; full OWS pipeline → MOSAIC index for AccGPT knowledge base
- **Nooon**: disability knowledge search (2.59M docs) for HR & Diversity & Inclusion





We should abolish the right to keep and bear arms

violence

All

Search

520 results

46 ms

Pro-Arguments

ranking score: 9.051

source

Argumentative Segment

Since the start of the year, there have been 80 mass shootings and more than 3,300 deaths due to gun violence nationwide, according to the Gun Violence Archive. We don't need any more thoughts and prayers from elected leaders. We shouldn't have any more moments of silence or vigils. We need action.

Key Point

Quality Score: 0.880

Gun ownership allows for mass-shootings/general gun violence

ranking score: 8.636

source

Argumentative Segment

"Any successful strategy to reduce gun violence requires preventing the diversion of lawful firearms into unlawful commerce," Steven Dettelbach said at the time. "Once there, these firearms end up in the hands of people who are sometimes violent criminals and intend to do harm to the people with whom we live, the innocent people who are victims and survivors of gun violence."

Key Point

Quality Score: -0.690

Guns can fall into the wrong hands

Con-Arguments

ranking score: 11.986

source

Argumentative Segment

Councilman Johnson, repeat after me: Violence is violent. The number one cause of violence is violent people. While semantically there is an argument that some people aren't violent until provoked, drunk and/or stoned, exposed to constant violence in entertainment media, video games etc., the fact remains that violence always, 100% of the time, involves at least one violent person. Blaming guns for violence is like blaming hammers for shoddy home construction.

Key Point

Quality Score: 0.660

ranking score: 8.489

source

Argumentative Segment

No matter how many times actual, real-life case studies are laid out for Cleveland's elected officials, they refuse to recognize that the only PROVEN method of reducing violence is removing violent persons from circulation, either through imprisonment or by vibrantly empowering the potential victims to remove the violent person when confronted by one.

Key Point

Quality Score: 0.870

Gun ownership promotes self protection

Demo: args.me/owi · args.me

MOSAIC

ows-lb-2.cern.ch/ui/

80%

MOSAIC Search

Search term: disability

Search

Index Info

Geo Filter:

West:

East:

North:

South:

Index:

Language:

Limit:

Keyword:

☐ default / all

☒ disability-index

☒ default / all

☐ English

☐ German

☒ default / 20

☐ 10 items

☐ 50 items

☐ 1,000,000

Search URL: https://ows-lb-2.cern.ch/mosaic/search?q=disability&index=disability-index

Search result for term: "disability"

Index: disability-index

Number of items: 20

Cockrell Hill Social Security Disability Attorneys | OneCLE Texas Attorney Directory

Kraft Dallas, TX Social Security Disability Lawyer with 54 years of experience (214) 999-9999 2777 N Stemmons Fwy Suite 1300 Dallas, TX 75207 Free ConsultationSocial Security Disability and Personal Injury Baylor Law School Show Preview View Website

Metadata: language:eng, word count:4273, index date:2025-08-10 00:13

Locations:

Keywords: Attorney • Attorneys • Social Security Disability • Cockrell Hill • Texas • https://lawyers.onecle.com/lawyers/social-security-disability/texas/cockrell-hill

Insurance Coordinator Manual

Partial disability must result from the same condition as the total disability. Proof of partial disability must be submitted within 31 days of the date the employee's total disability period ends. Pa

Metadata: language:eng, word count:20858, index date:2025-08-10 00:51

Locations:

Keywords: https://oklahoma.gov/egid/insurance-coordinators/manual.html

House of Lords/Commons - Joint Committee on the Draft Disability Discrimination Bill - First Report

Disability Rights Commission, DDB 1, para 10.6.1. British Council of Disabled People, DDB 6, para 9. Law Society, DDB 15, p7. 15(a): The Government agrees with the recommendations of the DRTF that MPs

Metadata: language:eng, word count:12308, index date:2025-08-10 00:25

Locations:

Keywords: https://publications.parliament.uk/pa/jt200304/jtselect/jtdisab/82/8222.htm


UKRI diversity data for funding applicants and awardees 2020 to 21 update – UKRI

After removing this opportunity, the award rate for fellows reporting a disability was 22% compared with

Third-Party Projects (FSTPs)

- **AKSE** — argumentation knowledge graph for structured deliberation
- **VERITAS / DEXAI** — verification-oriented RAG for sensitive geopolitical contexts (Ukraine)
- **OMMS** — enriching OpenStreetMap POI with OWI structured data
- **DTCommerce** — product extraction & enrichment pipeline for SME e-commerce
- **CIFFIL** — cross-collection index interoperability for municipal search (Dutch pilots)
- **TILDE** — trustworthy health search with bias detection and entity-level retrieval

Towards Federated RAG Systems



- RAG / Agentic systems require focused, high-quality collections
- Web information redundancy increases due to LLMs — no one-size-fits-all retrieval
- OpenWebSearch.eu supports low-cost creation of topic-specific RAG systems

Key ideas:

- Retrieve from well-designed RAG systems instead of generic search engines
- Topic-specific, deduplicated and curated content
- Domain-optimized ranking, aggregation and summarization
- Lightweight browser-based federation with personal data integration

OwlerLite (WWW 2026) — privacy-respecting browser extension for scope- and freshness-aware retrieval: user-defined **scopes** (trusted sources), semantic freshness detection via SimHash + vector similarity, Scope Fidelity improved from 0.64 → **0.83** (negligible NDCG loss)

Using the Open Web Index — Data & Analytics

Using the Open Web Index — Data & Analytics

Published Datasets

Dataset	Scale	Use case
WebFAQ	96M Q&A, 75 langs	QA training, cross-lingual IR
WebFAQ 2.0	198M pairs, 108 langs	Dense retrieval, hard negatives
Imprints	5.25M hosts, geocoded	Enterprise analytics, Eurostat
German Commons	154B tokens	German LLM pre-training
OWS-Curlie-2025	20.7M docs + qrels	IR evaluation benchmark

All datasets available at openwebindex.eu/corpora

Analytics Potential

- **Official Statistics:** enterprise data extraction from legal imprints (Eurostat, Destatis)
- **Media Monitoring:** supply chain analysis, news monitoring (JRC)
- **AI Training Data:** domain-specific corpora for LLM fine-tuning
- **Evaluation Benchmarks:** web-scale IR test collections with relevance assessments
- **Trend Analysis:** temporal web monitoring at Petabyte scale

The OWI is not just for search — it is a data infrastructure for the web-scale analytics ecosystem

Published Datasets

WebFAQ — Multilingual FAQ from the Web

- **96 million** Q&A pairs across **75+ languages**
- Extracted from structured FAQ markup (JSON-LD, Microdata) in the Open Web Index
- Fields: question, answer, URL, origin, topic, question type
- Parquet format, hosted on Hugging Face

Use cases: multilingual QA training, information retrieval, cross-lingual transfer learning

huggingface.co/datasets/PaDaS-Lab/webfaq

Top languages	Pairs
English	~30M
German	~8M
French	~6M
Spanish	~5M
...	75+ langs



WebFAQ 2.0 — Hard Negatives for Dense Retrieval

- **198 million** FAQ-based Q&A pairs across **108 languages**
- **14.3 million** bilingual aligned QA pairs — largest FAQ-based resource
- **1.25 million** queries with mined hard negatives (20 languages)
- 200 cross-encoder scored negatives per query

Key contributions:

1. Novel web crawling for diverse content extraction
2. Hard negatives for dense retriever training
3. Contrastive Learning & Knowledge Distillation
4. Open datasets and training scripts

Dinzinger, Caspari, Salman, Topi, Mitrović & Granitzer. "WebFAQ 2.0: A Multilingual QA Dataset with Mined Hard Negatives for Dense Retrieval." arXiv:2602.17327, 2025.

Imprints Dataset — Geospatially Grounded Web Curation

- **5.25 million** hosts, **47.8 GB**
- Each website annotated along two dimensions:
 - **Topical classification** via LLM (Foursquare category hierarchy)
 - **Geographic grounding** from legal imprints via geocoding
- **3.08 million** websites geocoded from extracted addresses
- Page types: main page, imprint, contact, about us, privacy, terms

Country	Geocoded	Share
Germany	2,636,913	85.5%
Austria	223,805	7.3%
Switzerland	160,683	5.2%

Top-Level Topic	Share
Business & Professional	41.6%
Community & Government	16.2%
Arts & Entertainment	12.2%



German Commons — 154B Tokens for German LLMs

- 154.56 billion tokens, 35.78 million documents
- Aggregated from 41 diverse sources under open licenses
- Quality-filtered, deduplicated, OCR scored

Domain	Tokens
News Commons	72.67B
Cultural Commons	54.49B
Web Commons	19.89B
Political Commons	3.57B
Legal Commons	2.99B
Scientific Commons	0.84B
Economics Commons	0.11B

Gienapp, Schröder, Schweter et al. "The German Commons — 154 Billion Tokens of Openly Licensed Text for German Language Models." arXiv:2510.13996, 2025.



OWS-Curlie-2025 — Web Search Test Collection

- **20.7 million** documents from 6 months of OWI crawl data (March–August 2025)
- English documents from the Curlie subcollection (websites labeled in the Curlie directory)
- Three evaluation splits with student-provided topics and relevance assessments
- Jina V3 embeddings, ClFF indexes for sparse retrieval
- Accessible via owilix CLI and ir-datasets-owi Python package

Split	Documents	Embeddings	Qrels
all	20,746,181	--	--
radboud-val	63,639	yes	yes
radboud	90,852	yes	--
jena-kassel	77,382	yes	--

openwebindex.eu/corpora · License: OWIL v1.0



Using the Open Web Index — Ideas

Potential Application Domains

Some application scenarios discussed at the start of OpenWebSearch.eu are now in reach:

Web-Analytics for Official Statistics

Eurostat: enterprise information extraction from legal imprints — who is doing business where?



Destatis: statistical data extraction from web sources to complement national statistics



Media Monitoring & Intelligence

JRC Media Monitor: monitoring media ecosystems and supply chain signals at European scale (*in start phase*)



A Live Use Case: EU Toy Safety & the Open Web Index

Can the OWI be used for EU product market surveillance? In preparation of the talk I ran a live exploratory analysis using the OUR²S structured-data API plus some AI Agents.

A Live Use Case: EU Toy Safety & the Open Web Index

Can the OWI be used for EU product market surveillance? In preparation of the talk I ran a live exploratory analysis using the OUR²S structured-data API plus some AI Agents.

"Make an analysis of the OURS API for the OWI corpus: find children's toys and assess how they appear in the dataset — particularly those potentially violating EU regulations."

Prompt used in JetBrains's Junie AI System with Gemini 3.1 Flash

A Live Use Case: EU Toy Safety & the Open Web Index

Can the OWI be used for EU product market surveillance? In preparation of the talk I ran a live exploratory analysis using the OUR²S structured-data API plus some AI Agents.

"Make an analysis of the OURS API for the OWI corpus: find children's toys and assess how they appear in the dataset — particularly those potentially violating EU regulations."

Prompt used in Jetbrain's Junie AI System with Gemini 3.1 Flash

The corpus

Schema	Documents
Web pages	310 M
Structured entities	137.5 M
Unique hosts	10 M

Entities Found

- **27 M Products** — largest single type
- 26 M Organisations
- 18 M Articles
- 3.8 M Local Businesses
- 2.8 M Events

→ Products are the single largest entity type — directly usable for market analysis.

A Live Use Case: Products at Scale — Toys in the Index

Pricing & geography

- ~15 M entities carry price data
- **3 M EUR-priced** products → significant EU market coverage
- Also: USD 2.2 M · UAH 1.9 M · BRL 1.2 M

Toy query results (live)

- "children toy" → **102,530 products**
- With pricing → **51,384** (50 % have structured prices)
- "Spielzeug Kinder" → **7,888** German results

Sample items found

- *"1:24 die-cast metal sports car"* — \$90.99
- *"Kids Microscope 1000× Magnifier"* — \$59.00
- *"48-teiliges Einhorn Prinzessin Set"* — €8.44

Structured fields available per product

- Name, brand, GTIN/EAN
- Price, currency, availability
- Description, images, ratings

~50 % of toy products include structured price data, reflecting real e-commerce markup adoption.

~92 % of all entities **lack location data** — a significant data-quality finding.

Whole analysis took 20 minutes in parallel to working on this presentation.

A Live Use Case: Conclusion

What the OWI provides today

- Structured Microdata with some features
 - GTIN/EAN presence assessment
 - Brand & manufacturer data
- Completeness and parsing ok, but should be improved

What would be missing for full EU compliance

- CE marking (*not in schema.org*)
- Age suitability / safety warnings
- Material composition
- ...

Could be added to parser

The OWI is a **foundation for such a use case** — Without an aggregated web index, such an analysis would not be possible. Of course you can use commercial search engines for such tasks, but they are not transparent and you don't know what data they use.

--

Other Research Directions

Next-Generation Retrieval

- Federated RAG across distributed OWI slices and 3rd-party corpora
- Agentic search orchestration for autonomous knowledge discovery
- Multimodal indexing: images, tables, charts alongside text
- Temporal web search: tracking how the web changes over time

AI & LLM Infrastructure

- Domain-specific corpora for sovereign European LLM training
- Retrieval-Augmented Generation benchmarks on OWI data
- Open evaluation infrastructure for IR and NLP research
- AI Agents & Search - actively pursued by Uni Passau
- **Multi-loop LLM orchestration:** classify incoming tasks, represent task types as vectors, and use cross-encoder-style scoring to route or sequence specialised LLM loops for the identified tasks
- **Stacked Kohonen embeddings:** map each sentence token to the top-n matching SOM codebook vectors, use their grid positions as explicit embedding coordinates, and condition the next token's SOM lookup on the previous token's match to form a joint sequence vector
- **Capability allocation across parameter space:** which capabilities — knowledge, reasoning, behaviour — should live in *model parameters* vs. *indexed memory* vs. *interaction traces*? Open design question; envisioned



figure maps capability types across the parameter / retrieval / trace space (parked from the TPC26 "Open System Question" slide)

Societal & Policy Applications

- **Information quality monitoring:** detecting coordinated inauthentic behaviour, SEO spam
- **Digital literacy tooling:** transparent, inspectable search for education
- **Regulatory compliance:** evidence-based auditing of search engine behaviour
- **Digital Sovereignty:** European alternative to US-dominated search infrastructure

The OWI is an open infrastructure — we invite researchers, businesses, and public institutions to build on it

We are looking to be part of interesting research projects to advance the OWI.

We hope to build an open ecosystem that puts open, fair and unbiased information exchange and societal values before commercial interest.

Conclusion

In Summary

What we built

- A **federated, open web index** — 1.1 PB raw data, 810+ daily datasets, running 24/7 across European HPC centres
- A **full open-source pipeline** — crawl → preprocess → index → distribute, all reproducible
- A **growing ecosystem** — search apps, RAG systems, datasets, and analytics tools built on the OWI

Why it matters

- Web data and web search are **critical infrastructure for Europe's digital sovereignty**
- The OWI **reduces upfront costs** for tapping the web as resource
- Open, transparent, and FAIR — enabling **reproducible research** and **open innovation**
- **Collaborative** approach to build and maintain the infrastructure

OpenWebSearch.eu proves that a European open web index is technically feasible, legally navigable, and economically viable

Before we come to the Questions

We are looking for ...

Helping Hands

- **Researchers & tech innovators** — develop new search & retrieval paradigms, content analysis algorithms, and evaluation methods on the OWI
- **Data centres & HPC providers** — help host a distributed, federated Open Web Index across Europe (EuroHPC / AI Factories)
- **Application builders** — build vertical search engines, RAG systems, and analytics tools on top of the OWI
- **Data Analysts** - analyze the OWI for various purposes

Helping Funds

- **Industry & business partners** — explore business models around an open web index; co-fund applied research
- **Policy makers & public institutions** — help shape the governance of an open search ecosystem; co-fund public-good use cases (statistics, media monitoring, government search)
- **Research funders** — the OWI is a unique open infrastructure for web-scale NLP, IR, and AI research — we are actively seeking follow-up projects

Get in touch

Website	openwebsearch.eu
Dashboard	openwebindex.eu
Code Repository	https://opencode.it4i.eu/openwebsearcheu-public/
Community	openwebsearch.eu/community
Join	join@openwebsearch.org
General	community@openwebsearch.eu

Thank you — Questions, Comments, Feedback very welcome!