

Countering Disinformation with
SOURCE
Crawling
Classification
Contextualization
Community-Building



From the Open Web to the Analyst's Desk

Scalable, Open and Comprehensive Recognition of Disinformation Campaigns on the Web

Prof. Dr. Michael Granitzer

Chair of Data Science, University of Passau · Coordinator SOURCE



Funded by

Partners



U N I K A S S E L
V E R S I T Ä T



1 · SOURCE in brief

DESINFORMATION ERKENNEN. DEMOKRATIE SCHÜTZEN.

SOURCE at a glance

SOURCE builds a freely accessible **early-warning platform** against disinformation: web and social-media content is found, analysed and prepared for countermeasures.

- analyses content by **where it comes from** and by **what it says**
- detects AI-generated material, manipulative language and factual errors
- builds on the **Open Web Index** of OpenWebSearch.eu
- open data, open interfaces, and a community around them



3 years

April 2026 – March 2029

€ 2.5 M


funded by the BMFTR

5 partners

research, infrastructure, civil society

source-project.de/projekt

Disinformation is a campaign, not a post



A single false post is easy to check and cheap to replace.

What does the damage is the **campaign** behind it:

- the same story, slightly changed, on many sites and channels
- accounts that amplify each other, across platforms
- timing that follows a plan, not the news

To counter a campaign you have to see the pattern, and describe it in a language others can act on.


The consortium




Partner	Focus in SOURCE
University of Passau · Data Science	coordination; source-based detection: finding and tracking content on the web, actors and networks
University of Kassel · Deep Semantic Learning	content-based detection: AI-generated text, stylometry and authorship, evaluation
Open Search Foundation	building and connecting the communities around the tools; science communication
LRZ · Leibniz Supercomputing Centre	compute, cloud and storage for the SOURCE database; federated infrastructure
Alliance4Europe	use cases and expert studies; the Counter Disinformation Network; integration into DISARM

2 · Planned System

FROM THE OPEN WEB TO THE ANALYST, AND BACK



From the open web to the analyst, and back



Infrastructure: compute, cloud and storage for the SOURCE database, on the OpenWebSearch.eu federation


Who works on which link

Component	What happens there	Partner
Open Web Index	daily web crawl, extended to social media; hashes that recognise reused text and images	Passau · LRZ · OSF
Actor analysis	where content comes from, how it spreads, which accounts act together	Passau · Kassel
Content analysis	AI-generated text, stylometry and authorship, manipulative rhetoric	Kassel
Indexing & search	web and social-media posts indexed with their analysis into the OURRS search service (ourrs.eu); user contributions are not indexed	Passau
Campaign analysis	the SOURCE campaign analysis: patterns, networks, trends across the index	Passau
Reverse image search	a separate service: find where an image appeared before	to be agreed
 SOURCE harness	the analyst's AI assistant: an LLM agent that reads social media and speaks DISARM/DISRUPT	Passau · Kassel · A4E

Component	What happens there	Partner
Analysts	use cases, expert studies, evaluation of what works	A4E
Community	contributions, shared tasks, outreach to fact-checkers and the public	OSF · all

3 · The SOURCE harness

A RESEARCH PROTOTYPE: AN AI ASSISTANT THAT WORKS FOR THE ANALYST, ON THE ANALYST'S MACHINE



What is a harness?

A language model on its own can only write text. A **harness** is everything around it that turns it into a useful and safe assistant:

- the **tools** it may use: search, a browser, DISARM
- the **rules** it must follow, and a guard that enforces them
- a **memory**: the research wiki it keeps
- a **person** who decides, reviews and approves

Think of the model as the engine, and the harness as the car around it: steering, brakes, and a driver.


Analyst SOURCE Desk, in the browser

Agent an open-weight model on university servers

Tools search · read-only browser · DISARM · wiki

Evidence every page read is kept as proof

An investigation, step by step



The assistant does the legwork: searching, reading, filing, cross-referencing. The **judgement** stays with the analyst.

DISARM / DISRUPT: a shared language for influence operations

DISARM is an open framework that describes disinformation campaigns the way cybersecurity describes attacks:

- **Red**: what the attacker does, in phases, tactics and techniques ("create fake accounts", "amplify with bots", ...)
- **Blue**: what defenders can do against each technique

The harness ships DISARM as data and tools: look a technique up, find the countermeasures against it, and find **past incidents** that showed the same pattern.

DISRUPT, Alliance4Europe's successor to DISARM, adds the operational side: a community toolbox of measures to stop an operation, from preparing evidence to disruption and mitigation.

391

techniques in 16 tactics

140

countermeasures

129

documented incidents

DISARM Foundation, DISARMframeworks (CC BY-SA 4.0) · Alliance4Europe, DISRUPT Framework, alliance4europe.eu/counter-disinfo/disruption-framework

4 · The SOURCE Desk

SCREENSHOTS OF THE RESEARCH PROTOTYPE, AND A DEMO

The SOURCE Desk

SOURCE Desk

YOUR RESEARCH HARNESS

ready

Restart SOURCE

INVESTIGATIONS

+ New

None yet. An investigation gets its own chats, goals and wiki area.

OTHER CHATS

+ New

New chat — no message yet

new

Monitoring Uni Passau

3 d ago

good morning. I want to continue working on the putin russia focus that osintrader analysed. Can you look that up in the wiki?

3 d ago

use owimed corpus to search social media entries about putin and rjussia

4 d ago

(an empty chat)

4 d ago

can you analyse recent news in ourrs?

4 d ago

Reply with exactly one word: ok

4 d ago

can you find recent social media data?

4 d ago

Reply with exactly: PI-OK

6 d ago

Chat

Wiki

Summary

Settings

New chat — no message yet

Rename

SOURCE

Your research harness. Ask it to look into a claim, a hashtag or an account — it searches owimed, reads the platforms read-only, and files what it finds in your wiki.

Ask SOURCE...

Enter to send - Shift+Enter for a new line

Stop

Steer now

Send

GOALS

No goal in this chat yet — SOURCE opens one when it starts an investigation.

Still open from earlier chats

Search owimed corpus for social media entries about Putin and Russia, with focus on @osintradar Bluesky account

WORKING

zauberberg analysis complete: profile captured, ~40 posts analyzed across 5 content categories, all topics cross-referenced to reliable news sources in wiki/actors/zauberberg.md. Network page updated. Next: user direction on coordination analysis, cross-platform comparison, or final hypothesis validation. - updated 3 d ago

The room's open goals not in a room yet

SETUP

ourrs — owimed search

SIGNED IN

signed in — connects on the first search

Sign in again

Relay

NOT SIGNED IN

logs in to yoars-api.ourrs.projects.innkube.fim.uni-passau.de (release)

Log in

Agent

NONE YET

the name other researchers see in the room

Room

NONE

an invitation link or a room key from the room's owner

Platform logins

YOUR ACCOUNTS

source: Refused: a login window is already open — the researcher finishes that one first

14

Discover: narratives from agentx

FOR AGENTS & ANALYSTS - API V1

Discover. Read the evidence. Follow the narrative.

Connect your agent harness to structured narratives, source evidence and saved annotations. Every result links back to the dossier so an analyst can examine the same material.

01 / CONNECT YOUR HARNESS

One contract, both ways to read it.

Import this OpenAPI 3.1 URL into a compatible harness or generate a client from its JSON schema. Operation IDs are stable: `discoverNarratives`, `getNarrative`, `getNarrativeEvidence` and `exportNarrative`.

https://agentx (dev)/api/openapi.json

Copy URL

[OpenAPI JSON ↗](#) [Discovery manifest ↗](#) [Agent instructions ↗](#)

PUBLIC READ ACCESS

No API key required

JSON over HTTPS on AgentX. The local dashboard serves the same API on its own address.

[Inspect snapshot & capabilities ↗](#)

A discovery workflow

1. Read snapshot status and available sources.
2. Find recent narratives, or filter by country, review status or keywords.
3. Follow the returned evidence and export links.
4. Pass the analyst link to a person for review.

agentx (development instance), discovery API page, October 2026.

agentx lets an agent discover narratives and read the evidence behind them. The harness uses it as a tool, `source_discover`.

- read the snapshot status and available sources
- find recent narratives, by country, review status or keywords
- follow the evidence and export links
- pass the analyst link to a person for review

Answers are **leads**, not findings: the harness treats them as untrusted.

What it writes: a research wiki

SCHEMA.md

EDITABLE

Edit

How this wiki works

This folder is a research wiki about possible disinformation campaigns, kept by an agent for the researcher who runs it. It follows Karpathy's "LLM wiki" pattern: **raw sources that never change, pages the agent maintains on top of them, and a check that keeps the two honest.** Open the folder in Obsidian (or any markdown editor) to read it; links, backlinks and the graph work there.

If you edit pages in Obsidian, set it to write the links this wiki uses: Settings → Files & Links → turn Use `[[Wikilinks]]` **off**, and set **"New link format" to "Relative path to file"**. `lint` reports a `[[wikilink]]`, because nothing outside Obsidian can follow one.

Layout

path	what	who writes it
<code>raw/</code>	captures: what was read, exactly, with URL, time and SHA-256	<code>source_capture / source_snapshot / source_search / source_fetch</code> only — never edited
<code>annotations /</code>	which DISARM technique a passage of a capture shows, and why — one file per capture, beside <code>raw/</code>	<code>disarm_annotate</code> only — never edited
<code>wiki/campaigns/</code>	one page per suspected coordinated campaign	the agent
<code>wiki/actors /</code>	accounts, channels, outlets, organisations	the agent
<code>wiki/narratives/</code>	recurring claims and framings	the agent
<code>wiki/incidents/</code>	one bounded event: a burst, a coordinated push, a hoax	the agent
<code>wiki/<section>/index.md</code>	one line per page in that section	the agent, with every new page
<code>disarm/</code>	the DISARM framework, imported at a pinned commit	nobody — re-imported, never edited
<code>index.md</code>	the entry point	the agent keeps the counts and highlights current

What it writes: a research wiki

wiki/actors/osintradar.md

EDITABLE

Edit

kind	actor
title	OSINTRadar (@osintradar.bsky.social)
status	hypothesis
updated	2026-09-27
disarm	["T0119.002", "T0143"]

OSINTRadar (@osintradar.bsky.social)

Summary

OSINTRadar is a Bluesky account with 11.7K followers and 12.6K posts that specializes in open-source intelligence reporting on the Russia-Ukraine war, Middle East, Europe, East Asia, and South Caucasus. The account posts frequent, time-stamped updates covering diplomatic statements, military operations, sanctions, and geopolitical analysis — primarily sourced from mainstream outlets (Reuters, Wall Street Journal, Washington Post, Politico, CBS News) and official government statements. The account also maintains a Telegram channel (@t.me/OSINTRadar05) and accepts donations via PayPal and Ko-fi.

Profile

Platform: Bluesky (@osintradar.bsky.social)
Followers: 11.7K
Following: 803
Total posts: 12.6K
Bio: "RussiaUkraine war | Middle east | Europe | East asia | South Caucasus"
Secondary channel: Telegram (t.me/OSINTRadar05)
Funding: Donations via PayPal (paypal.me/osintradar) and Ko-fi

Recent Posts (captured 2026-09-26)

The profile snapshot below shows posts from approximately the last 24–48 hours. Key Russia/Ukraine-related content:

Diplomatic & Political Statements

"Putin only needs to declare victory and go home" — attributed to Kit Kellogg, with analysis that Russia is not capable of capturing all of Ukraine ([capture](#))
Zelenskyy's call for Trump to activate the Lindsey Graham Act against Russia, plus additional European sanctions ([capture](#))
Sergey Lavrov stated nothing new was achieved during meeting with German Foreign Minister ([capture](#))
Lavrov at UN: "We see signs of neo-Nazism spreading among EU and NATO member states, including in Germany" ([capture](#))
Lavrov at UN demanding release of Maduro and his wife ([capture](#))



What it writes: a research wiki

wiki/narratives/osintradar-russia-ukraine-aggregation.md

EDITABLEEdit

kindnarrative

titleOSINTRadar — Russia/Ukraine Geopolitical News Aggregation

aliases[]

platforms["bluesky", "telegram"]

actorOSINTRadar

created_at2026-09-27

updated2026-09-27

statushypothesis

sources

DISARM[]

OSINTRadar — Russia/Ukraine Geopolitical News Aggregation

Summary

OSINTRadar (@osintradar.bsky.social / @OSINTRadar05 on Telegram) is a geopolitical news aggregation account with approximately 11,700 followers on Bluesky and an associated Telegram channel. It posts 100+ times per day, curating and reposting news from mainstream sources (Reuters, Wall Street Journal, Washington Post, CBS News, Politico) with minimal editorial commentary.

Content Profile

The account's content breaks down into several categories based on a 24-hour sample captured on 2026-09-26:

Category	Approx. Frequency	Examples
Russia-Ukraine military updates	Very high	Ukrainian drone strikes on Russian refineries, territorial advances, Russian strikes on Ukrainian cities
Russian official statements (quoted)	High	Putin quotes, Lavrov UN statements (neo-Nazism in EU, "cannon fodder" narrative)
Western government/military aid	High	USAID contracts, Patriot missiles, US-Poland base
Trump administration news	Medium	Iran ceasefire rejection, China-Xi diplomacy, G20 invite controversy
Other geopolitics	Medium	Japan-Russia tensions, Middle East (Houthis), Cuba
Fundraising solicitations	Low (1-2/day)	PayPal, Ko-fi donation links
Cross-platform promotion	Low	"Join me on Telegram" pinned post



It learns, under supervision

good morning. I want to continue working on the putin russia focus that osintrader analysed. Can you look that up in the wiki? [Rename](#)

Bluesky reads	source_search(account=...) works, source_fetch(profile) works	8s rate limit between calls; feeds are refused; only public profile/post pages	Rate-limited source_fetch on 3 accounts in one session
DISARM mapping	disarm_guide phase/tactic scoping, disarm_lookup definitions	Quote-matching fails when captures use smart quotes (") vs ASCII (")	6 annotation retries per technique before finding correct offset
Wiki writes	Schema is well-defined	Frontmatter fields (kind, title), capture paths (../depth), unsourced claims, invalid sections, DISARM code lint rules	9 lint findings on first try; needed 3 rounds of fixes
Cross-referencing	Manual topic extraction → source mapping was thorough	No tool to help match content topics to reliable news anchors	Purely manual

Skill Proposals

1. Update mcp-ourrs — Add Search Fallback Patterns

The current skill covers the serialization bug but doesn't address **relevance filtering** (nuggets returning semantically matched but topically irrelevant results) or **Bluesky-specific search patterns**.

2. disarm-annotation — Systematic DISARM Mapping

This skill codifies the DISARM evaluation → lookup → quote extraction → annotation workflow, with the offset-based quote matching pattern.

3. wiki-ingest — Wiki Page Creation & Lint-First Pattern

This skill codifies the correct frontmatter, section structure, and capture-path patterns that pass lint on first try.

4. bluesky-analysis — Bluesky Account Behavioral Analysis

This skill codifies the analysis patterns we developed for Bluesky accounts: posting cadence, engagement metrics, amplification networks, and cross-platform behavior.

The index has an invalid status for zauberberg. Let me verify and fix:

Analysis of Our Work & Skill Proposals

Here's what I've done:

Retrospective

Categor	What Worked Well	What Caused Retries	Frequencv
---------	------------------	---------------------	-----------

After an investigation the agent reviews its own work and proposes skills.

Appearance

The font, size and line spacing of this page. They apply at once, need no restart, and are kept in this browser only.

FONT

System

SIZE

Normal

LINE SPACING

Normal

Skills

What SOURCE may follow beyond its brief. SOURCE can draft a skill (skills/<name>/SKILL.md); only you approve it. An approval covers exactly what you saw: an edit makes it a draft again. An MCP override (mcp-<server>) replaces the harness's own help for that server. Changes apply when SOURCE restarts.

bluesky-analysis

APPROVED

Systematic analysis of Bluesky accounts: posting patterns, engagement metrics, amplification networks, cross-platform behavior, and critic mapping

Review

Edit

Revoke

disarm-annotation

APPROVED

Systematic DISARM technique mapping: phase evaluation, definition lookup, offset-based quote extraction, and annotation — avoiding Unicode/smart-quote mismatches

Review

Edit

Revoke

mcp-ourrs — MCP override for ourrs

APPROVED

Optimized search across SOURCE's ourrs MCP server with corpus routing, serialization-workaround, and relevance fallbacks

Review

Edit

Revoke

wiki-ingest

APPROVED

Wiki page creation that passes source_wiki_lint on first try — correct frontmatter, sections, capture paths, DISARM annotations, and section indices

Review

Edit

Revoke

The analyst reviews each skill and approves, edits or revokes it.



Live demo

Open an investigation · ask about an account or a claim · watch it search, read and file · open the wiki page it wrote

Where we are, and what comes next

A research prototype

A first technical idea, running in the lab:

- the Desk, a read-only browser and its guard
- search over web and social-media posts via ourrs.eu
- DISARM tools, annotations and a research wiki
- skills proposed by the agent, approved by people

Next: with the community

- discuss the approach and its rules with analysts; it needs the community's approval
- first trials with Alliance4Europe analysts
- DISARM/DISRUPT guiding questions reviewed by experts
- tools from the partners: coordination detection, image fingerprints, link provenance



source-project.de · info@source-projekt.de

Appendix

DESIGN RULES OF THE PROTOTYPE, AND A FIRST INTERNAL TEST

Design rules of the prototype

BUILT INTO THE HARNESS TODAY; TO BE DISCUSSED AND AGREED WITH THE COMMUNITY

- ✓ **Every claim cites a capture.** A check refuses wiki pages that do not.
- ✓ **DISARM codes come from the DISARM tools,** never from the model's memory.
- ✓ **Runs as you, on your machine.** What it collects stays there.
- ✓ **New skills need a human.** The agent proposes; only the analyst approves.
- ✗ **No posting, liking, following or messaging.** The browser is read-only.
- ✗ **No passwords.** The analyst logs in by hand; the agent never sees credentials.
- ✗ **No bulk scraping.** A reading pace per platform: at least 8 s between pages, 400 a day.
- ✗ **No attribution by resemblance.** Looking like a known incident is not proof.

SOURCE harness: agent brief (spec/brief.md), README, design decisions D1, D7, D32 (yoars, lane SOURCE, September 2026)

Does it help? A first internal test

An automated benchmark: DISARM documents example incidents with quoted source reporting. We gave the agent **62 of these descriptions**, with the technique codes removed, and asked which techniques each one shows.

Model (open weights)	Exact technique	Technique family
Gemma 4 31B	0.33	0.42
Qwen 3.6 35B	0.29	0.36
Qwen3-Next 80B	0.21	0.27

Micro F1 against 309 gold codes, 26 September 2026.

0

invented codes, in every run

Preliminary, not yet validated by analysts. About a third right is a **helper**, not a judge; the guiding questions go to experts next, then we measure again.